

Foreword

For more than twenty years, ontology engineering has carried a quiet reputation: powerful in theory, hard to reach in practice. Michael DeBellis' Pizza tutorial did more than almost any other resource to change that, by letting people learn semantics with their hands instead of only their heads. What Xiaoqi Zhao has done here is take that well-loved starting point and carry it all the way to where the work actually has to live, which is inside real systems, under real governance.

This book earns its subtitle. In a single progression it moves from the first class you build in Protégé to the question every enterprise eventually faces: how machine-readable meaning becomes something a system can act on safely. Along the way it never loses the teacher's discipline of one idea at a time.

I want to point readers in particular at Chapters 12 and 13, because their titles tell you something. They do not *only* say “object property characteristics” and “domain and range.” They say *governing* semantic relationship behavior and *governing* semantic boundaries. That word is the whole game. A functional or transitive property is not a checkbox in a tool. It is a rule about how a relationship is allowed to behave. A domain and a range are not decoration. They are a statement about what may legitimately connect to what. Xiaoqi teaches these as governance, and that is exactly the right approach.

Chapter 13 also does something I rarely see in an introductory text: it tells the truth about a sharp edge. OWL domain and range do not reject a bad assertion the way a database constraint would. The reasoner infers, under an open-world assumption, rather than refuses, and as the chapter notes, this surprises newcomers who expected enforcement. That surprise is where academic ontology meets the enterprise information system. The moment data leaves the editor and travels between systems, a boundary that only infers is not yet a boundary you can trust. It has to be able to refuse what does not belong, deterministically, with its rules traveling alongside the data rather than living in a separate document or a separate dashboard.

That is the work my colleagues and I have spent years on in the Semantic Data Charter. SDC treats meaning and the constraint definitions as shareable, bound assets. They are one inseparable data payload, so admissibility can be checked at the point of use, in the open, against published specifications. Readers will recognize this as related to the layer Xiaoqi names Γ in his EKA tuple: the governance layer that turns a clever model into a trustworthy one. EKA is the

architecture that orchestrates the pieces; a substrate like SDC is one open, standards-based way to make that governance layer real in production. The two meet exactly where this book spends its most careful chapters.

There is a larger reason to welcome a book like this now. Ontology engineering is becoming a core architectural capability, the way databases and APIs did before it, and the field is poorly served when it stays split between texts that are rigorous but unreachable and tutorials that are reachable but shallow. Xiaoqi has taken the harder middle path: technically honest, professionally structured, and grounded in the standards that real organizations have to answer to. That is the bridge between theory and practice, and we need more people building it.

Read this book with Protégé open. Do the exercises. And when you reach those governing chapters, remember that the discipline you are learning on a pizza is the same discipline that will one day decide whether a clinical, financial, or regulatory system is allowed to act. As Xiaoqi writes, meaningful intelligence depends on meaningful boundaries. This book teaches you how to build them.

Timothy W. Cook
Founder, Axius SDC, Inc.
Creator, the Semantic Data Charter (SDC)
June 2026