

Executable Knowledge Architecture (EKA)

Mastering Ontology Engineering with Protégé and Pizza.owl

Table of Content

- [From Ontology Learning to Executable Knowledge Architecture](#)
- [Why the Pizza.owl Tutorial Matters](#)
- [Why This Book?](#)
- [About This Book](#)
- [Why Ontology Matters Today](#)
- [Ontology Within the EKA Framework](#)
- [This Book Is NOT Only About Pizza](#)
- [Who This Book Is For](#)
- [What Makes This Book Different](#)
- [How to Use This Book - The Learning Journey](#)
- [Structural Integrity](#)
- [Acknowledgements](#)
 - [Our Intellectual Heritage & Licensing](#)
 - [Special Thanks](#)
- [Related Resources](#)
- [Final Thoughts](#)

From Ontology Learning to Executable Knowledge Architecture

The Semantic Web has existed for more than two decades, yet for many architects, engineers, analysts, and AI practitioners, ontology engineering still feels abstract, academic, or difficult to approach systematically.

One of the reasons is simple:

Most ontology learning materials focus either too heavily on theory or too narrowly on tool operations.

As a result, many learners can:

- remember OWL terminology,
- click through Protégé demonstrations,
- or understand isolated semantic concepts,

but still struggle to answer the most important question:

Why does ontology engineering actually matter in modern enterprise and AI systems?

This book was created to answer that question.

Why the Pizza.owl Tutorial Matters

Among all ontology learning materials ever created, Michael DeBellis' Pizza OWL tutorial remains one of the most influential and practical introductions to OWL ontology engineering.

Using a familiar pizza domain, the tutorial teaches learners how to:

- think semantically,
- construct ontology hierarchies,
- model relationships,
- apply OWL logic (rules), and
- and understand reasoning behavior inside Protégé

What makes the Pizza tutorial special is that it gradually introduces semantic complexity in a highly structured manner.

Rather than overwhelming learners immediately with advanced ontology theory, it guides readers step by step through the evolution of a real semantic model.

That educational philosophy strongly influenced the structure of this book.

This eBook serves as the comprehensive companion guide to the [Protégé OWL Pizza Tutorial Hands-on Series](#)

Why This Book?

Most ontology materials focus either on abstract theory or narrow tool operations. This book bridges that gap by connecting semantic engineering to the **Executable Knowledge Architecture (EKA)** framework.

It is designed to help you engineer machine-understandable meaning for AI systems, enterprise knowledge platforms, and digital twins.

About This Book

This eBook was written as a companion knowledge guide to the YouTube playlist:

****Protégé OWL Pizza Tutorial Hands-on Series****

(<https://www.youtube.com/playlist?list=PL6DEHvciXKeUx4P32B3hKMK1t6mC8RhsW>)

Since publishing the Protégé OWL Pizza Tutorial series in 2023, I've realized how valuable practical ontology learning remains for architects and AI practitioners. What started as a hands-on walkthrough of Michael DeBellis' tutorial gradually became an important foundation for my EKA (Executable Knowledge Architecture) thinking — connecting ontology, knowledge graphs, and executable intelligence. The playlist may use pizza as the example domain, but its real purpose is helping learners develop semantic thinking and understand how machine-readable knowledge is engineered.

However, the goal is not merely to transcribe the videos.

Instead, this book expands the tutorial into a much deeper professional learning experience by combining:

- ontology engineering theory,
- hands-on Protégé practice,
- semantic modeling methodology,
- enterprise architectural thinking,
- and modern knowledge graph perspectives.

Every chapter corresponds closely to the learning progression of the original video series so readers can learn both visually and conceptually in parallel.

This synchronization is intentional.

Ontology engineering is best learned progressively.

Why Ontology Matters Today

We are entering a new era of intelligent systems.

Modern enterprises increasingly rely on:

- AI systems,
- semantic search,
- enterprise knowledge graphs,
- digital twins,
- intelligent automation,
- contextual reasoning,
- and machine-readable knowledge.

Traditional architectures are built purely on:

- servers and databases,
- documents,
- diagrams
- and APIs

are no longer sufficient for advanced semantic intelligence.

The missing layer here is:

Meaning

Ontology engineering provides that layer.

OWL ontologies allow organizations to formally define:

- concepts,
- relationships,
- constraints,
- semantics,
- and inference logic

in a machine-readable and machine-processable form.

This transforms static information into executable knowledge.


Ontology Within the EKA Framework

This book also introduces ontology engineering through the lens of the EKA (Executable Knowledge Architecture) framework.

EKA represents an architectural approach for transforming enterprise knowledge into executable intelligence systems.

Within EKA, the implementation roadmap is:

The EKA Implementation Roadmap



Ontology occupies the critical semantic transformation layer.

This is where:

- visual architecture,
- conceptual models,
- and enterprise abstractions

become:

- machine-readable semantics,
- reasoning structures,
- and graph-ready knowledge.

Protégé therefore becomes more than just an ontology editor.

It becomes an engineering environment for executable semantics.

This perspective fundamentally changes how ontology engineering should be understood.

This Book Is NOT Only About Pizza

Although the tutorial domain uses pizza examples, the real purpose of the tutorial is much larger.

The Pizza ontology teaches foundational semantic engineering patterns that scale into real-world domains such as:

- enterprise architecture,
- healthcare,
- finance,
- manufacturing,
- government,
- cybersecurity,

- digital twins,
- and AI knowledge systems.

The same semantics principles used to model:

`(Pizza)-[:hasTopping]->(CheeseTopping)`

can later scale into:

`(Application)-[:supportsCapability]->(businessFunction)`

or:

`(Device)-[:connectedTo]->(Sensor)`

Ontology engineering is domain-independent!

The Pizza example simply provides an approachable learning environment.

Who This Book Is For

This book is intended for:

- Enterprise Architects
- Solution Architects
- Data Architects
- Knowledge Engineers
- AI Engineers
- Semantic Web Learners
- Knowledge Graph Practitioners
- Protégé Beginners
- Meta-Modeling Practitioners
- EKA Learners
- Digital Transformation Professionals

No prior ontology experience is required.

However, reader with backgrounds in:

- UML,
- ER Modeling,
- Enterprise Architecture,
- Databases,
- Graph Technologies,
- or Software Engineering

will often recognize important conceptual parallels throughout the book.

What Makes This Book Different

Most ontology books fall into one of two extremes:

Style	Problem
Highly academic	Difficult for practitioners
Pure tool tutorials	Lack conceptual depth

This book intentionally bridges both worlds.

The objective is to create A Practical Ontology Engineering Handbook that remains:

- technically rigorous,
- professionally structured,
- architecturally meaningful,
- and implementation-oriented.

The book therefore combines:

- OWL theory,
- Protégé operations,
- semantic engineering methodology,
- EKA architecture thinking,
- and knowledge graph perspectives

into a single progressive learning journey.

How to Use This Book - The Learning Journey

The recommended learning approach is:

1. **Watch:** The corresponding YouTube video chapter.
2. **Read:** The matching eBook chapter for deeper conceptual context.
3. **Practice:** Manual Protégé exercises and hands-on semantic modeling.
4. **Reflect:** Map these semantic principles to your own enterprise architecture.

Ontology engineering is not mastered through passive reading alone.

It requires semantic thinking practice.

The goal is not merely learning Protégé.

The goal is learning how to engineer machine-understandable meaning.

Structural Integrity

- **Foreword:** By Timothy W. Cook (Founder of SDC).
- **Core Curriculum:** Covers everything from basic class hierarchies and object properties to advanced SHACL, SPARQL, and reasoning behavior.
- **Error Tracker:** An included log to help you troubleshoot common modeling pitfalls.

Acknowledgements

This book has benefited greatly from the generous feedback, technical reviews, and professional discussions contributed by members of the ontology and semantic technologies community.

The full acknowledgements are maintained here: [ACKNOWLEDGEMENTS.md](#)

Our Intellectual Heritage & Licensing

This work is a direct descendant of the **Protégé 4 Tutorial (version 1.3)** by **Matthew Horridge**. We are deeply indebted to the original visionaries who developed the preceding versions and established the core teaching material: **Holger Knublauch, Alan Rector, Robert Stevens, Chris Wroe, Simon Jupp, Georgina Moulton, Nick Drummond, and Sebastian Brandt**.

In this modern adaptation, we have incorporated **Michael DeBellis's** revisions, which successfully bridged the transition to Protégé 5.5 and integrated advanced practices such as SWRL, SPARQL, and SHACL. Furthermore, we acknowledge the contributions of **Lorenz Buehmann, André Wolski, Dick Ooms, Colin Pilkington, Livia Pinera, Jans Aasman, Yan Xu, and the team at Franz Inc.**, whose work in applying these ontologies to real-world triple stores like AllegroGraph and Gruff has significantly shaped our current perspective on **Executable Knowledge Architecture (EKA)**.

All educational content in this eBook is licensed under **CC BY-SA 4.0**.

For full legal details, please see the `README.md` and `LICENSE.md` files in the project root.

Special Thanks

We extend our sincere gratitude to **Michael DeBellis** for creating one of the most influential practical OWL learning resources available to the ontology community. (Visit Michael's homepage here - <https://www.michaeldebellis.com/>). We also thank him for his meticulous technical audit for this series; his expert insights -- particularly regarding the crucial engineering distinction between modeling tool and persistent graph databases -- have been instrumental in ensuring this content meets the standards of modern production-grade environments.

We are also honored to feature a foreword and ongoing review by **Timothy Cook (Founder of SDC)**. His guidance and endorsement have been vital in validating our approach, helping to bridge the gap between traditional semantic modeling and the sophisticated demands of modern enterprise architecture / knowledge architecture.

Finally, we thank **Stanford University and the Protégé community** for their enduring commitment to advancing open ontology engineering tools and Semantic Web technologies.

By preserving this attribution chain, we honor the intellectual legacy of the individuals and organizations that have sustained this field. We stand on the shoulders of these giants, aiming to provide a bridge from traditional semantic modeling to the future of Executable Knowledge Architecture .

Related Resources

Official Resources:

- Protégé Official Website: <https://protege.stanford.edu/>
- Protégé Pizza Repository: https://github.com/yasenstar/protege_pizza/tree/main
- EKA Official Website: <https://xiaoqi.com/>

Learning Resources:

- YouTube Channel - Xiaoqi(Yasen) Zhao: <http://www.youtube.com/@yasenzhao>
- Udemy Course - Xiaoqi Zhao: <https://www.udemy.com/user/xiaoqi-zhao>

Final Thoughts

Ontology engineering is NO LONGER a niche academic discipline.

It is rapidly becoming part of the foundational infrastructure for:

- AI systems,
- enterprise knowledge platforms,
- semantic interoperability,
- and executable intelligence architectures.

Understanding ontology today is increasingly similar to understand databases or APIs twenty years ago.

It is becoming a core architectural capability.

The journey begins with a simple Pizza ontology.

But the destination is much larger:

Executable knowledge

Welcome to the World of Ontology Engineering!

"You" are the learner of this book, so while you're reading, I'll say to "you" directly instead of "learner" / "reader" from now on.

Last updated at: 2026-09-07