

From Semantic Foundations to
Executable Knowledge Architecture (EKA)



Mastering Ontology Engineering

with Protégé and Pizza.owl

VOLUME 1

Chapters 00-08

Ontology Foundations in
Protégé

XIAOQI ZHAO

Mastering Ontology Engineering with Protégé and Pizza.owl

– From Semantic Foundations to Executable Knowledge Architecture (EKA)



Table of Content

- [From Ontology Learning to Executable Knowledge Architecture](#)
- [Why the Pizza.owl Tutorial Matters](#)
- [Why This Book?](#)
- [About This Book](#)
- [Why Ontology Matters Today](#)
- [Ontology Within the EKA Framework](#)
- [This Book Is NOT Only About Pizza](#)
- [Who This Book Is For](#)
- [What Makes This Book Different](#)
- [How to Use This Book - The Learning Journey](#)

- [Structural Integrity](#)
- [Acknowledgements](#)
 - [Our Intellectual Heritage & Licensing](#)
 - [Special Thanks](#)
- [Related Resources](#)
- [Final Thoughts](#)

From Ontology Learning to Executable Knowledge Architecture

The Semantic Web has existed for more than two decades, yet for many architects, engineers, analysts, and AI practitioners, ontology engineering still feels abstract, academic, or difficult to approach systematically.

One of the reasons is simple:

Most ontology learning materials focus either too heavily on theory or too narrowly on tool operations.

As a result, many learners can:

- remember OWL terminology,
- click through Protégé demonstrations,
- or understand isolated semantic concepts,

but still struggle to answer the most important question:

Why does ontology engineering actually matter in modern enterprise and AI systems?

This book was created to answer that question.

Why the Pizza.owl Tutorial Matters

Among all ontology learning materials ever created, Michael DeBellis' Pizza OWL tutorial remains one of the most influential and practical introductions to OWL ontology engineering.

Using a familiar pizza domain, the tutorial teaches learners how to:

- think semantically,
- construct ontology hierarchies,

- model relationships,
- apply OWL logic (rules), and
- and understand reasoning behavior inside Protégé

What makes the Pizza tutorial special is that it gradually introduces semantic complexity in a highly structured manner.

Rather than overwhelming learners immediately with advanced ontology theory, it guides readers step by step through the evolution of a real semantic model.

That educational philosophy strongly influenced the structure of this book.

This eBook serves as the comprehensive companion guide to the [Protégé OWL Pizza Tutorial Hands-on Series](#)

Why This Book?

Most ontology materials focus either on abstract theory or narrow tool operations. This book bridges that gap by connecting semantic engineering to the **Executable Knowledge Architecture (EKA)** framework. It is designed to help you engineer machine-understandable meaning for AI systems, enterprise knowledge platforms, and digital twins.

About This Book

This eBook was written as a companion knowledge guide to the YouTube playlist:

****Protégé OWL Pizza Tutorial Hands-on Series****

(<https://www.youtube.com/playlist?list=PL6DEHvciXKeUx4P32B3hKMK1t6mC8RhsW>)

Since publishing the Protégé OWL Pizza Tutorial series in 2023, I've realized how valuable practical ontology learning remains for architects and AI practitioners. What started as a hands-on walkthrough of Michael DeBellis' tutorial gradually became an important foundation for my EKA (Executable Knowledge Architecture) thinking — connecting ontology, knowledge graphs, and executable intelligence. The playlist may use pizza as the example domain, but its real purpose is helping learners develop semantic thinking and understand how machine-readable knowledge is engineered.

However, the goal is not merely to transcribe the videos.

Instead, this book expands the tutorial into a much deeper professional learning experience by combining:

- ontology engineering theory,
- hands-on Protégé practice,
- semantic modeling methodology,
- enterprise architectural thinking,
- and modern knowledge graph perspectives.

Every chapter corresponds closely to the learning progression of the original video series so readers can learn both visually and conceptually in parallel.

This synchronization is intentional.

Ontology engineering is best learned progressively.

Why Ontology Matters Today

We are entering a new era of intelligent systems.

Modern enterprises increasingly rely on:

- AI systems,
- semantic search,
- enterprise knowledge graphs,
- digital twins,
- intelligent automation,
- contextual reasoning,
- and machine-readable knowledge.

Traditional architectures are built purely on:

- servers and databases,
- documents,
- diagrams
- and APIs

are no longer sufficient for advanced semantic intelligence.

The missing layer here is:

Meaning

Ontology engineering provides that layer.

OWL ontologies allow organizations to formally define:

- concepts,
- relationships,
- constraints,
- semantics,
- and inference logic

in a machine-readable and machine-processable form.

This transforms static information into executable knowledge.

Ontology Within the EKA Framework

This book also introduces ontology engineering through the lens of the EKA (Executable Knowledge Architecture) framework.

EKA represents an architectural approach for transforming enterprise knowledge into executable intelligence systems.

Within EKA, the implementation roadmap is:

The EKA Implementation Roadmap



Ontology occupies the critical semantic transformation layer.

This is where:

- visual architecture,
- conceptual models,
- and enterprise abstractions

become:

- machine-readable semantics,

- reasoning structures,
- and graph-ready knowledge.

Protégé therefore becomes more than just an ontology editor.

It becomes an engineering environment for executable semantics.

This perspective fundamentally changes how ontology engineering should be understood.

This Book Is NOT Only About Pizza

Although the tutorial domain uses pizza examples, the real purpose of the tutorial is much larger.

The Pizza ontology teaches foundational semantic engineering patterns that scale into real-world domains such as:

- enterprise architecture,
- healthcare,
- finance,
- manufacturing,
- government,
- cybersecurity,
- digital twins,
- and AI knowledge systems.

The same semantics principles used to model:

```
(Pizza)-[:hasTopping]->(CheeseTopping)
```

can later scale into:

```
(Application)-[:supportsCapability]->(businessFunction)
```

or:

```
(Device)-[:connectedTo]->(Sensor)
```

Ontology engineering is domain-independent!

The Pizza example simply provides an approachable learning environment.

Who This Book Is For

This book is intended for:

- Enterprise Architects
- Solution Architects
- Data Architects
- Knowledge Engineers
- AI Engineers
- Semantic Web Learners
- Knowledge Graph Practitioners
- Protégé Beginners
- Meta-Modeling Practitioners
- EKA Learners
- Digital Transformation Professionals

No prior ontology experience is required.

However, reader with backgrounds in:

- UML,
- ER Modeling,
- Enterprise Architecture,
- Databases,
- Graph Technologies,
- or Software Engineering

will often recognize important conceptual parallels throughout the book.

What Makes This Book Different

Most ontology books fall into one of two extremes:

Style	Problem
Highly academic	Difficult for practitioners
Pure tool tutorials	Lack conceptual depth

This book intentionally bridges both worlds.

The objective is to create A Practical Ontology Engineering Handbook that remains:

- technically rigorous,
- professionally structured,
- architecturally meaningful,
- and implementation-oriented.

The book therefore combines:

- OWL theory,
- Protégé operations,
- semantic engineering methodology,
- EKA architecture thinking,
- and knowledge graph perspectives

into a single progressive learning journey.

How to Use This Book - The Learning Journey

The recommended learning approach is:

1. **Watch:** The corresponding YouTube video chapter.
2. **Read:** The matching eBook chapter for deeper conceptual context.
3. **Practice:** Manual Protégé exercises and hands-on semantic modeling.
4. **Reflect:** Map these semantic principles to your own enterprise architecture.

Ontology engineering is not mastered through passive reading alone.

It requires semantic thinking practice.

The goal is not merely learning Protégé.

The goal is learning how to engineer machine-understandable meaning.

Structural Integrity

- **Foreword:** By Timothy W. Cook (Founder of SDC).
- **Core Curriculum:** Covers everything from basic class hierarchies and object properties to advanced SHACL, SPARQL, and reasoning behavior.
- **Error Tracker:** An included log to help you troubleshoot common modeling pitfalls.

Acknowledgements

This book has benefited greatly from the generous feedback, technical reviews, and professional discussions contributed by members of the ontology and semantic technologies community.

The full acknowledgements are maintained here: [ACKNOWLEDGEMENTS.md](#)

Our Intellectual Heritage & Licensing

This work is a direct descendant of the **Protégé 4 Tutorial (version 1.3)** by **Matthew Horridge**. We are deeply indebted to the original visionaries who developed the preceding versions and established the core teaching material: **Holger Knublauch, Alan Rector, Robert Stevens, Chris Wroe, Simon Jupp, Georgina Moulton, Nick Drummond, and Sebastian Brandt**.

In this modern adaptation, we have incorporated **Michael DeBellis's** revisions, which successfully bridged the transition to Protégé 5.5 and integrated advanced practices such as SWRL, SPARQL, and SHACL. Furthermore, we acknowledge the contributions of **Lorenz Buehmann, André Wolski, Dick Ooms, Colin Pilkington, Livia Pinera, Jans Aasman, Yan Xu, and the team at Franz Inc.**, whose work in applying these ontologies to real-world triple stores like AllegroGraph and Gruff has significantly shaped our current perspective on **Executable Knowledge Architecture (EKA)**.

All educational content in this eBook is licensed under **CC BY-SA 4.0**.

For full legal details, please see the `README.md` and `LICENSE.md` files in the project root.

Special Thanks

We extend our sincere gratitude to **Michael DeBellis** for creating one of the most influential practical OWL learning resources available to the ontology community. (Visit Michael's homepage here - <https://www.michaeldebellis.com/>). We also thank him for his meticulous technical audit for this series; his expert insights -- particularly regarding the crucial engineering distinction between modeling tool and persistent graph databases -- have been instrumental in ensuring this content meets the standards of modern production-grade environments.

We are also honored to feature a foreword and ongoing review by **Timothy Cook (Founder of SDC)**. His guidance and endorsement have been vital in validating our approach, helping to bridge the gap between traditional semantic modeling and the sophisticated demands of modern enterprise architecture / knowledge architecture.

Finally, we thank **Stanford University and the Protégé community** for their enduring commitment to advancing open ontology engineering tools and Semantic Web technologies.

By preserving this attribution chain, we honor the intellectual legacy of the individuals and organizations that have sustained this field. We stand on the shoulders of these giants, aiming to provide a bridge from traditional semantic modeling to the future of Executable Knowledge Architecture .

Related Resources

Official Resources:

- Protégé Official Website: <https://protege.stanford.edu/>
- Protégé Pizza Repository: https://github.com/yasenstar/protege_pizza/tree/main
- EKA Official Website: <https://xiaoqi.com/>

Learning Resources:

- YouTube Channel - Xiaoqi(Yasen) Zhao: <http://www.youtube.com/@yasenzhao>
- Udemy Course - Xiaoqi Zhao: <https://www.udemy.com/user/xiaoqi-zhao>

Final Thoughts

Ontology engineering is NO LONGER a niche academic discipline.

It is rapidly becoming part of the foundational infrastructure for:

- AI systems,
- enterprise knowledge platforms,
- semantic interoperability,
- and executable intelligence architectures.

Understanding ontology today is increasingly similar to understand databases or APIs twenty years ago.

It is becoming a core architectural capability.

The journey begins with a simple Pizza ontology.

But the destination is much larger:

Executable knowledge

Welcome to the World of Ontology Engineering!

"You" are the learner of this book, so while you're reading, I'll say to "you" directly instead of "learner" / "reader" from now on.

Last updated at: 2026-08-08

Acknowledgements

This project would not have reached its current level of maturity without the generosity of numerous members of the ontology, semantic web, enterprise architecture, and knowledge graph communities.

Many people contributed through technical discussions, detailed chapter review, engineering suggestions, issue reports, and encouragement throughout the writing process.

The following individuals deserve special recognition.

Michael DeBellis

Author of the original *Protégé 5 New OWL Pizza Tutorial*.

This book is built upon Michael's excellent tutorial, which has served for many years as one of the most practical introductions to OWL ontology engineering.

Beyond creating the original tutorial, Michael has kindly provided valuable feedback on the direction of this book and has been supportive of extending the tutorial into a broader engineering perspective through the Semantic Knowledge Development Lifecycle (SKDL) and Executable Knowledge Architecture (EKA).

Without his original educational work, this project would not exist.

Timothy W. Cook

Founder & CEO, Axius SDC, Inc.

Tim has generously reviewed multiple chapters of this book during its development, particularly those concerning semantic governance, ontology engineering practice, and enterprise-scale knowledge systems.

His feedback has helped strengthen several important engineering themes throughout the SKDL framework, including:

- treating ontology as a meaning-driven system rather than merely a validation system;
- distinguishing semantic inference from policy enforcement;
- clarifying the role of semantic governance within long-lived knowledge systems;
- improving the engineering narrative surrounding reusable semantic knowledge.

His perspective comes from extensive real-world work in semantic interoperability, healthcare data standards, and enterprise semantic governance.

I am deeply grateful for his thoughtful reviews, constructive criticism, and continued encouragement throughout the writing of this book.

The Ontology Community

Many members of the broader ontology and semantic web community have contributed indirectly through discussions, questions, publications, conference presentations, open-source software, and years of accumulated research.

Protégé, OWL, RDF, Description Logic, and the Semantic Web ecosystem are the result of decades of collaborative work by thousands of researches and practitioners.

This book stands on the collective foundation.

Readers - YOU

Finally, sincere thanks to every reader who as reported errors, suggested improvements, opened GitHub issues, starred the repository, shared the project with others, or simply spent time learning ontology engineering through this book.

Every question, correction, and discussion helps improve future editions.

Knowledge grows through reuse, collaboration, and continuous refinement -- the very principles promoted throughout this book.

Last updated at 2026-08-06

Legal & Licensing

This eBook is a derivative work of the "Pizza.owl Tutorial", a foundational resource for the ontology engineering community.

We are committed to transparency and the preservation of the intellectual legacy established by the original authors.

Licensing Terms

This work is governed by a dual-licencing strategy to balance educational openness with software best practices:

- **Documentation & Educational Content:** All textual content, diagrams, and educational explanations within this eBook are licensed under the **Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International (CC BY-NC-SA 4.0)** license.
 - *Requirement:* Any derivative works based on this educational content must maintain the non-commercial and Attribution-ShareAlike terms as required by the original authors.
- **Software & Scripts:** The associated code artifacts, scripts, and automation tools accompanying this eBook are provided under the **GNU General Public License v3.0 (GPL-3.0)**.

Attribution & Intellectual Heritage

This work stands on the shoulders of the community that created the original Protégé 4 Pizza Tutorial. We acknowledge the visionary contributions of **Matthew Horridge, Holger Knublauch, Alan Rector, Robert Stevens, Chris Wroe, Simon Jupp, Georgina Moulton, Nick Drummond, and Sebastian Brandt**.

We also extend our gratitude to **Michael DeBellis** for his invaluable modern revisions, technical audit, and his ongoing dedication to providing practical OWL learning resources for the ontology community.

For further details regarding the full attribution chain, the history of this project, or the specific licensing of individual files, please refer to the [README.md](#) and [LICENSE.md](#) files in the [official GitHub repository \(https://github.com/yasenstar/protege_pizza\)](https://github.com/yasenstar/protege_pizza).

Foreword from Timothy W. Cook

For more than twenty years, ontology engineering has carried a quiet reputation: powerful in theory, hard to reach in practice. Michael DeBellis' Pizza tutorial did more than almost any other resource to change that, by letting people learn semantics with their hands instead of only their heads. What Xiaoqi Zhao has done here is take that well-loved starting point and carry it all the way to where the work actually has to live, which is inside real systems, under real governance.

This book earns its subtitle. In a single progression it moves from the first class you build in Protégé to the question every enterprise eventually faces: how machine-readable meaning becomes something a system can act on safely. Along the way it never loses the teacher's discipline of one idea at a time.

I want to point readers in particular at Chapters 12 and 13, because their titles tell you something. They do not *only* say "object property characteristics" and "domain and range." They say *governing* semantic relationship behavior and *governing* semantic boundaries. That word is the whole game. A functional or transitive property is not a checkbox in a tool. It is a rule about how a relationship is allowed to behave. A domain and a range are not decoration. They are a statement about what may legitimately connect to what. Xiaoqi teaches these as governance, and that is exactly the right approach.

Chapter 13 also does something I rarely see in an introductory text: it tells the truth about a sharp edge. OWL domain and range do not reject a bad assertion the way a database constraint would. The reasoner infers, under an open-world assumption, rather than refuses, and as the chapter notes, this surprises newcomers who expected enforcement. That surprise is where academic ontology meets the enterprise information system. The moment data leaves the editor and travels between systems, a boundary that only infers is not yet a boundary you can trust. It has to be able to refuse what does not belong, deterministically, with its rules bound into the data rather than living in a separate document or a separate dashboard.

That is the work my colleagues and I have spent years on in the Semantic Data Charter. SDC treats meaning and the constraint definitions as shareable, bound assets. They are one inseparable data payload, so admissibility can be checked at the point of use, in the open, against published specifications. Readers will recognize this as related to the layer Xiaoqi names Γ in his EKA tuple: the governance layer that turns a clever model into a trustworthy one. EKA is the architecture that orchestrates the pieces; a substrate like SDC is one open, standards-based way to make that governance layer real in production. The two meet exactly where this book spends its most careful chapters.

There is a larger reason to welcome a book like this now. Ontology engineering is becoming a core architectural capability, the way databases and APIs did before it, and the field is poorly served when it stays split between texts that are rigorous but unreachable and tutorials that are reachable but shallow. Xiaoqi has taken the harder middle path: technically honest, professionally structured, and grounded in the standards that real organizations have to answer to. That is the bridge between theory and practice, and we need more people building it.

Read this book with Protégé open. Do the exercises. And when you reach those governing chapters, remember that the discipline you are learning on a pizza is the same discipline that will one day decide whether a clinical, financial, or regulatory system is allowed to act. As Xiaoqi writes, meaningful intelligence depends on meaningful boundaries. This book teaches you how to build them.

Timothy W. Cook

Founder, Axiom SDC, Inc.

Creator, the Semantic Data Charter (SDC)

June 2026

(Last Updated at 2026-09-04)

Preface from Michael DeBellis

I sometimes find it easy to be pessimistic about the human race. The world gives us plenty of reasons for discouragement. But one thing that illustrates some of the best in people is the open source movement.

It is remarkable how much of the software infrastructure that powers the modern internet is open source. The Linux operating system, Apache web servers, Kubernetes, Kafka, Docker containers, and countless programming languages, libraries, databases, and development tools were created as open-source projects or converted from internal projects to software freely available to all developers. Organizations and individuals understand that software becomes more valuable when it is shared openly so that the developer community can study it, improve it, and build on it. It's a very positive thing that corporations such as Google and LinkedIn are motivated to contribute software in which they've invested heavily into the open source infrastructure. Even more so that so many individuals will contribute their time to maintain and expand open source tools for the simple joy of seeing things work and providing tools that colleagues will find useful.

That generosity is not limited to code. Documentation, examples, tutorials, walkthroughs, bug reports, and teaching materials are also essential parts of the open-source ecosystem. They are just as important, because good software is wasted if no one can understand or learn to use it.

Tutorials are often underappreciated. As I said when I first created my revised version of the Protégé Pizza tutorial, no one gets a PhD for writing a tutorial. Tutorials do not usually receive the same professional recognition as research papers, software frameworks, or commercial products. Yet, for many learners, a good tutorial is the difference between a technology remaining abstract vs. becoming usable.

That is why I am pleased to see Xiaoqi Zhao's *Mastering Ontology Engineering with Protégé and Pizza.owl: From Semantic Foundations to Executable Knowledge Architecture (EKA)*. It continues a tradition that has always been central to the Semantic Web community: building on prior work, acknowledging intellectual lineage, and helping new learners move from abstract ideas to practical results.

The Pizza tutorial itself has a long history. The original tutorial was developed by members of the Protégé and ontology engineering community. It became one of the best-known introductions to OWL ontology development. Its genius was the choice of a domain that was familiar enough to get out of the learner's way. Almost everyone understands pizzas, toppings, bases, vegetarian choices, and spicy

ingredients. Because the domain is simple, learners can focus on the real subject: classes, properties, axioms, inference, queries, and rules.

My own revision of the Pizza tutorial was motivated by a practical teaching need. I was preparing a hands-on workshop and realized that the older tutorial no longer matched the current Protégé user interface closely enough for beginners. Especially since the majority of the students in that workshop were Library Science experts, not developers. Experienced OOP developers can translate instruction such as “add superclass” into revised user interfaces such as “add subclassOf”. New learners, especially those coming from fields such as library science and information architecture, should not have to do that translation while also trying to understand the concepts of OWL and other Semantic Web languages. The tutorial needed to match the actual tool.

What began as a small update became a much larger project. I updated the examples for Protégé 5, revised screenshots and instructions, and added material on topics such as SPARQL, SHACL, SWRL, IRIs, namespaces, and WebProtégé. My goal was not just to teach people where to click, but to help them understand why the clicks mattered.

Xiaoqi’s eBook takes that work in a new direction. It uses the Pizza tutorial as a foundation, but it is not simply a transcription or repetition of the older material. It connects OWL ontology engineering to a broader architectural vision that he calls Executable Knowledge Architecture (EKA). The eBook encourages learners to see ontology development not as an isolated academic exercise, but as part of a larger architecture for formal knowledge, knowledge graphs, reasoning, validation, and intelligent systems.

I’ve spent half of my career as a consultant and the other half doing research. One of the most important lessons I learned as a consultant is that the hard part isn’t writing software, the hard part is integration. That’s true whether we’re integrating rules with COBOL programs, LLMs with enterprise data, or ontologies and knowledge graphs into distributed computing environments.

That broader perspective is timely.

For a while, it seemed that the Semantic Web might remain a powerful idea that never achieved its hoped-for impact. The original vision was ambitious: not just a web of documents, but a web of data and meaning, where machines could understand and integrate information across sources. The challenge was that adding semantics requires work. It requires modeling. It requires shared identifiers. It requires explicit commitments about meaning. It requires different groups in large organizations to talk to each other and define shared concepts and domain boundaries.

Many organizations decided that this was too much effort. It was easier to publish documents, tables, and data dumps and let humans, search engines, and especially machine learning systems make sense

of them.

Ironically, the rise of Large Language Models (LLMs) has created renewed interest in the very technologies that some people thought machine learning would make unnecessary. LLMs are extraordinary tools. Used well, they can provide enormous productivity benefits. But they also have well-known limitations. Their reasoning is opaque. They can generate fluent but incorrect answers. They may struggle with provenance, controlled terminology, and domain-specific meaning.

These problems are not merely accidental bugs. They follow naturally from the fact that LLMs are statistical models rather than explicit knowledge models. Improvements in scale, training, and architecture will continue to make them more capable, but many practical applications still need complementary structures: curated data, explicit semantics, provenance, and constraints.

That is why knowledge graphs and ontologies are becoming more important. Retrieval Augmented Generation (RAG) was an early response to the problem of grounding LLMs in curated information. But as applications become more demanding, simple document retrieval is often not enough. Many domains require richly interconnected knowledge and automated reasoning. RDF, OWL, SPARQL, SHACL, and reusable ontologies such as Dublin Core, PROV-O, and SKOS provide mature standards for representing and working with exactly that kind of explicit knowledge representation. Explicit knowledge representation based on logic and set theory is the perfect complement to the powerful inductive statistical models that power LLMs.

This renewed interest also means that the word “ontology” is now used in many different ways. Some commercial systems use the term for models that are closer to enterprise object models than to ontologies in the Semantic Web sense. I understand why some researchers and ontology engineers find that frustrating. I sometimes do too. But I also see it as a sign of progress. I have seen this pattern since the early days of expert system adoption and later with the adoption of Object-Oriented Programming and Agile methods: when ideas move from research to industry, terms such as rule-based, knowledge base, *Rapid Application Development (RAD)*, object model, and ontology are stretched, simplified, and repurposed. That can create confusion, but it also shows that industry has recognized that the underlying ideas matter. Eventually as practitioners understand the core research ideas, they make educated decisions and can separate real technology from marketing hype. Tutorials such as this eBook are an essential part of that process. Learners need to understand not only that ontologies are useful, but what ontologies and other semantic technology provide that is new and why these new capabilities matter.

One of the things I appreciate about this eBook is that it tries to connect beginner-level Protégé modeling with larger architectural questions. A learner begins with classes such as *Pizza*, *PizzaTopping*, *CheeseTopping*, and *VegetableTopping*. But those beginner examples lead naturally to deeper questions. What is the difference between a class and an individual? When should a class be primitive,

and when should it be defined? What does a reasoner actually infer? Why does OWL's open-world assumption behave differently from a database constraint? When should we use OWL reasoning, and when should we use SHACL validation? How do ontologies relate to RDF graphs, triple stores, SPARQL queries, and larger distributed system architectures?

Those are not minor details. They are the conceptual foundation of ontology engineering.

My strongest advice to readers is to keep Protégé open while reading. Do the exercises. Run the reasoner frequently. That last part is really critical. For training ontologies, the reasoner executes immediately. Running it after every change means you'll see an error as soon as it is created. When that happens, the reasoner can almost always focus you on the specific problem. If you wait until there are several errors the reasoner feedback is usually much less helpful and finding errors can take hours instead of minutes.

When something surprising happens, slow down and understand why. In OWL, surprises are often the most valuable teaching moments. A reasoner that classifies something unexpectedly, fails to classify something you expected, or reports an inconsistency is showing you the difference between what you intended to model and what you actually modeled.

That distinction is at the heart of ontology engineering.

A good ontology is not just a diagram. It is not just a taxonomy. It is not just a data model. It is an explicit, formal model of meaning in a domain. Building one requires careful thought, iteration, testing, and patience. The Pizza tutorial remains useful because it gives learners a simple domain in which to practice those habits before applying them to medicine, finance, manufacturing, enterprise architecture, LLM integration and other business, engineering, and scientific domains.

I am glad to see this new contribution to the Pizza tutorial lineage. It reflects the best spirit of open educational work.

The journey begins with pizza. But the real subject is meaning.

Michael DeBellis

2 July 2026

San Francisco, CA

<https://www.michaeldebellis.com/blog>

Preface from Author

From Static Diagrams to Executable Knowledge: A Personal Journey

For nearly thirty years, I have lived and breathed enterprise IT. My career began long before “knowledge graph” was a common term, working across telecom, automotive, finance, and IT outsourcing consulting. For most of those years, my primary tool for making sense of complexity was static diagramming, mostly with Visio, draw.io, or even office suite like Excel and Powerpoint. It was powerful, but it was also... silent. The diagrams described architecture, but they could not execute it. They could not reason, validate, or adapt.

That was the before.

The shift (2021 – 2022) began when I moved from old toolbox to Archi, the open-source ArchiMate modeling tool. Suddenly, I was not just drawing; I was building a **repository**. I discovered the power of a shared, collaborative model – a single source of truth about an enterprise. But even Archi, as good as it is, had limits. The model was still primarily for human consumption. It was a rich, structured database, but it was not a reasoning engine.

It was around this same time that I discovered the **Essential EAS Open Source (Ontology)** – a glimpse of what happens when an EA repository is built on formal semantics. That was the first crack in the old paradigm, learnt and practiced with its modeling tool - Protégé - the first time I met this word.

The catalyst (2023) was the `Pizza.owl` tutorial. Through creating my first own video series on the topic, I got in touch with Michael DeBellis. More importantly, I got my hands on Protégé. For the first time, I was not just modeling structure; I was formalizing **meaning**. I was defining classes, properties, and rules that a machine could use to infer new knowledge, check for contradictions, and navigate relationships intelligently.

The pizza example was simple, but the potential was explosive. I realized that the future of enterprise architecture was not in prettier diagrams, but in **executable knowledge**. With that learning, back to the Essential EAS tool, I realized how powerful of treating enterprise architecture from semantic perspective, and how flexible the tool provided, e.g. if you need one meaning (triple) that EAS built-in meta-model doesn't exist, just create one new from it's slots and that's now fitting the knowledge of your enterprise reality, more and more.

The synthesis (2023 – 2026) became a period of intense exploration. My YouTube channel (and other video sharing platform, like Bilibili, DouYin, etc..) grew to host not just the Pizza.owl series, but courses on ArchiMate, meta-modeling, Python / C / C++ / Java, PlantUML (diagramming as code), Neo4j (graph

databases), and even 3D modeling with OpenSCAD. Each course was a piece of a larger puzzle. programming taught me logic, PlantUML taught me that diagrams could be generated from text (code), Neo4j showed me how to traverse relationships at scale, and OpenSCAD reinforced the power of parametric, declarative design.

All these threads eventually wove themselves into a single, coherent idea: **EKA – Executable Knowledge Architecture** (<https://xiaoqi.com>). EKA is the formalization of everything I have learned. It is the recognition that an ontology is not a document, but an executable specification. It is the understanding that $EKA = (K, R, \Theta, \Phi, \Gamma)$ – a framework where Knowledge, Reasoning, Triggers, Actions, and Governance become a single, integrated system.

Why this book? This book is not a transcript of my videos. It is the theory behind the practice. It is the result of years of hands-on work, distilled into a structured, pedagogical journey. I wrote it because I saw a gap:

countless tutorials show you how to click buttons in Protégé, but almost none explain the architectural thinking behind the clicks.

This book aims to bridge that gap, using the familiar `Pizza.owl` as a safe and repeatable domain to teach the foundational principles of EKA.

You will start with pizza. But if you follow the journey, you will be able to model anything: an enterprise architecture, a digital twin, a regulatory framework, or an AI knowledge system.

The goal is not just to teach you Protégé. The goal is to give you a new way to think about knowledge: as something that can be **formalized, reasoned over, governed, and eventually executed**.

Welcome to the journey.

Xiaoqi Zhao (Yasen)

Global Enterprise Architect, Founder of the EKA Framework

June, 2026

Chapter 00 -- EKA Formalized: The Executable Knowledge Architecture Definition

- [0.0 Why This Chapter Exists](#)
- [0.1 The Problem That EKA Solves](#)
- [0.2 EKA -- A Working Definition](#)
 - [0.2.1 Short Definition](#)
 - [0.2.2 Formal Definition](#)
- [0.3 The Five Layers of EKA \(Implementation Roadmap\)](#)
- [0.4 Why "Executable" Is Not the Same as "Reasoning"](#)
- [0.5 EKA in Relation to Existing Standards and Architectures](#)
- [0.6 EKA Maturity Levels](#)
- [0.7 EKA and the `Pizza.owl` Tutorial](#)
- [0.8 How to Use the EKA Definition in Your Work](#)
- [0.9 Key Concepts](#)
- [0.10 Chapter \(00\) Summary](#)
- [0.11 Reference](#)

0.0 Why This Chapter Exists

If you have already read the preface or browsed the book's outline, you have encountered the acronym **EKA** -- Executable Knowledge Architecture.

You have seen it in:

- The EKA Implementation Roadmap diagram (repeated across chapters)
- "EKA Connection" sections at the end of most chapters
- References to <https://xiaoqi.com>

But until now, we have not paused to define EKA **formally**.

This chapter changes that.

By the end of this chapter, you will not only understand what EKA is - you will be able to:

- Distinguish EKA from ordinary ontology projects, knowledge graphs, and rule engines.
- Identify the **five core layers** of an executable knowledge architecture
- Recognize why **executability** transforms semantic modeling into operational (executable) intelligence
- Position the `Pizza.owl` tutorial within the larger EKA maturity roadmap

If you are a practitioner, architect, or researcher, this chapter gives you the **conceptual vocabulary** to design, evaluate, and communicate EKA-based systems.

Let us begin.

0.1 The Problem That EKA Solves

Modern enterprises invest heavily in:

- Semantic modeling
 - OWL - Web Ontology Language,
 - RDF - Resource Description Framework,
 - SHACL - Shapes Constraint Language, and more
- Knowledge graphs (Neo4j as a property graph; GraphDB, Stardog as native RDF graph databases)
- Reasoning engines (Pellet, HermiT, RDFS)

Yet a common complaint remains:

"We have ontologies and graphs, but our knowledge does not *act*."

Models are validated.

Graphs are queried.

But **knowledge is not executed** as part of automated, trustworthy, and adaptive business processes.

This gap is not accidental. Most semantic projects stop at one of three stages below:

Stopping Point	What is Missing?
Ontology only	Execution engine, dynamic query, rules as actions
Graph only	Formal semantics, logical consistency, inheritance
Rules only	Domain structure, classification, semantic governance

EKA (Executable Knowledge Architecture) is a framework designed to close the gap.

EKA does not replace ontologies, knowledge graphs, or reasoners.

Instead,

it **orchestrates them** into a single, layered architecture where knowledge is not only represented but also executed.

0.2 EKA -- A Working Definition

0.2.1 Short Definition

EKA (Executable Knowledge Architecture) is an architectural framework in which formal knowledge (ontologies, rules, constraints) is transformed into machine-executable intelligence through a structured pipeline:

Diagrams → Meta-models → Ontologies → Knowledge Graphs → Executable Intelligence.

0.2.2 Formal Definition

Let us now provide a more precise, **formal definition** that can be used in research, tooling, and enterprise governance.

An **EKA system** is a tuple:

$$EKA = (K, R, \Theta, \Phi, \Gamma)$$

Where:

- K = Knowledge Graph layer - a set of entities (nodes) and semantic relationships (edges) conforming to an ontology.
- R = Reasoning & Rules layer - a set of inference rules (OWL, SWRL, or custom) that derive new facts from K .
- Θ = Trigger layer - a set of conditions (semantic events, queries, or time-based triggers) that initiate execution.
- Φ = Execution layer - a set of actions (API calls, process invocations, alerts, graph updates) that change the real world or the knowledge graph.
- Γ = Governance layer - constraints, validation rules (SHACL), and access policies that ensure semantic integrity.

An EKA system is **defined as executable if and only if (iff)** for every trigger $\theta \in \Theta$ that evaluates to `true`, at least one action $\phi \in \Phi$ is **automatically invoked** with guaranteed traceability.

In plain language: EKA is not a static ontology. It is a **semantic automation architecture**.

 **Note**

Note on Notation: Understanding $\in$

Throughout this book, you will occasionally see the symbol $\in$ (pronounced "is an element of" or "belongs to"). Borrowed from set theory, it is simply a shorthand way to show that a specific item is part of a larger collection or set. For example, writing $\theta \in \Theta$ simply means that a given trigger (θ) is a member of the broader set of all possible triggers (Θ). Think of it like saying a specific menu item belongs to the restaurant's master menu—no advanced mathematics required!

0.3 The Five Layers of EKA (Implementation Roadmap)

The EKA roadmap appears throughout this book's diagrams. Let us now explain **each layer** in operational terms.

Layers	Artifact	Tool example	Question it answers
1. Diagramming	Informal conceptual sketches	Draw.io , Visio®, whiteboard	<i>What are the key concerns?</i>
2. Meta-modeling	Formal structural rules among concepts	ArchiMate®, UML (Unified Modeling Language) class diagrams, custom DSL (Domain-Specific Languages)	<i>How are concepts legally related?</i>
3. Ontology	OWL ontology with classes, properties, restrictions, disjointness	Protégé, Fluent Editor	<i>What is the formal meaning?</i>
4. Knowledge Graph (Choice 1: Native RDF)	Populated graph with individuals and inferred relationships	AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune (RDF mode)	<i>What specific facts are true in a semantic, reasoner-compatible store?</i>
4. Knowledge Graph (Choice 2: Property Graph)	Transform graph for fast traversal	Neo4j, TigerGraph, JanusGraph, Amazon Neptune (property graph mode)	<i>How can we traverse relationships efficiently for specific queries?</i>
5. Executable Intelligence	Event-driven actions, queries, and decisions	<i>What should the system do now?</i>	

 Important

Critical Architectural Distinction

The "Knowledge Graph" layer in EKA can be implemented using two fundamentally different technologies:

- **Native RDF Graph Databases** (AllegroGraph, Stardog, GraphDB, RDFox) store data as RDF triples, use global IRIs as identifiers, support SPARQL queries, and be able to perform OWL reasoning natively. This is the **recommended choice** for preserving and executing the full logical semantics of your OWL ontology.
- **Property Graphs** (Neo4j, TigerGraph, JanusGraph) use internal, non-semantic element IDs and are optimized for fast graph traversal algorithms. An OWL ontology can be *transformed* into a property graph, but this is a **lossy transformation** that sacrifices OWL reasoning capabilities and semantic richness.

According to the formal EKA definition, the ***K* layer** can be implemented by either type, but the choice has significant architectural implications. For building a semantic layer over enterprise data that leverages reasoning and Linked Data, a Native RDF Graph Database is the correct architecture choice.

Furthermore, many projects claim to have an "EKA" because they have a knowledge graph. However, **without a trigger Θ and execution Φ , you do not have EKA** - you only have a semantic repository.

0.4 Why "Executable" Is Not the Same as "Reasoning"

A common misconception is that **OWL reasoning alone makes knowledge executable**.

Let us clarify.

Capability	OWL Reasoner	EKA Execution layer
Infers new class membership	✓ Yes	✓ Can use reasoner
Detects inconsistency	✓ Yes	✓ Yes
Triggers an external API	✗ No	✓ Yes
Sends an alert to a business system	✗ No	✓ Yes
Modifies a graph or database	Varies by system (see note below*)	✓ Yes
Operates over time (event-based)	✗ No	✓ Yes

 Note

***Important Note on Inference Materialization**

In Protégé’s default desktop workflow, reasoner results are normally shown as **inferred axioms** (displayed in the interface) rather than automatically saved as **asserted axioms** into the ontology file. This is a tool-specific behavior, not a universal principle of OWL reasoning or graph databases.

In production-grade semantic systems -- including triplestores, rule-enabled graph databases, and enterprise knowledge graph platforms -- inferred triples may be:

- **Materialized and persisted** (written back to storage),
- **Maintained as inferred closures** (updated incrementally), or
- **Exposed virtually** (computed on-the-fly at query time),

depending on the system architecture and configuration.

*The author thanks **Michael DeBellis** for his expert technical review and for clarifying this critical engineering distinction -- particularly the difference between Protégé’s UI behavior and broader enterprise-grade inference materialization practices. (Ref: GitHub Issue #9 (https://github.com/yasenstar/protege_pizza/issues/9))*

OWL reasoning answers: *What else is true?* (Note: as mentioned just now, the reasoner does not directly write inferred triples back into the original OWL file or persistent triple store, unless you run Drools and write to OWL manually in Protégé.)

EKA execution answers: *What should be done now, automatically, based on what is true?*

In an EKA system, the reasoner is a **component** - not the whole engine.

The execution layer Φ is what makes knowledge **operational**.

0.5 EKA in Relation to Existing Standards and Architectures

To avoid confusion, it is helpful to position EKA relative to well-known frameworks.

Framework	Primary focus	Relationship to EKA
TOGAF - The OpenGroup Architecture Framework (ADM - Architecture Development Method)	Enterprise architecture process	EKA can be used as a semantic automation pillar within all TOGAF’s ADM layers.
Semantic Web Stack (OWL, RDF, SPARQL - "SPARQL Protocol and RDF Query Language")	Knowledge representation and query	EKA extends the stack with an explicit execution layer.
Knowledge Graph (Native RDF: GraphDB, AllegroGraph, Stardog; Property Graph: Neo4j)	Graph storage and traversal	The <i>K</i> layer of EKA. Native RDF graphs are the recommended choice for preserving OWL semantics. Property graphs may be used for specialized, performance-intensive queries after transformation.

Framework	Primary focus	Relationship to EKA
Rule engines (Drools, DMN - Decision Model and Notation)	Business rule execution	EKA generalizes rules into semantic triggers + actions
Data Fabric / Data Mesh	Data integration and decentralization	EKA adds executable semantics on top of integrated data.

EKA is **not** a replacement for any of these frameworks.

It is an **architectural pattern** that assembles them into a coherent, executable knowledge pipeline.

0.6 EKA Maturity Levels

Not every ontology project needs full EKA.

We define four maturity levels to help you decide.

Level	Name	Characteristics	When to use
L0	Semantic modeling	Protégé + OWL only, no graph, no execution	Learning, academic exploration
L1	Knowledge graph	Ontology + populated graph, queryable (SPARQL/Cypher)	Enterprise knowledge repository
L2	Reactive knowledge	L1 + trigger (Θ) and notifications	Monitoring, compliance, alerting
L3	Executable intelligence (full EKA)	L2 + autonomous actions (Φ) + Governance (Γ)	Autonomous AI, closed-loop architecture, digital twins

This book:

- Chapters 1-7 primarily cover **L0 and L1** (ontology + reasoning)
- Chapters 8-11 introduce **graph readiness** (L1-L2)
- Later chapters on SWRL, SPARQL, and external integrations build toward **L2 and L3**
- Especially, chapters 17-24 bring the Semantic Knowledge Development Lifecycle (SKDL) in the systematically way as the implementation steps for reaching EKA

You are not expected to build a full L3 system after reading the `Pizza.owl` tutorial.

But you **will** understand how each layer contributes to executability.

0.7 EKA and the `Pizza.owl` Tutorial

You may now ask: *Why does the `Pizza.owl` tutorial appear in an EKA book?*, or *"Why we link to EKA in this `Pizza.owl` video reference book?"*

The answer to both of them is strategic.

The `Pizza.owl` tutorial teaches:

- Classes, properties, restrictions (Ontology layer)
- Reasoning (preparation for R)
- Disjointness and consistency (part of Γ)

But the original tutorial **does not** cover:

- Exporting to a graph database (K)
- Defining triggers (Θ)
- Automating actions (Φ)

Therefore, this book uses `Pizza.owl` as the **common, repeatable, beginning-friendly domain** to teach **EKA concepts layer by layer**.

By the end of this book, you will have built:

- 1. A correct `Pizza.owl` ontology (L0)
- 2. A populated knowledge graph from it (L1)
- 3. Example triggers and actions (simulated or via API) (L2-L3)

What you learn from `Pizza` ontology scales to:

- Enterprise risk ontologies
- Supply chain knowledge graphs
- Healthcare eligibility rules
- Digital twin automation

The domain changes. The EKA architecture does not.

0.8 How to Use the EKA Definition in Your Work

As a reader of this book, you are encouraged to apply the EKA definition to your own projects.

Here is a simple **EKA readiness checklist** for any semantic project:

- **Ontology (K schema)**: Do we have an OWL ontology with classes, properties, and constraints?
- **Reasoning (R)**: Can a reasoner infer new relationships and detect inconsistencies?
- **Graph population (K instance)**: Is the knowledge stored in a queryable graph? **(Critical: Is it a Native RDF Graph that preserves OWL semantics, or a Property Graph that requires transformation?)**
- **Triggers (Θ)**: Are these conditions (semantic, temporal, or event-based) that start automation?
- **Actions (Φ)**: Does the system automatically invoke APIs, workflows, or alerts based on triggers?
- **Governance (Γ)**: Are validation, access control, and auditability enforced?

If you answer "no" to **Triggers** or **Actions**, your system is a **knowledge graph** -- valuable, but not yet EKA.

The goal is not always full EKA.

The goal is to **choose the right level** for your problem.

0.9 Key Concepts

Term	Definition
Property Graph	A graph database that uses a property graph data model with nodes, edges, and key-value properties. Uses internal, non-semantic element IDs. Examples include Neo4j, TigerGraph, JanusGraph, and Amazon Neptune (property graph mode). Property graphs are optimized for fast graph traversal algorithms but do not natively support OWL reasoning.
Native RDF Graph Database (RDF Triplestore)	A graph database that stores data natively as RDF triples. Uses global IRIs as identifiers, supports SPARQL queries, and can perform OWL reasoning directly. Examples include AllegroGraph, GraphDB, RDFox, Stardog, and Amazon Neptune (RDF mode). This is the recommended choice for the EKA K layer when preserving OWL semantics is required.
Knowledge Graph (in EKA context)	The populated (K) layer of an EKA system, representing a set of entities (nodes) and semantic relationships (edges) conforming to an ontology. Can be implemented using either a Native RDF Graph Database (preserving OWL semantics) or a Property Graph (with lossy transformation and significant trade-off).
Neo4j (clarification)	A popular property graph database. While it can be used as part of an EKA implementation, it requires a lossy transformation from OWL/RDF and does not natively support OWL reasoning. Distinguished from Native RDF Graph Databases like GraphDB, Stardog, and AllegroGraph.

0.10 Chapter (00) Summary

In this chapter (00), we formally defined the **Executable Knowledge Architecture (EKA)** that appears throughout this book.

You learned:

- The **problem** EKA solves (knowledge is modeled but not executed)
- A **short and formal definition** of EKA as a tuple $(K, R, \Theta, \Phi, \Gamma)$
- The **five-layer implementation roadmap** from diagrams to executable intelligence
- The **difference** between OWL reasoning and EKA execution
- How EKA relates to TOGAF, Semantic Web Stack, knowledge graphs, and rule engines.
- **Four maturity levels** (L0-L3) to right-size EKA adoption
- Why the `Pizza.owl` tutorial is the **vehicle** - not the destination - for learning EKA

From this point forward, whenever you see the **EKA Connection** section at the end of a chapter, you will recognize it as a specific layer, component, or maturity level being strengthened.

The next chapter (01) resumes the ontology journey -- but now you see the full architectural horizon.

0.11 Reference

- EKA Official Website: <https://xiaoqi.com>
- EKA Community Site in Github: <https://github.com/EKA-Community/eka>
- Michael Debellis, *A Practical Guide to Building OWL Ontologies Using Protégé 5.5 and Plugins*
- W3C OWL 2 Specification: <https://www.w3.org/TR/owl2-syntax/>
- Neo4j Cypher Documentation (for EKA knowledge graph layer): <https://neo4j.com/docs/cypher-manual/>

Last updated at 2026-09-05

Chapter 01 -- Entering the World of Ontology Engineering with Protégé and Pizza.owl

The journey into ontology engineering often begins with a deceptively simple example: `pizza`. Behind this famous `pizza.owl` tutorial lies one of the most influential learning paths in the Semantic Web and Knowledge Graph ecosystem. In this first chapter, we establish the conceptual foundations of ontology modeling, introduce the Protégé ontology editor, and position ontology engineering within the broader vision of Executable Knowledge Architecture (EKA).

This chapter is based primarily on Michael DeBellis' excellent tutorial [A Practical Guide to Building OWL Ontologies Using Protégé 5.5 and Plugins](#), which modernized and extended the original Manchester Pizza Tutorial into a comprehensive hands-on learning experience.

At the same time, this book extends beyond the tutorial itself. The goal here is not only to learn Protégé as a standalone modeling tool, but to understand ontology as a critical layer in the transformation from diagrams to executable intelligence.

- [1.1 Why the Pizza Tutorial Became an Industry Classic](#)
- [1.2 Protégé: More Than an Ontology Editor](#)
- [1.3 The EKA Perspective: From Diagrams to Executable Intelligence](#)
- [1.4 Understanding OWL: The Semantic Core](#)
 - [1.4.1 Classes: Modeling Conceptual Categories](#)
 - [1.4.2 Individuals: Representing Concrete Instances](#)
 - [1.4.3 Properties: Defining Relationships](#)
- [1.5 Ontology vs Traditional Data Modeling](#)
- [1.6 Open World Assumption: A New Way of Thinking](#)
- [1.7 Reasoners: The Engine Behind Semantic Intelligence](#)
- [1.8 Why Ontology Matters in the AI Era](#)
- [1.9 What You Will Build Throughout This Book](#)
- [1.10 The Bigger Picture](#)
- [1.11 Chapter 01 Summary](#)
- [1.12 Next Chapter Preview](#)
- [1.13 Reference](#)
- [1.14 Key Concepts](#)
- [1.15 One More Thing...](#)

1.1 Why the Pizza Tutorial Became an Industry Classic

The `pizza.owl` tutorial has become the “Hello World” of ontology engineering for several reasons.

First, pizza is universally understandable. Nearly everyone intuitively understands toppings, ingredients, vegetarian classifications, spicy foods, and regional styles. This makes it easier to focus on semantic modeling concepts rather than domain complexity.

Second, the tutorial gradually introduces increasingly sophisticated OWL concepts in a controlled environment. Instead of overwhelming learners with abstract logic theory, it teaches ontology construction incrementally:

- Classes
- Subclasses
- Individuals
- Object Properties
- Restrictions
- Reasoners
- Logical Inference
- SWRL Rules
- SPARQL Queries
- SHACL Validation

This progression mirrors how real enterprise ontology projects evolve.

Third, the tutorial demonstrates something fundamentally important:

Ontologies are not just diagrams.

They are executable semantic systems.

This is the core idea that later becomes central in EKA.

1.2 Protégé: More Than an Ontology Editor

Protégé is often introduced simply as an ontology editor, but this description dramatically understates its importance.

Protégé is better understood as a **semantic modeling platform**.

It allows users to:

- Define domain concepts
- Encode formal semantics
- Build machine-readable knowledge structures
- Execute logical reasoning
- Validate constraints
- Query semantic relationships
- Prepare foundations for knowledge graphs

In traditional enterprise architecture tools, architects often create static diagrams that describe systems visually but lack executable semantics.

For example:

- UML diagrams describe structure
- BPMN diagrams describe processes
- ER diagrams describe data
- Capability maps describe organizations

However, these artifacts usually remain disconnected documentation.

Protégé introduces a fundamentally different paradigm.

Instead of merely drawing concepts, we formally define them using **logic-based semantics**. The ontology becomes interpretable by machines, enabling automated reasoning and inference.

This transition is foundational for modern AI-native architectures.

1.3 The EKA Perspective: From Diagrams to Executable Intelligence

Within the EKA (Executable Knowledge Architecture) framework introduced on xiaoqi.com, ontology engineering occupies a critical transformation layer.

The implementation roadmap of EKA can be summarized as:



🔥 Graph Database Loading Paths for OWL Ontologies

OWL ontologies can be loaded into graph databases, but the implementation path depends on how much semantic fidelity you need to preserve:

1. **Native RDF Graph Databases (Recommended for EKA):** AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune (RDF mode), and Apache Jena/Fuseki. These systems store data natively as RDF triples and can **load OWL ontologies directly**, preserving full formal semantics, global IRI identity, and reasoning capabilities. This is the recommended choice for the EKA *K* layer.
2. **Property Graphs (For Specialized Performance Scenarios):** Neo4j, TigerGraph, JanusGraph, and Amazon Neptune (property graph mode). These systems use internal, non-semantic element IDs. Loading an OWL ontology into a property graph is a **lossy transformation**, not a direct

mapping. You lose OWL semantics, reasoning capabilities, and IRI-based global interoperability. Use this path only for specialized, performance-intensive graph traversal scenarios.

The Pizza OWL tutorial specifically addresses the "Ontology" stage.

To understand why this matters, we must examine the limitations of traditional enterprise modeling.

Most enterprise architecture repositories contain thousands of diagrams:

- Application landscapes
- Process flows
- Integration maps
- Data models
- Organizational charts

But these assets - very valuable - are typically:

- Static
- Human-readable only
- Semantically ambiguous
- Non-executable
- Difficult to integrate

Ontology changes this situation fundamentally.

When enterprise concepts become ontological entities:

- Relationships gain formal meaning
- Constraints become machine-processable
- Inference becomes possible
- Knowledge graphs can be generated
- AI systems can consume structured semantics

In EKA terminology, ontology acts as the semantic bridge between abstract architecture and executable intelligence.

This is why learning Protégé is not merely learning another modeling tool.

IT IS learning how to formalize knowledge itself.

1.4 Understanding OWL: The Semantic Core

The tutorial centers around OWL, the Web Ontology Language.

[Web Ontology Language](#) is a W3C standard designed to represent knowledge in a machine-readable and logically rigorous form.

Unlike relational databases, OWL is not primarily designed for transaction processing.

Instead, OWL focuses on:

- Meaning
- Classification
- Inference
- Semantic relationships
- Knowledge interoperability

Michael DeBellis emphasizes that OWL ontologies are composed primarily of three core elements:

1. Classes
2. Individuals
3. Properties

These concepts form the semantic building blocks of all ontology systems.

1.4.1 Classes: Modeling Conceptual Categories


A class represents a category or type of thing.

Examples in the Pizza ontology include:

- Pizza
- PizzaTopping
- CheeseTopping
- VegetableTopping
- VegetarianPizza

Classes can be organized hierarchically.

For example:



You may also document above graph in this way:

```
PizzaTopping
|-- CheeseTopping
|-- MeatTopping
|-- SeafoodTopping
|-- VegetableTopping
```

This resembles inheritance structure in object-oriented modeling, but OWL classes are fundamentally semantic rather than implementation-oriented.

A class in OWL describes a set of possible individuals.

This distinction becomes important later when reasoning engines classify instances automatically.

1.4.2 Individuals: Representing Concrete Instances

Individuals are actual members of classes.

For example, members of class `Pizza` or `PizzaTopping` can be:

- MargheritaPizza
- MozzarellaTopping
- SpicyBeefPizza

An individual belongs to one or more classes.

However, unlike traditional systems where classification is often manually assigned, OWL reasoners can infer class membership automatically based on logical definitions.

This capability is one of the major breakthroughs of ontology systems.

1.4.3 Properties: Defining Relationships

Properties define how entities relate to one another.

In OWL, there are two major property types:

1. Object Properties: These connect individuals to other individuals.

Example:

```
hasTopping
```

2. Data Properties: These connect individuals to literal value.

Example:

```
hasCalories  
hasPrice  
hasSpicinessLevel
```

Properties become extraordinarily powerful when combined with restrictions and logical constraints.

1.5 Ontology vs Traditional Data Modeling

New learners often ask:

"How is ontology different from a database schema?"

This is one of the most important conceptual distinctions in semantic engineering.

Traditional databases focus on storing data efficiently.

Ontologies focus on representing meaning formally.

A relational schema might define:

```
Pizza(id, name)  
Topping(id, name)  
PizzaTopping(pizza_id, topping_id)
```

But the schema itself does not formally express semantic truths such as:

- Vegetarian pizzas cannot contain meat toppings
- Spicy pizzas contain spicy ingredients
- Certain toppings are subclasses of vegetable toppings
- Seafood toppings are disjoint from meat toppings

OWL allows these semantics to be explicitly modeled.

This is why ontologies become foundational for AI and knowledge graph systems.

1.6 Open World Assumption: A New Way of Thinking

One of the most difficult concepts for beginners is the Open World Assumption (OWA).

Traditional databases generally operate under a Closed World Assumption (CWA):

If something is not stored, it is considered false.

OWL works differently!

Under OWA:

Absence of knowledge does not imply falsehood.

It just means the knowledge that is unknown yet.

For example:

If an ontology does not specify whether a pizza contains meat, we cannot conclude that it is vegetarian.

This principle aligns much more closely with real-world knowledge representation, where information is often incomplete.

Understanding OWA is essential because it fundamentally changes how semantic systems behave compared to traditional enterprise systems.

Michael DeBellis highlights this repeatedly throughout the tutorial because many modeling errors originate from misunderstanding **open-world reasoning**.

1.7 Reasoners: The Engine Behind Semantic Intelligence

The real magic of ontology engineering appears when we introduce reasoners.

A reasoner is a logic engine that analyzes ontology axioms and derives new knowledge automatically.

For example, if we define:

```
VegetarianPizza ≡ Pizza AND (hasTopping only (CheeseTopping or VegetableTopping))
```

and later create a pizza that only contains cheese and vegetables, the reasoner can automatically classify it as a `VegetarianPizza`.

No manual tagging is required!

This is a profound shift.

The ontology stops being passive documentation and becomes a **reasoning-enabled semantic system**. The reasoner can automatically infer classifications and detect inconsistencies. This is an essential step toward **Executable Knowledge Architecture (EKA)**, but true executability -- where reasoning outcomes automatically trigger actions and workflows -- requires additional layers: **trigger (Θ)** and **execution action (Φ)**, which we will explore in later chapters.

This is precisely the transition EKA describes as moving toward executable intelligence.

1.8 Why Ontology Matters in the AI Era

Modern AI systems increasingly require structured semantic context.

Large Language Models (LLMs) are powerful, but without explicit knowledge structures they often struggle with:

- Consistency
- Explainability
- Governance
- Enterprise semantics
- Controlled reasoning

Poor structures for LLMs normally lead to the "beautiful garbage" to be generated by AI.

Ontology provides the semantic backbone that complements probabilistic AI systems.

This is why ontology engineering is re-emerging as a critical discipline in:

- Knowledge Graphs
- Semantic Search
- AI Agents
- Digital Twins
- Enterprise Knowledge Management

- Data Fabric architectures

The Pizza tutorial may appear simple, but the underlying modeling principles scale into enterprise-grade semantic systems.

1.9 What You Will Build Throughout This Book

As this eBook progresses, we will gradually build increasingly sophisticated ontology capabilities using Protégé.

The journey includes:

- Constructing class hierarchies
- Modeling object properties
- Defining logical restrictions
- Using automated reasoners
- Performing semantic queries
- Applying SWRL rules
- Validating ontologies with SHACL
- Connecting ontology concepts to knowledge graph thinking
- Positioning ontology inside the EKA execution pipeline

By the end, readers should understand not only how to use Protégé, but why ontology engineering matters strategically for modern AI-native enterprises.

1.10 The Bigger Picture

The Pizza ontology is not really about pizza.

It is about learning how to formalize knowledge.

Once you understand:

- classes,
- relationships,
- restrictions,
- inference, and
- semantic reasoning,

you can model almost anything:

- enterprise architectures,
- business capabilities,
- manufacturing systems,
- healthcare vocabularies,
- supply chains,
- digital twins,
- AI agent memory systems, and
- knowledge graphs.

See https://github.com/yasenstar/ArchiMate_Ontology as one practical sample that I'm applying to use Protégé to build ontology for ArchiMate (EA Modeling language).

This is where ontology engineering evolves from academic theory into enterprise intelligence infrastructure.

The future of architecture is not static diagrams.

The future is **reasoning-enabled semantics** that can evolve into **executable knowledge** -- where semantic insights automatically trigger actions, workflows, and intelligent system responses..

And Protégé is one of the gateways into that world.


1.11 Chapter 01 Summary

In this chapter, we established the conceptual and strategic foundations of ontology engineering using the famous `pizza.owl` tutorial as the entry point. Rather than treating Protégé as merely a standalone modeling application (ontology editor), we positioned it as a semantic engineering platform capable of transforming static knowledge into **reasoning-enabled semantic models** -- the essential foundation for **Executable Knowledge Architecture (EKA)**, where triggers and execution actions ultimately operationalize knowledge.

We explored why the Pizza Tutorial became one of the most influential ontology learning examples in the Semantic Web community and examined how Protégé enables users to formally define concepts, relationships, restrictions, and semantic logic through OWL ontologies.

The chapter also introduced the EKA (Executable Knowledge Architecture) perspective, where ontology acts as the semantic bridge between enterprise modeling and knowledge-driven intelligence systems.

Within the EKA implementation roadmap:



the Pizza OWL tutorial specifically focuses on the ontology layer, helping learners understand how formal semantic structures are created and reasoned upon.

In addition, we introduced several core OWL concepts including:

- Classes
- Individuals
- Properties
- Open World Assumption
- Reasoners
- Semantic Inference

These concepts form the foundation for all future ontology engineering activities throughout this book.

Finally, we connected ontology engineering to the broader AI era, explaining why semantic modeling, knowledge graphs, and executable semantics are becoming increasingly important in modern enterprise architecture and intelligence systems.

1.12 Next Chapter Preview

In the next chapter, we will begin working directly inside Protégé and start constructing the foundational structure of the Pizza ontology.

Readers will learn how to:

- Create ontology projects
- Understand Protégé workspace components
- Build class hierarchies
- Create subclasses
- Define semantic categories
- Organize ontology structures properly

This marks the beginning of hands-on ontology engineering using Protégé.

The transition from theory to implementation starts there.

1.13 Reference

- Michael DeBellis, A Practical Guide to Building OWL Ontologies Using Protégé 5.5 and Plugins
- Protégé official resources and tutorials
- Protégé Pizza Ontology learning repository: https://github.com/yasenstar/protege_pizza
- EKA Official Website: <https://xiaoji.com>

1.14 Key Concepts

Term	Definition
Ontology	A formal, machine-readable model that defines concepts, relationships, constraints, and inference logic within a domain.
OWL (Web Ontology Language)	A W3C standard language for representing semantic knowledge in a logically rigorous, machine-interpretable form.
Protégé	An open-source semantic modeling platform and ontology editor used to build, reason over, and validate OWL ontologies.
Class	A category or type of thing in an ontology (e.g., Pizza, CheeseTopping).
Individual	A concrete instance of a class (e.g., MargheritaPizza, MozzarellaTopping).
Object Property	A relationship that connects one individual to another (e.g., hasTopping).
Data Property	A relationship that connects an individual to a literal value (e.g., hasCalories).
Reasoner	A logic engine that checks consistency, infers implicit relationships, and automatically classifies individuals based on defined axioms.
Open World Assumption (OWA)	The principle that absence of information does not imply falsity; knowledge may be incomplete.
Closed World Assumption (CWA)	The principle that anything not explicitly known is assumed false (typical of relational databases).
Executable Knowledge Architecture (EKA)	An architectural framework that transforms formal knowledge into machine-executable intelligence through a structured pipeline of diagrams, meta-models, ontologies, knowledge graphs, and executable intelligence.
Semantic Inference	The process by which a reasoner derives new knowledge from existing axioms without manual intervention.

1.15 One More Thing...

Small tip, how to type Protégé in Markdown with those French é ?

Answer: 😊 , just using your keyboard's number pad, hold Alt key and key in 0233 or 233 , like Alt + 0233 or Alt + 233 , then you see the magic. (This is called Alt Code.) Enjoy!

 **Note**

The Alt Code is Windows only. This should be noted. Mac OS and Linux have their own methods for creating diacritical characters. (*Thanks for pointed out by @sarcanon, refer to [issue log #43](#)*)

Last updated at: 2026-09-05

Chapter 02 -- Building Your First Ontology in Protégé

After establishing the conceptual foundations of ontology engineering in Chapter 01, we now move into the practical world of ontology construction inside Protégé. This chapter marks the transition from semantic theory into hands-on ontology modeling.

The goal of this chapter is not merely to "click through a tool," but to understand how semantic structures are engineered systematically. In traditional software engineering, developers write executable code. In ontology engineering, architects build executable meaning.

This distinction is fundamental.

The Pizza ontology may appear simple on the surface, but it introduces modeling patterns used throughout enterprise semantic systems, knowledge graphs, and AI-driven architectures.

This chapter is primarily based on the ontology construction sections from Michael DeBellis' revised [Pizza Tutorial for Protégé 5.5](#).


- [2.1 From Theory to Construction](#)
- [2.2 Understanding the Protégé Workspace](#)
- [2.3 Creating a New OWL Ontology](#)
- [2.4 The Importance of Ontology Naming Conventions](#)
- [2.5 Understanding `Thing` : The Root of All Classes](#)
- [2.6 Creating the Core `Pizza` Taxonomy](#)
- [2.7 Building Class Hierarchies Systematically](#)
- [2.8 Disjoint Classes: Preventing Semantic Contradictions](#)
- [2.9 Primitive vs Defined Classes](#)
 - [2.9.1 Primitive Class](#)
 - [2.9.2 Defined Class](#)
- [2.10 Why Automated Classification Matters](#)
- [2.11 Ontology Engineering vs Diagramming](#)
- [2.12 Early Modeling Discipline](#)
- [2.13 The Emerging Semantic Architecture Mindset](#)
- [2.14 Chapter 02 Summary](#)
- [2.15 Protégé Operations \(Skills\) Learned](#)
- [2.16 EKA Connection](#)
- [2.17 Knowledge Graph Perspective](#)
- [2.18 Key Concepts](#)
- [2.19 Next Chapter \(03\) Preview](#)
- [2.20 Reference](#)

2.1 From Theory to Construction

In Chapter 01, we introduced the core OWL concepts:

- Classes
- Individuals
- Properties
- Reasoners
- Open World Assumption

A quick AI-generated picture below brings them together systematically:



Now we begin building actual semantic structures.

This is where ontology engineering starts to feel very different from traditional modeling disciplines.

In a traditional diagramming tool, creating a class hierarchy is often just a visual activity. In Protégé, however, every modeling action contributes to a machine-processable semantic structure.

This means that:

- class hierarchies have logical implications
- property definitions affect reasoning behavior
- restrictions become reasoning-enabled constraints, and
- inference engines can derive new knowledge automatically

As Michael DeBellis emphasizes throughout the tutorial, ontology engineering should be approached as semantic knowledge modeling rather than merely graphical editing.

2.2 Understanding the Protégé Workspace

When Protégé is launched for the first time, many new users feel overwhelmed by the interface.

This is normal.

Protégé is not designed as a lightweight diagramming application. On the contrary, it is a **semantic engineering environment**.

The workspace contains multiple perspectives for editing different ontology components:

- Classes
- Object Properties
- Data Properties
- Individuals
- Annotations
- Reasoner views

- Query views
- Rule views

One of the most important concepts for beginners to understand is that Protégé does not separate "diagramming" from "logic".

Everything visible in the interface ultimately contributes to OWL axioms underneath.


This is why ontology engineering feels closer to knowledge programming than traditional modeling.

2.3 Creating a New OWL Ontology


The first practical step is creating a new ontology project.

During this chapter's writing, I'm using Protégé v5.6.9.

Inside Protégé: File > New... :




If you're having existing ontology file opened in Protégé, you will get below pop-up with question:



If you choose **Yes** (default), your current ontology file will be closed and your new ontology project will occupy current window, ensure you save any previous changes before losing them.

Or, you may choose **No** so a new window will be opened for your new ontology project, your existing opened ontology will not be impacted.

You then can see the new ontology project screen:



When you want to save the project, Protégé then asks for:

- Ontology IRI (Internationalized Resource Identifier)
- Document storage location
- RDF/XML or OWL serialization format

At this stage, beginners often underestimate the importance of ontology [IRIs - Internationalized Resource Identifier](#).

As defined by W3C, "Each ontology may have an ontology IRI, which is used to identify an ontology." An ontology IRI acts as the global semantic identifier for the ontology.


Unlike filenames or database table names, ontology IRIs are intended to be globally unique identifiers in semantic ecosystems.

For example:

```
# IRI for pizza.owl
http://xiaoqi.com/ontologies/pizza.owl

# My new ontology project IRI (local)
http://www.semanticweb.org/xiaoqi/ontologies/2026/4/untitled-ontology-121
```

Using [WebVOWL - the OWL Visualization Tool](#), you can paste above pizza.owl IRI then visualize it for now:



This is conceptually similar to namespaces in programming languages, but with web-scale semantic interoperability.

In enterprise environments, ontology naming strategies become critically important because ontologies may later integrate across:

- business domains,
- data catalogs,
- APIs,
- knowledge graphs,
- AI agents, and
- semantic search platforms.

This is one of the first points where ontology engineering begins to intersect with enterprise architecture governance.

2.4 The Importance of Ontology Naming Conventions

One of the subtle but important lessons from the Pizza tutorial is disciplined naming.

Ontology projects quickly become unmanageable if naming conventions are inconsistent.

Michael DeBellis repeatedly emphasizes semantic clarity and hierarchy organization throughout the tutorial.

Good ontology naming should satisfy several criteria, as below:

- Human-readable
- Machine-processable
- Semantically meaningful
- Consistent across domains
- Stable over time

For example in our `Pizza.owl` :

```
Pizza
PizzaTopping
VegetarianPizza
SpicyPizza
```

These names are concise and semantically clear.

In enterprise ontology projects, naming conventions become even more important because ontologies may evolve into shared organizational semantic standards.

Within EKA, naming discipline is essential because ontology entities later become:


- graph nodes
- semantic APIs
- AI context objects
- digital twin entities
- and executable business semantics

Poor naming eventually creates **semantic debt** and chaos across the organization.

This is analogous to **technical debt** in software engineering.

2.5 Understanding Thing : The Root of All Classes

Every OWL ontology begins with a universal root class: `owl:Thing`



All classes ultimately inherit from `Thing` .

This concept initially appears trivial, but it represents an important philosophical principle in ontology engineering.

OWL assumes that all modeled entities exist within a universal semantic space.

For example:

```
Thing
|-- Pizza
|-- PizzaTopping
|-- PizzaBase
|-- Spiciness
```

Unlike relational databases where tables are independent structures, OWL organizes concepts into a globally connected semantic universe.

This is one reason ontology systems integrate naturally with knowledge graphs.

2.6 Creating the Core Pizza Taxonomy

The next step involves building the foundational class hierarchy.

From Wikipedia, **Taxonomy** is a practice and science concerned with classification or categorization.

The Pizza ontology begins with several top-level conceptual categories:

```
Pizza
PizzaTopping
PizzaBase
Spiciness
```

These represent the primary semantic domains within the ontology.

At this stage, ontology engineering resembles taxonomy construction.

However, ontologies extend far beyond simple taxonomies.

An ontology additionally defines:

- semantics
- constraints
- properties
- logical restrictions
- inference behavior
- and reasoning rules


This distinction is critical.

Many enterprise "ontologies" are actually only taxonomies because they lack formal semantics.

2.7 Building Class Hierarchies Systematically

One of the most powerful productivity features demonstrated in the Pizza tutorial is the "Create Class Hierarchy" wizard.

In Protégé, from menu [Tools] > [Create class hierarchy...]



Instead of manually creating classes one by one, Protégé allows ontology engineers to rapidly construct semantic hierarchies using indentation-based structure.

For example:

```
PizzaTopping
  CheeseTopping
  MeatTopping
  SeafoodTopping
  VegetableTopping
```

Tip: you may create those structure in one TXT file upfront, then paste into the wizard.

This approach provides several advantages:

- Faster ontology construction
- Better structural consistency
- Easier semantic organization
- Able to see `disjoint` in one go (see 2.8 below)
- Reduced modeling errors

In enterprise-scale ontology projects containing thousands of concepts, hierarchical discipline becomes essential.

Without strong hierarchy governance:

- semantic duplication emerges
- inheritance becomes inconsistent
- and reasoning performance degrades

Ontology engineering therefore requires architectural thinking, not merely modeling skills.

2.8 Disjoint Classes: Preventing Semantic Contradictions

An extremely important concept introduced during hierarchy creation is **class disjointness**.

For example:

```
MeatTopping
VegetableTopping
SeafoodTopping
```

These categories should logically be mutually exclusive.

A topping cannot simultaneously be both:

- MeatTopping
- and VegetableTopping

Protégé allows classes to be declared as disjoint.

This enables reasoners to detect contradictions automatically.

For example, if an individual were incorrectly classified as both:

```
HamTopping
AND
VegetableTopping
```

Since `HamTopping` is one type of `MeatTopping`, the reasoner would identify an inconsistency.

This capability becomes enormously important in enterprise environments.

The Disjointness supports:

- data quality enforcement
- semantic governance
- AI validation
- and knowledge consistency checking

Within EKA, this is one of the first places where architecture becomes executable.

The ontology no longer merely describes the domain, it actively validates semantic correctness.

2.9 Primitive vs Defined Classes

One of the most intellectually important concepts in OWL modeling is the distinction between:

- Primitive Classes
- Defined Classes

2.9.1 Primitive Class

A primitive class only specifies necessary conditions.

For example:

```
Pizza
```

may simply be declared as a subclass of:

```
Food
```

without fully defining what constitutes a pizza.

2.9.2 Defined Class

A defined class, however, specifies both:

- necessary conditions
- and sufficient conditions

You may find similar terms in Sets Theory, correct, they're inherited from that theory indeed.

For example:

```
VegetarianPizza
  = Pizza AND (hasTopping only (CheeseTopping or VegetableTopping))
```

This means the reasoner can automatically infer membership.

If a pizza satisfies the logical conditions, it becomes classified automatically.

This is one of the defining characteristics of semantic systems.

Traditional systems rely heavily on manual categorization, while ontology systems support computational classification.

2.10 Why Automated Classification Matters

This concept may initially seem academic, but it becomes transformative at enterprise scale.


Consider a large organization containing:

- thousands of applications,
- millions of data entities,
- thousands of APIs, and
- hundreds (or possible thousands as well) of business capabilities.

Manually classifying all semantic relationships becomes impossible.

Ontology reasoning introduces automation into semantic governance.

Within EKA, automated classification enables:



This is one of the major strategic advantages of semantic architecture over static repositories.

Automated classification makes the ontology **reasoning-enabled**. However, reasoning alone does not make the ontology **executable**. True executability within the EKA framework requires **triggers (Θ)** and **execution action (Φ)** that respond to semantic outcomes.

2.11 Ontology Engineering vs Diagramming

A common beginner's mistake is treating Protégé like Visio or UML modeling software.

This is fundamentally incorrect.

In diagramming tools:

- relationships are often visual only
- semantics remain informal
- and diagrams rarely execute

While in ontology engineering:

- semantics are formalized
- relationships are machine-readable
- and logic becomes executable

This difference is profound.

Ontology engineering sits much closer to software engineering and knowledge programming than to traditional architecture diagramming.

Protégé therefore represents a semantic compiler environment rather than merely a drawing tool. (In fact, you don't have actual "drawing" activities in Protégé.)

This distinction becomes central within EKA because the ultimate goal is not documentation.

The goal is executable knowledge.

2.12 Early Modeling Discipline

One of the most valuable lessons from the Pizza tutorial is modeling discipline.

Good ontology engineering requires:

- semantic clarity,
- naming consistency,
- hierarchy governance,
- logical rigor,
- and iterative refinement

New ontology engineers often attempt to model everything immediately.

Experienced ontology architects instead:

1. establish clean top-level structures,
2. define semantic boundaries,
3. create reusable patterns,
4. and evolve the ontology incrementally.

This mirrors good enterprise architecture practices.

Ontology engineering is therefore both:

- a technical discipline,

- and an architectural discipline.

2.13 The Emerging Semantic Architecture Mindset

At this stage in the learning journey, you should begin recognizing a major conceptual transition.

Traditional enterprise systems are largely:

Data-centric

Ontology systems are

Meaning-centric

This shift has enormous implications for AI systems.

Modern intelligent systems increasingly require:

- semantic interoperability,
- contextual understanding,
- knowledge reasoning,
- and machine-readable meaning

Ontology engineering provides the formal semantic layer necessary for this evolution.

This is why ontology, knowledge graphs, and semantic AI are rapidly converging across modern enterprise architecture.

2.14 Chapter 02 Summary

In this chapter, we began constructing the foundational structure of the Pizza ontology inside Protégé.

We have explored together:

- the Protégé workspace,
- ontology project initialization,
- ontology IRIs,
- semantic naming conventions,
- class hierarchy creation,
- disjoint classes,
- primitive vs. defined classes,
- and automated classification concepts.

More importantly, we examined how ontology engineering fundamentally differs from traditional diagramming approaches.

Protégé was introduced not simply as a modeling tool, but as a semantic engineering environment capable of creating executable knowledge structures.

The chapter also reinforced the EKA perspective that ontology serves as the semantic transformation layer between enterprise architecture and executable intelligence systems.

2.15 Protégé Operations (Skills) Learned

During this chapter, you learned how to:

- Create a new ontology
- Configure ontology IRIs

- Navigate the Protégé interface
- Create top-level classes
- Build class hierarchies
- Use the Create Class Hierarchy wizard
- Define semantic taxonomies
- Apply disjointness
- Organize ontology structures systematically

2.16 EKA Connection

Within EKA, this chapter represents the transition from static architecture modeling toward executable semantic structure.

Traditional enterprise repositories often contain disconnected diagrams.

Ontology introduces:

- formal semantics,
- machine readability,
- automated reasoning,
- and semantic consistency validation.

The class hierarchies created in this chapter are not merely documentation artifacts.

They are the foundational semantic structures that later evolve into:

- enterprise knowledge graphs,
- AI semantic layers,
- digital twins,
- and reasoning-enabled architectural intelligence -- the prerequisite for truly executable knowledge systems.

2.17 Knowledge Graph Perspective

Knowledge graphs require semantically meaningful entities and relationships.

The ontology structures created in this chapter provide the semantic backbone necessary for graph intelligence.

Graph Database Loading Paths for OWL Ontologies

OWL ontologies can be loaded into graph databases, but the implementation path depends on how much semantic fidelity you need to preserve:

1. **Native RDF Graph Databases (Recommended for EKA):** AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune (RDF mode), and Apache Jena/Fuseki. These systems store data natively as RDF triples and can **load OWL ontologies directly**, preserving full formal semantics, global IRI identity, and reasoning capabilities. This is the recommended choice for the EKA *K* layer.
2. **Property Graphs (For Specialized Performance Scenarios):** Neo4j, TigerGraph, JanusGraph, and Amazon Neptune (property graph mode). These systems use internal, non-semantic element IDs. Loading an OWL ontology into a property graph is a **lossy transformation**, not a direct mapping. You lose OWL semantics, reasoning capabilities, and IRI-based global interoperability. Use this path only for specialized, performance-intensive graph traversal scenarios.

For example:

```
Pizza
  ↓ hasTopping
CheeseTopping
```

can later become graph nodes (entities) and graph edges (relationships) enriched within formal semantics.

Ontology therefore serves as the semantic schema layer for knowledge graphs.

Without ontology:

- graph semantics become ambiguous,
- relationships lose formal meaning,
- and reasoning capability becomes limited.

2.18 Key Concepts

Term	Definition
OWL Ontology	A formal semantic model built using the Web Ontology Language (OWL), consisting of classes, properties, and individuals that represent knowledge in a machine-readable and logically rigorous form.
Protégé	An open-source ontology editor and semantic modeling platform used to create, visualize, and reason over OWL ontologies. It serves as the primary tool for hands-on ontology engineering in this book.
owl:Thing	The universal superclass from which all OWL classes inherit.
Taxonomy	A hierarchical classification structure organizing concepts into parent-child relationships.
Ontology	A semantic model that includes taxonomy, properties, restrictions, semantics, and reasoning logic.
Class	A category or type of thing in an ontology (e.g., <code>Pizza</code> , <code>PizzaTopping</code>). Classes can be organized into hierarchies to represent semantic inheritance.
Subclass	A class that inherits characteristics from a parent class. For example, <code>VegetarianPizza</code> is a subclass of <code>Pizza</code> .
Disjoint Class	Classes that are declared mutually exclusive — an individual cannot belong to two disjoint classes simultaneously (e.g., <code>MeatTopping</code> and <code>VegetableTopping</code>).
Primitive Class	A class that defines only necessary conditions. The reasoner cannot automatically classify individuals into primitive classes based on logical definitions alone.
Defined Class	A class that defines both necessary and sufficient conditions. The reasoner can automatically classify individuals into defined classes when the conditions are satisfied.
Ontology IRI	An Internationalized Resource Identifier that uniquely identifies an ontology. It serves as the global semantic identifier for the ontology, enabling interoperability and reuse across systems.
Taxonomy	A hierarchical classification structure that organizes concepts into parent-child relationships (e.g., <code>PizzaTopping</code> → <code>CheeseTopping</code>). A taxonomy provides the structural foundation for an ontology.
Necessary Condition	A condition that must be true for an individual to belong to a class. For example, all <code>Pizza</code> must have a base.
Sufficient Condition	A condition that, if satisfied, is enough to classify an individual into a class. For example, if a pizza contains only vegetable toppings, it can be inferred to be <code>VegetarianPizza</code> .
Automated Classification	The process by which a reasoner automatically assigns individuals to defined classes based on logical conditions, without manual intervention.
Semantic Debt	The cumulative consequences of poor ontology design decisions, such as inconsistent naming, missing constraints, or improper hierarchy. Analogous to technical debt in software engineering.
EKA (Executable Knowledge Architecture)	An architectural framework that transforms formal knowledge into machine-executable intelligence through a structured pipeline: Diagrams → Meta-Models → Ontologies → Knowledge Graphs → Executable Intelligence.

Term	Definition
Knowledge Graph	A graph-structured knowledge base that represents entities as nodes and relationships as edges, enriched with formal semantics from an ontology.
Reasoner	A logic engine that checks consistency, infers implicit relationships, and automatically classifies individuals based on defined axioms.

2.19 Next Chapter (03) Preview

In the next chapter, we will install Protégé and become familiar with the ontology engineering workspace used throughout the Pizza OWL tutorial series.

Readers will learn how to:

- install Protégé properly,
- navigate the Protégé interface,
- understand the major workspace areas,
- explore ontology editing perspectives,
- and begin developing the mindset of semantic engineering.

This step establishes the operational foundation for all future ontology modeling activities in the book.

Protégé is not simply a modeling application.

It is the primary semantic engineering environment where ontology structures are transformed into machine-readable knowledge models. Within the EKA framework, this represents the beginning of hands-on ontology engineering practice and the transition from conceptual understanding into executable semantic modeling.

2.20 Reference

- Protégé Pizza Repository: https://github.com/yasenstar/protege_pizza
- WebVOWL project in GitHub: <https://github.com/VisualDataWeb/WebVOWL>
- EKA Official Website: <https://xiaoqi.com>

Last updated: 2026-09-05

Chapter 03 -- Installing Protégé and Understanding the Ontology Engineering Workspace

In the previous chapters, we explored the conceptual foundations of ontology engineering and introduced the semantic thinking behind OWL and the Pizza ontology. We also discussed how ontology fits within the broader EKA (Executable Knowledge Architecture) framework as the critical semantic layer between meta-modeling and knowledge graph intelligence.

This chapter now moves into the practical environment where ontology engineering actually happens:

Protégé

Before building ontologies, you must first understand the ontology engineering workspace itself. Just as software developers require familiarity with their IDEs, ontology engineers must understand the tools, perspectives, editors, and semantic workflow provided by Protégé

This chapter focuses on two major objectives:

1. Installing Protégé properly
2. Understanding the Protégé user interface and semantic engineering environment

Although installation may initially appear straightforward, the deeper goal of this chapter is much more important:

Understanding Protégé not simply as software, but as a semantic engineering platform.

This chapter is aligned with the hands-on installation and UI walkthrough demonstrated in the video tutorial and builds upon the foundational theory introduced in Michael DeBellis' Pizza OWL tutorial.

- [3.1 Why Protégé Matters in Ontology Engineering](#)
- [3.2 Installing Protégé](#)
- [3.3 Protégé as an Ontology IDE](#)
- [3.4 Understanding the Protégé Workspace](#)
- [3.5 The Class Tab](#)
- [3.6 Understanding the Entity Tree](#)
- [3.7 Ontology Editing vs. Diagram Drawing](#)
- [3.8 OWL Serialization and Ontology Files](#)

- [3.9 Plugins and Extensibility](#)
- [3.10 The Importance of Semantic Workspace Familiarity](#)
- [3.11 Protégé Within the EKA Vision](#)
- [3.12 Developing the Ontology Engineering Mindset](#)
- [3.13 Chapter \(03\) Summary](#)
- [3.14 Protégé Operations Learned](#)
- [3.15 EKA Connection](#)
- [3.16 Knowledge Graph Perspective](#)
- [3.17 Key Concepts](#)
- [3.18 Next Chapter Preview](#)
- [3.19 Extended Reading - *Bonus*](#)
 - [3.19.1 Fix the](#) Unable to check for plugins [Problem](#)
- [3.20 Reference](#)

3.1 Why Protégé Matters in Ontology Engineering

Protégé is one of the most widely used ontology engineering platforms - far more than an editor - in the Semantic Web and Knowledge Graph ecosystem.

Originally developed at Stanford University, Protégé has become a foundational tool for:

- OWL ontology modeling
- Semantic Web development
- Knowledge graph engineering
- AI semantic modeling
- Enterprise ontology architecture
- Biomedical ontology research
- Linked data systems

Unlike traditional modeling tools, Protégé is designed specifically for semantic engineering.

This distinction is extremely important.

Most architecture or diagramming tools focus primarily on visualization:

- UML tools visualize structure
- BPMN tools visualize process
- ERD tools visualize data relationships

Protégé, however, focuses on:

Formal semantic meaning

Everything created inside Protégé ultimately contributes to executable semantic structures.

This is why Protégé plays such an important role with EKA.

In the EKA implementations roadmap:

EKA Implementation Roadmap



Protégé becomes the primary engineering environment for the "Ontology" stage.

It is where semantic structures are formally defined before later evolving into knowledge graphs and executable intelligence systems.

3.2 Installing Protégé

Protégé is distributed as a desktop application and supports multiple operating systems including:

- Windows
- macOS
- Linux

The official Protégé website provides downloadable installation packages:

<https://protege.stanford.edu>

You may choose to use WebProtégé at <https://protege.stanford.edu/software#web-protege> to experience its features before installing desktop.

FREE & OPEN SOURCE

Protégé is a free, open-source OWL ontology editor

Used by researchers, organisations, and governments worldwide to build and manage ontologies. Available as Protégé Desktop for local use and WebProtégé for collaborative, browser-based editing.

[DOWNLOAD NOW](#)[USE WEBPROTÉGÉ](#)

The tutorial series uses Protégé 5.x, which supports modern OWL 2 ontology modeling capabilities.

Installation is generally straightforward:

1. Download the appropriate installer
2. Install Java if required
3. Launch Protégé
4. Configure workspace preferences if necessary

Although the installation process itself is relatively simple, you should understand an important architectural point:

 **Protégé is not a lightweight note-taking or diagramming tool.**

It is a semantic reasoning environment.

As ontology projects grow larger, Protégé becomes capable of managing:

- thousands of classes,
- semantic restrictions,
- logical inference,
- reasoning engines,
- ontology plugins,
- and graph-scale semantic models.

This is why understanding the workspace properly is essential from the beginning.

3.3 Protégé as an Ontology IDE

One of the most useful ways to understand Protégé is to think of it as an IDE (Integrated Development Environment) for ontology engineering.

Software developers use environments such as:

- Visual Studio Code
- IntelliJ IDEA
- Eclipse IDE

to develop executable software.

Ontology engineers use Protégé to develop executable semantics.

This comparison helps explain why Protégé contains:

- editors,
- views,
- semantic structures,
- validation tools,
- plugins,
- reasoning engines,
- and ontology management capabilities.

Protégé is therefore not merely an editor.

It is a semantic engineering workspace.

Within EKA, Protégé can be viewed as:

An IDE for Executable Knowledge

This framing is extremely important because it changes how learners perceive ontology engineering itself.

Ontology engineering is not documentation work.

It is semantic system engineering.

3.4 Understanding the Protégé Workspace

When Protégé launches for the first time, the interface may appear complex to new users.

This is normal because Protégé exposes multiple semantic engineering perspectives simultaneously.

The workspace is designed around ontology components rather than visual diagrams.

Major areas of the interface include:

- Class hierarchy panel
- Entity editors
- Annotation views
- Object property tabs
- Data property tabs
- Individual management views
- Reasoner controls
- Ontology metadata panels

Unlike ordinary modeling tools, Protégé is organized around semantic structures.

This reflects a fundamental principle of ontology engineering:

Meaning comes before visualization

If you're familiar with ArchiMate and used to Archi (the ArchiMate modeling tool), the difference in their philosophy in modeling is significant:

- Archi (ArchiMate): Diagramming First, then Relationships Second
- Protégé (Ontology): Relationships First, then Visualization Second


Protégé prioritizes semantic precision over graphical appearance. (see more in [3.7](#))

This may initially feel unfamiliar to users coming from traditional enterprise architecture tools.

However, this semantic-first design is one reason ontology systems scale effectively into AI and knowledge graph ecosystems.

3.5 The Class Tab

One of the first and most important areas you will encounter is the `Classes` tab.



This is where ontology taxonomies are constructed and managed.

Inside the Classes tab, users can:

- Create classes,
- organize hierarchies,
- define subclasses,
- add annotations,
- and later define logical restrictions.

At the top of the hierarchy sits `owl:Thing` which represents the universal root class in OWL.

All ontology classes ultimately inherit from `Thing`.

This is one of the first places where you begin experiencing semantic hierarchy construction directly.

Even simple class creation contributes to formal OWL semantics underneath.

3.6 Understanding the Entity Tree

The entity tree displayed inside Protégé is more than a visual navigation structure.

It represents semantic inheritance relationships.

For example:

```
PizzaTopping
|-- CheeseTopping
|-- MeatTopping
|-- VegetableTopping
```

is not simply a visual grouping.

It formally means: $\text{CheeseTopping} \subseteq \text{PizzaTopping}$.

This distinction is essential.

Protégé is continuously generating semantic axioms behind the interface.

This is one reason ontology engineering differs fundamentally from ordinary taxonomy management.

3.7 Ontology Editing vs. Diagram Drawing

Many new users initially expect Protégé to behave like a graphical modeling application such as:

- Microsoft Visio
- Sparx Systems Enterprise Architect
- Lucidchart
- Archi Tool
- draw.io, etc

However, Protégé operates differently.

Traditional modeling tools often prioritize:

- layout,
- shapes,
- visual connectors,
- and presentations.

Protégé prioritizes:

- semantic logic,
- ontology structures,
- formal meaning,
- and reasoning consistency.

This distinction becomes one of the most important mindset transitions in ontology engineering.

The ontology itself is the executable artifact.

Visualization is secondary.

Within EKA, this reflects the shift from:

Static Architecture Documentation

to:

Executable Semantic Architecture

3.8 OWL Serialization and Ontology Files

Another important concept introduced during installation and project creation is **ontology serialization**.

Protégé stores ontologies using standardized semantic formats such as:

- RDF/XML
- Turtle
- OWL/XML

These formats allow ontologies to become portable semantic assets.

Unlike proprietary diagram files, OWL ontologies can later integrate into:

- semantic APIs,
- graph databases,
- Linked Data systems,
- AI semantic layers,
- and enterprise knowledge graphs.


This portability is one reason ontology engineering has become increasingly important in modern AI ecosystems.

Semantic knowledge becomes reusable infrastructure.


3.9 Plugins and Extensibility

Protégé also supports a plugin ecosystem.

From menu `File > Check for plugin...`, you can open the plugins window:



The plugin list will be shown and you can check the version and detail information:



This is important because ontology engineering often evolves beyond simple class modeling.

Plugins may provide capabilities such as:

- reasoning engines,
- graph visualization,

- query execution,
- ontology validation,
- semantic analysis,
- and knowledge graph integration.

As you progress further into ontology engineering, Protégé gradually transforms from a simple editor into a semantic engineering platform.

This extensibility aligns closely with EKA thinking, where ontology eventually becomes part of larger executable intelligence pipelines.

If you see the error `Unable to check for plugins`, check the [tip](#) in the end of this chapter - Extended Reading.

3.10 The Importance of Semantic Workspace Familiarity

One of the most underestimated aspects of ontology engineering is **workspace fluency**.

Beginners often focus entirely on OWL theory while neglecting tool mastery.

However, productive ontology engineering requires both:

- semantic understanding,
- and engineering workflow familiarity.

Experienced ontology engineers develop fluency in:

- navigating ontology structures,
- managing semantic hierarchies,
- organizing entities,
- interpreting inferred structures,
- and maintaining semantic consistency.

This is very similar to software development.

Knowing programming theory alone is insufficient without development environment proficiency.

Protégé therefore becomes part of the ontology engineer's cognitive workflow.

3.11 Protégé Within the EKA Vision

Within EKA, Protégé occupies a strategically important role.

Traditional enterprise architecture often produces:

- static diagrams,
- disconnected repositories,
- and non-executable documentations.

Protégé enables the transition into:

Machine-readable semantics

This is a foundational shift.

The ontology models built inside Protégé later become candidates for:

- semantic APIs,
- knowledge graphs,
- AI context layers,
- digital twins,
- and executable enterprise intelligence.

Graph Database Loading Paths for OWL Ontologies

OWL ontologies can be loaded into graph databases, but the implementation path depends on how much semantic fidelity you need to preserve:

1. **Native RDF Graph Databases (Recommended for EKA):** AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune (RDF mode), and Apache Jena/Fuseki. These systems store data natively as RDF triples and can **load OWL ontologies directly**, preserving full formal semantics, global IRI identity, and reasoning capabilities. This is the recommended choice for the EKA *K* layer.
2. **Property Graphs (For Specialized Performance Scenarios):** Neo4j, TigerGraph, JanusGraph, and Amazon Neptune (property graph mode). These systems use internal, non-semantic element IDs. Loading an OWL ontology into a property graph is a **lossy transformation**, not a direct mapping. You lose OWL semantics, reasoning capabilities, and

IRI-based global interoperability. Use this path only for specialized, performance-intensive graph traversal scenarios.

This is why ontology engineering matters far beyond academic Semantic Web research.

Protégé is not just a modeling application.

It is part of the infrastructure stack for intelligent architecture systems.

3.12 Developing the Ontology Engineering Mindset

At this stage, you should begin shifting from:

Diagram Thinking

toward:

Semantic Thinking

This transition is one of the MOST IMPORTANT conceptual evolutions in the entire Pizza tutorial journey.

Ontology engineering asks fundamentally different questions:

Instead of:

"What should this diagram look like?"

ontology engineering asks:

"What does this concept mean semantically?"

This semantic-first mindset later becomes essential for:

- knowledge graphs,
- semantic AI,
- enterprise reasoning,
- and executable intelligence systems.

3.13 Chapter (03) Summary

In this chapter, we installed Protégé and explored the ontology engineering workspace for the first time.

We examined:

- Protégé installation,
- the semantic engineering environment,
- ontology editing perspectives,
- class hierarchy navigation,
- OWL serialization formats,
- plugin extensibility,
- and the conceptual difference between ontology engineering and traditional diagramming.

Most importantly, we introduced Protégé as an ontology IDE and semantic engineering platform rather than merely a modeling application.

This chapter also reinforced Protégé's role within the EKA framework as the primary environment for transforming conceptual architecture into executable semantic structures.

3.14 Protégé Operations Learned

During this chapter, you learned how to:

- Install Protégé
- Launch the Protégé environment
- Understand the workspace layout
- Navigate ontology panels
- Explore class hierarchies
- Understand ontology file structures
- Recognize Protégé perspectives and editors

3.15 EKA Connection

Within EKA, Protégé represents the primary semantic engineering environment for ontology construction.

It enables the transition from:

Meta-Model Thinking

to:

Executable Semantic Modeling

The workspace familiarity developed in this chapter forms the operational foundation for all future ontology engineering activities inside the EKA execution pipeline.

3.16 Knowledge Graph Perspective

Protégé itself is not a graph database.

However, the semantic structures created within Protégé form the foundational layer that later enables knowledge graph generation.

The ontology classes, hierarchies, and semantic definitions managed inside Protégé eventually become graph-ready semantic assets.

Ontology engineering therefore acts as a precursor to enterprise knowledge graph engineering.

3.17 Key Concepts

Term	Definition
Ontology Editor	A software environment used to create and manage semantic ontologies.
OWL Serialization	The standardized storage formats used to represent OWL ontologies, including RDF/XML, Turtle, and OWL/XML. These formats enable portability and interoperability across systems.
Semantic Hierarchy	A class structure representing formal inheritance relationships.
Ontology Workspace	The engineering environment used for semantic modeling and ontology management.
Semantic Engineering	The discipline of formally modeling machine-readable meaning.

Term	Definition
Ontology IDE	The concept of Protégé as an Integrated Development Environment for semantic engineering, analogous to how Visual Studio or Eclipse serves software developers.
Semantic Workspace	The Protégé interface environment organized around ontology components (classes, properties, individuals) rather than visual diagrams, emphasizing semantic structures over graphical representation.
Class Hierarchy	A tree-structured organization of ontology classes showing parent-child relationships (e.g., <code>PizzaTopping</code> → <code>CheeseTopping</code>). This represents semantic inheritance in OWL.
Entity Tree	The hierarchical display in Protégé that visually represents class inheritance relationships, where each node corresponds to an OWL class.
Semantic Engineering	The discipline of formally modeling machine-readable meaning using logic-based languages like OWL, enabling automated reasoning and inference by machines.
Ontology Editor	A software tool used to create, edit, and manage semantic ontologies. Protégé is the most prominent example.
Semantic Hierarchy	A class structure representing formal inheritance relationships in an ontology, where subclasses inherit properties and meaning from parent classes.
Plugin Ecosystem	The extensible architecture of Protégé that allows third-party plugins for reasoning engines, graph visualization, query execution, ontology validation, and knowledge graph integration.
EKA (Executable Knowledge Architecture)	An architectural framework that transforms formal knowledge into machine-executable intelligence through a structured pipeline: Diagrams → Meta-Models → Ontologies → Knowledge Graphs → Executable Intelligence.
Semantic Reasoning Environment	A software platform (like Protégé) that supports logical inference, consistency checking, and automated classification of semantic models.
WebProtégé	The browser-based collaborative version of Protégé that allows multiple users to edit and manage ontologies simultaneously without installing desktop software.

3.18 Next Chapter Preview

In the next chapter, we will begin building the Pizza ontology from the ground up inside Protégé.

You will learn how to create classes, establish subclass hierarchies, and annotate your ontology with meaningful descriptions.


This chapter marks the first step in transforming your understanding of semantic structures into hands-on ontology modeling.

By the end, you will gain practical experience in organizing conceptual categories, preparing the foundation for more advanced semantic modeling in subsequent chapters.

3.19 Extended Reading - *Bonus*

3.19.1 Fix the Unable to check for plugins Problem

If you want to check plugins, but get below popup error:



No worries!

This error message, specifically the `SunCertPathBuilderException`, usually indicates that the Java environment Protégé is running on doesn't trust the SSL certificate of the plugin registry.

This is very common when you're working behind a corporate firewall or proxy that uses SSL inspection (man-in-the-middle decryption).

I've struggled with this problem in my company assigned machine, and below steps are for your reference and hope they can solve issues on your side as well:

Since my system uses a custom root CA certificate, it is needed to manually import that certificate into the Java TrustStore used by Protégé.

1. Identify the Java Runtime

Protégé often comes with its own bundled Java Runtime Environment (JRE). You need to find where it's located:

- Look in the Protégé installation folder for a directory named `jre`.
- If it's not there, Protégé is likely using your system's default Java (check your `JAVA_HOME` environment variable).

2. Export the Corporate Root Certificate

- Open your web browser (Chrome or Edge).
- Go to the URL the registry is trying to hit (usually `https://protege.stanford.edu`).
- Click the **Padlock icon** in the address bar -> **Connection is secure** -> **Certificate is valid**.
- Go to the **Details** or **Certification Path** tab.
- Select the **Root Certificate** (the one at the very top of the hierarchy).
- Click **Export** or **Copy to File** and save it as a `.cer` or `.crt` file (e.g. `company_root.cer`).

3. Import the Certificate into the KeyStore

Open a Command Prompt or Terminal as **Administrator** and run the `keytool` command. You will need to point it to the `cacerts` file inside the Protégé JRE.

Below command assumes you're running in Windows machine.

```
keytool -import -trustcacerts -alias corporate-root -file "C:\path\to\company_root.cer" -keystore
```

- **Default Password:** The default password for the Java keystore is `changeit`.
- **Trust:** When asked "Trust this certificate?", type `yes`.

4. Alternative: Manual Plugin Installation

If you can't bypass the proxy/SSL issues, you may download plugins manually:

- 1) Visit the "Protégé Plugin Library" (`https://protege.stanford.edu/ontologies.php`)
- 2) Download the `.jar` file for the plugin you need.
- 3) Drop the file into the `plugins` folder inside your Protégé installation directory.
- 4) Restart Protégé.

3.20 Reference

- Protégé Official Website: <https://protege.stanford.edu>

- Protégé Wiki: <https://protegewiki.stanford.edu>
- Michael DeBellis, *A Practical Guide to Building OWL Ontologies Using Protégé 5.5 and Plugins*:
https://www.researchgate.net/publication/351037551_A_Practical_Guide_to_Building_OWL_Ontologies_Using_Protege_55_and_Plugins
- Protégé Pizza Repository: https://github.com/yasenstar/protege_pizza
- EKA Official Website: <https://xiaoqi.com/>

Last updated at: 2026-09-05

Chapter 04 -- Creating Classes and Building the Pizza Ontology Skeleton in Protégé

In the previous chapter, we explored the Protégé workspace, navigating the ontology editing environment and understanding its key panels and perspectives. We also began to develop the mindset of a semantic engineer -- thinking in terms of conceptual categories, machine-readable meaning, and structured knowledge.

Now, we move from understanding the tool to **actively modeling domain concepts**. This chapter focuses on the first practical step in ontology engineering: creating classes and defining a structured hierarchy. Classes are the core of OWL ontologies -- they represent abstract concepts in a domain, define the structure for semantic relationships, and form the foundation for reasoning and knowledge graph construction.

From the perspective of the **Executable Knowledge Architecture (EKA)**, creating classes represents the first concrete step in translating enterprise knowledge from **meta-models and diagrams** into **formal, machine-understandable semantics**. Classes are the backbone nodes of your knowledge model, linking conceptual design to the eventual executable intelligence layer.

By the end of this chapter, you will be able to:

- Understand the role of classes in ontology modeling
- Create root and subclass hierarchies in Protégé
- Apply meaningful annotations to each class
- Organize the `pizza` ontology skeleton for scalability and reasoning
- Connect your class structures to the broader EKA framework

Table of Content

- [4.1 Understanding Classes in Ontology](#)
- [4.2 Creating Classes in Protégé](#)
- [4.3 Adding Annotations](#)
- [4.4 Organizing Subclass Hierarchies](#)
- [4.5 Practical Exercise: Building the Skeleton Ontology](#)
- [4.6 EKA Perspective](#)
- [4.7 Chapter Summary](#)

- [4.8 Key Concepts](#)
- [4.9 Protégé Skills Learned](#)
- [4.10 Next Chapter \(05\) Preview](#)

4.1 Understanding Classes in Ontology

In OWL, **classes** represent **conceptual categories** -- abstract groupings of things that share common characteristics. They answer the typical question:

"What types of entities exist in this domain?"

In the Pizza ontology:

- `Pizza` represents the overarching domain concepts for all pizzas.
- `Topping` represents ingredients or components added to pizzas.
- `VegetarianPizza` and `NonVegetarianPizza` represent specialized subclasses of `Pizza`, reflecting the dietary classification.
- `CheeseTopping`, `MeatTopping`, and `VegetableTopping` are subclasses of `Topping`, further organizing the ingredient concepts.

Classes are the **semantic anchors** for your ontology. They enable logical inheritance, reasoning, and relationships to be defined later.

In EKA terms, classes transform conceptual diagrams into **formal semantic nodes**, bridging the gap between enterprise knowledge design and executable intelligence.


Without a solid class hierarchy, downstream ontology operations like object property assignment, instance creation, and reasoning can become inconsistent or unreliable.

4.2 Creating Classes in Protégé

The process (step) of creating classes in Protégé introduces you to **hands-on ontology engineering**:


1. Open the Classes Tab:

Protégé organizes ontology classes hierarchically in this tab. This interface allows for creation, subclassing, and annotation management.



2. Create the Root Class:

While `owl:Thing` exists by default as the universal root, creating a domain-specific root, `Pizza`, gives semantic clarity and establishes a class starting point for the ontology.



3. Create Topping Class:

We will have Topping at the same level as Pizza class, so depending on your viewpoint, Topping can be treated as a subclass of root `owl:Thing` class, or a sibling class of just created Pizza class. If you select `owl:Thing`, then click Add subclass icon:

 ch04_08_topping-as-subclass

If you select Pizza, then click Add sibling class icon (only now it's clickable):

Active ontology x Entities x Individuals by class x OWLViz x DL Query x OntoGraf x SPARQL Query x

Class hierarchy: VegetarianPizza

Annotations Usage

Annotations: VegetarianPizza1

Annotations +

Description: VegetarianPizza1

Types +

VegetarianPizza

Same Individual As +

Different Individuals +

The pop-up dialogs are same, just type `Topping` and click `ok` , your ontology should be looked as below:

Classes Object properties Data properties

Class hierarchy: Topping

Annotations Usage

Annotations: Topping1

Annotations +

Description: Topping1

Types +


Topping

Same Individual As +

Different Individuals +

4. Add Subclasses:

- Under `Pizza` , create `VegetarianPizza` and `NonVegetarianPizza` .
- Under `Topping` , create `CheeseTopping` , `MeatTopping` , and `VegetableTopping`



5. Manage the Hierarchy:

- Ensure logical "is-a" relationships.
- $\text{VegetarianPizza} \subseteq \text{Pizza}$ means every vegetarian pizza is also a `Pizza`.
- Subclassing defines semantic inheritance, which is foundational for future reasoning.

6. Validate and Save:

Regularly saving and checking hierarchy consistency prevents structural errors. Protégé's visual hierarchy aids in spotting potential misclassifications.

This process demonstrates how **conceptual knowledge is formalized**. It also highlights the EKA principle that **well-structured ontology classes serve as the nodes of a future knowledge graph**, capable of supporting reasoning and executable intelligence workflows.


4.3 Adding Annotations

Annotations provide human-readable metadata for classes, making the ontology understandable to humans while remaining machine-readable.

- **Label** (`rdfs:label`): A description name of the class.
- **Comment** (`rdfs:comment`): An explanation of the class' purpose or semantic meaning.

For example:


- Class: `VegetarianPizza`
- Label: "Vegetarian Pizza"
- Comment: "A pizza containing only non-meat toppings, suitable for vegetarian diets."



The screenshot shows the Protege IDE interface. The main window displays the 'VegetarianPizza' class hierarchy, with 'VegetarianPizza' highlighted. A red arrow points from this class to the 'Annotation' dialog box. The dialog box is titled 'VegetarianPizza' and shows the 'rdfs:comment' property selected. The 'Value' field contains the text: 'A pizza containing only non-meat toppings, suitable for vegetarian diets.' The 'Datatype' dropdown is set to 'xsd:string'. The 'Language Tag' field is empty. The 'OK' button is visible at the bottom of the dialog.

 You have to click **OK** after Label setting then add Comment annotation, every time only one annotation can be set

After adding, two annotations are as below:



4.4 Organizing Subclass Hierarchies

Hierarchy is more than a visual tree -- it defines **semantic inheritance**, which is the foundation of reasoning:

- `VegetarianPizza` inherits all properties and constraints from `Pizza`.
- `CheeseTopping` inherit from `Topping`.

Best practices:

1. Maintain a clear root-to-leaf path to avoid confusion.
2. Avoid duplicate classes; semantic ambiguity undermines reasoning.
3. Use comments and labels to clarify distinctions.
4. Ensure consistency with your meta-model diagrams.

This step mirrors the **Diagramming** → **Meta-Model** → **Ontology** pathway in EKA. Your diagrams represent conceptual ideas, meta-models define the structural rules, and ontology classes encode these as executable semantic constructs.

4.5 Practical Exercise: Building the Skeleton Ontology

By the end of this chapter, you should have a **working skeleton ontology**:

- **Root Classes:** `Pizza`, `Topping`
- **Subclasses:** `VegetarianPizza`, `NonVegetarianPizza`, `CheeseTopping`, `MeatTopping`, `VegetableTopping`

- **Annotations:** Labels and comments for all classes
- **Clean Hierarchy:** Logical parent-child relationships aligned with semantic inheritance principles

This skeleton forms the base upon which object properties, individuals, and reasoning rules will later be applied. The exercise reinforces the **progressive semantic difficulty principle**, gradually preparing you for more complex modeling.

4.6 EKA Perspective

Creating classes is a **critical step in the EKA roadmap**:

1. **Diagramming:** Conceptual pizza categories visualized in diagrams.
2. **Meta-Model:** Defining rules and semantic categories for structured knowledge.
3. **Ontology:** Translating classes into machine-readable semantic nodes.
4. **Knowledge Graph:** Classes will serve as nodes connected by properties, forming the backbone of semantic reasoning.
5. **Executable Intelligence:** Accurate class structures ensure that reasoning engines can infer, validate, and execute knowledge.

By properly designing classes, you are establishing **semantic scaffolding** that allows enterprise knowledge to become executable and AI-ready, bridging traditional modeling with cutting-edge knowledge engineering.

4.7 Chapter Summary

In this chapter (04), you have:

- Explored the central role of classes in ontology modeling
- Created root and subclass hierarchies in Protégé
- Added annotations to improve human and machine interoperability
- Built the first skeleton of the Pizza ontology
- Connected practical exercises to EKA, understanding how classes become semantic nodes and eventually executable knowledge.

This exercise forms the essential foundation for future ontology enhancements, including object properties, individuals, and reasoning logic.

4.8 Key Concepts

Concept	Description
Class	A conceptual category representing domain entities.
Subclass	A specialization inheriting semantic properties.
Annotation	Human-readable metadata supporting clarity and documentation.
Hierarchy	Structured parent-child relationships defining semantic inheritance.
EKA Connection	Classes link diagrams and meta-models to knowledge graphs and executable intelligence.

4.9 Protégé Skills Learned

- Navigating the Classes tab
- Creating root and subclass hierarchies
- Adding labels and comments (annotations)
- Organizing semantic hierarchies for reasoning readiness

4.10 Next Chapter (05) Preview

In the next chapter, we will explore **Named Classes**, the building blocks of formal ontology definition in OWL.

You will learn how to:

- Define clear and semantically meaningful class names
- Organize class hierarchies for scalability and clarity
- Apply naming conversions to support reasoning and knowledge integration
- Begin understanding how naming decisions impact ontology maintenance and interoperability

This step continues the progressive construction of the Pizza ontology, transforming your skeleton into a **structured, well-defined semantic model**, ready for further enrichment and reasoning.

Last updated at: 2026-09-04

Chapter 05 -- Defining Named Classes in the Pizza Ontology

In the previous chapter, we constructed the **skeleton of the Pizza ontology**, creating root classes and a preliminary hierarchy in Protégé. We established the foundational structure upon which the semantic richness of the ontology can be built. However, these initial classes, while necessary, are only placeholders without explicit definitions.

This chapter introduces **Named Classes**, which formalize each concept with a unique identifier, annotations, and clear semantic meaning. Named classes are critical because they allow the ontology to move from a conceptual structure into a **machine-readable semantic model**, laying the groundwork for relationships, reasoning, and knowledge graph construction.

Named classes serve as the **primary semantic nodes** in any OWL ontology. Each named class encapsulates a domain concept in a way that can be referenced consistently throughout the ontology. Unlike temporary or anonymous classes, named classes are explicitly identifiable, enabling them to participate in object properties, class restrictions, and individual instances. In the Pizza ontology, examples include `Pizza`, `VegetarianPizza`, `NonVegetarianPizza`, and various topping classes such as `CheeseTopping` and `MeatTopping`. Each of these named classes represents a real-world concept, formalized in a way that machines and humans can both understand and utilize.

From the perspective of **Executable Knowledge Architecture (EKA)**, named classes occupy a central position. In the EKA roadmap, knowledge moves from **diagrams and conceptual models** to **meta-models**, then to **formal ontologies**, and finally into **knowledge graphs and executable intelligence**. Named classes mark the transition from abstract conceptualization to structured ontology nodes. These nodes are the semantic backbone of knowledge graphs, forming the vertices that will later be linked via object properties and populated with instances to support reasoning and AI-driven inference. The quality and clarity of named classes directly affect the effectiveness of the ontology in generating **reliable, executable intelligence**.

- [5.1 The Concept of Named Classes](#)
- [5.2 Creating Named Classes in Protégé](#)
- [5.3 Best Practices for Named Class Design](#)
- [5.4 Linking Named Classes to the EKA Framework](#)
- [5.5 Practical Exercise: Completing the Named Class Structure](#)
- [5.6 Advanced Considerations](#)
- [5.7 Chapter \(05\) Summary](#)

- [5.8 Key Concepts](#)
- [5.9 Protégé Skills Learned](#)
- [5.10 Next Chapter \(06\) Preview](#)

5.1 The Concept of Named Classes

A **Named Class** in OWL is more than a label. It is a formally defined entity within a **unique URI**, a descriptive label, and semantic annotations. These classes serve as reference points in the ontology, providing the structure necessary for logical consistency and automated reasoning.

For instance, a class like `VegetarianPizza` is formally defined as:

```
VegetarianPizza ≡ Pizza and (hasTopping only (CheeseTopping or VegetableTopping))
```

it is not only a subclass of `Pizza` but also a node that can be referenced when defining constraints, object properties, or when populating the ontology with individual pizza.

Defining Named Classes requires a careful balance between **semantic clarity** and **practical usability**. Each class name must accurately reflect the concept it represents and be easily understood by both humans and automated systems. Poorly named classes can introduce confusion, compromise reasoning accuracy, and hinder the ontology's utility in enterprise-scale knowledge systems.

Named classes, therefore, are the first step toward transforming abstract domain knowledge into an ontology that is both **maintainable** and **executable**, perfectly aligning with EKA principles.

5.2 Creating Named Classes in Protégé

Creating named classes in Protégé is an intuitive but deliberate process.


After opening the `classes` tab, you begin by adding each concept as a named class. Unlike earlier placeholder classes, each Named Class is given a **unique, meaningful identifier**, which establishes it as a semantic entity in the ontology. In practice, this involves entering a class name, positioning it appropriately in the hierarchy, and adding labels and comments that clarify its purpose.

For example, the class `VegetarianPizza` is placed under `Pizza` in the hierarchy. Its label is "Vegetarian Pizza", and a comment might read: "A pizza containing only non-meat toppings, suitable for vegetarian diets."

Such annotations serve multiple purposes, they

- clarify intent for humans,
- guide ontology maintenance,
- and provide context for reasoning engines.

Similarly, CheeseTopping under Topping is labeled "Cheese Topping" with a comment like "Any pizza topping derived from cheese products."



Each annotation ensures that the semantic intent is explicit, facilitating both **knowledge reuse** and **interoperability**.

When adding named classes, it is crucial to maintain **hierarchical integrity**. Subclasses must logically inherit from their parents. For instance, `VegetarianPizza` is a subclass of `Pizza`, meaning all properties or restrictions applied to `Pizza` will propagate to `VegetarianPizza`. Maintaining this hierarchy ensures that the ontology can support **reasoning engines**, which depend on inheritance relationships to infer new knowledge.

5.3 Best Practices for Named Class Design

Creating effective named classes involves more than simply typing names into Protégé. One of the most important practices is **naming consistency**. All classes should follow a clear naming convention that reflects the domain logic and is easy for both humans and automated systems (machines) to interpret.

For example, all pizza types might use the suffix "Pizza", while toppings can consistently use the suffix "Topping".

Such conventions reduce ambiguity and improve the maintainability of the ontology.

Another critical practice is **annotation discipline**. Every named class should include descriptive labels and comments. This documentation provides semantic context, which becomes increasingly important as the ontology grows and as multiple engineers collaborate on its development. Well-annotated classes are also easier to integrate into **knowledge graphs**, where clear semantic definitions are essential for AI-driven inference.

Finally, class hierarchies must reflect **real-world domain logic**. Avoid unnecessary duplication and ensure that every subclass meaningfully specializes its parent.

For instance, `VegetarianPizza` should include all vegetarian varieties without overlapping with `NonVegetarianPizza`. This discipline preserves both semantic clarity and reasoning accuracy, essential for enterprise-grade ontologies.

5.4 Linking Named Classes to the EKA Framework

Within the **EKA framework**, named classes provide a direct link between conceptual design and executable intelligence.

The process can be seen as a multi-layer transformation:

1. **Diagramming Layer**: Conceptual pizza types and ingredient categories are visualized
2. **Meta-Model Layer**: Rules governing the relationships between pizzas, toppings, and dietary classifications are defined
3. **Ontology Layer**: Named classes formalize these concepts into OWL classes with clear hierarchy, labels, and annotations.
4. **Knowledge Graph Layer**: Each named class becomes a node that can participate in relationships and reasoning operations.
5. **Executable Intelligence Layer**: Reasoning engines and AI systems can now infer new knowledge, validate consistency, and support enterprise decision-making.

See below the Version 1.1 (as of when book is written) of EKA Framework:

 EKA1.1-Framework-en

You may visit <https://xiaoqi.com> for more detail about EKA

By defining named classes meticulously, you ensure that the ontology is not only correct but also **ready for AI-driven inference**, a key goal of EKA. In practice, this means that any subsequent steps -- such as defining object properties, adding restrictions, or creating individuals -- operate on a **semantically sound foundation**.

5.5 Practical Exercise: Completing the Named Class Structure

To solidify your understanding, this chapter's exercises guide learners through:

- Reviewing the Chapter 04 class skeleton
- Defining named classes for all key pizza types and toppings
- Applying consistent labels and comments for semantic clarity
- Organizing the hierarchy to ensure logical inheritance and scalability
- Documenting naming conventions and design rationale for maintainability

Upon completion, the Pizza ontology will have a **full articulated class layer**, forming a robust semantic foundation. This layer is now ready for the addition of relationships, property restrictions, and reasoning in subsequent chapters.

5.6 Advanced Considerations

As your ontology grows, consider:

- **Scalability:** Ensure that naming and hierarchy choices support the future addition of new pizza varieties, toppings, or dietary categories.
- **Interoperability:** Named classes should align with external standards when integrating with other ontologies or knowledge graphs.
- **Reasoning Readiness:** Clear class definitions prevent logical inconsistencies, enabling the reasoner to infer meaningful conclusions.
- **Enterprise Knowledge Alignment:** Named classes formalize organizational knowledge, making it reusable and actionable across business systems.

5.7 Chapter (05) Summary

In this chapter, you have:

- Defined and structured **Named Classes** in Protégé
- Applied annotations and labels for semantic clarity
- Organized subclass hierarchies to preserve inheritance
- Connected named classes to the **EKA framework**, ensuring they serve as nodes in the knowledge graph and form the basis for executable intelligence
- Prepared the ontology for reasoning and advanced semantic modeling

The Pizza ontology is now a **semantically robust structure**, ready to incorporate relationships, reasoning, and instance data.

5.8 Key Concepts

Concept	Description
Named Class	A formally defined, uniquely identifiable concept in OWL
Hierarchy	Parent-child relationships that define semantic inheritance
Annotation	Labels and comments clarifying meaning for humans and machines
EKA Integration	Named classes bridge diagrams and meta-models to knowledge graphs and executable intelligence

5.9 Protégé Skills Learned

- Creating and naming classes explicitly
- Organizing hierarchical structures for semantic inheritance
- Applying labels and comments consistently
- Preparing the ontology for reasoning and knowledge graph development

5.10 Next Chapter (06) Preview

In the next chapter, we will introduce **using a reasoner** in Protégé. This chapter will demonstrate how to:

- Apply a reasoning engine to your ontology

- Detect logical inconsistencies in class hierarchies
- Infer implicit relationships between classes and individuals
- Validate the ontology against OWL semantic rules

With the reasoning engine, your structured named classes will no longer be static; they will form the basis of a **dynamic, executable ontology** capable of supporting intelligent knowledge inference decision-making, and integration into enterprise knowledge systems under the EKA framework.

Last updated at 2026-09-04, thanks for @000wahab000's pull request

Chapter 06 -- Applying a Reasoner to the Pizza Ontology

Having established a robust semantic foundation in Chapter 05 with **Named Classes**, we are now prepared to advance to a pivotal step in ontology engineering: **reasoning**.

Reasoners are software engines that can automatically **infer knowledge**, detect inconsistencies, and validate the logical structure of an ontology. They transform your ontology from a static hierarchy of classes into a **dynamic, intelligent knowledge model**, capable of supporting AI-driven insights and enterprise decision-making.

Within the **EKA framework**, reasoners operate at the intersection of the **Ontology layer** and the **Knowledge Graph layer**. Named classes, which were previously defined and annotated, serve as the nodes, while the reasoner evaluates their relationships, constraints, and hierarchy to produce new knowledge or detect logical errors.

In essence, reasoning allows an ontology to become **executable intelligence**, aligning directly with EKA's goal of turning formal knowledge structures into actionable insights.

- [6.1 Understanding Reasoners: What They Are and Why They Matter](#)
- [6.2 Preparing the Ontology for Reasoning](#)
- [6.3 Running a Reasoner in Protégé](#)
- [6.4 Interpreting Reasoner Outputs](#)
- [6.5 Practical Exercise: Applying Reasoners to the Pizza Ontology](#)
- [6.6 One More Hands-on Sample](#)
- [6.7 EKA Perspective: Reasoners as Enablers of Executable Knowledge](#)
- [6.8 Common Challenges and Best Practices](#)
- [6.9 Chapter \(06\) Summary](#)
- [6.10 Key Concepts](#)
- [6.11 Protégé Skills Learned](#)
- [6.12 Next Chapter \(07\) Preview](#)
- [6.13 Demo Video for this Chapter \(06\)](#)

6.1 Understanding Reasoners: What They Are and Why They Matter

A **reasoner** is a software engine that evaluates an ontology's logical consistency and derives implicit knowledge. Think of it as a semantic "auditor" that ensures every class, subclass, and property aligns with the rules defined in the ontology. Unlike humans, reasoners can process thousands of axioms instantaneously, making them indispensable in large-scale or enterprise ontologies.

Reasoners perform the following three major functions:

1. **Consistency Checking:** Detecting logical contradictions, such as a pizza being both `VegetarianPizza` and containing `MeatTopping` .
2. **Classification:** Automatically computing subclass relationships that were not explicitly defined.
3. **Inference:** Deductively identifying implicit knowledge, e.g., a pizza containing only `VegetableTopping` can be inferred to be a `VegetarianPizza` .

In EKA, reasoners bridge the **ontology layer** and the **knowledge graph layer**, allowing named classes to become **dynamic nodes**. They ensure that enterprise semantic models are **cohesive, reliable, and ready for automated inference**, which is the essence of executable intelligence.

6.2 Preparing the Ontology for Reasoning

Before running a reasoner, the ontology must be **reasoner-ready**.

This includes ensuring:


- **Complete Named Class Hierarchy:** Chapter 05 provided this foundation. Every named class, such as `VegetarianPizza` , `NonVegetarianPizza` , and `Topping` subclasses, should be well-defined and logically positioned.
- **Annotation Completeness:** Labels (`rdfs:label`) and comments (`rdfs:comment`) help the reasoner interpret the ontology and improve maintainability for human collaborators.
- **Preliminary Logical Validation:** Manual checks for obvious contradictions in subclass relationships, e.g.,
no pizza should simultaneously inherit conflicting classes without explicit design .

These preparations ensure that the reasoner can **maximize its inference capabilities**, generating meaningful insights without false positives or errors.

6.3 Running a Reasoner in Protégé

Protégé supports several high-quality reasoners, including **HermiT**, **Pellet**, and **Fact++**.

Here is my view from Reasoner menu:



It is fine that you may have different Reasoners installed, I'm normally use `Pellet` reasoner. The `Start reasoner` is in gray means the reasoner has been started.

Using a reasoner is straightforward but requires careful observation of its outputs.

The steps are:

1. **Select a Reasoner:** From the Protégé toolbar, select `Pellet` (or your preferred reasoner, like `HermiT`).
2. **Start the Reasoner:** Activating it will initiate logical evaluation across all classes and properties
3. **Observe Inferred Hierarchies:** The reasoner updates subclass relationships to reflect implicit connections. For example, pizza defined only by their toppings may be automatically categorized into `VegetarianPizza` or `NonVegetarianPizza`.
4. **Examine Warnings and Errors:** Unsatisfiable classes or inconsistencies will be flagged. This provides immediate feedback on areas requiring correction.

Unlike manual validation, reasoners can handle complex ontologies with **thousands of classes**, uncovering subtle contradictions that might be missed by even experienced engineers.

6.4 Interpreting Reasoner Outputs

Once the reasoner completes, Protégé provides a variety of outputs:

- **Inferred Class Hierarchy:** Subclasses that were previously undefined but are logically implied are now visible.
- **Unsatisfiable Classes:** Classes that cannot logically exist given current axioms, such as a pizza defined as both `VegetarianPizza` and containing `MeatTopping`.
- **Inconsistency Warnings:** Highlight potential logical conflicts that may compromise reasoning.
- **Suggested Inferences:** Implicit relationships that can now be exploited in knowledge graphs.

Understanding these outputs is critical for refining the ontology. Each flagged issue is an opportunity to **improve semantic clarity**, enforce domain logic, and ensure that the ontology can serve as a **reliable knowledge graph backbone** in enterprise applications.

Important Note on Inference Materialization

In Protégé's default desktop workflow, the inferences you see in the interface are **displayed as inferred axioms** -- they are not automatically written back (materialized) as asserted triples into the saved ontology file. This is a common design choice for ontology editing environments, where the goal is to allow modelers to review and validate reasoning results before deciding whether to accept them.

However, this behavior should not be mistaken for a universal limitation of OWL reasoners or semantic systems -- while that happens on myself when I only learnt Protégé first. In production-grade environments -- such as **Native RDF Graph Databases (triplestores)**, rule-enabled graph databases, and enterprise knowledge graph platforms -- inferred triples may be:

- **Materialized and persisted** (written back to storage for performance),
- **Maintained as inferred closures** (updated incrementally as new data arrives), or
- **Exposed virtually** (computed dynamically at query time),

depending on the system's architecture and configuration.

*The author here gratefully acknowledges **Michael DeBellis** for his expert technical review and for clarifying this critical engineering distinction. The corrected wording in this note borrowed directly from his suggestions. (Ref: GitHub Issue #9 (https://github.com/yasenstar/protege_pizza/issues/9))*

Important

Graph Technology Clarification: This distinction between Protégé's UI behavior and production-grade inference materialization is particularly relevant when choosing your graph technology.

- **Native RDF Graph Databases** (AllegroGraph, GraphDB, Stardog, RDFox) can materialize inferred triples, maintain inferred closures, or expose inferences virtually.
- **Property Graphs** (Noe4j, TigerGraph) do not natively support OWL reasoning or inference materialization -- any reasoning must be done before transforming the data into the property graph model.

6.5 Practical Exercise: Applying Reasoners to the Pizza Ontology

To gain hands-on experience, you should:

1. Open the Pizza ontology with all named classes and hierarchy defined in Chapter 05.
2. Select the Pellet reasoner and start reasoning.
3. Observe the inferred class hierarchy and note any changes in subclass relationships.
4. Investigate unsatisfiable classes, and if necessary, adjust annotations or hierarchy to resolve inconsistencies.
5. Document inferred knowledge, noting how implicit relationships enrich the ontology.

Through this exercise, you gain insight into how **logical consistency and inference** transform static class hierarchies into **dynamic, actionable semantic models**, directly supporting EKA's executable intelligence layer.

6.6 One More Hands-on Sample

It is possible that from above 6.5's steps, you don't see any warning or errors, that's fine since your ontology is still clean enough and no error.


To give you some tangible feeling, follow me - using the ontology file from Chapter 05 - and let's create the inconsistency.

Objective: we will have one pizza instance under `VegetarianPizza` subclass and let it contain a `MeatTopping`, and that should trigger inconsistency warning with the rule `VegetarianPizza` should only

contain CheeseTopping or VegetableToppings .

Just practice, some knowledge will be mentioned in later chapters, so if you don't understanding for now, no worry at all.

Firstly, let's create one new Object Property called `hasTopping` , to keep simple, we leave the `Domain` and `Range` no defined.



Secondly, let's create two instances:

- `VegetarianPizza1` within class `VegetarianPizza`

Active ontology × Entities × Individuals by class × OWLViz × DL Query × OntoGraf × SPARQL Query

Class hierarchy: VegetarianPizza

Annotations Usage

Annotations: VegetarianPizza1

Annotations +

Description: VegetarianPizza1

Types +

VegetarianPizza

Same Individual As +

Different Individuals +

Individuals Individuals (Inferred)

Direct instances: VegetarianPizza1

For: VegetarianPizza

VegetarianPizza1

```
graph TD
    owlThing[owl:Thing] --> Pizza[Pizza]
    Pizza --> NonVegetarianPizza[NonVegetarianPizza]
    Pizza --> VegetarianPizza[VegetarianPizza]
    VegetarianPizza --> Topping[Topping]
    Topping --> CheeseTopping[CheeseTopping]
    Topping --> MeatTopping[MeatTopping]
    Topping --> VegetableTopping[VegetableTopping]
```


- MeatTopping1 within class MeatTopping

The screenshot shows the OWL Viz interface with the following components:

- Top Bar:** Active ontology, Entities, Individuals by class, OWLViz, DL Query, OntoGraf, SPARQL Query.
- Class hierarchy: MeatTopping:** A tree view showing the hierarchy. The **MeatTopping** class is highlighted with a red box. The hierarchy is: **owl:Thing** (parent) → **Pizza** (child) → **NonVegetarianPizza** (child) → **VegetarianPizza** (child) → **Topping** (child) → **CheeseTopping** (child) → **MeatTopping** (child) → **VegetableTopping** (child).
- Annotations:** A tab showing annotations for **MeatTopping1**. It includes a section for **Annotations** with a plus sign.
- Description: MeatTopping1:** A tab showing the description of **MeatTopping1**. It includes a section for **Types** with a plus sign, and a section for **Same Individual As** and **Different Individuals** with plus signs.
- Individuals:** A tab showing individuals. It includes a section for **Direct instances: MeatTopping1** with a plus sign. Below this, there is a section for **For: MeatTopping** with a plus sign. The **MeatTopping1** instance is highlighted with a red box.


Thirdly, we tell ontology our expected rule, to do this:

1. Switch to Tab Entities then Classes
2. Expand class Pizza and click VegetarianPizza (Note: select the class, not instance!)
3. In Description area, click + sign in the right of Equivalent To
4. In tab Class expression editor, write down the rule
Pizza and (hasTopping only (CheeseTopping or VegetableTopping))
5. Click OK




Now the fourth, let's create the inconsistency in reality, with following steps to **force** set the MeatTopping1 to VegetarianPizza1 :

1. Switch to tab Individuals by class
2. Click class VegetarianPizza
3. In the Direct instances area below class hierarchy, click VegetarianPizza1 (Note: you may only have this one for now)
4. In the right Property assertions area, click "+" sign to add Object property assertions
5. Ensure you're in English language input, the left key in the object property hasTopping , the right key in target instance MeatTopping1 , click OK
6. You should see hasTopping MeatTopping1 under Object property assertions




After above four steps, ensure your reasoner is running already, now do **Synchronize reasoner**, the **Magic** warning should be shown as below




Error message:

InconsistentOntologyException: Cannot do reasoning with inconsistent ontologies!
Reason for inconsistency: Individual <http://www.semanticweb.org/xiaoqi/ontologies/2026/4/pizza-ebook#MeatTopping1> is forced to belong to class <http://www.semanticweb.org/xiaoqi/ontologies/2026/4/pizza-ebook#VegetableTopping> and its complement

This is what we expected, so no need be afraid of. Simply close that, and you may click the  triangle in the bottom right corner to open the detail log window and read more detail information:



Once the reasoner detects such kind of issues, your ontology may have many items are shown in RED color for warning purpose:



To fix this inconsistency is fairly easy, you now get noticed that the `MeatTopping1` is wrongly added to one `VegetarianPizza1`, just delete that `Object` property assertions (in reality, you may also choose the proper non-MeatToppings), then Synchronize Reasoner again, you can see the warning is disappeared.

Exciting to see the power from reasoner? Congratulation for you've done this and let's continue.

6.7 EKA Perspective: Reasoners as Enablers of Executable Knowledge

In EKA, reasoning is not merely a technical task -- it is a **strategic enabler** of executable intelligence. Reasoners convert the ontology into a **living semantic network**, ensuring that knowledge is **consistent, inferable, and actionable**.

- **Ontology Layer:** Named classes serve as nodes; reasoners validate their integrity.
- **Knowledge Graph Layer:** Inferred subclass relationships and logical connections establish edges between nodes.
- **Executable Intelligence:** Reasoners empower AI-driven applications to make decisions based on robust, verified knowledge structures.

By incorporating reasoners, enterprises gain the ability to **automatically validate knowledge**, detect contradictions, and ensure that semantic models scale without human error -- a key requirement for knowledge-driven organizations.

6.8 Common Challenges and Best Practices

While reasoners are powerful, ontology engineers must be mindful of:

1. **Complexity:** Large ontologies may slow reasoning. Use modularization if necessary.
2. **Unsatisfiable Classes:** Always investigate why a class is flagged to avoid misinterpretation.
3. **Annotation Discipline:** Clear labels and comments improve both human and machine understanding.
4. **Hierarchy Integrity:** Ensure subclass relationships remain logical, particularly as new classes or properties are added.

Following these best practices ensures the reasoning remains **efficient, reliable, and aligned with enterprise-scale semantic modeling goals**.

6.9 Chapter (06) Summary

In this chapter 06, you have:

- Introduced reasoning as a core methodology for validating and enriching ontologies.

- Learned to run a reasoner in Protégé and interpret outputs such as inferred hierarchies and unsatisfiable classes
- Observed how reasoning transforms named classes into a **dynamic semantic model**.
- Connected reasoning processes to the **EKA framework**, linking ontology validation with knowledge graph construction and executable intelligence.
- Prepared the ontology for more advanced modeling concepts, including `Disjoint Classes` in Chapter 07

With reasoning applied, the Pizza ontology evolves from a `static hierarchy of named classes` into a `living semantic structure` capable of supporting intelligent inference and enterprise knowledge operations.

6.10 Key Concepts

Concept	Description
Reasoner	Software that checks consistency, infers subclass relationships, and detects logic contradictions.
Consistency Checking	Validates that all ontology elements follow defined rules.
Inference	Deduction of implicit knowledge from explicit axioms.
EKA Integration	Reasoners convert ontologies into actionable, intelligent knowledge graphs.

6.11 Protégé Skills Learned

- Selecting and running reasoner
- Interpreting inferred class hierarchies and warnings
- Identifying and resolving unsatisfiable classes
- Documenting inferred relationships for knowledge graph integration

6.12 Next Chapter (07) Preview

The next chapter introduces **Disjoint Classes**, a mechanism for explicitly declaring that certain classes cannot overlap. For example, `VegetarianPizza` and `NonVegetarianPizza` will be formally disjoint to prevent logical conflict. Understanding disjoint classes is essential for maintaining **semantic integrity** and ensuring that the reasoner can operate accurately on your ontology.

6.13 Demo Video for this Chapter (06)

Demo video reference: YouTube - Chapter 06 (<https://youtu.be/TKMW5udKzIM>)

Last updated at 2026-08-29

Chapter 07 -- Ensuring Semantic Integrity with Disjoint Classes

After learning to structure your Pizza ontology and applying a reasoner in Chapter 06, your ontology now possesses a **dynamic semantic hierarchy** and the ability to detect inconsistencies. However, even with reasoning in place, logical errors can persist if classes that should not overlap are not explicitly constrained. This is where **Disjoint Classes** come into play. By defining disjoint classes, you as ontology engineers can enforce **semantic exclusivity**, preventing logically incompatible classifications and ensuring the integrity of inferred knowledge.

Disjoint classes are critical in enterprise ontologies because they formalize domain rules that may otherwise rely solely on reasoning inference. In our Pizza ontology, for example, `VegetarianPizza` and `NonVegetarianPizza` are mutually exclusive. Without declaring them disjoint, a reasoner might not flag subtle modeling errors, potentially allowing a pizza to be categorized as both. Disjointness introduces **explicit semantic constraints**, which enhance automated reasoning and maintain **knowledge graph accuracy** in the EKA framework.

- [7.1 Understanding Disjoint Classes](#)
- [7.2 Defining Disjoint Classes in Protégé](#)
- [7.3 Linking Disjoint Classes to the EKA Framework](#)
- [7.4 Practical Implications and Examples](#)
- [7.5 Running the Reasoner with Disjoint Classes](#)
- [7.6 Practical Exercise: Implementing Disjoint Classes](#)
- [7.7 Best Practices for Disjoint Classes](#)
- [7.8 Chapter \(07\) Summary](#)
- [7.9 Key Concepts](#)
- [7.10 Protégé Skills Learned](#)
- [7.11 Next Chapter \(08\) Preview](#)

7.1 Understanding Disjoint Classes

A **disjoint class** is an OWL concept indicating that two or more classes cannot share instances. If an individual is asserted to belong to one class, it **CANNOT belong to any of the disjoint classes**.

Disjoint declarations help:

- Prevent **logical contradictions** in the ontology
- Aid reasoners in **consistency checking**
- Ensure that **inferred relationships** are semantically sound

In the Pizza ontology, some of the key disjoint relationships include:

- VegetarianPizza and NonVegetarianPizza
- CheeseTopping and MeatTopping
- SpicyTopping and NonSpicyTopping

Note: for Class names, we use standard Camel convention.

By defining these disjoint relationships, we provide the reasoner with **rules that cannot be violated**, which helps maintain the ontology's **semantic correctness** as it grows in complexity.

Note

VegetarianPizza is formally defined as `VegetarianPizza ≡ Pizza and (hasTopping only (CheeseTopping or VegetableTopping))`

This means a vegetarian pizza must have a pizza base and can only have cheese or vegetable toppings. It cannot have meat or seafood (or other) toppings.

7.2 Defining Disjoint Classes in Protégé

Creating disjoint classes in Protégé is straightforward but requires careful planning:


1. Navigate to the **classes** tab and select the class you wish to make disjoint with others.

The screenshot displays the Protege OWL editor interface. The top navigation bar includes tabs for 'Active ontology', 'Entities', 'Individuals by class', 'OWLViz', 'DL Query', 'OntoGraf', and 'SPARQL Query'. The 'Entities' tab is active, showing a left sidebar with 'Classes', 'Object properties', and 'Data properties'. The 'Classes' tab is selected, displaying a class hierarchy for 'SpicyTopping'. The hierarchy shows 'owl:Thing' as the root, with 'Pizza' as a child. 'Pizza' has a child 'Topping', which in turn has children 'NonSpicyTopping', 'SpicyTopping', 'CheeseTopping', 'MeatTopping', and 'VegetableTopping'. The 'SpicyTopping' class is highlighted with a red box. A tooltip 'Scroll to selection' is visible near the 'Topping' class. The main area on the right shows the 'Annotations: SpicyTopping' and 'Description: SpicyTopping' panels. The 'Description' panel lists various class relationships: 'Equivalent To', 'SubClass Of' (with 'Topping' selected), 'General class axioms', 'SubClass Of (Anonymous Ancestor)', 'Instances', 'Target for Key', 'Disjoint With', and 'Disjoint Union Of'.

2. Open the **Disjoint With** section and select the classes that should be mutually exclusive.

The screenshot shows the Protege software interface for editing the 'pizza-ebook' ontology. The main window displays the class hierarchy for 'SpicyTopping', which includes 'Pizza', 'Topping', 'CheeseTopping', 'MeatTopping', 'NonSpicyTopping', 'SpicyTopping', and 'VegetableTopping'. The 'NonSpicyTopping' class is highlighted with a red box. A red arrow points from this box to the 'Disjoint With' button in the 'Description: SpicyTopping' panel. Another red arrow points from the 'SpicyTopping' class in the hierarchy to the same 'Disjoint With' button. The 'Disjoint With' button is also highlighted with a red box. The 'OK' button in the 'SpicyTopping' dialog is also highlighted with a red box.

3. Save changes and run the reasoner to check for conflict



From above screens sample, when declaring `NonSpicyTopping` disjoint with `SpicyTopping`, the reasoner immediately flags any pizza mistakenly classified under both classes. This ensures that **all individuals are consistently assigned**, reducing errors in semantic inference and knowledge graph construction.

As below, simply create one topping instance called `myTopping` and add its Types to both `SpicyTopping` and `NonSpicyTopping`, you may see the two types are highlighted in red after synchronizing the reasoner.

The screenshot shows the Protege OWL editor interface. The top menu bar includes File, Edit, View, Reasoner, Tools, Refactor, Window, and Help. The address bar shows the URL: <http://www.semanticweb.org/xiaoqi/ontologies/2026/4/pizza-ebook>. The main window has tabs for Active ontology, Entities, Individuals by class, OWLViz, DL Query, OntoGraf, and SPARQL Query. The left pane shows the Class hierarchy for SpicyTopping, with a tree structure including owl:Thing, Pizza, Topping, CheeseTopping, MeatTopping, Non SpicyTopping, SpicyTopping, and VegetableTopping. The right pane shows the Description of MyTopping, with a red box highlighting the Types section, which lists Non SpicyTopping and SpicyTopping. Below the types, there are sections for Same Individual As, Different Individuals, Object property assertions, Data property assertions, Negative object property assertions, and Negative data property assertions. A red arrow points from the 'For: SpicyTopping' section in the left pane to the 'Types' section in the right pane.

The error popup is as below:

The screenshot shows an error message popup in the Protege editor. The message reads: "An error occurred during reasoning. InconsistentOntologyException: Cannot do reasoning with inconsistent ontologies! Reason for inconsistency: Individual <http://www.semanticweb.org/xiaoqi/ontologies/2026/4/pizza-ebook#MyTopping> is forced to belong to class <http://www.semanticweb.org/xiaoqi/ontologies/2026/4/pizza-ebook#NonSpicyTopping> and its complement". The popup has an OK button. The background shows the same Protege interface as the previous screenshot, with the error message overlaying the main content.

Error message is:

InconsistentOntologyException: Cannot do reasoning with inconsistent ontologies!

Reason for inconsistency: Individual <http://www.semanticweb.org/xiaoqi/ontologies/2026/4/pizza-ebook#MyTopping>

To make the ontology working, you can simply remove `MyTopping` from either `SpiceTopping` or `NonSpicyTopping`, re-synchronize reasoner, the type left will be turned from red to black. -- conflict is fixed and your ontology is back to consistency.

In practice, disjoint classes often reflect **real-world domain constraints**. In enterprise ontology modeling, failing to define such relationships can result in **incorrect inferences** that propagate throughout downstream knowledge graphs, potentially affecting AI-driven decision-making.

7.3 Linking Disjoint Classes to the EKA Framework

Within EKA, disjoint classes enforce **semantic rigor** between the ontology and the resulting knowledge graph.

Important

Critical Distinction: The "resulting knowledge graph" in EKA can be implemented as either a Native RDF Graph Database or a Property Graph.

- In a **Native RDF Graph**, disjointness constraints are preserved and can be enforced by the reasoner.
- In a **Property Graph**, disjointness becomes a property or label constraint that the database does not inherently understand as a logical rule.

for full EKA implementations that rely on automated reasoning and semantic integrity, a Native RDF Graph is the recommended choice for the *K* layer.

The impact is multi-layered:

- **Ontology Layer:** Disjoint declarations define hard rules between semantic nodes (classes).
- **Knowledge Graph Layer:** Edges representing relationships between individuals cannot violate these constraints.
- **Executable Intelligence Layer:** AI and reasoning engines operate on a consistent, validated knowledge structure, ensuring trustworthy insights.

In other words, disjoint classes prevent ambiguity, preserve data integrity, and allow **enterprise knowledge models to scale** without introducing conflicts -- essential in professional, large-scale EKA deployments.

7.4 Practical Implications and Examples

Consider the following scenarios in the Pizza ontology:

- If `VegetarianPizza` and `NonVegetarianPizza` were not declared disjoint, a pizza containing both vegetables and meat might erroneously be inferred to belong to both classes.
- By making them disjoint, the reasoner flags this as inconsistent, prompting corrective action.

- Similarly, disjointing `CheeseTopping` and `MeatTopping` ensures that inferred dietary classifications remain accurate, which is critical for applications like menu recommendation engines or AI dietary planning tools.

Explicitly modeling disjointness **reduces semantic ambiguity** and strengthens reasoning outputs. In enterprise contexts, this discipline ensures that the **knowledge graph reflects true business rules** rather than merely inferred relationships.

7.5 Running the Reasoner with Disjoint Classes

Once disjoint classes are defined, the next step is to **re-run the reasoner** in Protégé:

1. Activate your preferred reasoner (e.g., Pellet or HermiT).
2. Observe any newly detected inconsistencies flagged due to disjoint constraints.
3. Examine the **inferred class hierarchy** to see if any individuals or subclasses conflict with the disjoint rules.
4. Correct any errors to ensure consistency across the ontology.

By combining named classes, reasoning, and disjoint declarations, the ontology achieves a **robust semantic framework** capable of supporting complex knowledge representation, ready for eventual integration into knowledge graphs and AI applications.

7.6 Practical Exercise: Implementing Disjoint Classes

1. Open the Pizza ontology with reasoned hierarchy from Chapter 06.
2. Identify pairs of classes that should never share instances (e.g., `VegetarianPizza` VS. `NonVegetarianPizza`).
3. Use the **Disjoint With** feature in Protégé to enforce mutual exclusivity.
4. Run the reasoner and confirms that conflicts are flagged appropriately.
5. Document how these constraints improve consistency and reliability.

This exercise reinforces the importance of **explicit semantic constraints** in ontology design, bridging the gap between structured knowledge and executable intelligence in EKA.

7.7 Best Practices for Disjoint Classes

- **Plan Before Declaring:** Ensure that disjoint classes reflect true domain logic to avoid unnecessary conflicts.
- **Use Sparingly but Strategically:** Excessive disjoint declarations may complicate reasoning, so focus on critical semantic exclusivity.
 - This is especially valid when you create classes hierarchy bulk from `Tools` menu, simply check the `Disjoint` will create lots of disjoint by default!
- **Maintain Hierarchical Integrity:** Disjoint declarations should not contradict subclass relationships.
- ****Integrate with EKA Principles:** Always consider how disjoint classes affect the downstream knowledge graph and reasoning outcomes.

By following these principles, ontologies remain **clean, scalable, and ready for advanced AI integration**.

7.8 Chapter (07) Summary

In Chapter (07), you have:

- Learned the theory and purpose of **Disjoint Classes** in OWL ontologies.
- Applied disjoint declarations in Protégé to enforce semantic exclusivity.
- Observed how disjoint classes enhance reasoning outputs and prevent logical conflicts.
- Linked disjoint modeling to the **EKA framework**, ensuring semantic integrity in knowledge graphs and executable intelligence.
- Prepared the ontology for advanced structural inspection, including RDF file understanding in Chapter 08 (next one).

With disjoint classes in place, the Pizza ontology is now a **highly consistent and reliable semantic model**, ready for **export, integration, and further refinement**.

7.9 Key Concepts

Concept	Description
Disjoint Class	A class explicitly defined to have no instances in common with other classes.
Semantic Exclusivity	Ensures logical separation between incompatible concepts.
Consistency Enforcement	Strengthens reasoning reliability and knowledge graph integrity.
EKA Integration	Disjoint classes support executable intelligence by maintaining accurate semantic relationships.

7.10 Protégé Skills Learned

- Identifying and defining disjoint classes
- Enforcing semantic exclusivity in the ontology
- Running the reasoner to validate disjoint constraints
- Documenting and maintaining logical integrity for enterprise-scale ontologies

7.11 Next Chapter (08) Preview

The next chapter (08) explores **understanding the RDF file structure**, the format in which OWL ontologies is serialized. You will see how the Pizza ontology is represented in **RDF/XML**, how classes, individuals, and

properties are encoded, and how this structure facilitates **sharing, integration, and reasoning** across systems.

Last updated at 2026-09-05

Chapter 08 -- Understanding the RDF File Structure: Looking Beneath Protégé into the Language of Semantic Knowledge

Before we begin this chapter, it is important to clarify something to you.

This chapter intentionally goes **beyond the original scope of Michael DeBellis' Pizza.owl tutorial**.

The Pizza tutorial primarily focuses on helping you understand ontology engineering through practical, hands-on exercises inside Protégé. Its emphasis is naturally centered on learning OWL concepts, building classes, using reasoners, and understanding semantic modeling fundamentals.

However, through years of practical work in enterprise architecture, ontology engineering, meta-modeling, graph databases, and knowledge-driven systems, I have repeatedly observed a common challenge among learners:

Many people learn how to **use the tool**, but far fewer truly understand **what the tool is generating underneath**.

This distinction matters!

Previously, while teaching enterprise architecture and meta-modeling, I introduced a similar perspective through analysis of **ArchiMate model structure and exchange formats**. Many enterprise architects become proficient at drawing ArchiMate diagrams but never examine the underlying model exchange specification that makes architecture portable, interoperable, and executable across repositories and tools.

Eventually, I realize ontology learning suffers from a similar challenge.

Protégé is simply the editor. (→ mapping to Archi, the ArchiMate Modeling Tool)

The ontology is the language. (→ mapping to ArchiMate, the language)

For this reason, Chapter 08 intentionally introduces a more theoretical perspective. Instead of treating ontology merely as a modeling exercise, we will examine ontology from the viewpoint of **RDF language specification**, semantic representation, and practical implementation into **Knowledge Graphs**.

This perspective becomes especially important if your long-term ambition extends beyond Pizza.owl toward:

- Enterprise semantic modeling
- Knowledge graph engineering
- AI-ready knowledge systems
- Executable Knowledge Architecture (EKA)

Because eventually ontology engineers must answer a deeper question:

What exactly is an ontology made of?

The answer begins with:


RDF -- Resource Description Framework.

This chapter therefore marks an important transition in the EKA roadmap:



Up until now, we have focused primarily on the **Ontology phase**.

Beginning here, we start preparing for the next transformation:



And RDF is the bridge that makes this transition possible.

- [8.1 Protégé Is the Editor, RDF Is the Ontology](#)
- [8.2 Understanding RDF from a Language Specification Perspective](#)
- [8.3 Understanding RDF Core Concepts Through `Pizza.owl`](#)
- [8.4 Reading RDF/XML Inside `Pizza.owl`](#)
- [8.5 Why RDF Is Already Graph-Shaped](#)
- [8.6 From Ontology to Knowledge Graph: Why RDF Naturally Evolves into Graph Databases](#)
- [8.7 Understanding the Relationship Between Ontology and Graph Database](#)
- [8.8 The Practical EKA Pipeline: From `Pizza.owl` to Knowledge Graph](#)
- [8.9 Practical Graph Integration: Two Paths from `Pizza.owl`](#)
 - [8.9.1 Path 1: Native RDF Graph Database \(Recommended for EKA\)](#)
 - [8.9.2 Path 2: Property Graph](#)
 - [8.9.3 Choosing the Right Path for Your EKA Implementation](#)
- [8.10 EKA Perspective -- RDF as the Missing Bridge](#)
- [8.11 Chapter \(08\) Summary](#)
- [8.12 Next Chapter \(09\) Preview](#)
- [8.13 Reference Demo Video](#)

8.1 Protégé Is the Editor, RDF Is the Ontology

One of the most important conceptual transitions in ontology engineering is recognizing a subtle but fundamental truth:

When you see in Protégé is not the ontology itself.

Rather, what you see is a **visual abstraction of the ontology**.

This idea may initially feel surprising because you spend most of your time interacting with:

- The Classes tab
- The Object Properties tab
- The Individuals tab
- The Reasoner interface

Everything feels highly visual and tool-oriented.

However, underneath Protégé exists something far more important.

A standards-based semantic representation language.

Whenever Protégé saves an ontology, it serializes knowledge into formal semantic structures using standards based on RDF and OWL.

In other words:

Protégé is simply an **ontology authoring environment**.

The real ontology exists independently of the tool.

This mirrors an important lesson from enterprise architecture.

When architects create ArchiMate diagrams, the diagram itself is not the architecture model. The real model becomes portable through a formal exchange structure.

Likewise:

Protégé diagrams are not the ontology.

The RDF/OWL representation is.

This serialization fundamentally changes how ontology engineers think.

You stop thinking:

"I built something in Protégé."

And begin thinking:

"I authored a semantic knowledge specification."

That is a major maturity leap.

8.2 Understanding RDF from a Language Specification

Perspective

Many beginner's explanations describe RDF as:

"A graph data format."

While technically correct, this explanation is incomplete.

To understand RDF professionally, we should instead view it as:

A formal language specification for representing semantic assertions.

RDF was standardized to solve an important problem:

How can machines understand meaning and relationships?

Traditional systems primarily store data.

RDF stores **meaning**.

This distinction is profound.

Relational databases organize information through:

- Tables
- Columns
- Primary keys
- Foreign key relationships

But RDF organizes knowledge differently.

Instead of tables, RDF models everything as **statements**.

Every semantic statement follows the same structure - triple:

Subject --> Predicate --> Object

For example:

VegetarianPizza --> subClassOf --> Pizza

In semantic terms:

- Subject = VegetarianPizza
- Predicate = subClassOf
- Object = Pizza

Similarly:

```
CheesePizza --> hasTopping --> CheeseTopping
VegetarianPizza --> hasTopping --> VegetableTopping
```

These are not simply data records.

They are **formal semantic assertions**.

This distinction becomes extremely important for reasoning.

Reasoners do not interpret diagrams.

Reasoners interpret semantic statements.

And RDF is the language they understand.

8.3 Understanding RDF Core Concepts Through `Pizza.owl`

To understand RDF deeply, you should begin thinking like language designers.

Resource

Everything meaningful in RDF is modeled as a **resource**.

Examples from `Pizza.owl` include:

- `Pizza`
- `VegetarianPizza`
- `CheeseTopping`
- `TomatoTopping`

Each concept receives a globally identifiable reference.

This solves ambiguity and enables semantic interoperability.

For enterprise systems, this becomes extremely important because concepts can be reused across repositories, APIs, metadata platforms, and knowledge graphs.

Property

Relationships themselves are also explicitly defined.

Examples include:

- `subClassOf`
- `disjointWith`
- `hasTopping`

Unlike traditional systems where relationships may feel secondary, RDF treats relationships as **first-class semantic meaning**.

With my sayings in certain EA tooling practices:

- Traditional EA Tool: diagramming first, relationships follow
- RDF: relationships first, visualization (diagrams) follows

Triple

The triple becomes RDF's atomic unit of knowledge.

Everything eventually becomes:

Subject --> Predicate --> Object

Like sentences in human language.

This elegant simplicity explains why RDF scales effectively for knowledge systems.

Complex ontologies are ultimately collections of semantic statements.

8.4 Reading RDF/XML Inside Pizza.owl

At first glance, RDF/XML syntax may appear intimidating.

You may encounter elements such as:

- `rdf:RDF`
- `owl:Class`
- `rdf:about`
- `rdfs:subClassOf`

Initially, many learners assume they must understand XML deeply.

In reality, ontology engineers should instead learn to read RDF/XML as **semantic meaning**.

For example:

```
<owl:Class rdf:about="#VegetarianPizza">
  <rdfs:subClassOf rdf:resource="#Pizza"/>
</owl:Class>
```

Rather than viewing this as technical syntax, translate it semantically:

VegetarianPizza is a subclass of Pizza.

Similarly:

```
<owl:disjointWith rdf:resource="#NonVegetarianPizza"/>
```

means:

VegetarianPizza cannot be NonVegetarianPizza.

This perspective is transformative.

You stop reading RDF/XML as code.

You begin reading RDF/XML as **machine-readable meaning**.

That is precisely how semantic technologies should be understood.

8.5 Why RDF Is Already Graph-Shaped

One of the most important realizations in ontology engineering is this:

RDF is already a graph.

Every RDF triple naturally forms:

Node --> Relationship --> Node

Example:

VegetarianPizza ↓ subClassOf Pizza

From ontology perspective:

VegetarianPizza is a subclass of Pizza.

From graph perspective:

VegetarianPizza node connected to Pizza node through SUBCLASS_OF relationship.

The transformation is almost direct.

This insight becomes strategically important.

When ontology engineers understand RDF deeply, they begin realizing:

We are already modeling graph structures.

Protégé may appear to be an ontology editor.

But underneath, you are constructing graph-ready semantic structures.

This realization often represents a major maturity leap for practitioners, especially enterprise architects transitioning toward **knowledge engineering**.

8.6 From Ontology to Knowledge Graph: Why RDF Naturally Evolves into Graph Databases

At this point, you may naturally ask:

If ontology already represents semantic knowledge, why do we still need a graph database?

Good catch! This question deserves a deeper answer.

Ontology and graph databases are related, but they solve different problems. However, it is critical to distinguish between two types of graph databases, as this choice profoundly affects your EKA implementation.

	Ontology	Graph Databases
Focus on	Semantic meaning and formal logic	Operational querying, scalable relationship traversal, and executable intelligence
Answer questions such as	• What does a concept mean? • What rules govern relationships? • What constraints must remain true?	• How are concepts connected? • What paths exist between entities? • What hidden patterns emerge? • What knowledge can be operationalized?
Provides	Semantic truth	Operational execution

Important

Critical Distinction: Native RDF vs. Property Graph

When we say "RDF naturally evolves into graph databases", we are specifically referring to **Native RDF Graph Databases** (also called RDF triplestores). These databases:

- Store data as RDF triples natively
- Use global IRIs as identifiers (preserving Linked Data compatibility)
- Support SPARQL queries
- Can perform OWL reasoning directly

Examples include GraphDB, Startdog, AllegroGraph, RDFox, and Amazon Neptune (in RDF mode).

Property Graphs (such as Neo4j, JanusGraph, and TigerGraph) are different. They:

- Use internal, non-semantic element IDs

- Are optimized for fast graph traversal algorithms
- Do not natively support OWL reasoning
- Require a lossy transformation to load OWL data

While an OWL ontology *can* be loaded into a property graph, it is a **transformation**, not a direct mapping. You lose the explicit semantics, IRI-based identity, standardized query compatibility, and reasoning capabilities that make OWL powerful.

For an EKA system that aims to preserve and execute the full logical semantics of its OWL ontology, a Native RDF Graph is the appropriate foundational choice.

Back to the table, this ontology vs. graph database distinction is fundamental in EKA.

Many organizations successfully document enterprise knowledge.

Far fewer operationalize it.

The problem is often not lack of modeling.

The problem is lack of execution.

This is exactly why the EKA roadmap exists.

Each layer increases knowledge maturity.

8.7 Understanding the Relationship Between Ontology and Graph Database

One of the biggest misconceptions in knowledge engineering is assuming ontology and graph databases compete with one another.

In fact, they do not!

They complement one another.

A useful way to understand this relationship is:

- Ontology = Semantic Brain
- Native RDF Graph Database = Semantic Execution Engine
- Property Graph = Performance-Optimized Traversal Engine



For EKA practitioners, this distinction is critical:

- A **Native RDF Graph Database** understands the ontology directly. It can run reasoners, answer SPARQL queries, and maintain semantic consistency without losing meaning.
- A **Property Graph** requires you to *transform* your ontology into a different data model. It may be faster for certain queries (like pathfinding), but it no longer "understands" the OWL semantics.

Ontology defines meaning.

For example:

- What qualifies as `VegetarianPizza` ?
- What topologies are allowed?
- What concepts are disjoint?

Ontology answers:

What should be true?

Meanwhile, graph databases operationalize these semantics.

For example:

- Which pizzas contain mozzarella?
- Which pizzas violate vegetarian rules?
- What toppings commonly co-occur?

Graph databases answer:

What is happening?

This mirrors enterprise architecture practice.

Meta-models define meaning.

Repositories operationalize knowledge.

This is precisely why ontology and graph databases work naturally together.

8.8 The Practical EKA Pipeline: From `Pizza.owl` to Knowledge Graph

Inside EKA, ontology is not the final destination.

It is an intermediate maturity stage.

The real goal is:

Executable Intelligence

A practical implementation path may look like this.

Step 1 -- Conceptual Modeling

Architects begin with diagrams and conceptual structures.

Knowledge remains visual.

Step 2 -- Meta-Modeling

Relationships become formalized.

Meaning becomes structured.

Step 3 -- Ontology Engineering in Protégé

Concepts become formal semantic definitions.

Now we define:

- Classes
- Restrictions
- Reasoning logic
- Disjointness
- Semantic inheritance

Step 4 -- RDF Serialization

Protégé exports semantic knowledge as:

- RDF/XML
- OWL
- Turtle (.ttl)

At this stage, the ontology becomes portable.

This is where many organizations stop.

However, semantic knowledge still remains underutilized.

Step 5 -- Knowledge Graph Implementation

RDF now becomes operational.

Semantic structures can be imported into graph databases. The choice depends on your goals:

- For a **Native RDF Graph** (e.g., AllegroGraph, GraphDB, Stardog), the imports is a direct mapping from RDF triples to the database's native storage model.
- For a **Property Graph** (e.g., Neo4j), the import requires a **lossy transformation** (see [section 8.9](#))

Classes become nodes.

Properties become relationships.

Semantic triples become traversable graph structures.

The ontology becomes queryable.

Step 6 -- Executable Intelligence

Once implemented as a graph, organizations can begin asking intelligent questions.

For example:

- What business capabilities depend on a deprecated application?
- What systems are impacted by a regulation change?
- What architecture elements violate governance rules?
- What semantic dependencies exist across enterprise domains?

Suddenly:

Knowledge becomes executable.

This is the vision of EKA.

8.9 Practical Graph Integration: Two Paths from `Pizza.owl`

A natural next question becomes:

Can `Pizza.owl` be imported into a graph database?

The answer is:

Yes -- but the "how" depends on what you want to preserve.

8.9.1 Path 1: Native RDF Graph Database (Recommended for EKA)

Tool examples: AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune (RDF mode)

Process:

1. Export from Protégé as RDF/XML, OWL, or Turtle.
2. Load the RDF file directly into the triplestore.

3. The ontology becomes immediately queryable via SPARQL
4. OWL reasoning can be performed natively.

What you preserve:

- All OWL semantics (classes, individuals, properties, restrictions, disjointness)
- Global IRIs for all entities
- Reasoning capabilities
- Compatibility with Linked Data

Result:

Pizza.owl → Native RDF Graph → SPARQL queries + Reasoning/Rules

8.9.2 Path 2: Property Graph

Tool examples: Neo4j, TigerGraph, JanusGraph, Amazon Neptune (property graph mode)

Process:

1. Export from Protégé to RDF/XML, OWL, or Turtle.
2. Transform RDF triples into required graph structure (node, edges, and properties) (see below).
3. Load the transformed data into the property graph.
4. Query using Cypher or Gremlin (depending on different tool).

What you lose (trade-offs):

- **IRI-based identity** → Property graphs use internal IDs, losing web-scale interoperability.
- **OWL reasoning** → The reasoner cannot operate directly on a property graph.
- **Explicit semantics** → OWL restrictions become just properties, losing their logical meaning.
- **SPARQL queries** → You must use the property graph's native query language.

Practical Transformation Example:

The RDF triple:

```
VegetarianPizza rdfs:subClassOf Pizza
Pizza :HAS_TOPPING CheeseTopping
```

becomes a property graph node relationship (e.g., in Neo4j):

```
(vegetarianPizzaNode)-[:SUBCLASS_OF]->(pizzaNode)
(pizzaNode)-[:HAS_TOPPING]->(cheeseToppingNode)
```

But `SUBCLASS_OF` or `HAS_TOPPING` are now just labels -- the property graph does not "know" what subclass or topping means semantically.

8.9.3 Choosing the Right Path for Your EKA Implementation

Requirement	Recommended Graph Technology
Preserve OWL semantics for reasoning	Native RDF Graph (AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune RDF mode, Apache Jena/Fuseki)
Execute complex, relationship-heavy queries (e.g., pathfinding)	Property Graph (Neo4j, TigerGraph, JanusGraph, Amazon Neptune property graph mode)
Integrate with Linked Data and the Semantic Web	Native RDF Graph
High-performance analytics on internal graph data	Property Graph
Full EKA (L2-L3) with triggers(Θ) and actions(Φ)	Native RDF Graph for semantic reasoning; Property Graph as optional performance layer

Important

EKA Recommendation: OWL ontologies can be loaded directly into Native RDF Graph Databases (AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune RDF mode, Apache Jena/Fuseki). These systems preserve full OWL semantics, reasoning capabilities, and IRI-based global identity — making them the correct architectural choice for the EKA K layer.

A Property Graph (Neo4j, TigerGraph, JanusGraph, Amazon Neptune property graph mode) requires a **lossy transformation** that sacrifices OWL semantics and reasoning. It can be used as a downstream, specialized cache or execution layer for specific, performance-intensive queries, but it should not be the primary K layer for a system that depends on OWL semantics.

8.10 EKA Perspective -- RDF as the Missing Bridge

Many organizations already have:

Diagrams

but lack semantics.

Others have:

Meta-models

but lack execution.

Some have:

Ontologies

but still lack operational intelligence.

The missing bridge is often:

Knowledge Graph realization

And RDF is the enabling layer.

RDF transforms ontology from:

Semantic specification

into:

Executable graph knowledge - but **only when deployed in a Native RDF Graph Database** that understands and preserves OWL semantics.

The complete EKA maturity path becomes visible:



This is perhaps the most important message of Chapter 08.

Because at this stage, you are no longer merely learning `Pizza.owl`.

You are learning how to engineer knowledge itself.

Key Insight for EKA Practitioners:

The bridge RDF builds is to **Native RDF Graph Databases**. This is where semantic knowledge becomes truly operational while preserving the meaning encoded in the ontology.

Property graphs like Neo4j are powerful tools for different problems. They are excellent for:

- Fast pathfinding and graph algorithms
- Social network analysis
- Recommendation engines

However, for an EKA system where the goal is **executable semantics** -- where the ontology's constraints, reasoning, and logical rules must be preserved and acted upon -- the correct choice is a Native RDF Graph Database.

8.11 Chapter (08) Summary

In this chapter, we intentionally moved beyond the standard `Pizza.owl` tutorial to examine ontology from a language specification perspective.

You explored:

- Why Protégé is only the editor while RDF is the true ontology artifact
- How RDF represents semantic meaning through triples
- How RDF/XML stores classes, relationships, and semantic logic
- Why RDF is inherently graph-shaped
- How ontology naturally evolves into graph databases
- How **Native RDF Graph Databases** (GraphDB, AllegroGraph, Stardog) preserve and operationalize ontology semantics
- How **Property Graph** (Neo4j) can be used but require a lossy transformation with significant trade-offs
- The critical architectural distinction between Native RDF Graph Databases and Property Graphs
- Why choosing the right graph technology is essential for preserving OWL reasoning in EKA

For the first time in this book, ontology becomes more than modeling.

It becomes:

Portable, Executable, and Operational Knowledge.

8.12 Next Chapter (09) Preview

In the next chapter (09), we continue exploring ontology implementation in Protégé by deepening our understanding of OWL structures and semantic modeling practices. Building upon the RDF foundation introduced here, you will continue progressing toward more advanced ontology engineering capabilities.

8.13 Reference Demo Video

Demo Video Reference on YouTube: Chapter 08 - (<https://youtu.be/qjer-vEKMNg>)

Last updated at 2026-08-29

Appendix A -- Chapter Mapping with Exercises

Volume #	Chapter	Exercise	Tutorial Section	Video	Notes
Volume-1	Chapter 00~03	N/A	N/A	N/A	Foundations & EKA Introduction
Volume-1	Chapter 04	Exercise 1, 2, 3	Section 4.0	04	Creating Named Classes
Volume-1	Chapter 05	Exercise 4	Section 4.1	05	Using a Reasoner
Volume-1	Chapter 06	Exercise 5	Section 4.2	06	Disjoint Classes
Volume-1	Chapter 07	Exercise 6	Section 4.3	07	Create Class Hierarchy
Volume-1	Chapter 08	N/A	N/A	N/A	Conceptual Modeling Foundations
Volume-2	Chapter 09	Exercise 7	Section 4.4	09	Object Properties
Volume-2	Chapter 09	Exercise 8	Section 4.5	09	Creating Property Hierarchies
Volume-2	Chapter 10	Exercise 9	Section 4.6	10	Inverse Properties
Volume-2	Chapter 11	Exercise 10	Section 4.7	11	Property Characteristics
Volume-2	Chapter 12	(Added by Author)	N/A	N/A	Property Characteristics Deep Dive

Volume #	Chapter	Exercise	Tutorial Section	Video	Notes
Volume-2	Chapter 13	Exercise 11, 12	Section 4.9	13	Domain & Range
Volume-3	Chapter 14	Exercise 13	Section 4.10.2	14	Existential Restrictions
Volume-3	Chapter 15	N/A	Section 4.10.2	14	Semantic Description (Theory)
Volume-3	Chapter 16	N/A	Section 4.10.2	14	Description Logic Foundations
Volume-4	Chapter 17	N/A	N/A	N/A	SKDL Introduction
Volume-4	Chapter 18	Exercise 14	Section 4.10.3	15	Stage 1 — Conceptual Modeling
Volume-4	Chapter 19	Exercise 15	Section 4.10.3	16	Stage 2 — Semantic Description
Volume-4	Chapter 20	Exercise 16	Section 4.10.3	17	Stage 3 — Knowledge Reuse (Cloning Pattern)
Volume-4	Chapter 20	Exercise 17	Section 4.10.3	18	Stage 3 — Knowledge Reuse (Practicing Specialization)
Volume-4	Chapter 21	Exercise 18	Section 4.10.3	19	Stage 4 — Semantic Governance (Disjointness)
Volume-4	Chapter 22	Exercise 19	Section 4.10.4	20	Stage 5 — Semantic Validation (Probe Class)
Volume-4	Chapter 23	Exercise 20	Section 4.11	21	Stage 6 — Semantic Reasoning (Primitive Classes)
Volume-4	Chapter 23	Exercise 21	Section 4.11		Stage 6 — Semantic Reasoning (Defined Classes)

Volume #	Chapter	Exercise	Tutorial Section	Video	Notes
Volume-4	Chapter 24	N/A	N/A		Stage 7 — Executable Knowledge (Introduction)
Volume-5	Chapter 25	Exercise 22	Section 4.12		Universal Restrictions
Volume-5		Exercise 23	Section 4.13		Open World Assumption & Closure Axioms
Volume-5		Exercise 24	Section 4.14		Enumerated Classes
Volume-5		Exercise 25	Section 4.15		hasValue Restrictions
Volume-5		Exercise 26	Section 4.15		Cardinality Restrictions
Volume-5		Exercise 27	Section 5.1		Datatype Properties
Volume-5		Exercise 28, 29, 30, 31, 32	Section 5.2		Data Property Values & Classification
Volume-5		Exercise 33	Section 9.1		DL Queries
Volume-5		Exercise 34, 35, 36	Section 10.0		SWRL & SQwRL Rules
Volume-6					Advanced Topics
Volume-7					EKA & AI Integration

Contributors

Name / GitHub Handle	Contribution Type	Issue/Discussion	Chapter(s)	Date	Memo
000wahab000	Pull Request Contribution	PR #57	Various	2026-09-04	Contributed updates and improvements via Pull Request #57.
Jah-yee	Pull Request Contribution	PR #54	Various	2026-09-04	Submitted the repository's first community-contributed pull request, helping enhance project collaboration and open-source engagement.
sarcanon	Question / Issue Report	#39-#50	Various	2026-08-28	Many detail issue reports
LaraAcuna	Question / Content Improvement	#26	13	2026-08-12	Raised a thoughtful question about modeling OR semantics in Domain and Range (Chapter 13). Her question prompted the addition of a complete guide to owl:unionOf ,

Name / GitHub Handle	Contribution Type	Issue/Discussion	Chapter(s)	Date	Memo
					separate properties, and local restriction patterns, including new screenshots, RDF examples, and a comparison table. This contribution significantly improved the practical guidance of Chapter 13. Fixed in v0.44.0.
Michael DeBellis	Refinement Suggestion	#19	15	2026-07-01	Proposed revised definition for VegetarianPizza pattern to align with common real-world interpretation (include cheese/dairy). Pending.
Michael DeBellis	Terminology Clarification	#20	Multiple	2026-07-01	Highlighted need to clearly distinguish "Annotation Properties"

Name / GitHub Handle	Contribution Type	Issue/Discussion	Chapter(s)	Date	Memo
					(reasoner-invisible) from logical property axioms. Pending.
Michael DeBellis	Critical Distinction / Error Report	#21	00, 08	2026-07-01	Identified fundamental conflation between Property Graphs (e.g., Neo4j) and Native RDF Graphs (e.g., AllegroGraph, GraphDB, Stardog). Major architectural correction required. Pending.
Michael DeBellis	Question / Technical Review	#9	00, 06, 14	2026-06-13	Clarified reasoner inference materialization distinction (Protégé UI behavior vs. production triplestores). Fixed in v0.39.0.
Michael DeBellis	Error Report	#10, #11, #12, #13, #14, #15	Various	2026-06-16	Multiple error reports including OWL modeling errors,

Name / GitHub Handle	Contribution Type	Issue/Discussion	Chapter(s)	Date	Memo
					contributor attribution, proofreading, and other corrections. See individual issues for details.
mlungsta89	Bug Report	#3	—	2025-09-15	Reported reasoner getting stuck issue. Closed.
nikokaoya	Feedback	#2	—	2024-06-26	Positive feedback and encouragement. Closed.
Your name here	Error report / Pull request / Question	—	X	YYYY-MM-DD	Brief Memo