

From Semantic Foundations to
Executable Knowledge Architecture (EKA)



Mastering Ontology Engineering

with Protégé and Pizza.owl

VOLUME 4

Chapters 17-24

SKDL — Semantic Knowledge
Development Lifecycle

XIAOQI ZHAO

Mastering Ontology Engineering with Protégé and Pizza.owl

– *From Semantic Foundations to Executable Knowledge Architecture (EKA)*




Table of Content

- [From Ontology Learning to Executable Knowledge Architecture](#)
- [Why the Pizza.owl Tutorial Matters](#)
- [Why This Book?](#)
- [About This Book](#)
- [Why Ontology Matters Today](#)
- [Ontology Within the EKA Framework](#)
- [This Book Is NOT Only About Pizza](#)
- [Who This Book Is For](#)
- [What Makes This Book Different](#)
- [How to Use This Book - The Learning Journey](#)
- [Structural Integrity](#)
- [Acknowledgements](#)
 - [Our Intellectual Heritage & Licensing](#)
 - [Special Thanks](#)
- [Related Resources](#)
- [Final Thoughts](#)

From Ontology Learning to Executable Knowledge Architecture

The Semantic Web has existed for more than two decades, yet for many architects, engineers, analysts, and AI practitioners, ontology engineering still feels abstract, academic, or difficult to approach systematically.

One of the reasons is simple:

Most ontology learning materials focus either too heavily on theory or too narrowly on tool operations.

As a result, many learners can:

- memorize OWL terminology,
- click through Protégé demonstrations,
- or understand isolated semantic concepts,

but still struggle to answer the most important question:

Why does ontology engineering actually matter in modern enterprise and AI systems?

This book was created to answer that question.

Why the Pizza.owl Tutorial Matters

Among all ontology learning materials ever created, Michael DeBellis' Pizza OWL tutorial remains one of the most influential and practical introductions to OWL ontology engineering.

Using a familiar pizza domain, the tutorial teaches learners how to:

- think semantically,
- construct ontology hierarchies,
- model relationships,
- apply OWL logic (rules), and
- and understand reasoning behavior inside Protégé

What makes the Pizza tutorial special is that it gradually introduces semantic complexity in a highly structured manner.

Rather than overwhelming learners immediately with advanced ontology theory, it guides readers step by step through the evolution of a real semantic model.

That educational philosophy strongly influenced the structure of this book.

This eBook serves as the comprehensive companion guide to the [Protégé OWL Pizza Tutorial Hands-on Series](#)

Why This Book?

Most ontology materials focus either on abstract theory or narrow tool operations. This book bridges that gap by connecting semantic engineering to the **Executable Knowledge Architecture (EKA)** framework. It is designed to help you engineer machine-understandable meaning for AI systems, enterprise knowledge platforms, and digital twins.

About This Book

This eBook was written as a companion knowledge guide to the YouTube playlist:

****Protégé OWL Pizza Tutorial Hands-on Series****
(<https://www.youtube.com/playlist?list=PL6DEHvciXKeUx4P32B3hKMK1t6mC8RhsW>)

Since publishing the Protégé OWL Pizza Tutorial series in 2023, I've realized how valuable practical ontology learning remains for architects and AI practitioners. What started as a hands-on walkthrough of Michael DeBellis' tutorial gradually became an important foundation for my EKA (Executable Knowledge Architecture) thinking — connecting ontology, knowledge graphs, and executable intelligence. The playlist may use pizza as the example domain, but its real purpose is helping learners develop semantic thinking and understand how machine-readable knowledge is engineered.

However, the goal is not merely to transcribe the videos.

Instead, this book expands the tutorial into a much deeper professional learning experience by combining:

- ontology engineering theory,
- hands-on Protégé practice,
- semantic modeling methodology,
- enterprise architectural thinking,
- and modern knowledge graph perspectives.

Every chapter corresponds closely to the learning progression of the original video series so readers can learn both visually and conceptually in parallel.

This synchronization is intentional.

Ontology engineering is best learned progressively.

Why Ontology Matters Today

We are entering a new era of intelligent systems.

Modern enterprises increasingly rely on:

- AI systems,
- semantic search,
- enterprise knowledge graphs,
- digital twins,
- intelligent automation,
- contextual reasoning,
- and machine-readable knowledge.

Traditional architectures are built purely on:

- servers and databases,
- documents,
- diagrams
- and APIs

are no longer sufficient for advanced semantic intelligence.

The missing layer here is:

Meaning

Ontology engineering provides that layer.

OWL ontologies allow organizations to formally define:

- concepts,
- relationships,
- constraints,
- semantics,
- and inference logic

in a machine-readable and machine-processable form.


This transforms static information into executable knowledge.

Ontology Within the EKA Framework

This book also introduces ontology engineering through the lens of the EKA (Executable Knowledge Architecture) framework.

EKA represents an architectural approach for transforming enterprise knowledge into executable intelligence systems.

Within EKA, the implementation roadmap is:



Ontology occupies the critical semantic transformation layer.

This is where:

- visual architecture,
- conceptual models,
- and enterprise abstractions

become:

- machine-readable semantics,
- reasoning structures,
- and graph-ready knowledge.

Protégé therefore becomes more than just an ontology editor.

It becomes an engineering environment for executable semantics.

This perspective fundamentally changes how ontology engineering should be understood.

This Book Is NOT Only About Pizza

Although the tutorial domain uses pizza examples, the real purpose of the tutorial is much larger.

The Pizza ontology teaches foundational semantic engineering patterns that scale into real-world domains such as:

- enterprise architecture,
- healthcare,
- finance,
- manufacturing,
- government,
- cybersecurity,
- digital twins,
- and AI knowledge systems.

The same semantics principles used to model:

```
(Pizza)-[:hasTopping]->(CheeseTopping)
```

can later scale into:

```
(Application)-[:supportsCapability]->(businessFunction)
```

or:

```
(Device)-[:connectedTo]->(Sensor)
```

Ontology engineering is domain-independent!

The Pizza example simply provides an approachable learning environment.

Who This Book Is For

This book is intended for:

- Enterprise Architects
- Solution Architects
- Data Architects
- Knowledge Engineers
- AI Engineers

- Semantic Web Learners
- Knowledge Graph Practitioners
- Protégé Beginners
- Meta-Modeling Practitioners
- EKA Learners
- Digital Transformation Professionals

No prior ontology experience is required.

However, reader with backgrounds in:

- UML,
- ER Modeling,
- Enterprise Architecture,
- Databases,
- Graph Technologies,
- or Software Engineering

will often recognize important conceptual parallels throughout the book.

What Makes This Book Different

Most ontology books fall into one of two extremes:

Style	Problem
Highly academic	Difficult for practitioners
Pure tool tutorials	Lack conceptual depth

This book intentionally bridges both worlds.

The objective is to create **A Practical Ontology Engineering Handbook** that remains:

- technically rigorous,
- professionally structured,
- architecturally meaningful,
- and implementation-oriented.

The book therefore combines:

- OWL theory,
- Protégé operations,
- semantic engineering methodology,
- EKA architecture thinking,
- and knowledge graph perspectives

into a single progressive learning journey.

How to Use This Book - The Learning Journey

The recommended learning approach is:

1. **Watch:** The corresponding YouTube video chapter.
2. **Read:** The matching eBook chapter for deeper conceptual context.
3. **Practice:** Manual Protégé exercises and hands-on semantic modeling.
4. **Reflect:** Map these semantic principles to your own enterprise architecture.

Ontology engineering is not mastered through passive reading alone.

It requires semantic thinking practice.

The goal is not merely learning Protégé.

The goal is learning how to engineer machine-understandable meaning.

Structural Integrity

- **Foreword:** By Timothy W. Cook (Founder of SDC).
- **Core Curriculum:** Covers everything from basic class hierarchies and object properties to advanced SHACL, SPARQL, and reasoning behavior.
- **Error Tracker:** An included log to help you troubleshoot common modeling pitfalls.

Acknowledgements

This book has benefited greatly from the generous feedback, technical reviews, and professional discussions contributed by members of the ontology and semantic technologies community.

The full acknowledgements are maintained here: [ACKNOWLEDGEMENTS.md](#)

Our Intellectual Heritage & Licensing

This work is a direct descendant of the **Protégé 4 Tutorial (version 1.3) by Matthew Horridge**. We are deeply indebted to the original visionaries who developed the preceding versions and established the core teaching material: **Holger Knublauch, Alan Rector, Robert Stevens, Chris Wroe, Simon Jupp, Georgina Moulton, Nick Drummond, and Sebastian Brandt**.

In this modern adaptation, we have incorporated **Michael DeBellis's** revisions, which successfully bridged the transition to Protégé 5.5 and integrated advanced practices such as SWRL, SPARQL, and SHACL. Furthermore, we acknowledge the contributions of **Lorenz Buehmann, André Wolski, Dick Ooms, Colin Pilkington, Livia Pinera, Jans Aasman, Yan Xu, and the team at Franz Inc.**, whose work in applying these ontologies to real-world triple stores like AllegroGraph and Gruff has significantly shaped our current perspective on **Executable Knowledge Architecture (EKA)**.

All educational content in this eBook is licensed under **CC BY-SA 4.0**.

For full legal details, please see the [README.md](#) and [LICENSE.md](#) files in the project root.

Special Thanks

We extend our sincere gratitude to **Michael DeBellis** for creating one of the most influential practical OWL learning resources available to the ontology community. (Visit Michael's homepage here - <https://www.michaeldebellis.com/>). We also thank him for his meticulous technical audit for this series; his expert insights -- particularly regarding the crucial engineering distinction between modeling tool and

persistent graph databases -- have been instrumental in ensuring this content meets the standards of modern production-grade environments.

We are also honored to feature a foreword and ongoing review by **Timothy Cook (Founder of SDC)**. His guidance and endorsement have been vital in validating our approach, helping to bridge the gap between traditional semantic modeling and the sophisticated demands of modern enterprise architecture / knowledge architecture.

Finally, we thank **Stanford University and the Protégé community** for their enduring commitment to advancing open ontology engineering tools and Semantic Web technologies.

By preserving this attribution chain, we honor the intellectual legacy of the individuals and organizations that have sustained this field. We stand on the shoulders of these giants, aiming to provide a bridge from traditional semantic modeling to the future of **Executable Knowledge Architecture**.

Related Resources

Official Resources:

- Protégé Official Website: <https://protege.stanford.edu/>
- Protégé Pizza Repository: https://github.com/yasenstar/protege_pizza/tree/main
- EKA Official Website: <https://xiaoqi.com/>

Learning Resources:

- YouTube Channel - Xiaoqi(Yasen) Zhao: <http://www.youtube.com/@yasenzhao>
- Udemy Course - Xiaoqi Zhao: <https://www.udemy.com/user/xiaoqi-zhao>

Final Thoughts

Ontology engineering is NO LONGER a niche academic discipline.

It is rapidly becoming part of the foundational infrastructure for:

- AI systems,
- enterprise knowledge platforms,
- semantic interoperability,
- and executable intelligence architectures.

Understanding ontology today is increasingly similar to understand databases or APIs twenty years ago.

It is becoming a core architectural capability.

The journey begins with a simple Pizza ontology.

But the destination is much larger:

Executable knowledge

Welcome to the World of Ontology Engineering!

"You" are the learner of this book, so while you're reading, I'll say to "you" directly instead of "learner" / "reader" from now on.

Last updated at: 2026-08-08

Acknowledgements

This project would not have reached its current level of maturity without the generosity of numerous members of the ontology, semantic web, enterprise architecture, and knowledge graph communities.

Many people contributed through technical discussions, detailed chapter review, engineering suggestions, issue reports, and encouragement throughout the writing process.

The following individuals deserve special recognition.

Michael DeBellis

Author of the original *Protégé 5 New OWL Pizza Tutorial*.

This book is built upon Michael's excellent tutorial, which has served for many years as one of the most practical introductions to OWL ontology engineering.

Beyond creating the original tutorial, Michael has kindly provided valuable feedback on the direction of this book and has been supportive of extending the tutorial into a broader engineering perspective through the Semantic Knowledge Development Lifecycle (SKDL) and Executable Knowledge Architecture (EKA).

Without his original educational work, this project would not exist.

Timothy W. Cook

Founder & CEO, Axius SDC, Inc.

Tim has generously reviewed multiple chapters of this book during its development, particularly those concerning semantic governance, ontology engineering practice, and enterprise-scale knowledge systems.

His feedback has helped strengthen several important engineering themes throughout the SKDL framework, including:

- treating ontology as a meaning-driven system rather than merely a validation system;
- distinguishing semantic inference from policy enforcement;
- clarifying the role of semantic governance within long-lived knowledge systems;
- improving the engineering narrative surrounding reusable semantic knowledge.

His perspective comes from extensive real-world work in semantic interoperability, healthcare data standards, and enterprise semantic governance.

I am deeply grateful for his thoughtful reviews, constructive criticism, and continued encouragement throughout the writing of this book.

The Ontology Community

Many members of the broader ontology and semantic web community have contributed indirectly through discussions, questions, publications, conference presentations, open-source software, and years of accumulated research.

Protégé, OWL, RDF, Description Logic, and the Semantic Web ecosystem are the result of decades of collaborative work by thousands of researches and practitioners.

This book stands on the collective foundation.

Readers - YOU

Finally, sincere thanks to every reader who as reported errors, suggested improvements, opened GitHub issues, starred the repository, shared the project with others, or simply spent time learning ontology engineering through this book.

Every question, correction, and discussion helps improve future editions.

Knowledge grows through reuse, collaboration, and continuous refinement -- the very principles promoted throughout this book.

Last updated at 2026-08-06

Chapters List -- Volume 4

- [Cover Page](#)
- [Cover Chapter \(README\)](#)
- [ACKNOWLEDGEMENTS](#)
- [Chapter List -- This File](#)
- [Legal & Licensing](#)
- [Foreword from Timothy W. Cook](#)
- [Foreword from Michael DeBellis](#)
- [Preface from the Author](#)

This Volume's Chapters (Volume 4 of 5)

- [Chapter 17 -- Semantic Knowledge Development Lifecycle \(SKDL\)](#)
- [Chapter 18 -- Conceptual Modeling: Stage 1 of the Semantic Knowledge Development Lifecycle](#)
- [Chapter 19 -- Semantic Description: Stage 2 of the Semantic Knowledge Development Lifecycle](#)
- [Chapter 20 -- Knowledge Reuse: Stage 3 of the Semantic Knowledge Development Lifecycle](#)
- [Chapter 21 -- Semantic Governance: Stage 4 of the Semantic Knowledge Development Lifecycle](#)
- [Chapter 22 -- Semantic Validation: Stage 5 of the Semantic Knowledge Development Lifecycle](#)
- [Chapter 23 -- Semantic Reasoning: Stage 6 of the Semantic Knowledge Development Lifecycle](#)
- [Chapter 24 -- Executable Knowledge: Stage 7 of the Semantic Knowledge Development Lifecycle](#)

Full Book Roadmap (for reference -- other volumes sold separately)

- Volume 1 -- Chapters 00-08: Ontology foundations in Protégé
- Volume 2 -- Chapters 09-13: Object properties, characteristics, domain & range
- Volume 3 -- Chapters 14-16: Semantic requirements (restrictions) and EKA governance
- **Volume 4 (this book)** -- Chapters 17-24: Semantic Knowledge Development Lifecycle (SKDL), Stages 1-7 — complete methodology
- Volume 5 -- Chapters 25-XX: Executable Knowledge Architecture (EKA) and AI integration
- [Appendix A - Chapter mapping with Exercises](#)
- [Contributors](#)
- [Error Tracker](#)
- [Error Hunting](#)

Last Updated at 2026-08-15

Legal & Licensing

This eBook is a derivative work of the "[Pizza.owl](#) Tutorial", a foundational resource for the ontology engineering community.

We are committed to transparency and the preservation of the intellectual legacy established by the original authors.

Licensing Terms

This work is governed by a dual-licencing strategy to balance educational openness with software best practices:

- **Documentation & Educational Content:** All textual content, diagrams, and educational explanations within this eBook are licensed under the **Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0)** license.
 - *Requirement:* Any derivative works based on this educational content must maintain the Attribution-ShareAlike terms as required by the original authors.
- **Software & Scripts:** The associated code artifacts, scripts, and automation tools accompanying this eBook are provided under the **GNU General Public License v3.0 (GPL-3.0)**.

Attribution & Intellectual Heritage

This work stands on the shoulders of the community that created the original Protégé 4 Pizza Tutorial. We acknowledge the visionary contributions of **Matthew Horridge, Holger Knublauch, Alan Rector, Robert Stevens, Chris Wroe, Simon Jupp, Georgina Moulton, Nick Drummond, and Sebastian Brandt**.

We also extend our gratitude to **Michael DeBellis** for his invaluable modern revisions, technical audit, and his ongoing dedication to providing practical OWL learning resources for the ontology community.

For further details regarding the full attribution chain, the history of this project, or the specific licensing of individual files, please refer to the [README.md](#) and [LICENSE.md](#) files in the [official GitHub repository](https://github.com/yasenstar/protege_pizza) (https://github.com/yasenstar/protege_pizza).

Foreword from Timothy W. Cook

For more than twenty years, ontology engineering has carried a quiet reputation: powerful in theory, hard to reach in practice. Michael DeBellis' Pizza tutorial did more than almost any other resource to change that, by letting people learn semantics with their hands instead of only their heads. What Xiaoqi Zhao has done here is take that well-loved starting point and carry it all the way to where the work actually has to live, which is inside real systems, under real governance.

This book earns its subtitle. In a single progression it moves from the first class you build in Protégé to the question every enterprise eventually faces: how machine-readable meaning becomes something a system can act on safely. Along the way it never loses the teacher's discipline of one idea at a time.

I want to point readers in particular at Chapters 12 and 13, because their titles tell you something. They do not *only* say "object property characteristics" and "domain and range." They say *governing* semantic relationship behavior and *governing* semantic boundaries. That word is the whole game. A functional or transitive property is not a checkbox in a tool. It is a rule about how a relationship is allowed to behave. A domain and a range are not decoration. They are a statement about what may legitimately connect to what. Xiaoqi teaches these as governance, and that is exactly the right approach.

Chapter 13 also does something I rarely see in an introductory text: it tells the truth about a sharp edge. OWL domain and range do not reject a bad assertion the way a database constraint would. The reasoner infers, under an open-world assumption, rather than refuses, and as the chapter notes, this surprises newcomers who expected enforcement. That surprise is where academic ontology meets the enterprise information system. The moment data leaves the editor and travels between systems, a boundary that only infers is not yet a boundary you can trust. It has to be able to refuse what does not belong, deterministically, with its rules traveling alongside the data rather than living in a separate document or a separate dashboard.

That is the work my colleagues and I have spent years on in the Semantic Data Charter. SDC treats meaning and the constraint definitions as shareable, bound assets. They are one inseparable data payload, so admissibility can be checked at the point of use, in the open, against published specifications. Readers will recognize this as related to the layer Xiaoqi names Γ in his EKA tuple: the governance layer that turns a clever model into a trustworthy one. EKA is the architecture that orchestrates the pieces; a substrate like SDC is one open, standards-based way to make that governance layer real in production. The two meet exactly where this book spends its most careful chapters.

There is a larger reason to welcome a book like this now. Ontology engineering is becoming a core architectural capability, the way databases and APIs did before it, and the field is poorly served when it stays split between texts that are rigorous but unreachable and tutorials that are reachable but shallow. Xiaoqi has taken the harder middle path: technically honest, professionally structured, and grounded in the standards that real organizations have to answer to. That is the bridge between theory and practice, and we need more people building it.

Read this book with Protégé open. Do the exercises. And when you reach those governing chapters, remember that the discipline you are learning on a pizza is the same discipline that will one day decide whether a clinical, financial, or regulatory system is allowed to act. As Xiaoqi writes, meaningful intelligence depends on meaningful boundaries. This book teaches you how to build them.

Timothy W. Cook
Founder, Axius SDC, Inc.
Creator, the Semantic Data Charter (SDC)
June 2026

Preface from Michael DeBellis

I sometimes find it easy to be pessimistic about the human race. The world gives us plenty of reasons for discouragement. But one thing that illustrates some of the best in people is the open source movement.

It is remarkable how much of the software infrastructure that powers the modern internet is open source. The Linux operating system, Apache web servers, Kubernetes, Kafka, Docker containers, and countless programming languages, libraries, databases, and development tools were created as open-source projects or converted from internal projects to software freely available to all developers. Organizations and individuals understand that software becomes more valuable when it is shared openly so that the developer community can study it, improve it, and build on it. It's a very positive thing that corporations such as Google and LinkedIn are motivated to contribute software in which they've invested heavily into the open source infrastructure. Even more so that so many individuals will contribute their time to maintain and expand open source tools for the simple joy of seeing things work and providing tools that colleagues will find useful.

That generosity is not limited to code. Documentation, examples, tutorials, walkthroughs, bug reports, and teaching materials are also essential parts of the open-source ecosystem. They are just as important, because good software is wasted if no one can understand or learn to use it.

Tutorials are often underappreciated. As I said when I first created my revised version of the Protégé Pizza tutorial, no one gets a PhD for writing a tutorial. Tutorials do not usually receive the same professional recognition as research papers, software frameworks, or commercial products. Yet, for many learners, a good tutorial is the difference between a technology remaining abstract vs. becoming usable.

That is why I am pleased to see Xiaoqi Zhao's *Mastering Ontology Engineering with Protégé and Pizza.owl: From Semantic Foundations to Executable Knowledge Architecture (EKA)*. It continues a tradition that has always been central to the Semantic Web community: building on prior work, acknowledging intellectual lineage, and helping new learners move from abstract ideas to practical results.

The Pizza tutorial itself has a long history. The original tutorial was developed by members of the Protégé and ontology engineering community. It became one of the best-known introductions to OWL ontology development. Its genius was the choice of a domain that was familiar enough to get out of the learner's way. Almost everyone understands pizzas, toppings, bases, vegetarian choices, and spicy ingredients. Because the domain is simple, learners can focus on the real subject: classes, properties, axioms, inference, queries, and rules.

My own revision of the Pizza tutorial was motivated by a practical teaching need. I was preparing a hands-on workshop and realized that the older tutorial no longer matched the current Protégé user interface closely enough for beginners. Especially since the majority of the students in that workshop were Library Science experts, not developers. Experienced OOP developers can translate instruction such as "add superclass" into revised user interfaces such as "add subclassOf". New learners, especially those coming from fields such as library science and information architecture, should not have to do that translation while also trying to understand the concepts of OWL and other Semantic Web languages. The tutorial needed to match the actual tool.

What began as a small update became a much larger project. I updated the examples for Protégé 5, revised screenshots and instructions, and added material on topics such as SPARQL, SHACL, SWRL, IRIs, namespaces,

and WebProtégé. My goal was not just to teach people where to click, but to help them understand why the clicks mattered.

Xiaoqi's eBook takes that work in a new direction. It uses the Pizza tutorial as a foundation, but it is not simply a transcription or repetition of the older material. It connects OWL ontology engineering to a broader architectural vision that he calls Executable Knowledge Architecture (EKA). The eBook encourages learners to see ontology development not as an isolated academic exercise, but as part of a larger architecture for formal knowledge, knowledge graphs, reasoning, validation, and intelligent systems.

I've spent half of my career as a consultant and the other half doing research. One of the most important lessons I learned as a consultant is that the hard part isn't writing software, the hard part is integration. That's true whether we're integrating rules with COBOL programs, LLMs with enterprise data, or ontologies and knowledge graphs into distributed computing environments.

That broader perspective is timely.

For a while, it seemed that the Semantic Web might remain a powerful idea that never achieved its hoped-for impact. The original vision was ambitious: not just a web of documents, but a web of data and meaning, where machines could understand and integrate information across sources. The challenge was that adding semantics requires work. It requires modeling. It requires shared identifiers. It requires explicit commitments about meaning. It requires different groups in large organizations to talk to each other and define shared concepts and domain boundaries.

Many organizations decided that this was too much effort. It was easier to publish documents, tables, and data dumps and let humans, search engines, and especially machine learning systems make sense of them.

Ironically, the rise of Large Language Models (LLMs) has created renewed interest in the very technologies that some people thought machine learning would make unnecessary. LLMs are extraordinary tools. Used well, they can provide enormous productivity benefits. But they also have well-known limitations. Their reasoning is opaque. They can generate fluent but incorrect answers. They may struggle with provenance, controlled terminology, and domain-specific meaning.

These problems are not merely accidental bugs. They follow naturally from the fact that LLMs are statistical models rather than explicit knowledge models. Improvements in scale, training, and architecture will continue to make them more capable, but many practical applications still need complementary structures: curated data, explicit semantics, provenance, and constraints.

That is why knowledge graphs and ontologies are becoming more important. Retrieval Augmented Generation (RAG) was an early response to the problem of grounding LLMs in curated information. But as applications become more demanding, simple document retrieval is often not enough. Many domains require richly interconnected knowledge and automated reasoning. RDF, OWL, SPARQL, SHACL, and reusable ontologies such as Dublin Core, PROV-O, and SKOS provide mature standards for representing and working with exactly that kind of explicit knowledge representation. Explicit knowledge representation based on logic and set theory is the perfect complement to the powerful inductive statistical models that power LLMs.

This renewed interest also means that the word "ontology" is now used in many different ways. Some commercial systems use the term for models that are closer to enterprise object models than to ontologies in the Semantic Web sense. I understand why some researchers and ontology engineers find that frustrating. I sometimes do too. But I also see it as a sign of progress. I have seen this pattern since the early days of expert system adoption and later with the adoption of Object-Oriented Programming and Agile methods: when ideas

move from research to industry, terms such as rule-based, knowledge base, *Rapid Application Development (RAD)*, object model, and ontology are stretched, simplified, and repurposed. That can create confusion, but it also shows that industry has recognized that the underlying ideas matter. Eventually as practitioners understand the core research ideas, they make educated decisions and can separate real technology from marketing hype. Tutorials such as this eBook are an essential part of that process. Learners need to understand not only that ontologies are useful, but what ontologies and other semantic technology provide that is new and why these new capabilities matter.

One of the things I appreciate about this eBook is that it tries to connect beginner-level Protégé modeling with larger architectural questions. A learner begins with classes such as *Pizza*, *PizzaTopping*, *CheeseTopping*, and *VegetableTopping*. But those beginner examples lead naturally to deeper questions. What is the difference between a class and an individual? When should a class be primitive, and when should it be defined? What does a reasoner actually infer? Why does OWL's open-world assumption behave differently from a database constraint? When should we use OWL reasoning, and when should we use SHACL validation? How do ontologies relate to RDF graphs, triple stores, SPARQL queries, and larger distributed system architectures?

Those are not minor details. They are the conceptual foundation of ontology engineering.

My strongest advice to readers is to keep Protégé open while reading. Do the exercises. Run the reasoner frequently. That last part is really critical. For training ontologies, the reasoner executes immediately. Running it after every change means you'll see an error as soon as it is created. When that happens, the reasoner can almost always focus you on the specific problem. If you wait until there are several errors the reasoner feedback is usually much less helpful and finding errors can take hours instead of minutes.

When something surprising happens, slow down and understand why. In OWL, surprises are often the most valuable teaching moments. A reasoner that classifies something unexpectedly, fails to classify something you expected, or reports an inconsistency is showing you the difference between what you intended to model and what you actually modeled.

That distinction is at the heart of ontology engineering.

A good ontology is not just a diagram. It is not just a taxonomy. It is not just a data model. It is an explicit, formal model of meaning in a domain. Building one requires careful thought, iteration, testing, and patience. The Pizza tutorial remains useful because it gives learners a simple domain in which to practice those habits before applying them to medicine, finance, manufacturing, enterprise architecture, LLM integration and other business, engineering, and scientific domains.

I am glad to see this new contribution to the Pizza tutorial lineage. It reflects the best spirit of open educational work.

The journey begins with pizza. But the real subject is meaning.

Michael DeBellis

2 July 2026

San Francisco, CA

<https://www.michaeldebellis.com/blog>

Preface from Author

From Static Diagrams to Executable Knowledge: A Personal Journey

For nearly thirty years, I have lived and breathed enterprise IT. My career began long before “knowledge graph” was a common term, working across telecom, automotive, finance, and IT outsourcing consulting. For most of those years, my primary tool for making sense of complexity was static diagramming, mostly with Visio, draw.io, or even office suite like Excel and Powerpoint. It was powerful, but it was also... silent. The diagrams described architecture, but they could not execute it. They could not reason, validate, or adapt.

That was the before.

The shift (2021 – 2022) began when I moved from old toolbox to Archi, the open-source ArchiMate modeling tool. Suddenly, I was not just drawing; I was building a **repository**. I discovered the power of a shared, collaborative model – a single source of truth about an enterprise. But even Archi, as good as it is, had limits. The model was still primarily for human consumption. It was a rich, structured database, but it was not a reasoning engine.

It was around this same time that I discovered the **Essential EAS Open Source (Ontology)** – a glimpse of what happens when an EA repository is built on formal semantics. That was the first crack in the old paradigm, learnt and practiced with its modeling tool - Protégé - the first time I met this word.

The catalyst (2023) was the [Pizza.owl](#) tutorial. Through creating my first own video series on the topic, I got in touch with Michael DeBellis. More importantly, I got my hands on Protégé. For the first time, I was not just modeling structure; I was formalizing **meaning**. I was defining classes, properties, and rules that a machine could use to infer new knowledge, check for contradictions, and navigate relationships intelligently.

The pizza example was simple, but the potential was explosive. I realized that the future of enterprise architecture was not in prettier diagrams, but in **executable knowledge**. With that learning, back to the Essential EAS tool, I realized how powerful of treating enterprise architecture from semantic perspective, and how flexible the tool provided, e.g. if you need one meaning (triple) that EAS built-in meta-model doesn't exist, just create one new from it's slots and that's now fitting the knowledge of your enterprise reality, more and more.

The synthesis (2023 – 2026) became a period of intense exploration. My YouTube channel (and other video sharing platform, like BiliBili, DouYin, etc..) grew to host not just the Pizza.owl series, but courses on ArchiMate, meta-modeling, Python / C / C++ / Java, PlantUML (diagramming as code), Neo4j (graph databases), and even 3D modeling with OpenSCAD. Each course was a piece of a larger puzzle. programming taught me logic, PlantUML taught me that diagrams could be generated from text (code), Neo4j showed me how to traverse relationships at scale, and OpenSCAD reinforced the power of parametric, declarative design.

All these threads eventually wove themselves into a single, coherent idea: **EKA – Executable Knowledge Architecture** (<https://xiaoqi.com>). EKA is the formalization of everything I have learned. It is the recognition that an ontology is not a document, but an executable specification. It is the understanding that $EKA = (K, R, \Theta, \Phi, \Gamma)$ – a framework where Knowledge, Reasoning, Triggers, Actions, and Governance become a single, integrated system.

Why this book? This book is not a transcript of my videos. It is the theory behind the practice. It is the result of years of hands-on work, distilled into a structured, pedagogical journey. I wrote it because I saw a gap:

countless tutorials show you how to click buttons in Protégé, but almost none explain the architectural thinking behind the clicks.

This book aims to bridge that gap, using the familiar [Pizza.owl](#) as a safe and repeatable domain to teach the foundational principles of EKA.

You will start with pizza. But if you follow the journey, you will be able to model anything: an enterprise architecture, a digital twin, a regulatory framework, or an AI knowledge system.

The goal is not just to teach you Protégé. The goal is to give you a new way to think about knowledge: as something that can be **formalized, reasoned over, governed, and eventually executed**.

Welcome to the journey.

Xiaoqi Zhao (Yasen)

Global Enterprise Architect, Founder of the EKA Framework

June, 2026

Chapter 17 -- Semantic Knowledge Development Lifecycle (SKDL)

From Semantic Modeling to Executable Knowledge

- [17.1 Congratulations -- You Have learned the Language of OWL](#)
- [17.2 From Learning OWL to Practicing Ontology Engineering](#)
- [17.3 What Is the Semantic Knowledge Development Lifecycle \(SKDL\)?](#)
- [17.4 The Semantic Knowledge Development Lifecycle \(SKDL\)](#)
- [17.5 Understanding Each Lifecycle Stage](#)
- [17.6 Relationship Between SKDL and EKA Framework](#)
- [17.7 Mapping Michael DeBellis' Tutorial to the Lifecycle -- Preview of the Remaining Chapters](#)
- [17.8 Interesting Reading -- The Lifecycle of Knowledge Development](#)
- [17.9 Preview of the Next Chapters](#)
- [17.10 Chapter Summary](#)

17.1 Congratulations -- You Have learned the Language of OWL

By completing previous Chapters (14-16) - I know it's a long chapter - you have reached an important milestone in your ontology engineering journey.

So far, this book has introduced the fundamental building blocks of **OWL**, including:

- ontology classes
- object properties
- inverse properties
- property characteristics
- domain and range
- existential restrictions (**some**)
- universal restrictions (**only**)
- logical reasoning
- semantic governance
- the Open World Assumption

These concepts form the logical language used to describe knowledge in an ontology.

At this point, you already understand **how OWL expresses semantic knowledge**.

The focus now changes.

Instead of learning individual OWL constructs one by one, the remaining chapters will show how these constructs work together to build complete semantic systems.

In other words, you are moving from

learning a language

toward:

learning an engineering discipline.

Congratulations on reaching this milestone!

Before continuing, consider taking a short break, grab a cup of coffee or tea.

Reflect on how far you have progressed -- from creating your first ontology in Protégé to understanding semantic restrictions and automated reasoning.

The chapters that follow will build upon everything you have learned so far.

17.2 From Learning OWL to Practicing Ontology Engineering

Learning OWL syntax is similar to learning the vocabulary and grammar of a spoken language.

Knowing individual words does not automatically enable someone to write a novel.

Likewise, understanding individual OWL constructs DOES NOT immediately translate into designing scalable enterprise ontologies.

Ontology engineering is a systematic process.

Each modeling decision influences future **reasoning**, **validation**, **governance**, and **knowledge reuse**.

Professional ontology engineers therefore think beyond individual classes or properties.

They consider questions such as:

- *How should domain concepts be organized?*
- *Which concepts should become reusable semantic patterns?*
- *How can semantic consistency be maintained as the ontology grows?*
- *Which axioms enable automatic reasoning?*
- *How can ontology models evolve without becoming difficult to maintain?*

These questions shift ontology engineering from software operation to architectural design.

The remaining chapters of this book introduce a structured engineering workflow that answers these questions.

17.3 What Is the Semantic Knowledge Development Lifecycle (SKDL)?


Throughout the remainder of this book, we introduce the **Semantic Knowledge Development Lifecycle (SKDL)**.

The SKDL describes the progressive stages through which semantic knowledge evolves -- from an **initial conceptual model** into **executable knowledge** capable of supporting intelligent systems.

Unlike a software development lifecycle (SDLC), which primarily transforms requirements into executable code, the SKDL transforms domain knowledge into executable semantics.

This lifecycle provides a practical roadmap for ontology engineers, enterprise architects, and knowledge graph practitioners.

17.4 The Semantic Knowledge Development Lifecycle (SKDL)



Each stage builds upon the previous one.

1. Conceptual models provide the structure.
2. Semantic descriptions enrich those structures with meaning.
3. Reusable patterns accelerate ontology development.
4. Governance ensures consistency.
5. Validation detects logical problems.
6. Reasoning derives new knowledge automatically.

Together, these stages transform static data models into intelligent semantic systems.

17.5 Understanding Each Lifecycle Stage

The **Semantic Knowledge Development Lifecycle (SKDL)** divides ontology engineering into seven (7) progressive stages. Each stage addresses a distinct engineering objective while building upon the semantic artifacts produced by the previous stage.

Rather than viewing ontology development as a collection of isolated modeling techniques, SKDL organizes the process into a structured lifecycle in which semantic knowledge gradually evolves -- from conceptual abstraction to executable intelligence.

Each stage contributes a specific capability that ultimately supports the vision of **Executable Knowledge Architecture (EKA)** introduced in Chapter (00).

The seven stages are summarized below:

Stage	Purpose	Typical Activities
1. Conceptual Modeling	Identify and organize domain concepts into coherent semantic taxonomies.	Create classes, subclasses, abstraction hierarchies, and taxonomy structures.
2. Semantic Description	Enrich concepts with formal semantic definitions, properties, and logical restrictions.	Define object properties, existential and universal restrictions, equivalent classes, and semantic axioms.
3. Knowledge Reuse	Extend existing ontology components through semantic reuse and incremental refinement.	Duplicate, specialize, refactor, and extend existing ontology structures while minimizing redundancy.
4. Semantic Governance	Maintain semantic integrity through modeling standards, constraints, and governance policies.	Apply ontology design principles, naming conventions, modularization, documentation, and semantic governance practices.

Stage	Purpose	Typical Activities
5. Semantic Validation	Verify ontology correctness before deployment by ensuring logical consistency and semantic quality.	Execute reasoners, detect inconsistencies, validate constraints, review inferred hierarchies, and perform ontology quality assurance.
6. Semantic Reasoning	Discover implicit knowledge through logical inference and automated semantic classification.	Execute OWL reasoning, infer new knowledge, classify concepts, and enrich the knowledge graph automatically.
7. Executable Knowledge	Connect semantic knowledge with enterprise systems, automation, AI and intelligent decision-making.	Integrate ontologies with EKA workflows, knowledge graphs, APIs, enterprise applications, intelligent agents, and semantic automation.

Although these stages are presented sequentially, ontology engineering is rarely a strictly linear process. As new requirements emerge or domain knowledge evolves, ontology engineers frequently revisit earlier stages to refine conceptual models, improve semantic definitions, strengthen governance, or validate newly introduced knowledge.

For this reason, SKDL should be viewed as an iterative engineering lifecycle rather than a one-time development methodology. Each iteration progressively improves the semantic model while preserving consistency with the knowledge already established.

Most importantly, each stage contributes directly to the realization of the EKA framework. As the lifecycle progresses, the ontology evolves from a static conceptual model into an executable semantic asset capable of supporting reasoning, governance, automation, and intelligent decision-making.

17.6 Relationship Between SKDL and EKA Framework

One of the distinguishing characteristics of the Semantic Knowledge Development Lifecycle is that it does not exist independently of EKA. Instead, SKDL provides the practical engineering methodology for progressively constructing the five components of the **Executable Knowledge Architecture (EKA)**.

$$EKA = (K, R, \Theta, \Phi, \Gamma)$$

- K : Knowledge Graph Layer
- R : Reasoning & Rules Layer
- Θ : Trigger Layer
- Φ : Execution Layer
- Γ : Governance Layer

Rather than attempting to build all EKA components simultaneously, ontology engineers develop them incrementally as the semantic model matures.

The relationship between SKDL and EKA is illustrated below:

SKDL Stage	Primary EKA Components	Contribution
1. Conceptual Modeling	K -- Knowledge Graph	Establishes the conceptual vocabulary and semantic taxonomy of the domain.
2. Semantic Description	K -- Knowledge Graph	Enriches concepts with formal semantic definitions and logical meaning.
3. Knowledge Reuse	K -- Knowledge Graph	Expands and refines the knowledge graph through reusable semantic patterns.
4. Semantic Governance	Γ -- Governance	Introduces semantic quality controls, modeling standards, and governance policies.
5. Semantic Validation	$\Gamma + K$	Verifies semantic correctness, logical consistency, and ontology quality before deployment.
6. Semantic Reasoning	R -- Reasoning & Rules	Enables automated inference and semantic classification using formal logic.
7. Executable Knowledge	$\Theta + \Phi$	Connects semantic knowledge to triggers, enterprise workflows, APIs, AI agents, and executable actions.

This mapping also illustrates an important engineering observation.

The K - **Knowledge Graph** is not constructed in a single activity. Instead, it evolves progressively across the first three stages SKDL. Conceptual Modeling establishes the structural foundation, Semantic Description enriches that structure with formal meaning, and Knowledge Reuse expands the semantic model through incremental refinement.

Only after a sufficiently rich and stable knowledge graph has been established do governance, validation, and automated reasoning become truly effective.

Finally, the semantic knowledge produced by these stages can be operationalized through triggers and executable actions, completing the vision of **Executable Knowledge Architecture**.

17.7 Mapping Michael DeBellis' Tutorial to the Lifecycle -- Preview of the Remaining Chapters

The remainder of this book follows the Semantic Knowledge Development Lifecycle introduced in this chapter. Rather than treating the **Pizza** ontology as a sequence of isolated Protégé exercises, each chapter explores one stage of semantic knowledge engineering using Michael DeBellis' **Pizza** tutorial as a continuous case study.

Chapter	Tutorial Exercise	SKDL Stage	Engineering Focus
16	Exercise 14	Conceptual Modeling	Build semantic taxonomies through classes and subclass hierarchies.

Chapter	Tutorial Exercise	SKDL Stage	Engineering Focus
17	Exercise 15	Semantic Description	Enrich concepts with formal semantic definitions and logical restrictions.
18	Exercise 16	Knowledge Reuse	Reuse and refine semantic patterns through incremental ontology development.
19	Exercise 17-18	Semantic Governance	Improve semantic quality through disciplined ontology engineering practices.
20	Exercise 19	Semantic Validation	Verify logical consistency and semantic correctness using OWL reasoning.
21	Exercise 20	Semantic Reasoning	Discover implicit knowledge through automated inference and semantic classification.
	Later Chapters	Executable Knowledge	Extend semantic models into knowledge graphs, EKA, automation, and AI-driven applications.

- Viewed individually, these exercises appear to be simple Protégé operations.
- Viewed collectively, they demonstrate the lifecycle followed by professional ontology engineers.

This progression reflects a central philosophy of ontology engineering:

Semantic knowledge should evolve systematically through disciplined engineering practices rather than ad hoc modeling.

By following the Semantic Knowledge Development Lifecycle, readers will gradually progress from constructing conceptual models to developing intelligent semantic systems capable of supporting reasoning, governance, and enterprise execution.

17.8 Interesting Reading -- The Lifecycle of Knowledge Development

The idea that knowledge develops through progressive stages is not new. Throughout history, philosophers, scientists, and management scholars have proposed different ways of understanding how knowledge is created, refined, and applied.

One of the earliest discussions can be found in **Plato's *Theaetetus***, where Socrates explores the question:

What is knowledge?

Among the ideas discussed is the famous concept that knowledge can be understood as **justified true belief** -- a belief that is true and supported by sufficient justification. Although modern epistemology has refined and challenged this definition, it established an enduring principle:

Knowledge is more than information; it requires structure, meaning, and justification.

More than two thousand years later, knowledge management researchers continued exploring how knowledge evolves within organization. A particularly influential example is the **SECI Model**, proposed by Ikujiro Nonaka and Hirotaka Takeuchi. The model describes organizational knowledge creation as a continuous cycle of:

- **Socialization** -- sharing tacit knowledge through experience.
- **Externalization** -- expressing tacit knowledge in explicit forms.
- **Combination** -- integrating explicit knowledge into more complex structures.
- **Internalization** -- absorbing explicit knowledge back into individual understanding.


The Semantic Knowledge Development Lifecycle (SKDL) introduced in this book is inspired by the same fundamental observation:

Knowledge is not created in a single step -- it evolves through successive stages.

However, SKDL addresses a different problem.

Rather than describing how people create organizational knowledge, SKDL describes how **semantic knowledge** is engineered into machine-interpretable models capable of supporting automated reasoning and intelligent systems.

The progression can be summarized as:



Each stage enriches the semantic model with additional structure, precision, and computational capability.

From a mathematical perspective, this evolution can be viewed as the gradual enrichment of a formal knowledge system.

Initially, we define a set of concepts:

$$C = c_1, c_2, \dots, c_n$$

Semantic relationships organize these concepts into a graph:

$$G = (V, E)$$

Logical axioms further constrain and enrich the graph:

$$A = a_1, a_2, \dots, a_m$$

Finally, reasoning derives new knowledge:

$$K_{\text{new}} = K \cup \text{Infra}(A)$$

Each stage increases the expressive power of the knowledge base without changing its underlying domain.

Seen from this perspective, ontology engineering is not merely about creating classes and properties. It is a disciplines process of progressively transforming domain knowledge into **executable semantic intelligence**.

The Semantic Knowledge Development Lifecycle provides a practical engineering roadmap for managing that transformation, bridging classical theories of knowledge with modern ontology engineering and the vision of **Executable Knowledge Architecture (EKA)**

17.9 Preview of the Next Chapters

The next several chapters correspond directly to the stages of the Semantic Knowledge Development Lifecycle.

Chapter	Lifecycle Stage	Focus
Chapter 16	Conceptual Modeling	Building the first semantic taxonomy
Chapter 17	Semantic Description	Adding meaning through logical characteristics
Chapter 18	Knowledge Reuse	Reusing and refining ontology design patterns
Chapter 19	Semantic Governance	Protecting semantic consistency
Chapter 20	Validation	Detecting logical inconsistencies
Chapter 21	Reasoning	Enabling automatic semantic inference

Rather than viewing these chapters as isolated tutorials, think of them as successive stages in a single engineering process.

Each chapter contributes another capability required to build robust, scalable, and executable semantic knowledge systems.

17.10 Chapter Summary

This chapter introduced the **Semantic Knowledge Development Lifecycle (SKDL)**, a practical engineering roadmap that will guide the remainder of this book.

Rather than treating ontology development as a collection of unrelated modeling techniques, the lifecycle presents ontology engineering as a progressive process in which semantic knowledge evolves from conceptual structures into executable intelligence.

As you continue through the following chapters, keep this lifecycle in mind.

Every new modeling techniques you lean represents another step toward building semantic systems that are not only logically consistent, but also capable of supporting automated reasoning, enterprise governance, and ultimately the vision of **Executable Knowledge Architecture (EKA)**

Last updated at 2026-08-08, renamed from original Chapter 15

Chapter 18 -- Conceptual Modeling: Stage 1 of the Semantic Knowledge Development Lifecycle

Building Semantic Taxonomies Through Subclass Hierarchies

- 18.1 Introduction -- Conceptual Modeling: The Foundation of Semantic Knowledge
- 18.2 Learning Objectives
- 18.3 Position Within the Semantic Knowledge Development Lifecycle
- 18.4 Exercise 14 -- Creating the First Conceptual Hierarchy
 - 18.4.1 **Pizza** -- The General Concept
 - 18.4.2 **NamedPizza** -- Introducing an Intermediate Abstraction
 - 18.4.3 **MargheritaPizza** -- The First Concrete Domain Concept
 - 18.4.4 An Engineering Perspective
- 18.5 Why Conceptual Modeling Comes Before Semantic Modeling
 - 18.5.1 Concept Identification Before Semantic Definition
 - 18.5.2 The Engineering Pattern Across Disciplines
 - 18.5.3 Why Modeling Semantics Too Early Creates Problems
 - 18.5.4 Conceptual Modeling as the Foundation of EKA Knowledge
 - 18.5.5 From Structure to Meaning
- 18.6 Building a Semantic Taxonomy
 - 18.6.1 Generalization and Specialization
 - 18.6.2 Designing for Semantic Growth
 - 18.6.3 More Than Classification
 - 18.6.4 The Conceptual Backbone of Semantic Knowledge
- 18.7 From Taxonomy to Semantic Inheritance
 - 18.7.1 From Organization to Meaning
 - 18.7.2 The Meaning of **SubClassOf**
 - 18.7.3 Semantic Inheritance in Practice
 - 18.7.4 Semantic Inheritance Is Not Software Inheritance
 - 18.7.5 Why Semantic Inheritance Matters
 - 18.7.6 From Conceptual Hierarchies to Logical Reasoning
- 18.8 Interesting Reading -- Mathematical Foundations of Semantic Taxonomy
 - 18.8.1 Classes as Sets
 - 18.8.2 Why Inheritance Happens Automatically
 - 18.8.3 A Taxonomy as a *Partially Ordered Set*
 - 18.8.4 Why Ontology Engineers Should Care
 - 18.8.5 Transition
- 18.9 Engineering Perspective -- Modeling Before Meaning
 - 18.9.1 Separating Conceptual Structure from Semantic Meaning
 - 18.9.2 Exercise 14 as an Engineering Case Study
 - 18.9.3 Engineering Before Implementation
 - 18.9.4 The Role of the Semantic Knowledge Development Lifecycle
 - 18.9.5 Why Incremental Modeling Matters
 - 18.9.6 Engineering Perspective

- 18.9.7 Transition to the Next Perspective
- 18.10 Interesting Reading -- Two Great Classification Traditions
 - 18.10.1 Biological Taxonomy: Nature's Original Ontology
 - 18.10.1.1 Linnaean Classification System
 - 18.10.1.2 Biological Taxonomy and the **Pizza** Ontology
 - 18.10.1.3 From Classification to Ontology Engineering
 - 18.10.2 Library Classification: Organizing Human Knowledge
 - 18.10.2.1 Dewey Decimal Classification
 - 18.10.2.2 Library of Congress Classification
 - 18.10.2.3 Library Classification and Ontology Engineering
 - 18.10.3 What Both Traditions Teach Ontology Engineers
- 18.11 EKA Perspective -- Conceptual Modeling Establishes the Knowledge Layer (*K*)
 - 18.11.1 Conceptual Modeling and the EKA Tuple
 - 18.11.2 Mapping Chapter 16 to the EKA Framework
 - 18.11.3 Why the Knowledge Layer Comes First
 - 18.11.4 From Conceptual Organization to Executable Knowledge
 - 18.11.5 An Architectural Perspective
 - 18.11.6 Engineering Takeaway
- 18.12 Engineering Guidelines for Conceptual Modeling
 - 18.12.1 Think Semantically Rather Than Visually
 - 18.12.2 Model Concepts Before Properties
 - 18.12.3 Prefer Generalization Before Duplication
 - 18.12.4 Design for Future Growth
 - 18.12.5 Maintain Consistent Levels of Abstraction
 - 18.12.6 Distinguish Classes from Individuals
 - 18.12.7 Building Stable Taxonomies Before Introducing Semantics
 - 18.12.8 Engineering Takeaway
- 18.13 Key Concepts
- 18.14 Chapter Summary
- 18.15 Looking Ahead -- From Concepts to Meaning

18.1 Introduction -- Conceptual Modeling: The Foundation of Semantic Knowledge

Every engineering and architectural discipline begins with **abstraction**.

- Before software engineers design classes and implement algorithms, they first identify the fundamental concepts that define the problem domain.
- Before database designers normalize tables and establish integrity constraints, they first determine the entities, attributes, and relationships that represent the information being managed.
- Before enterprise architects analyze business capabilities, information flows, and technology landscapes, they first construct conceptual models that describe how an enterprise is organized and how its components interact.

Although these disciplines employ different methodologies, modeling languages, and engineering practices, they all share a common principle:

A sound conceptual model must precede implementation!

Ontology engineering follows exactly the same principle.

Before semantic relationships can be established, before logical restrictions can be defined, before automated reasoning can derive new knowledge, and before intelligent systems can interpret semantic information, an ontology must first establish a coherent conceptual representation of the domain it intends to describe.

As introduced in Chapter (15), the **Semantic Knowledge Development Lifecycle (SKDL)** provides a structured engineering methodology for developing semantic knowledge systems. Rather than viewing ontology development as a collection of isolated Protégé operations, SKDL organizes the entire engineering process into seven progressive stages that gradually transform conceptual knowledge into executable intelligence.

This chapter (16) begins **Stage 1 -- Conceptual Modeling**, the foundation upon which every subsequent stage of the lifecycle depends.

Conceptual Modeling addresses the first and arguably the most fundamental question in ontology engineering:

What concepts exist within the domain?

Although deceptively simple, this question determines the quality of every subsequent stage of semantic knowledge development.

A poorly designed conceptual model inevitably leads to ambiguous semantics, inconsistent reasoning, difficult maintenance, and limited opportunities for semantic reuse. Conversely, a well-designed conceptual hierarchy provides a stable semantic foundation upon which semantic definitions, reusable knowledge patterns, governance policies, validation mechanisms, reasoning capabilities, and ultimately executable knowledge can be progressively constructed.

This engineering philosophy is clearly reflected in the structure of **Michael DeBellis' Pizza tutorial**.

Throughout the previous chapters, the tutorial introduced the fundamental building block of OWL, including classes, object properties, property characteristics, inverse properties, domains, ranges, and semantic restrictions. Having established this semantic vocabulary, the tutorial now shifts its focus from learning individual OWL constructs to a coherent semantic model.

Rather than immediately introducing increasingly sophisticated logical expressions, **Exercise 14** begins by extending the **Pizza** ontology through the creation of its first meaningful subclasses hierarchy. While the modeling activity itself appears straightforward, its engineering significance is far greater than simply creating additional classes.

The exercise demonstrates a principle that applies to every ontology, regardless of its domain:

Semantic meaning should be built upon conceptual organization -- not the other way around!

Concepts must first be identified before they can be described. They must be organized before they can be governed. They must be formally defined before they can be validated. Only after these foundations have been established can automated reasoning discover new knowledge and intelligent systems transform semantic knowledge into executable behavior.

Throughout this chapter, we will examine Conceptual Modeling from multiple complementary perspectives. Beginning with the practical implementation of subclass hierarchies in Protégé, we will progressively explore the underlying principles of ontology engineering, the mathematical foundations of semantic taxonomies, parallels with biological classification, and the role of conceptual modeling within the **Executable Knowledge Architecture (EKA)** framework.

By the end of this chapter, you will recognize that subclass hierarchies are far more than graphical structures displayed within Protégé. They constitute the conceptual backbone of every ontology and establish the semantic vocabulary upon which scalable, governable, and executable knowledge systems are built.

18.2 Learning Objectives

After completing this chapter, you should be able to achieve the following learning outcomes:

Knowledge

- Explain why **Conceptual Modeling** is the first stage of the Semantic Knowledge Development Lifecycle (SKDL).
- Describe the purpose of semantic taxonomies in ontology engineering.
- Understand the relationship between abstraction, generalization, and specialization.
- Explain the semantic meaning of subclass (**SubClassOf**) relationships in OWL

Practical Skills

- Create semantic class hierarchies in Protégé.
- Model subclass relationships correctly using OWL.
- Organize domain concepts into reusable conceptual taxonomies.
- Distinguish conceptual organization from semantic definition.

Engineering Perspective


- Explain why conceptual modeling precedes semantic description.
- Understand how conceptual taxonomies contribute to the **$\mathcal{K}$ - Knowledge Graph** layer within the Executable Knowledge Architecture (EKA) framework.
- Recognize Conceptual Modeling as the engineering foundation for semantic reuse, governance, validation, reasoning, and executable knowledge.

18.3 Position Within the Semantic Knowledge Development Lifecycle

As introduced in Chapter (15), ontology engineering progresses through seven stages of the **Semantic Knowledge Development Lifecycle (SKDL)**.

This chapter focuses on **Stage 1 -- Conceptual Modeling**, highlighted below:

Semantic Knowledge Development Lifecycle



The objective of this stage is straightforward yet fundamental.

Rather than describing the semantic characteristics of concepts or performing automated reasoning, **Conceptual Modeling** establishes the conceptual vocabulary of the domain by identifying its principal concepts and organizing them into meaningful abstraction hierarchies.

The deliverable produced during this stage is therefore not logical inference or executable intelligence, but a **semantic taxonomy** that provides the structural backbone of the ontology.

Within the EKA framework, this stage primarily establishes the conceptual foundation of the *K* - **Knowledge Graph** by defining the concepts that later stages will progressively enrich with semantic meaning.

18.4 Exercise 14 -- Creating the First Conceptual Hierarchy

Having established the engineering motivation for Conceptual Modeling, we now turn to its practical implementation through **Exercise 14** in Michael DeBellis' *Protégé 5 New OWL Pizza Tutorial*.

This exercise represents the first hands-on activity of **Stage 1 - Conceptual Modeling** within the Semantic Knowledge Development Lifecycle (SKDL).

Unlike the previous exercises, which introduced individual OWL language constructs and semantic modeling principles, Exercise 14 focuses on building the first meaningful conceptual hierarchy of the *Pizza* ontology.

The practical modeling task is intentionally modest. Rather than constructing an extensive taxonomy containing numerous pizza varieties, the exercise introduces only two new classes:

- **NamedPizza**
- **MargheritaPizza**

At first glance, this may appear surprisingly simple. However, this simplicity is intentional and reflects a sound ontology engineering methodology.

Instead of overwhelming you with a large number of concepts, the tutorial introduces only the minimum semantic structure required to demonstrate the principles of abstraction and specialization. Once these principles have been mastered, the ontology will gradually expand throughout the subsequent exercises.

The resulting hierarchy is shown conceptually below:

```
./
└─ Thing
    └─ Food
        └─ Pizza
            └─ NamedPizza
                └─ MargheritaPizza
```

Although only a single concrete pizza has been introduced, this hierarchy already illustrates one of the most fundamental principles of ontology engineering.

Rather than treating every pizza as a direct specialization of **Pizza**, the ontology now introduces multiple levels of semantic abstraction.

NamedPizza represents an abstract semantic category rather than an individual pizza variety. It establishes a reusable conceptual container intended to accommodate many named pizzas that will be introduced later in the tutorial. In contrast, **MargheritaPizza** represents the first concrete domain concept belonging to that category.

This design provides a clear separation between general concepts and increasingly specialized domain concepts.

18.4.1 **Pizza** -- The General Concept

The class **Pizza** represents the most generic semantic concepts of a pizza within the ontology.

At this level, the ontology intentionally captures only those semantic characteristics that are common to every pizza, leaving specialized characteristics to be introduced by subclass.

By maintaining this high level of abstraction, future subclasses can inherit these common semantics without unnecessary duplication.

This approach follows one of the central principles of ontology engineering:

Common knowledge should be modeled once and inherited wherever possible.

This principle is widely recognized across many engineering disciplines.

Software engineers avoid duplicating code through inheritance and abstraction.

Database designers normalize schemas to eliminate redundant information.

Ontology engineers apply exactly the same philosophy by organizing concepts into semantic hierarchies, allowing common knowledge to be defined once and inherited consistently by more specialized concepts.

18.4.2 **NamedPizza** -- Introducing an Intermediate Abstraction

Many beginners initially question why the tutorial introduces **NamedPizza** rather than placing **MargheritaPizza** directly beneath **Pizza**.

The answer lies in the importance of **abstraction**.

Rather than representing a single pizza variety, **NamedPizza** represents an **abstraction** that groups pizzas identified by their recognized names. At this stage of the tutorial, only **MargheritaPizza** belongs to this category. However, the abstraction has been intentionally introduced in anticipation of future semantic expansion.

Introducing this intermediate concept provides several engineering advantages.

1. It groups all commercially recognized ("Named") pizzas into a coherent semantic category.
2. It prevents the **Pizza** hierarchy from becoming unnecessarily crowded as additional pizza varieties are introduced.
3. It creates a reusable abstraction upon which future semantic definitions can be consistently applied.

As ontologies continue to evolve, intermediate abstractions such as **NamedPizza** become increasingly valuable because they **reduce redundancy** while **improving maintainability**.

18.4.3 **MargheritaPizza** -- The First Concrete Domain Concept

The class **MargheritaPizza** represents the first concrete pizza type introduced within the ontology.

Unlike **NamedPizza**, which functions primarily as an organizational abstraction, **MargheritaPizza** corresponds to a specific pizza variety that people recognize in everyday life.

Every instance of **MargheritaPizza** is, by definition, also an instance of **NamedPizza**, while every instance of **NamedPizza** is also a **Pizza**. Knowledge therefore flows naturally through the hierarchy via inheritance of semantic meaning.

Importantly, at this stage **MargheritaPizza** possesses only its conceptual identity.

It does not yet specify:

- which toppings it contains;
- which ingredients distinguish it from other pizzas; or
- which logical restrictions formally define it.

For this reason, **MargheritaPizza** serves as the first **instantiable semantic concept** within the **Pizza** ontology, providing a concrete example that can progressively acquire richer semantic descriptions throughout the remaining stages of the Semantic Knowledge Development Lifecycle.

For now, the objective is considerably simpler:

Identify the concept before describing its meaning.

18.4.4 An Engineering Perspective

Separating **conceptual identification** and **semantic definition** is a hallmark of professional ontology engineering. Rather than attempting to model every aspect of a concept simultaneously, ontology engineers first establish *what* the concept is before specifying *how* it should be interpreted by machines.

Exercise 14 therefore illustrates an important distinction between learning an ontology editor and learning ontology engineering.

From the perspective of Protégé, you simply create two new classes and establish a subclass relationship.

From the perspective of ontology engineering, however, you have established the first conceptual taxonomy that will organize every semantic definition introduced throughout the remainder of the ontology.

The modeling activity may be small, but its architectural significance is substantial.

Rather than expanding through isolated concepts, the ontology now begins evolving as an organized conceptual system.

Each newly introduced concept becomes part of a coherent semantic hierarchy that can subsequently support semantic description, knowledge reuse, governance, validation, automated reasoning, and ultimately executable knowledge.

- It allows the conceptual model to stabilize before additional semantic complexity is introduced.
- It reduces modeling errors by separating structural design from logical specification.
- It encourages reuse of conceptual abstractions as the ontology grows.
- It produces an ontology that is easier to maintain, validate, and evolve over time.

Consequently, Exercise 14 should not be viewed merely as an exercise in creating subclasses within Protégé. Instead, it represents the first engineering milestone in transforming an OWL vocabulary into a structured semantic knowledge model.

From the perspective of the Semantic Knowledge Development Lifecycle, the deliverables produced by this exercise can be summarized as follows.

Engineering Perspective	Contribution of Exercise 14
SKDL Stage	Stage 1 -- Conceptual Modeling
Primary Objective	Establish the initial conceptual hierarchy of the Pizza domain.
Modeling Activity	Create NamedPizza and its first concrete subclass, MargheritaPizza .
Engineering Outcome	Define the first semantic taxonomy that will support future semantic enrichment.
Preparation for Later Chapters	Provides the conceptual foundation upon which semantic descriptions, reusable modeling patterns, governance, validation, reasoning, and executable knowledge will progressively be developed.

Although Exercise 14 contains only a single concrete concept, it establishes the conceptual framework upon which the remainder of the **Pizza** ontology will evolve. Every subsequent named pizza introduced throughout the tutorial will become a specialization of the abstraction established in this chapter, demonstrating how carefully designed conceptual hierarchies enable semantic knowledge systems to evolve incrementally while remaining **coherent, maintainable, and reusable**.

18.5 Why Conceptual Modeling Comes Before Semantic Modeling

A common misconception among newcomers to ontology engineering is that the primary purpose of an ontology is to define logical rules, restrictions, and inference mechanisms.

Although these capabilities are important, they represent later stages of semantic knowledge development.

Before an ontology can express formal meaning, it must first establish a stable conceptual foundation.

In other words:

An ontology cannot formally describe concepts that have not yet been conceptually identified.

This principle is the reason why the Semantic Knowledge Development Lifecycle begins with **Stage 1 -- Conceptual Modeling**.

Conceptual Modeling establishes the vocabulary of the knowledge domain by answering:

What concepts exist?

Semantic Modeling, introduced in the next stage of SKDL, builds upon this foundation by answering:

What do these concepts mean?

These two activities are closely related, but they serve different engineering purposes.

18.5.1 Concept Identification Before Semantic Definition

Consider the **Pizza** ontology developed throughout this tutorial.

After Exercise 14, the ontology contains the following conceptual structure:



At this stage, the ontology understands that:

- **MargheritaPizza** is a type of **NamedPizza**.
- **NamedPizza** is a type of **Pizza**.

However, the ontology does not yet understand **why** a **MargheritaPizza** is different from other pizzas.

It does not yet know:

- which toppings characterize a Margherita pizza;
- which ingredients are required;
- which ingredients are prohibited;
- how it differs logically from other pizza concepts.

These questions belong to **Semantic Description**, not Conceptual Modeling.

The conceptual identity must exist before its semantic definition can be expressed.

This separation creates a disciplined progression:

SKDL Stage	Fundamental Question	Example
Stage 1 -- Conceptual Modeling	What concepts exist?	MargheritaPizza is a type of pizza.

SKDL Stage	Fundamental Question	Example
Stage 2 -- Semantic Description	What defines these concepts?	A Margherita pizza has tomato, mozzarella, and basil.
Stage 3 -- Knowledge Reuse	How can existing knowledge be reused?	Reuse existing pizza toppings and ingredient concepts.
Stage 4 -- Semantic Governance	How do we maintain semantic quality?	Ensure naming, consistency, and modeling standards.
Stage 5 -- Semantic Validation	Is the ontology logically correct?	Detect inconsistent concepts.
Stage 6 -- Semantic Reasoning	What new knowledge can be inferred?	Derive classifications automatically.
Stage 7 -- Executable Knowledge	How can knowledge drive actions?	Trigger intelligent decisions or automation.

The lifecycle therefore moves progressively from **concept existence** toward **machine-understandable meaning and execution**.

18.5.2 The Engineering Pattern Across Disciplines

The separation between Conceptual Modeling and Semantic Modeling is not unique to ontology engineering. It reflects a fundamental engineering pattern shared by many disciplines.

Although the terminology differs, the underlying engineering principle remains remarkably consistent.

Every discipline begins by identifying **what exists** before attempting to describe **how it behaves**.

Ontology engineering follows exactly the same progression.

Discipline	Conceptual Question	Semantic / Implementation Question	Initial Modeling Activity
Software Engineering	What objects exist?	How do objects behave?	Class Modeling
Database Design	What entities exist?	What constraints govern the data?	Entity-Relationship Modeling
Business Architecture	What business capabilities exist?	How do capabilities describe the business?	Capability Mapping
Enterprise Architecture	What architectural elements exist?	How do they interact and operate?	Architecture Modeling
Ontology Engineering	What concepts exist?	What logical meaning defines them?	Conceptual Modeling

In each case, engineers first establish the **structure of the domain** before defining the detailed behavior, constraints, or semantics.

For example:

A software architect does not begin by writing complex algorithms before identifying the classes and responsibilities within the system.

A database architect does not begin by creating constraints before determining the entities that the database must represent.

A business architect does not begin by modeling detailed capability relationships before first identifying the business capabilities themselves.

An enterprise architect does not begin by analyzing operational processes before understanding the capabilities and architectural elements of the enterprise.

Ontology engineering follows the same principle.

18.5.3 Why Modeling Semantics Too Early Creates Problems

Attempting to introduce semantic definitions before establishing a conceptual hierarchy often leads to several common problems.

1. Ambiguous Concepts

Without clear conceptual boundaries, the same term may represent different meanings.

For example, if an ontology immediately defines "**Pizza**" through detailed restrictions without first identifying subclasses, it becomes difficult to determine which characteristics belong to all pizzas and which belong only to specific pizza types.

2. Duplicate Knowledge

Without abstraction, common semantic definitions are repeatedly recreated.

For example, if every pizza variety independently defines its base, toppings, and ingredients, the ontology becomes difficult to maintain.

A conceptual hierarchy allows common knowledge to be defined once and inherited by specialized concepts.

3. Difficult Evolution

Domains continuously evolve.

New concepts will inevitably appear as knowledge grows.

A well-designed conceptual taxonomy provides stable extension points where new concepts can naturally be introduced without restructuring existing semantic definitions.

4. Reduced Reasoning Quality

Automated reasoning depends on well-structured concepts.

Reasoners can infer new knowledge only when concepts, relationships, and constraints are represented consistently.

A weak conceptual foundation limits the effectiveness of later reasoning capabilities.

18.5.4 Conceptual Modeling as the Foundation of EKA Knowledge

From the perspective of **Executable Knowledge Architecture (EKA)** framework, Conceptual Modeling establishes the initial foundation of the *K* - **Knowledge Graph** layer.

Conceptual Modeling answers only one fundamental questions:

What concepts exist within the knowledge domain?

At this stage, the ontology provides:

- a shared vocabulary of domain concepts;
- a structured hierarchy of knowledge categories;
- a foundation for future semantic enrichment.


However, the knowledge graph is not yet fully executable.

Later stages of the Semantic Knowledge Development Lifecycle answer progressively model sophisticated questions.

For example:

- How are concepts semantically related?
- Which logical constraints govern them?
- Which facts can be inferred automatically?
- How should semantic quality be validated?
- How can semantic knowledge drive enterprise execution?

The remaining EKA components depend on future stages of semantic development:



Without the conceptual foundation established by Stage 1, later components cannot reliably operate.

- Reasoning requires concepts to reason about.
- Triggers require meaningful semantic events.
- Execution requires knowledge that can be interpreted correctly.
- Governance requires stable semantic structures that can be maintained over time.

Therefore:

Conceptual Modeling is the starting point where knowledge becomes structured, but not yet executable.


18.5.5 From Structure to Meaning

Exercise 14 represents a small but important milestone.

The **Pizza** ontology has moved beyond a collection of isolated OWL classes and has become an organized conceptual system.

The ontology now has a semantic vocabulary upon which future knowledge can be progressively constructed.

The engineering progression is therefore:



This progression reflects a fundamental principle of professional ontology engineering:

First establish what exists. Then define what it means. Finally determine what can be done with that knowledge.

Chapter (16) focuses on the first step of this journey.

The next chapter begins the transition from **conceptual to meaning** by introducing formal semantic descriptions using OWL axioms and logical expressions.

18.6 Building a Semantic Taxonomy

Once the primary concepts within a domain have been identified, the next task is to organize them into a coherent conceptual hierarchy.

This structure is known as a **semantic taxonomy**.

A semantic taxonomy is more than a simple list of concepts. It organizes domain knowledge according to levels of **abstraction**, allowing broad concepts to be progressively refined into increasingly specialized concepts while preserving their conceptual relationships.

Within the **Pizza** ontology, Exercise 14 established the first **meaningful** taxonomy:

```
. /
└ Thing
  └ Food
    └ Pizza
      └ NamedPizza
        └ MargheritaPizza
```

Although this hierarchy currently contains only a few concepts, it already illustrates the fundamental organization strategy used by virtually every ontology, regardless of its domain.

Each child concept represents a more specialized interpretation of its parent while inheriting all of the semantic meaning already established at higher levels.

Each successive level introduces additional semantic specificity.

Moving downward through the hierarchy does not create unrelated concepts.

Instead, every specialization preserves the semantic identity of its ancestors.

Consequently:

- Every **MargheritaPizza** is a **NamedPizza**.
- Every **NamedPizza** is a **Pizza**.
- Every **Pizza** is **Food**. (although we don't create this layer in our working model)
- Every **Food** is ultimately a **Thing**.

This hierarchical organization provides numerous engineering benefits.

Rather than viewing concepts as isolated definitions, ontology engineers organize them into semantic hierarchies that support inheritance, reuse, extensibility, and ultimately automated reasoning.

The taxonomy therefore becomes the conceptual backbone of the ontology.

18.6.1 Generalization and Specialization

A semantic taxonomy is constructed through two complementary modeling activities:


- **Generalization**, which identifies the common characteristics shared by multiple concepts and represents them through a more abstract parent concept.
- **Specialization**, which refines a general concept into increasingly specific concepts that inherit the characteristics of their ancestors while introducing additional distinctions.

These two perspectives describe the same hierarchy from opposite directions.

Moving upward through the hierarchy emphasizes **generalization**, where individual concepts are abstracted into broader categories.

Moving downward emphasizes **specialization**, where broad concepts are progressively refined into more specific semantic concepts.

For the **Pizza** ontology, this progression can be viewed as follows:



Viewed from the top downward, the hierarchy represents increasing specialization.

Viewed from the bottom upward, it represents increasing abstraction (generalization).

A semantic taxonomy captures the **generalization-specialization relationships** between concepts.

Professional ontology engineers constantly alternate between these two viewpoints while designing an ontology.

They ask questions such as:

- Which concepts share common characteristics?

- Should these characteristics be represented once at a higher level?
- Does a new concept represent a genuinely different type, or merely another instance of an existing concept?

The answer determine whether a new subclass should be introduced or an existing abstraction should be reused.

18.6.2 Designing for Semantic Growth

Once of the primary objectives of conceptual modeling is not simply to represent today's knowledge, but to accommodate tomorrow's knowledge.

Real-world domains continuously evolve.

- New products are introduced
- Business processes change.
- Scientific discoveries redefine existing concepts.
- Enterprise architectures expand as organizations grow.

Consequently, a well-designed taxonomy should remain stable even as new concepts are added.

The introduction of `NamedPizza` in Exercise 14 illustrates this principle.

Although only `MargheritaPizza` currently belongs to this category, the intermediate abstraction has been intentionally created to support future expansion.

As additional pizza varieties are introduced, the hierarchy can naturally evolve:

```
. /
  └─ Pizza
    └─ NamedPizza
      ├── MargheritaPizza
      ├── AmericanaPizza
      ├── CapricciosaPizza
      ├── FourSeasonsPizza
      └─ QuattroFormaggiPizza
```

No restructuring of the ontology is required.

The conceptual model already provides a stable extension point.

This illustrates an important engineering principle:

Good conceptual models anticipate future knowledge rather than merely describing current knowledge.

18.6.3 More Than Classification

It is tempting to think of a taxonomy as nothing more than a classification tree.

While organization is certainly one of its purposes, a semantic taxonomy serves a much broader role within ontology engineering.

A well-designed taxonomy provides several important engineering benefits:

1. First, it improves **readability** by organizing concepts into a logical and intuitive structure that is easier for both humans and machines to understand.
2. Second, it promotes **reuse** by encouraging common knowledge to be modeled once at higher levels of abstraction rather than repeatedly duplicated across multiple concepts.
3. Third, it enhances **maintainability** because changes to general concepts automatically benefits more specialized concepts throughout the hierarchy.
4. Finally, it improves **extensibility** by enabling new concepts to be incorporated without disrupting the overall conceptual organization.

These characteristics make semantic taxonomies one of the most important design artifacts produced during Conceptual Modeling.

18.6.4 The Conceptual Backbone of Semantic Knowledge

The taxonomy established during Stage 1 of the **Semantic Knowledge Development Lifecycle** does not yet contain detailed semantic definitions.

At this stage, the ontology simply answers the question:

What concepts exist within the domain, and how are they conceptually related?

The answers to more sophisticated questions -- such as which properties characterized those concepts, which logical restrictions distinguish them, and what new knowledge can be inferred -- will be introduced during later stages of the lifecycle.

Nevertheless, every one of those future semantic capabilities depends upon the conceptual structure established here.

Without a robust and coherent conceptual hierarchy (taxonomy):

- semantic definitions become inconsistent;
- common knowledge cannot be effectively reused;
- logical restrictions become difficult to manage;
- automated reasoning becomes less reliable;
- governance becomes harder to enforce; and
- ontology evolution becomes increasingly expensive.

For this reason, ontology engineers regard a semantic taxonomy not as a preliminary modeling exercise, but as the **conceptual backbone** upon which semantic knowledge systems are progressively constructed.

By the completion of Stage 1, the ontology possesses a stable conceptual vocabulary organized into a reusable hierarchy.

The next step is to understand why this hierarchy carries formal semantic meaning rather than serving merely as a visual organization of concepts.

This transition leads naturally to the next section:

From Taxonomy to Semantic Inheritance.

18.7 From Taxonomy to Semantic Inheritance

By the completion of **Stage 1 -- Conceptual Modeling**, the **Pizza** ontology now contains a coherent semantic taxonomy.

At first glance, this taxonomy appears to be little more than a graphical hierarchy displayed within Protégé.

Many beginners therefore assume that subclass relationships exist primarily to organize concepts visually, much like folders in a file system.

This interpretation is understandable -- but fundamentally **incorrect**.

A folder hierarchy is primarily an organizational mechanism. It helps humans navigate information but carries little semantic meaning. Moving a document from one folder to another changes its location, but not its intrinsic meaning.

A semantic taxonomy is **not** an organizational convenience.

It is **fundamentally different** since it's a **formal semantic model**.

Every subclass in an ontology represents a **logical specification** of its parent class, and the subclass relationship contributes logical meaning to the ontology and establishes how semantic knowledge is inherited throughout the conceptual hierarchy.

Understanding this distinction marks an important milestone in learning ontology engineering.

Rather than seeing the class hierarchy as a navigation tree, we begin to recognize it as one of the ontology's primary mechanisms for representing knowledge.

18.7.1 From Organization to Meaning

Consider the conceptual hierarchy introduced in Exercise 14:

```
. /  
  └ Thing  
    └ Pizza  
      └ NamedPizza  
        └ MargheritaPizza
```

At first glance, this hierarchy resembles the directory structure of an operating system.

Many newcomers naturally interpret it as:

"MargheritaPizza is stored under NamedPizza."

However, this interpretation confuses **organization** with **semantics**.

A folder hierarchy merely groups objects together for convenience.

Moving a document from one folder to another folder changes its location, but it does not change the meaning of the document itself.

Ontology hierarchies works very differently.

Every subclass relationship changes the semantic interpretation of the concepts involved.

Instead of expressing containment, the hierarchy expresses **specialization**.

It therefore should be interpreted as:

Every MargheritaPizza is a NamedPizza, and every NamedPizza is a Pizza.

This seemingly small distinction transforms the hierarchy from a visual diagram into a logical specification of the domain.

18.7.2 The Meaning of SubClassOf

OWL represents specialization through the **SubClassOf** axiom.

Unlike a folder hierarchy, the **SubClassOf** relationship expresses an **IS-A** relationship.


In other words:

- MargheritaPizza **IS A** NamedPizza
- NamedPizza **IS A** Pizza

rather than

- MargheritaPizza is stored inside NamedPizza.

The distinction is illustrated directly in Protégé.



The significance of this relationship extends far beyond visual organization.

Once a class is declared as a subclass of another class, it automatically inherits the semantic commitments established by its parent unless those commitments are explicitly refined or constrained in later axioms.

This mechanism is known as **semantic inheritance**.


18.7.3 Semantic Inheritance in Practice

The practical consequences of semantic inheritance become immediately apparent once semantic restrictions are introduced.

Suppose the ontology defines the following restriction for the class **Pizza**:

```
Pizza
  SubClassOf
    hasBase some PizzaBase
```

This states that every pizza must have at least one pizza base:



No additional semantic restrictions have yet been defined for either **NamedPizza** or **MargheritaPizza**.



Nevertheless, Protégé already understands that both subclasses satisfy this restriction through inheritance.

The ontology therefore recognizes that:

- every **NamedPizza** has a pizza base; and
- every **MargheritaPizza** has a pizza base.

This behavior requires no additional modeling effort.

It follows automatically from the subclass hierarchy established during **Conceptual Modeling**.



These three screenshots together illustrate one of the most important principles in ontology engineering:

Semantic knowledge defined once can be inherited many times.

This dramatically reduces duplication while ensuring that semantic knowledge remains consistent throughout the ontology.

18.7.4 Semantic Inheritance Is Not Software Inheritance

Readers familiar with object-oriented programming often compare ontology subclasses with software class inheritance.

Although the terminology is similar, their purposes are fundamentally different.

In object-oriented programming, subclasses primarily inherit implementation details, executable methods, and software behavior.

The goal is **code reuse**.

Ontology engineering has a different objective.

Ontology subclasses inherit **logical meaning**.

The goal is **semantic consistency** rather than software implementation.

An ontology class does not contain executable code.

Instead, it represents:

a formal description of a concept within the domain.

Consequently, semantic inheritance propagates knowledge rather than program behavior.

This distinction explains why ontology reasoners can derive new knowledge from subclass hierarchies even though the ontology contains no executable algorithms.

18.7.5 Why Semantic Inheritance Matters

Semantic inheritance is one of the key reasons why ontologies scale effectively.

Without inheritance, every specialized concept would need to redefine the characteristics already shared with its parent concepts.

As an ontology grows, this quickly leads to

- duplicated knowledge,
- inconsistent semantics, and
- increasingly difficult maintenance.

By contrast, a well-designed semantic taxonomy allows common knowledge to be modeled once and automatically reused throughout the hierarchy.

As additional pizza varieties are introduced in later chapters, they will inherit existing semantic knowledge while contributing only the characteristics that distinguish them from other pizzas.

This approach keeps the ontology both concise and extensible.

More importantly, it establishes the foundation upon which automated reasoning will operate in later stages of the **Semantic Knowledge Development Lifecycle**.

18.7.6 From Conceptual Hierarchies to Logical Reasoning

Semantic inheritance represents the bridge between **Conceptual Modeling** and **Semantic Reasoning**.

At the current stage of the lifecycle, the taxonomy has organized the conceptual vocabulary of the domain.

As future chapters introduce logical restrictions, equivalence axioms, and reasoning rules, those semantic will naturally propagate throughout the hierarchy via inheritance.

In this sense, the taxonomy established during Stage 1 is already an active component of the ontology's semantic architecture.

It is not simply a diagram used for navigation.

It is the mechanism through which semantic knowledge becomes **reusable, maintainable, and machine-processable**.

This naturally raises a deeper question:

Why can ontology reasoners rely on subclass hierarchies with such confidence?

The answer lies not merely in OWL itself, but in the mathematical foundations upon which OWL is built.

The next section explores those foundations by examining semantic taxonomies through the lens of **set theory** and **partially ordered sets**.

18.8 Interesting Reading -- Mathematical Foundations of Semantic Taxonomy

One of the distinguishing characteristics of ontology engineering is that its conceptual structure are not based solely on intuition or visual organization.

Behind every subclass relationship lies a rigorous mathematical foundation that enables ontology reasoners to derive knowledge **consistently** and **predictably**.

Unlike many modeling techniques, OWL is grounded in formal logic and set theory. Consequently, the semantic taxonomy constructed during **Conceptual Modeling** is more than an engineering artifact -- it is also a mathematical structure.

Understanding these mathematical foundations is not required to build simple ontologies, but it provides valuable insight into **why automated reasoning works** and **why semantic inheritance can be trusted**.

18.8.1 Classes as Sets

From the perspective of *set theory*, every OWL class may be interpreted as a **set** whose members are the individuals belonging to that class.

Suppose we define the following sets for the **Pizza** ontology:

$$P = \{\text{all pizzas}\}$$
$$N = \{\text{all named pizzas}\}$$
$$M = \{\text{all Margherita pizzas}\}$$

Exercise 14 establishes the following conceptual hierarchy:

```
Pizza
  NamedPizza
    MargheritaPizza
```

Mathematically, this hierarchy corresponds directly to the subset relationships:

$$M \subseteq N \subseteq P$$

This expression states that every member of the set M also belongs to (a member of) N , and every member of N also belongs to P .

Note: M , N and P may denote exactly the same set of individuals. In OWL, this situation is modeled using `EquivalentClasses` rather than `SubClassOf`, although mathematically the corresponding sets are equal.

In ontology terminology (OWL), these same relationships are expressed using the `SubClassOf` axiom:

```
MargheritaPizza SubClassOf NamedPizza
NamedPizza SubClassOf Pizza
```

Although one notation is mathematical and the other is ontological, they express exactly the same conceptual relationship.

18.8.2 Why Inheritance Happens Automatically

The mathematical interpretation immediately explains the semantic inheritance discussed in the previous section.

If: $M \subseteq N$

and: $N \subseteq P$

then it follows logically that: $M \subseteq P$

without introducing any additional axioms.

Using `Pizza` ontology's Exercise 14 case:

If every Margherita pizza belongs to the set of Named pizzas, and every Named pizza belongs to the set of pizzas, then it follows logically that Margherita pizza belongs to the set of pizzas without requiring an additional axiom.

This conclusion follows directly from the **transitivity of subset inclusion**, one of the fundamental properties of set theory.

Consequently, semantic inheritance is not a special feature invented by Protégé.

Rather, Protégé and OWL reasoners apply well-established mathematical principles to semantic knowledge.

The ontology therefore inherits semantic restrictions because the underlying mathematics guarantees that the subclass relationships remain logical consistent.

18.8.3 A Taxonomy as a *Partially Ordered Set*

The mathematical structure of a semantic taxonomy extends beyond simple subset relationships.

The collection of OWL classes connected through `SubClassOf` relationships form what mathematicians call a **partially ordered set**, or **poset**.

A partial order satisfies three important properties:

Property	Mathematical Expression
Reflexivity	Every class is a subclass of itself: $A \subseteq A$
Anti-symmetry	If $A \subseteq B$ and $B \subseteq A$, then $A = B^*$
Transitivity	If $A \subseteq B$ and $B \subseteq C$, then $A \subseteq C$

*: the two classes represents the same set of individuals.

These mathematical properties guarantee that semantic hierarchies remain logically consistent as they expand.

From this perspective, OWL does not invent a new hierarchy mechanism. Instead, it formalizes centuries of mathematical reasoning about sets, inclusion, and logical specification into a machine-understandable and machine-processable representation (meaning).

For ontology engineers, the most important of these is **transitivity**, because it enables semantic knowledge to propagate automatically through the hierarchy.

18.8.4 Why Ontology Engineers Should Care

Fortunately, ontology engineers do not need to perform mathematical proofs while building ontologies.

Reasoners perform this work automatically.

Nevertheless, understanding the mathematical foundations of semantic taxonomies provides several important insights.

1. First, it explains why subclass relationships are more than graphical connections.
2. Second, it demonstrates why semantic inheritance is logically sound rather than an implementation convenience.
3. Third, it reveals why ontology reasoners can infer new knowledge with mathematical rigor instead of relying on heuristic algorithms.

This mathematical foundation is one of the reasons why ontologies have become an important technology for knowledge representation, scientific data integration, enterprise architecture, and artificial intelligence.

18.8.5 Transition

The mathematical perspective explains **why semantic inheritance is logically valid**.

Interestingly, long before formal logic and computer science existed, humans had already developed remarkably sophisticated hierarchical classification systems.

The next section explores perhaps the most influential example -- **biological taxonomy**, whose principles continue to inspire modern ontology engineering.

18.9 Engineering Perspective -- Modeling Before Meaning

Exercise 14 deliberately postpones one question that many beginners naturally ask:

Why doesn't the tutorial define the toppings of a *MargheritaPizza* immediately?

After all, once the class has been created, wouldn't it be more efficient to immediately specify its ingredients, logical restrictions, equivalent classes, and reasoning rules?

Surprisingly, the answer is **no**, and this answer lies in the engineering discipline of *incremental semantic modeling*.

Professional ontology engineering deliberately resists this temptation.

Instead of attempting to describe every aspect of a concept simultaneously, ontology engineers follow a disciplined engineering methodology that separate **conceptual organization** from **semantic specification**.

This approach may initially appear slower, but it consistently produces ontologies that are more understandable, maintainable, extensible, and reusable.

The underlying principle can be summarized as follows:

Model the concepts first. Describe their meaning afterwards.

This philosophy is one of the defining characteristics of the **Semantic Knowledge Development Lifecycle (SKDL)** and distinguishes professional ontology engineering from ad hoc ontology construction.

18.9.1 Separating Conceptual Structure from Semantic Meaning

One of the fundamental principles of engineering is the **separation of concerns**.

Complex systems become easier to understand and maintain when different responsibilities are addressed independently rather than simultaneously.

Ontology engineering follows exactly this principle.

During **Conceptual Modeling**, the objective is intentionally limited.

Ontology engineers identify the concepts that exist within the domain and organize them into a coherent semantic taxonomy.

At this stage, they deliberately avoid introducing unnecessary semantic complexity.

Questions such as:

- Which toppings define a Margherita pizza?

- Which pizzas are vegetarian?
- Which logical restrictions distinguish one pizza from another?
- Which classes should be inferred automatically?

are intentionally postponed.

Instead, **Stage 1** ask only one question:

What concepts exist within the domain, and how are they related through abstraction and specialization?

Only after this conceptual foundation has stabilized do later stages progressively enrich the ontology with formal semantic definitions.

This separation allows the conceptual structure to evolve independently of the semantic rules that will later operate upon it.

18.9.2 Exercise 14 as an Engineering Case Study

Exercise 14 provides an excellent example of this methodology.

The entire exercise introduces only two new classes:

- `NamedPizza`
- `MargheritaPizza`

No toppings are specified.

No equivalent classes are created.

No necessary and sufficient conditions are defined.

No reasoning rules are introduced.

From the perspective of someone learning Protégé, this may seem like a surprisingly small amount of work.

From the perspective of ontology engineering, however, something much more significant has been accomplished.

The ontology now possesses its first **meaningful conceptual taxonomy**.

Rather than expanding through isolated classes, the ontology has begun evolving as a structured conceptual system.

Every future named pizza introduced throughout the tutorial will naturally become part of this hierarchy.

As the ontology grows, semantic descriptions, logical restrictions, and reasoning rules can all be attached to concepts that already occupy a well-defined place within the conceptual model.

This incremental approach greatly simplifies both **ontology development** and **long-term maintenance**.

18.9.3 Engineering Before Implementation

This staged methodology -- SKDL -- is by no means unique to ontology engineering.

It reflects a broader engineering practice adopted across many disciplines.

When architects design a building, they first create the structural blueprint before specifying electrical wiring, plumbing, or interior decoration.

Software application architects design class structures before implementing algorithms.

Database designers identify entities before defining integrity constraints, indexes, and optimization strategies.

Enterprise architects establish capability maps before modeling application interactions or technology deployments.

Although these disciplines employ different modeling languages and engineering practices, they all follow the same underlying methodology:

Stabilize the structure before increasing the complexity.

Ontology engineering applies exactly the same discipline.

Rather than:

constructing semantic knowledge in a single step,

professional ontology engineers:

progressively refine the ontology through successive stages of increasing semantic sophistication.


18.9.4 The Role of the Semantic Knowledge Development Lifecycle

The **Semantic Knowledge Development Lifecycle (SKDL)** formalizes this incremental engineering process.

Instead of viewing ontology development as a collection of unrelated modeling activities, SKDL organizes development into seven complementary stages.

1. **Conceptual Modeling** identifies the concepts and establishes conceptual.
2. **Semantic Description** defines the meaning of those concepts.
3. **Knowledge Reuse** promotes semantic consistency and extensibility.
4. **Semantic Governance** establishes quality standards and modeling policies.
5. **Semantic Validation** verifies semantic correctness.
6. **Semantic Reasoning** derives new knowledge through logical inference.
7. **Executable Knowledge** enables semantic knowledge to drive intelligent systems and enterprise execution.

Each stage builds directly upon the work completed during the previous stage.



Rather than representing isolated modeling activities, these stages collectively describe the gradual maturation of semantic knowledge.

18.9.5 Why Incremental Modeling Matters

Attempting to model every semantic detail at the beginning of an ontology project often creates unnecessary complexity.

Concepts that have not yet stabilized may require repeated redesign.

Semantic definitions become tightly coupled to evolving conceptual structure.

Common knowledge is frequently duplicated because appropriate abstractions have not yet been identified.

As the ontology grows, these problems become increasingly difficult to correct.

Incremental semantic modeling avoids these issues, and prevents ontologies from becoming overly complex during the early phases of development.

By establishing a stable conceptual hierarchy first, ontology engineers create a framework that supports continuous refinement without disrupting the overall architecture.

New concepts can be introduced naturally.

Shared semantic can be inherited rather than duplicated.

Reasoning rules can operate over a consistent conceptual vocabulary.

Maintenance becomes significantly simpler because structural changes and semantic changes remain largely independent.

This disciplined progression enables ontologies to evolve over time while preserving semantic consistency.

18.9.6 Engineering Perspective

Viewed from an engineering perspective, Exercise 14 represents far more than the creation of two subclasses in Protégé.

It demonstrates a professional methodology for developing **semantic knowledge systems**.

Rather than attempting to model every aspect of the domain immediately, ontology engineers first establish a stable conceptual vocabulary.

That vocabulary is then progressively enriched with

- semantic definitions,
- reusable modeling patterns,
- governance mechanisms,
- validation strategies,
- reasoning capabilities, and
- ultimately executable knowledge.

This incremental refinement transforms ontology development from a collection of isolated editing activities into a disciplined engineering process.

It also explains why mature ontologies remain understandable, extensible, and maintainable even as they grow to contain thousands -- or even millions -- of concepts and relationships.

The conceptual taxonomy established during **Stage 1 -- Conceptual Modeling** therefore represents much more than the beginning of the ontology.

It provides the architectural foundation upon which every subsequent semantic capability depends.

18.9.7 Transition to the Next Perspective

The engineering methodology described in this section demonstrates how ontology engineers organize the development of semantic knowledge.

Interestingly, the underlying idea of organizing concepts into progressively more specialized hierarchies is far older than computer science itself.

Centuries before the invention of ontologies, scientists were already developing hierarchical classification systems to organize the natural world.

The next section explores **biological taxonomy**, one of history's most successful knowledge organization systems, and examines how its principles continue to influence modern ontology engineering.

18.10 Interesting Reading -- Two Great Classification Traditions

Throughout this chapter, we have explored **Conceptual Modeling** from multiple perspectives.

From ontology engineering, we learned how concepts are identified and organized into semantic taxonomies.

From mathematics, we examined why subclass relationships support logical inheritance.

From engineering practice, we explored why professional knowledge systems are developed incrementally.

However, the fundamental idea behind **conceptual modeling** is far older than computers and semantic technologies.

Long before ontology engineering existed, humanity had already developed sophisticated methods for organizing knowledge.

Two particularly influential classification traditions demonstrate this principle:

1. **Biological Taxonomy** -- organizing the natural world.
2. **Library Classification** -- organizing human knowledge.

Although these two traditions classify very different things, they share a common objective:

Transforming complexity into a structured conceptual system that humans can understand and navigate.

These historical examples provide valuable insight into why conceptual modeling remains one of the most fundamental activities in knowledge engineering.

18.10.1 Biological Taxonomy: Nature's Original Ontology

Long before the emergence of computer science, artificial intelligence, and semantic technologies, **biologists** faced a challenge remarkably similar to that encountered by **ontology engineers** today:

How can an enormous collection of entities be organized into a coherent, reusable, and universally understandable knowledge system?

The answer became one of the most influential classification systems in scientific history:

biological taxonomy.

Biological taxonomy represents one of humanity's earliest examples of systematic conceptual modeling.


Rather than treating every living organism as an isolated entity, scientists organized life into hierarchical categories based on shared characteristics and increasing levels of specialization.

This approach created a structured representation of biological knowledge that continues to support scientific discovery today.

18.10.1.1 Linnaean Classification System

In the eighteenth century, the Swedish naturalist **Carl Linnaeus** introduced a hierarchical classification framework that became the foundation of modern biological taxonomy.

Living organisms are classified into progressively more specialized categories:



Each level represents a different degree of abstraction.

- Higher levels describe broad conceptual categories containing many organisms
- Lower levels describe increasingly specific concepts.

For example, the domestic cat is classified as:

Classification Level (Rank)	Classification
Kingdom	<i>Animalia</i>
Phylum	<i>Chordata</i>
Class	<i>Mammalia</i>
Order	<i>Carnivora</i>

Classification Level (Rank)	Classification
Family	<i>Felidae</i>
Genus	<i>Felis</i>
Species	<i>Felis catus</i>

Each successive level introduces greater specialization while preserving characteristics inherited from broader categories.

A domestic cat is not simply placed below the category "Mammalia".

The relationship has semantic meaning:

Every *Felis catus* is a *Mammalia*, and every *Mammalia* is an *Animalia*.

This is conceptually identical to the meaning of an ontology subclass relationship:

```
Felis catus SubClassOf Felis
Felis SubClassOf Felidae
Felidae SubClassOf Mammalia
```

The hierarchy therefore represents more than classification.

It represents **inherited knowledge**.


18.10.1.2 Biological Taxonomy and the **Pizza** Ontology

The relationship between biological classification and ontology engineering becomes especially clear when comparing their structure.

Biological taxonomy:

```
. /
└─ Animalia
    └─ Mamalia
        └─ Felidae
            └─ Felis
                └─ Felis catus
```

Pizza ontology:



Both structures follow the same conceptual principle:

Conceptual Principle	Biological Taxonomy	Pizza Ontology
General concept	Animalia	Pizza
Intermediate abstraction	Mammalia	NamedPizza
Specific concept	Felis catus	MargheritaPizza
Relationship	Type of organism	Type of concept
Inherited meaning	Biological characteristics	Semantic axioms

The similarity is not accidental.

Both systems solve the same fundamental knowledge organization problem:

How can general knowledge be defined once and inherited by more specific concepts?

This is exactly the purpose of semantic inheritance in ontology engineering.

18.10.1.3 From Classification to Ontology Engineering

Although biological taxonomy strongly resembles ontology engineering, modern ontologies extend beyond traditional classification.

Biological taxonomy primarily answers:

How should living organisms be categorized?

Ontology engineering addresses a broader question:

What concepts exist, how are they related, and what knowledge can be logically derived from those relationships?

Modern ontologies add capabilities that traditional taxonomy does not provide, including:

- formal semantic relationships;
- logical restrictions;
- equivalence definitions;
- domain and range constraints; and
- automated reasoning.

Therefore, ontology engineering can be viewed as an evolution of classical taxonomy.

It preserves the hierarchical organization developed by centuries of scientific classification while adding formal semantics that enable **machines** to interpret, validate, and reason over knowledge.

Within the **Executable Knowledge Architecture (EKA)** framework, biological taxonomy provides an intuitive analogy for the first stage of knowledge construction.

Before knowledge can be executed, it must first be organized.

Before reasoning can occur, concepts must first exist.

Before intelligent automation can operate, semantic structures must first be established.

This is why **Conceptual Modeling** establishes the foundation of the *K* - **Knowledge Graph** layer.

18.10.2 Library Classification: Organizing Human Knowledge

Biological taxonomy demonstrates how humans organize the natural world.

However, humanity has also faced another equally challenging classification problem:

How can the enormous amount of human knowledge itself be organized, stored, and retrieved?

Libraries represent one of the oldest and most successful examples of knowledge organization systems.

Unlike biological taxonomy, which attempts to discover relationships that exist in nature, library classification systems are intentionally designed by humans to organize intellectual concepts.

Their purpose is not to describe physical entities, but to create a navigable structure for human knowledge.

18.10.2.1 Dewey Decimal Classification

One of the most widely recognized library classification systems is the **Dewey Decimal Classification (DDC)** system.

Developed by Melvil Dewey in 1876, DDC organizes human knowledge into ten major categories:

Range	Knowledge Domain
000	Computer Science, Information & General Works
100	Philosophy & Psychology
200	Religion
300	Social Sciences
400	Language
500	Science
600	Technology
700	Arts & Recreation

Range	Knowledge Domain
800	Literature
900	History & Geography

Each broad category is progressively refined into more specific concepts.

For example:

```
 /
└ 500 Science
   └ 510 Mathematics
      └ 512 Algebra
```

The structure allows people to navigate from broad conceptual domains toward increasingly specialized knowledge areas.

18.10.2.2 Library of Congress Classification


Another major classification system is the **Library of Congress Classification (LCC)** system.

Unlike DDC's numerical structure, LCC uses alphabetic categories.

It organizes knowledge into 21 major classes, represented primarily by letters:

Class	Domain
A	General Works
B	Philosophy, Psychology, Religion
C	Auxiliary Sciences of History
D	World History
Q	Science
QA	Mathematics and Computer Science
T	Technology

For example:



Similar to biological taxonomy and ontology hierarchies, LCC moves from broader conceptual categories toward increasingly specific concepts.

18.10.2.3 Library Classification and Ontology Engineering

Library classification provides another important perspective on conceptual modeling.

Although a library classification systems is not an ontology in the formal OWL sense, it demonstrates many principles shared with ontology engineering:

- identifying concepts;
- organizing concepts hierarchically;
- creating meaningful abstraction levels;
- enabling navigation and discovery.

The key difference is the basis of classification.

Dimension	Biological Taxonomy	Library Classification
Classification Object	Living organisms	Human Knowledge
Classification Basis	natural characteristics and evolutionary relationships	Intellectual domains and human design
Primary Purpose	Scientific understanding	Knowledge organization and retrieval
Hierarchy Example	Species → Genus → Family	Topic → Subject Area → Domain
Ontology Connection	Inspires inheritance and specialization	Inspires conceptual organization

Biological taxonomy discovers patterns in the natural world.

Library classification creates structures to help humans navigate the conceptual world.

Both, however, demonstrate the same fundamental idea:

Complex knowledge becomes manageable when organized through meaningful concepts and relationships.

18.10.3 What Both Traditions Teach Ontology Engineers

Biological taxonomy and library classification represent two great traditions of human knowledge organization.

One emerged from the natural sciences, and organizes **entities that exist in nature**.

The other emerged from librarianship and information science, and organizes **knowledge created by humanity**.

Despite their differences, both demonstrate a universal principle:

Conceptual modeling is a fundamental strategy for managing complexity.


Ontology engineering builds upon this long history by combining the strengths of both traditions.

Like biological taxonomy, ontologies use hierarchical specialization to represent concepts and inheritance.

Like library classification, ontologies organize knowledge into structures that support discovery, navigation, and reuse.

However, ontology engineering extends these traditions further by adding formal semantics and automated reasoning.

This progressive can be summarized as:



Within the **Semantic Knowledge Development Lifecycle (SKDL)**, this progression begins with **Conceptual Modeling**.

Before semantic descriptions can be written, before reasoning can derive new knowledge, and before executable intelligence can emerge, concepts must first be identified, organized, and connected through a coherent conceptual hierarchy.

The simple taxonomy created in Exercise 14 therefore represents much more than a Protégé class hierarchy.

It reflects a knowledge organization strategy that has guided scientific discovery, library science, and information management for centuries, while laying the conceptual foundation for modern ontology engineering and the Executable Knowledge Architecture (EKA).

18.11 EKA Perspective -- Conceptual Modeling Establishes the Knowledge Layer (K)

The two great classification traditions explored in the previous section teach us an important lesson:

- Biological taxonomy organizes the natural world.
- Library classification organizes human knowledge.

Although they were developed independently and for different purposes, both demonstrate the same fundamental principle:

Organizing knowledge into coherent conceptual structures is one of humanity's most enduring and effective strategies for managing complexity.

Neither system begins with detailed descriptions, logical rules, or sophisticated reasoning.

Instead, both begin by identifying concepts, organizing them into meaningful hierarchies, and establishing relationships that enable knowledge to be understood, navigated, and progressively expanded.

Ontology engineering follows exactly the same philosophy.

The **Executable Knowledge Architecture (EKA)** framework formalizes this long-established strategy for modern semantic knowledge systems, extending conceptual organization into machine-understandable and ultimately machine-executable knowledge.

Throughout this chapter, we have focused exclusively on the first step of that journey.

By identifying concepts such as **Pizza**, **NamedPizza**, and **MargheritaPizza** and organizing them into a semantic taxonomy, we have established the conceptual vocabulary of the domain.

Although the ontology remains relatively simple, it has already completed the first -- and arguably the most fundamental -- stage of the **Executable Knowledge Architecture (EKA)**.

18.11.1 Conceptual Modeling and the EKA Tuple

As introduced in Chapter (00), EKA represents executable knowledge through five complementary components:

- K - **Knowledge Graph layer**
- R - **Reasoning & Rules layer**
- Θ - **Trigger layer**

- Φ - **Execution layer**
- Γ - **Governance layer**

These components are not independent technologies.

Rather, they represent progressively richer layers of semantic capability built upon a common conceptual foundation.

Within this architecture, the work completed in Exercise 14 contributes almost entirely to the K - **Knowledge Graph** layer.

At this stage, the ontology has not yet introduced sophisticated logical restrictions, reasoning rules, event processing, or executable behavior

Instead, it has established something even more fundamental:

A shared conceptual model describing what exists within the domain.

This conceptual model provides the semantic vocabulary upon which every subsequent stage of semantic knowledge development depends.

18.11.2 Mapping Chapter 16 to the EKA Framework

The contribution of this chapter to each component of the EKA tuple can be summarized as follows:

EKA Component	Contribution of Chapter (16)
K - Knowledge Graph	Establishes the conceptual vocabulary and semantic taxonomy of the Pizza domain through class hierarchies and subclass relationships.
R - Reasoning & Rules	Not yet introduced. The conceptual taxonomy provides the structure upon which logical axioms, restriction, and inference rules will later operate.
Θ - Trigger	Not involved at this stage. Semantic events require formal semantics descriptions that will be developed in later chapters.
Φ - Execution	No executable behavior has been defined. The ontology currently models knowledge rather than operational processes.
Γ - Governance	Concept naming, abstraction, and hierarchy organization establish the first elements of semantic governance, supporting long-term consistency, maintainability, and evolution.

This mapping illustrates an important characteristic of EKA:

Executable intelligence cannot emerge without a well-structured knowledge foundation.

The five components are **progressive rather than parallel**.

Each layer depends upon the successful completion of the layers beneath it.

18.11.3 Why the Knowledge Layer Comes First

Many newcomers to ontology engineering are naturally attracted to the more advanced capabilities of semantic technologies.

Automated reasoning, intelligent querying, knowledge validation, and AI-driven decision support are all compelling examples of what ontologies can achieve.

However, none of these capabilities can exist without a well-organized conceptual foundation.

- Reasoning engines cannot infer relationships between concepts that have never been identified.
- Validation mechanisms cannot enforce constraints upon concepts that have not yet been modeled.
- Semantic queries cannot retrieve knowledge that lacks a coherent conceptual structure.
- Likewise, intelligent automation cannot operationalize knowledge that has not first been represented semantically.

For this reason, the *K* - **Knowledge Graph** layer is not simply one component of EKA.

It is the architectural foundation upon which every other component depends.

18.11.4 From Conceptual Organization to Executable Knowledge


The significance of Exercise 14 therefore extends far beyond creating two additional class in Protégé.

By introducing **NamedPizza** and **MargheritaPizza**, the ontology has established a reusable conceptual taxonomy capable of supporting further semantic enrichment.

Over the remaining stages of **Semantic Knowledge Development Lifecycle (SKDL)**, this taxonomy will progressively acquire:

- formal semantic description;
- reusable modeling patterns;
- governance policies;
- validation mechanisms;
- logical inference; and
- executable behaviors.

In other words, the ontology evolves from:



Every subsequent capability introduced throughout this book ultimately depends upon the conceptual foundation established in this chapter.

18.11.5 An Architectural Perspective

Biological taxonomy provides the conceptual structure for understanding the natural world.

Library classification provides the conceptual structure for organizing human knowledge.

In exactly the same way, **Conceptual Modeling** provides the conceptual structure for semantic knowledge systems.

The comparison is more than an analogy.

All three systems begin with the same architectural principle:

Structure precedes meaning, and meaning precedes intelligence.

Within the EKA framework, this principle explains why K - **Knowledge Graph** occupies the first position in the EKA tuple.

- Without concepts, there can be no semantics.
- Without semantics, there can be no reasoning.
- Without reasoning, there can be no executable knowledge.

18.11.6 Engineering Takeaway

Stage 1 -- **Conceptual Modeling** should therefore never be viewed as merely creating classes within Protégé.

Instead, it establishes the semantic vocabulary that defines the scope of the ontology, organizes knowledge into reusable conceptual structures, and provides the architectural foundation for every subsequent stage of semantic knowledge development.

Just as biological taxonomy and library classification have enabled generations of scientists and scholars to manage increasingly complex **bodies of knowledge**, Conceptual Modeling enables ontology engineers to construct semantic knowledge systems that are scalable, maintainable, governable, and ultimately executable.

For this reason, Stage 1 is not simply the beginning of the **Semantic Knowledge Development Lifecycle**.

It is the indispensable foundation upon which the entire **Executable Knowledge Architecture (EKA)** is progressively constructed.

18.12 Engineering Guidelines for Conceptual Modeling

Conceptual Modeling is far more than simply creating classes within Protégé.

Although the mechanics of defining subclasses are straightforward, designing a semantic taxonomy that remains **understandable, reusable, and maintainable** over time requires careful engineering judgment.

Experienced ontology engineers recognize that conceptual modeling establishes the foundation upon which every subsequent stage of ontology development depends. Decisions made during this stage influence semantic descriptions, logical reasoning, governance policies, validation strategies, and future extensibility.

The following engineering guidelines summarize widely accepted practices for building robust conceptual hierarchies.

18.12.1 Think Semantically Rather Than Visually

One of the most common beginner mistakes is designing an ontology based primarily on how the class hierarchy appears in Protégé.

Although Protégé displays concepts as a tree, the objective is not to produce an attractive diagram. Every node represents a formal semantic concept, and every edge represents a logical **IS-A** relationship.

Each node represents a formal semantic concept (with IRI), not merely a visual element.

This distinction need to be carefully understood, especially when transitioning from visual diagramming tools (such as Microsoft Visio or draw.io) to Protégé.

When deciding where a class belongs, always ask:

Is every instance of this class also an instance of its parent?

- If the answer is yes, the subclass relationship is probably appropriate.
- If not, the hierarchy should be reconsidered.

The objective is not to construct an attractive diagram but to capture the **conceptual structure** of domain knowledge accurately.

For example, the following hierarchy is semantically correct:

```
Pizza
  NamedPizza
    MargheritaPizza
```

Every **MargheritaPizza** **is a** **NamedPizza**, and every **NamedPizza** is a **Pizza**.

A common mistake is organizing classes according to convenience rather than meaning.

For example, placing **PizzaBase** or **PizzaTopping** beneath **Pizza** simply because they are "related" to pizzas would incorrectly imply that every pizza base **is a** pizza.

```
Pizza
  PizzaBase
  PizzaTopping
```

Remember:

Semantic correctness is always more important than visual appearance.

18.12.2 Model Concepts Before Properties

Another frequent mistake is attempting to define object properties, restrictions, and logical axioms before the conceptual hierarchy has stabilized.

Professional ontology development follows a deliberate progression.

1. First identify the concepts.
2. Then organize them into a coherent taxonomy.

3. Only after the conceptual structure has become sufficiently stable should semantic relationships and logical constraints be introduced.

This staged approach reduces unnecessary redesign as the ontology evolves.

Exercise 14 deliberately follows this philosophy.

At this point in the tutorial, **MargheritaPizza** has been introduced as a concept, but it still contains no semantic definition regarding its toppings, ingredients, or distinguishing characteristics.

Those semantic descriptions belong to the next stage of the Semantic Knowledge Development Lifecycle:

Identify concepts first. Describe their meaning later.

Separating conceptual organization from semantic specification produces ontologies that are easier to **understand**, **extend**, and **maintain**.

18.12.3 Prefer Generalization Before Duplication

Whenever several concepts share common characteristics, introduce an appropriate higher-level abstraction instead of repeatedly modeling similar structures.

The introduction of **NamedPizza** provides an excellent example.

Instead of placing every pizza variety directly beneath **Pizza**, the ontology introduces an intermediate abstraction that represents commercially recognized pizzas.

As additional pizzas are introduced in later chapters, they naturally become subclasses of **NamedPizza**, avoiding unnecessary duplication while keeping the hierarchy organized.

Without this abstraction, the taxonomy would gradually become increasingly crowded as new pizza varieties were added.

Generalization therefore improves:

- semantic reuse;
- maintainability;
- consistency; and
- future extensibility.

This principle closely parallels inheritance in software engineering and normalization in database design:

Model common knowledge once, then allow specialized concepts to inherit it.

18.12.4 Design for Future Growth

An ontology should never be designed solely to satisfy current requirements.

Instead, it should provide a conceptual framework capable of accommodating future expansion without major structural redesign.

Exercise 14 illustrates this principle perfectly.

Although only one concrete pizza (`MargheritaPizza`) currently exists beneath `NamedPizza`, the abstraction anticipates many future subclasses.

Later chapters will introduce additional pizzas such as `AmericanaPizza`, `CapricciosaPizza`, `FourSeasonsPizza`, and others.

Because the conceptual hierarchy has already been established, these additions require no architectural changes.

A scalable taxonomy should allow new concepts to fit naturally within the existing structure.

Professional ontology engineers therefore design today's hierarchy with tomorrow's ontology in mind.

18.12.5 Maintain Consistent Levels of Abstraction

Sibling classes should represent concepts at the same level of **generality**.

Mixing broad conceptual categories with highly specific concepts within the same hierarchy often leads to semantic ambiguity and inconsistent reasoning.

For example, the following hierarchy maintains a consistent level of abstraction:

```
Pizza
  NamedPizza
  VegetarianPizza
  NonVegetarianPizza
  PizzaBase
  PizzaTopping
```

Each child represents a major conceptual category within the ontology. (Top level concepts are subClassOf `owl:Thing`)

By contrast, the following organization mixes different abstraction levels:

```
Pizza
  NamedPizza
  MargheritaPizza
  MozzarellaTopping
```

Here, `NamedPizza` represents a broad abstraction, `MargheritaPizza` represents a concrete pizza type, and `MozzarellaTopping` belongs to an entirely different conceptual branch.

Although Protégé allows such structures, they reduce readability and often indicate weaknesses in conceptual design

Maintaining consistent abstraction levels improves both human understanding and automated reasoning.

18.12.6 Distinguish Classes from Individuals

Another common source of confusion is failing to distinguish between **types** and **instances**.

Remember That Classes Describe Types, Not Individuals

Ontology classes describe categories of things rather than specific objects.

Individuals represent concrete objects belonging to those categories

For example:

- `MargheritaPizza` is a class
- "The pizza served on Table 5" is an individual
- `NamedPizza` describes a category.
- `'Pizza #1045'` represents a particular instance.

This distinction mirrors the relationship between a blueprint and the building constructed from it.

- The blueprint defines a type.
- The completed building is an instance of that type.

Maintaining a clear distinction between classes and individual **prevent conceptual confusion** and ensures that ontologies accurately represent domain knowledge rather than operational data.

18.12.7 Building Stable Taxonomies Before Introducing Semantics

Perhaps the most important lesson of this chapter is that conceptual organization should stabilize before detailed modeling begins.

Many beginners are eager to define restrictions such as:

```
hasTopping some MozzarellaTopping
```

or

```
hasBase some ThinAndCrispyBase
```

before establishing whether the conceptual hierarchy itself is appropriate.

Professional ontology engineers work in the opposite order.

They first determine whether concepts such as `NamedPizza`, `VegetarianPizza`, or `SpicyPizza` represent meaningful abstractions.

Only after the conceptual model has matured do they begin enriching it with semantic descriptions.

This incremental approach greatly reduces restructuring later in the development process and produces ontologies that remain coherent as they grow.

It also reflects the philosophy of the **Semantic Knowledge Development Lifecycle (SKDL)**.

- Stage 1 establishes the conceptual foundation.
- Stage 2 enriches that foundation with semantic meaning.

Subsequent stages progressively introduce reuse, governance, validation, reasoning, and executable knowledge.

A stable conceptual taxonomy is the prerequisite for every subsequent stage of ontology engineering.

18.12.8 Engineering Takeaway

Although creating subclasses in Protégé requires only a few mouse clicks, constructing an effective conceptual model requires considerably more thought.

A well-designed taxonomy should accurately represent the conceptual structure of the domain, support semantic inheritance, encourage reuse, accommodate future growth, and provide a stable foundation for semantic enrichment.

The engineering guidelines presented in this section are not specific to the **Pizza** ontology. They reflect general principles that apply equally to enterprise ontologies, biomedical knowledge graphs, digital libraries, manufacturing systems, and large-scale semantic knowledge bases.

By following these principles, ontology engineers create conceptual models that remain understandable, maintainable, and extensible long after the initial ontology has been completed.

18.13 Key Concepts

Throughout this chapter, we explored Conceptual Modeling from practical, mathematical, engineering, historical and architectural perspectives. The following concepts summarize the essential terminology introduced during **Stage 1 -- Conceptual Modeling** of the Semantic Knowledge Development Lifecycle.

Concept	Description
Conceptual Modeling	The process of identifying, abstracting, and organizing domain concepts into a coherent semantic structure before introducing formal semantic definitions or logical axioms.
Conceptual Vocabulary	The complete collection of concepts that define the scope of an ontology and provide the shared language of a knowledge domain.
Taxonomy	A structured classification system that organizes concepts according to IS-A relationships, enabling semantic organization and reuse.
Semantic Taxonomy	A hierarchical organization of concepts based on generalization and specialization, forming the conceptual backbone of an ontology.
Generalization	The process of identifying common characteristics shared by multiple concepts and representing them through a higher-level abstraction.
Specialization	The process of refining a general concept into more specific semantic categories while preserving inherited semantic meaning.
Abstraction	The engineering practice of modeling common knowledge at an appropriate conceptual level to improve reuse, maintainability, and scalability within the ontology.

Concept	Description
Subclass (SubClassOf)	An OWL construct expressing that every instance of one class is also an instance of another class, representing an IS-A relationship.
Semantic Inheritance	The automatic logical inheritance of semantic meaning from parent classes to their subclasses through the ontology hierarchy.
Conceptual Hierarchy	A semantic structure that organizes concepts into multiple levels of abstraction, supporting reuse, extensibility, and automated reasoning.
<i>K</i> - Knowledge Graph Layer	The first component of the Executable Knowledge Architecture (EKA), establishing the conceptual vocabulary and semantic structure upon which reasoning, governance, and execution are progressively built.
Stage 1 - Conceptual Modeling	The first stage of the Semantic Knowledge Development Lifecycle (SKDL), responsible for establishing the conceptual foundation of an ontology before semantic enrichment begins.

18.14 Chapter Summary

This chapter introduced **Stage 1 -- Conceptual Modeling**, the first stage of the **Semantic Knowledge Development Lifecycle (SKDL)** and the essential standing point of professional ontology engineering.

Using **Exercise 14** from Michael DeBellis' *Protégé 5 New OWL Pizza Tutorial* as a practical case study, we explored how ontology engineers begin by identifying and organizing concepts before introducing formal semantics. Although the modeling exercise only the creation of two new classes -- **NamedPizza** and **MargheritaPizza** -- its engineering significance extends far beyond adding nodes to a class hierarchy.

Throughout the chapter, we demonstrated that a subclass hierarchy is not simply a visual tree displayed within Protégé. Instead, it represents a **semantic taxonomy** in which every subclass expresses as **IS-A** relationships with its parent and inherits its conceptual meaning. This semantic organization provides the stable foundation upon which later stages of ontology development progressively introduce richer semantic descriptions, reusable modeling patterns, governance policies, validation mechanisms, automated reasoning, and ultimately executable knowledge.

To place Conceptual Modeling in a broader context, we examined the discipline from several complementary perspectives.

From an **engineering perspective**, we emphasized that every mature engineering discipline begins by identifying the concepts that define its problem domain before describing their behavior or implementation.

From a **mathematical perspective**, we interpreted subclass hierarchies as subset relationships within set theory and showed how semantic inheritance naturally follows from the transitivity of subset inclusion.

From a **historical perspective**, we explored two of humanity's most influential classification traditions. Biological taxonomy demonstrated how the natural world can be organized through hierarchical specialization, while library classification illustrated how human knowledge can be systematically structured into navigable conceptual hierarchies. Although developed for different purposes, both traditions reinforce the same fundamental principle:

Complex knowledge becomes understandable only after it has been organized conceptually.

Finally, we related Conceptual Modeling to the **Executable Knowledge Architecture (EKA)** framework. Within EKA, the work completed in this chapter establishes the *K* -- **Knowledge Graph** layer by defining the conceptual vocabulary and semantic taxonomy that support every subsequent architectural capability.

Without this foundation:

- reasoning has no concepts to infer over,
- governance has no structure to manage,
- validation has no model to verify, and
- execution has no knowledge to operationalize.

By the end of this chapter, the **Pizza** ontology has progressed from a collection of isolated OWL constructs to a coherent conceptual model.

Although the ontology still contains relatively little semantic detail, it now possesses the stable conceptual foundation required for the next stage of semantic development.

18.15 Looking Ahead -- From Concepts to Meaning

At the conclusion of this chapter, the **Pizza** ontology contains a coherent conceptual taxonomy.

Concepts such as **Pizza**, **NamedPizza**, and **MargheritaPizza** have been identified, organized into a meaningful hierarchy, and connected through formal **SubClassOf** relationships.

The ontology therefore knows **what concepts exist** and **how they are conceptually related**.

However, an important question remains unanswered:

What actually distinguishes one concept from another?

At present, the ontology understands that a **MargheritaPizza** is a specialized type of **NamedPizza**, but it does not yet understand **why** it is a Margherita pizza. The hierarchy identifies the concept, but it does not yet describe its semantic meaning.

For example, the ontology cannot yet answer questions such as:

- Which toppings defines a **MargheritaPizza**?
- Which ingredients distinguish it from other named pizzas?
- Which logical restrictions characterize this class?
- how can an OWL reasoner automatically recognize a pizza as a **MargheritaPizza**?

Answering these questions requires more than conceptual organization. It requires **formal semantic descriptions** expressed through OWL axioms, property restrictions, and logical class definitions.

This transition marks the beginning of **stage 2 -- Semantic Description** within the Semantic Knowledge Development Lifecycle.

Whereas **Stage 1** established the conceptual vocabulary of the domain, **Stage 2** progressively enriches those concepts with machine-interpretable meaning. Rather than merely stating that concepts exist, ontology

engineers begin specifying the semantic characteristics that distinguish one concept from another and enable automated reasoning.

The journey therefore continues from answering:

What concepts exist within the domain?

to answer the equally important question:

What do those concepts actually mean?

Last updated at: 2026-08-08, renamed from original Chapter 16

Chapter 19 — Semantic Description: Stage 2 of the Semantic Knowledge Development Lifecycle

Defining Concepts Through OWL Restrictions

- 19.1 Introduction -- From Concepts to Meaning
- 19.2 Learning Objectives
- 19.3 Position Within SKDL -- Stage 2 Highlighted
- 19.4 Exercise 15 -- Creating the First Semantic Description
 - 19.4.1 Engineering Objective
 - 19.4.2 Opening the Class Expression Editor
 - 19.4.3 Creating the First Semantic Restriction
 - 19.4.4 Adding the Second Restriction
 - 19.4.5 Synchronizing the Reasoner
 - 19.4.6 Understanding What Was Actually Added
 - 19.4.7 From Concept to Definition
 - 19.4.8 Engineering Perspective
- 19.5 Why Semantic Description Follows Conceptual Modeling
 - 19.5.1 The Universal Engineering Principle: Structure Before Detail
 - 19.5.1.1 Software Engineering
 - 19.5.1.2 Database Design
 - 19.5.1.3 Enterprise Architecture
 - 19.5.1.4 Ontology Engineering
 - 19.5.2 Why Introducing Semantics Too Early Creates Problems
 - 19.5.3 Exercise 14 and Exercise 15 Demonstrate This Principle
 - 19.5.4 Concept Maturity Before Semantic Commitment
 - 19.5.5 The Connection with "Delay Decisions Until Understanding Improves"
 - 19.5.6 Engineering Takeaway
- 19.6 Introduction to Description Logic (DL) -- The Formal Language Behind OWL
 - 19.6.1 What Is Description Logic?
 - 19.6.2 Why Does OWL Need Description Logic?
 - 19.6.3 The Relationship Between OWL and Description Logic
 - 19.6.4 The Three Fundamental Building Blocks of Description Logic
 - 19.6.5 Decidability -- Why Description Logic Matters
 - 19.6.6 Description Logic in Exercise 15
 - 19.6.7 Preparing for the First Description Logic Constructor
- 19.7 Existential Restrictions (some) -- The First Description Logic Constructor
 - 19.7.1 From Naming Concepts to Description Concepts
 - 19.7.2 Understanding "At Least One"
 - 19.7.3 How Existential Restrictions Supports Automated Classification
 - 19.7.4 Semantic Descriptions Grow Incrementally
 - 19.7.5 Common Misunderstandings About Existential Restrictions
 - 19.7.6 Engineering Takeaway
- 19.8 Class Expressions -- Building Semantic Meaning with Description Logic

- 19.8.1 From Individual Restrictions to Complete Semantic Descriptions
- 19.8.2 What Is a Class Expression?
- 19.8.3 Named Classes and Anonymous Classes
- 19.8.4 Description Logic as a Language of Composition
- 19.8.5 Semantic Composition in the *Pizza* Ontology
- 19.8.6 Necessary Conditions Versus Complete Definitions
- 19.8.7 Open World Assumption Enables Incremental Semantic Modeling
- 19.8.8 Engineering Perspective -- Building Meaning Incrementally
- 19.8.9 Preparing for Future Chapters
- 19.9 Mathematical Perspective -- Semantic Description as Formal Knowledge Representation
 - 19.9.1 From Conceptual Sets to Semantic Sets
 - 19.9.2 Object Properties as Binary Relations
 - 19.9.3 Existential Restrictions as Set constructors
 - 19.9.4 Building Meaning Through Set Operations
 - 19.9.5 Interpretation Functions Give Meaning to Syntax
 - 19.9.6 Consistency, Satisfiability, and Logical Soundness
 - 19.9.7 Decidability Makes Automated Reasoning Possible
 - 19.9.8 From Mathematics to Engineering
- 19.10 Interesting Reading -- From Aristotle to Description Logic
 - 19.10.1 Aristotle -- the Birth of Conceptual Definition
 - 19.10.2 Porphyry's Tree -- Visualizing Hierarchical Knowledge
 - 19.10.3 Leibniz -- The Dream of a Universal Knowledge Language
 - 19.10.4 Frege -- Modern Predicate Logic
 - 19.10.5 Description Logic -- Making Logic Computable
 - 19.10.6 OWL -- The Modern Realization of a Two-Thousand-Year Journey
 - 19.10.7 Reflection -- Why Ontology Engineering Is Different
- 19.11 Engineering Perspective -- Meaning Before Reasoning
 - 19.11.1 Semantics Description Provides the Raw Material for Reasoning
 - 19.11.2 Every Restriction Should Capture One Independent Semantic Fact
 - 19.11.3 Engineering Analogy -- Blueprint Before Construction
 - 19.11.4 Incremental Semantic Maturity
 - 19.11.5 Design Semantic Models for Humans First
 - 19.11.6 Let the Reasoner Become Your Design Assistant
 - 19.11.7 Engineering Discipline Throughout the Semantic Knowledge Development Lifecycle
 - 19.11.8 Engineering Takeaway
- 19.12 EKA Perspective -- Stage 2 Enriches the Knowledge Layer K and Prepares Reasoning R
- 19.13 Engineering Guidelines for Semantic Description
 - 19.13.1 Model Meaning Incrementally
 - 19.13.2 Prefer Necessary Conditions Before Complete Definitions
 - 19.13.3 Write Readable Class Expressions
 - 19.13.4 Reuse Semantic Patterns
 - 19.13.5 Validate Frequently with the Reasoner
 - 19.13.6 Avoid Over-Constraining Early Models
 - 19.13.7 Separate Conceptual Organization from Semantic Description
 - 19.13.8 Engineering Takeaway
- 19.14 Key Concepts

- [19.15 Chapter Summary](#)
- [19.16 Looking Ahead -- Toward Knowledge Reuse](#)

19.1 Introduction -- From Concepts to Meaning

Every successful knowledge system begins by answering two fundamental questions:

What concepts exist?

and

What do those concepts actually mean?

Although these questions appear closely related, they represent two distinct stages of ontology engineering.

In **Chapter (18)**, we completed **Stage 1 -- Conceptual Modeling** of the **Semantic Knowledge Development Lifecycle (SKDL)**. Rather than describing concepts in detail, we focused on identifying the conceptual vocabulary of the domain and organizing it into a coherent semantic taxonomy. Through Exercise 14, the **Pizza** ontology established the first meaningful class hierarchy:

```
./
├─ Pizza
│   └─ NamedPizza
│       └─ MargheritaPizza
```

This hierarchy answered the first question:

What concepts exist within the **Pizza domain?**

At this point, however, the ontology possessed only **conceptual structure**.

Although the ontology knew that **MargheritaPizza** was a specialized type of **NamedPizza**, it still had no understanding of **why** a pizza should be considered a Margherita pizza.

The class existed, but its **semantic meaning** had not yet been defined.

This chapter begins **Stage 2 -- Semantic Description**, where concepts evolve from simple names into formally defined semantic entities.

Rather than merely organizing concepts into a hierarchy, we begin describing the characteristics that distinguish one concept from another. In OWL, these semantic descriptions are expressed through **class expressions** and **logical restrictions**, enabling machines to interpret concepts according to their formal meaning rather than their names alone.

Exercise 15 introduces this important transition by defining the first semantic characteristics of **MargheritaPizza**.

Instead of:

simply stating that a Margherita pizza exists,

we begin:

specifying the conditions that describe it, such as the toppings that characterize this familiar pizza variety.

Although the practical exercise requires adding only two OWL restrictions, its engineering significance extends far beyond the Protégé user interface.

For the first time, the ontology begins expressing **machine-understandable semantics**.

This distinction marks an important milestone in ontology engineering.

A conceptual hierarchy alone provides organization, but it cannot explain the meaning of the concepts it contains.

Semantic description transforms that hierarchy into a formal knowledge model capable of supporting **validation, inference, reuse**, and ultimately **intelligent behavior**.

This progression reflects a principle found throughout engineering.

- Software architects first identify classes before implementing their behavior.
- Database designers first identify entities before defining integrity constraints.
- Enterprise architects first establish architectural elements before specifying their relationships and responsibilities.
- Ontology engineering follows exactly the same philosophy: concepts should first be identified, only then should their semantics be defined.

The transition from **Conceptual Modeling** to **Semantic Description** therefore represents far more than the addition of OWL syntax. It marks the point at which an ontology begins transforming from a conceptual classification system (taxonomy) into a formal semantic knowledge model.

Throughout this chapter, we will examine semantic description from multiple complementary perspectives. Beginning with the practical creation of OWL restrictions in Protégé, we will progressively explore existential restrictions, class expressions, the philosophical foundations of semantic definition, and the mathematical principles that ultimately evolved into modern **Description Logic**, the formal language upon which OWL is built.

By the end of this chapter, you will recognize that a semantic restriction is far more than a line of OWL syntax. It represents the first formal statement of meaning within an ontology and establishes the semantic foundation upon which automated reasoning, knowledge reuse, governance, and executable intelligence will be progressively constructed throughout the remaining stages of the **Semantic Knowledge Development Lifecycle (SKDL)**.

19.2 Learning Objectives

After completing this chapter, you should be able to achieve the following learning outcomes.

Knowledge

- Explain why **Semantic Description** is the second stage of the **Semantic Knowledge Development Lifecycle (SKDL)**.
- Describe the purpose of **OWL restrictions** in defining the meaning of ontology concepts.
- Understand the difference between **conceptual organization** and **semantic description**.
- Explain the semantic meaning of **existential restrictions (some)** in **OWL**.
- Understand how **class expressions** describe concepts through logical conditions rather than hierarchical structure.

Practical Skills

- Create **existential restrictions** using the Protégé Class Expression Editor.
- Define concepts using **OWL class expressions**.
- Apply multiple semantic restrictions to describe a domain concept.
- Use **auto-completion (Ctrl + Space)** to efficiently construct OWL expressions.
- Synchronize the reasoner to verify the consistency of newly added semantic descriptions.

Engineering Perspective


- Explain why semantic description follows conceptual modeling within the SKDL.
- Recognize semantic restrictions as **formal descriptions of meaning** rather than merely Protégé syntax.
- Understand how semantic descriptions prepare an ontology for **knowledge reuse, validation, reasoning, and machine interpretation**.
- Explain how Stage 2 enriches the *K* — **Knowledge Graph** layer and establishes the foundation for the *R* — **Reasoning & Rules** layer within the **Executable Knowledge Architecture (EKA)**.
- Appreciate that semantic description represents the transition from a **conceptual taxonomy** to a **machine-understandable semantic knowledge model**.

19.3 Position Within SKDL -- Stage 2 Highlighted

As introduced in **Chapter (17)**, ontology engineering progresses through the seven stages of the **Semantic Knowledge Development Lifecycle (SKDL)**, each stage contributing a distinct engineering objective toward the construction of executable semantic knowledge systems.

Having completed **Stage 1 -- Conceptual Modeling** in the previous chapter, this chapter advances to **Stage 2 -- Semantic Description**, highlighted below:

Semantic Knowledge Development Lifecycle



Although the ontology already contains a conceptual hierarchy, concepts alone are insufficient for machine understanding.

At the completion of **Stage 1**, the ontology could answer questions such as:

- Which concepts exist?
- How are they organized?
- Which concepts specialize others?

However, it still could not answer a much more important question:

What does each concept actually mean?

Providing formal answers to that questions is the primary objective of **Stage 2 -- Semantic Description**.

Rather than introducing new concepts, this stage enriches the existing conceptual vocabulary by attaching formal semantic definitions to classes through **OWL class expressions, property restrictions**, and other logical constructs.

For example, in the previous chapter the ontology introduced the concept:

MargheritaPizza

At that stage, the class served only as a conceptual placeholder within the taxonomy.

During this chapter, the ontology begins to describe its semantic meaning by stating that a **MargheritaPizza** has toppings such as **MozzarellaTopping** and **TomatoTopping**.

These statements are not merely annotations for human readers.

They are formal logical expressions that become part of the ontology's machine-interpretable semantics, allowing automated reasoners to understand, validate, and eventually infer new knowledge.

Consequently, the primary deliverable of **Stage 2** is not longer a conceptual hierarchy, but a **semantic model** in which concepts are enriched with logical meaning.

This distinction is fundamental:

- Stage 1 establishes **what** exists.
- Stage 2 begins defining **what those concepts mean**.

Within the **Executable Knowledge Architecture (EKA)** framework, this stage continues enriching the **K -- Knowledge Graph** layer by transforming conceptual nodes into semantically described knowledge objects.

Although reasoning has not yet become the primary focus of development, the semantic descriptions introduced during this stage establish the logical foundation upon which the **R -- Reasoning & Rules** layer will later operate.

The relationship between the first two stages of SKDL can therefore be summarized as follows:

Stage	Primary Engineering Objective	Primary Deliverable
Stage 1 -- Conceptual Modeling	Identify and organize domain concepts into a coherent taxonomy.	Conceptual hierarchy (semantic taxonomy).

Stage	Primary Engineering Objective	Primary Deliverable
Stage 2 -- Semantic Description	Define the formal meaning of concepts using OWL class expressions and logical restrictions.	Semantically enriched ontology capable of supporting logical interpretation.

Together, these two stages establish the semantic foundation of the ontology.

First, the ontology determines **what concepts exist**.

Then, it specifies **what those concepts mean**.

Only after both the conceptual structure and semantic descriptions have been established can subsequent stages introduce knowledge reuse, governance, validation, automated reasoning, and ultimately executable knowledge.

19.4 Exercise 15 -- Creating the First Semantic Description

Having established the engineering motivation for **Semantic Description**, we now turn to its practical implementation through **Exercise 15** in Michael DeBellis' *Protégé 5 New OWL Pizza Tutorial*.

This exercise represents the first hands-on activity of **Stage 2 -- Semantic Description** within the **Semantic Knowledge Development Lifecycle (SKDL)**.

Unlike the previous exercise, which focused on organizing concepts into a conceptual hierarchy, Exercise 15 introduces the first formal semantic descriptions of those concepts.

The practical modeling task remains intentionally simple.


Rather than constructing a complete semantic definition for **MargheritaPizza**, the exercise introduces only two existential restrictions:

- `hasTopping some MozzarellaTopping`
- `hasTopping some TomatoTopping`

Although these additions appear modest, they represent a significant milestone in ontology engineering.

For the first time, the ontology begins expressing **machine-understandable meaning** rather than merely identifying concepts.

The overall workflow of Exercise 15 can be summarized as follows:



Unlike Chapter (18), where the primary objective was to determine **what concepts exist**, the objective of this exercise is fundamentally different.

Rather than introducing new classes, we enrich an existing concepts by describing the semantic characteristics that distinguish it from other pizzas.

This distinction marks the transition from **conceptual organization** toward **semantic definition**.

19.4.1 Engineering Objective


At the completion of **Exercise 14**, the ontology already contained the following conceptual hierarchy:

```

/
├ Pizza
│   └ NamedPizza
│       └ MargheritaPizza

```

Those 3 concepts are highlighted in below screen in Protégé:



This hierarchy successfully established that a **MargheritaPizza** is a specialized type of **NamedPizza**, which in turn is a specialized type of **Pizza**.

However, the ontology still lacked an answer to an obvious semantic question:

Why should a pizza be considered a Margherita pizza?

The class existed only as a conceptual category.

Nothing within the ontology explained the characteristics that distinguished a Margherita pizza from every other named pizza.

Exercise 15 begins answering that question by introducing formal semantic descriptions using **OWL class expressions**.


Rather than relying on the class name itself, the ontology starts defining the concept through logical conditions that software systems can interpret and reason about.

Consequently, the engineering objective of this exercise is not simply to add OWL syntax.

Its objective is to transform a conceptual class into a semantically described concept.

19.4.2 Opening the Class Expression Editor

To begin the exercise, select **MargheritaPizza** from the class hierarchy within the **Classes** tab.



In the **Description** view, locate the **SubClassOf** section and click the **Add (+)** button.

Instead of using the Object Restriction Creator introduced in earlier exercises, Exercise 15 now employs the **Class Expression Editor**.

This editor allows ontology engineers to construct complete OWL class expressions directly using Description Logic syntax.


It therefore represents one of the most important editing interfaces within Protégé, as many advanced semantic definitions are created through this editor.

19.4.3 Creating the First Semantic Restriction

After opening the dialog, switch to the **Class Expression Editor** tab.


Enter the following expression:

`hasTopping some MozzarellaTopping`



Rather than typing the complete topping name manually, Protégé provides a highly useful productivity feature.

After typing `hasTopping some Mo`, press `Ctrl + Space`, Protégé automatically searches the ontology for matching entities.



[!Note] You need ensure you're in the English input mode for these shortcut working.

If only one candidate exists, the editor completes the remaining text automatically.

If multiple candidates match the partial name, Protégé presents a list of possible completions from which you may select the desired concept.

Although this shortcut appears minor, experienced ontology engineers use it continuously when constructing large ontologies containing hundreds of even thousands of classes and properties.

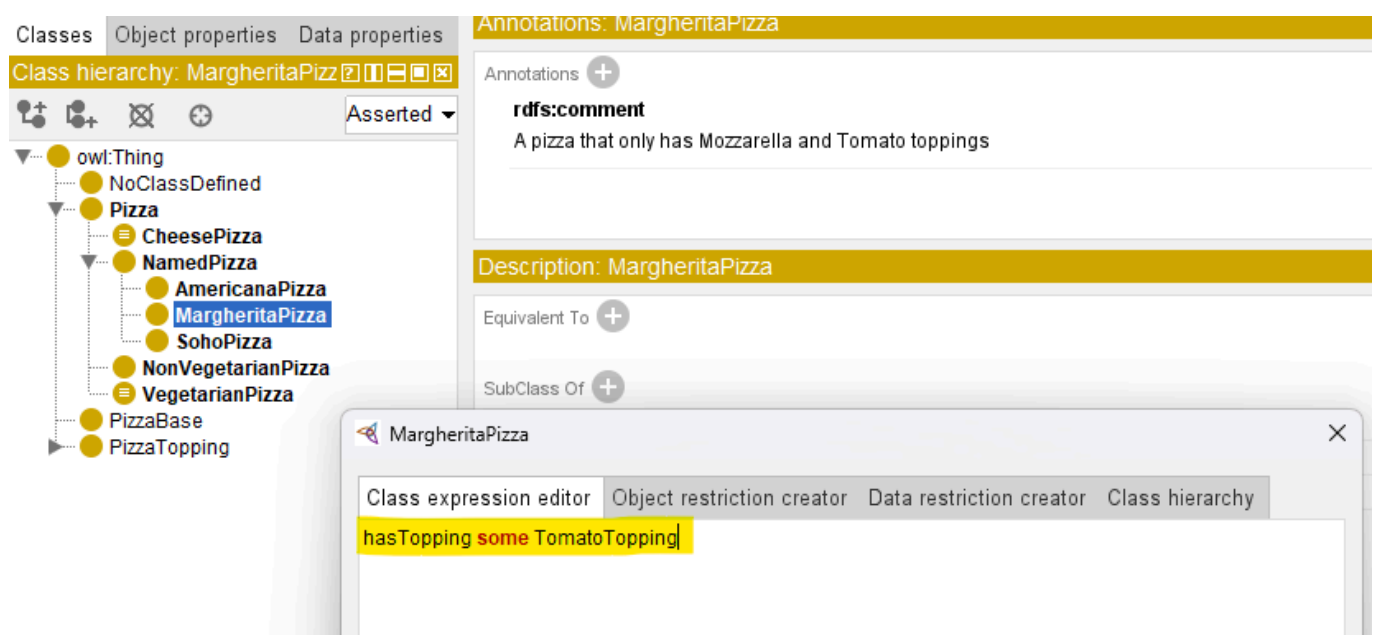
Using auto-completion:

- reduces typing errors,
- improves modeling efficiency, and
- encourages the consistent reuse of existing ontology entities.

19.4.4 Adding the Second Restriction

After confirming the first restriction, repeat the same procedure to add a second semantic description:

`hasTopping some TomatoTopping`



At this point, **MargheritaPizza** now contains two independent existential restrictions.

The screenshot shows the Protégé ontology editor interface. On the left, the 'Classes' tab is active, displaying a class hierarchy. The hierarchy starts with 'owl:Thing' at the top, followed by 'NoClassDefined', 'Pizza', 'CheesePizza', 'NamedPizza', 'AmericanaPizza', 'MargheritaPizza' (highlighted in blue), 'SohoPizza', 'NonVegetarianPizza', 'VegetarianPizza', 'PizzaBase', 'PizzaTopping', 'CheeseTopping', and 'MozzarellaTopping'. On the right, the 'Annotations: MargheritaPizza' section is visible. It shows an 'rdfs:comment' annotation: 'A pizza that only has Mozzarella and Tomato toppings'. Below this, the 'Description: MargheritaPizza' section is shown, which includes a 'SubClass Of' section. This section contains two blue boxes: 'hasTopping some MozzarellaTopping' and 'hasTopping some TomatoTopping', both of which are highlighted with a red rectangle. Below these boxes is a 'NamedPizza' entry.

Unlike programming languages, where multiple statements are executed sequentially, OWL treats these restrictions as logical conditions that collectively contribute to the semantic description of the class.

Each restriction captures one aspect of the concept's meaning.

Together, they begin forming a richer semantic definition of **MargheritaPizza**.

Although the ontology still does not provide a complete definition of a Margherita pizza, it has now taken its first beyond conceptual classification.

19.4.5 Synchronizing the Reasoner

After both restrictions have been added, synchronize the reasoner.

Running the reasoners ensures that the newly introduced semantic description remain logically consistent with the existing ontology.

Although no new inferred classification appear at this stage, regularly synchronizing the reasoner is considered good engineering practice.

Frequent validation allows modeling errors to be detected early, when they are still relatively easy to correct.

Professional ontology development follows an incremental cycle:

Model → **Validate** → **Refine**

rather than postponing validation until the ontology has become substantially larger.

19.4.6 Understanding What Was Actually Added

From the perspective of Protégé, Exercise 15 appears to involve only two additional entries beneath the **SubClass Of** section.

From the perspective of ontology engineering, however, something much more significant has occurred.

Before this exercise, the ontology contains only the conceptual statement:

```
MargheritaPizza
  SubClassOf NamedPizza
```

After completing the exercise, the semantic description has evolved into:

```
MargheritaPizza
  SubClassOf NamedPizza

  hasTopping some MozzarellaTopping

  hasTopping some TomatoTopping
```

The ontology no longer identifies only **where** `MargheritaPizza` belongs within the taxonomy.

It also begins expressing **what characteristics distinguish it**.

This represents the ontology's first formal semantic description.

The concept has progressed from

being merely a named category

toward

becoming a **logically described** class.

19.4.7 From Concept to Definition

Exercise 15 illustrates one of the central principles of ontology engineering.

Stage 1 established **conceptual identity**.

Stage 2 begins establishing **semantic meaning**.

This distinction is fundamental.

Conceptual Modeling determines that a concept exists.

Semantic Description specifies the logical conditions that characterize that concept.

The ontology therefore progresses from asking:

What concepts exist?

to asking:

What formally defines those concepts?

This progression transforms a conceptual taxonomy into a semantic knowledge model capable of supporting automated interpretation.

19.4.8 Engineering Perspective

Exercise 15 demonstrates an important distinction between **learning Protégé** and **learning ontology engineering**.

From the perspective of Protégé, you simply just enter two class expressions and then synchronize the reasoner.

From the perspective of ontology engineering, however, you have introduced the ontology's **first formal semantic definitions**.

Although only two restrictions have been added, they fundamentally change the role of the ontology.

Previously, the ontology organized concepts.

Now, it begins describing their **meaning**.

This marks the transition from a **semantic taxonomy** to a **machine-interpretable semantic model**.

Every subsequent semantic definition introduced throughout the remainder of the tutorial will build upon this same modeling approach.

The engineering contribution of **Exercise 15** can be summarized as below table:

Engineering Perspective	Contribution from Exercise 15
SKDL Stage	Stage 2 -- Semantic Description
Primary Objective	Define the semantic meaning of ontology concepts using OWL class expressions and logical restrictions.
Modeling Activity	Add existential restrictions describing the toppings associated with MargheritaPizza .
Engineering Outcome	Transform a conceptual class into a semantically described ontology concept.
Preparation for Later Chapters	Establish the semantic foundation required for knowledge reuse, validation, reasoning, and ultimately executable intelligence.

Although **Exercise 15** introduces only two semantic restrictions, it represents one of the most important transitions within the **Semantic Knowledge Development Lifecycle (SKDL)**.

From this point onward, the ontology is no longer concerned solely with identifying concepts. It begins constructing a formal semantic model in which concepts are defined through **logical meaning**, enabling machines to **interpret**, **validate**, and eventually **reason** over the knowledge they represent.

19.5 Why Semantic Description Follows Conceptual Modeling

The transition from **Conceptual Modeling** to **Semantic Description** represents one of the most important engineering principles in ontology development:

A concept must be identified before its meaning can be formally defined.

Although conceptual models and semantic descriptions are closely connected, they solve fundamentally different problems.

Conceptual Modeling answers:

What concepts exist in the domain?


While, Semantic Description answers:

What conditions define the meaning of those concepts?

Confusing these two stages often leads:

- unstable ontologies,
- unnecessary redesign, and
- semantic inconsistencies.

Professional ontology engineering therefore follows a disciplined progression:



The purpose of this section is to understand why semantic descriptions should be introduced only after the conceptual foundation has become sufficiently stable.

19.5.1 The Universal Engineering Principle: Structure Before Detail

The relationship between conceptual modeling and semantic description reflect a broader principle found across many engineering disciplines:

Establish the structure first, then define the detailed behavior and constraints.


This principle is not unique to ontology engineering.

It appears repeatedly in software architecture, database design, enterprise architecture, and many other engineering practices.

19.5.1.1 Software Engineering

Software engineers normally do not begin by implementing complex business logic before identifying the fundamental software entities.

A typical progression is:



For example, before implementing methods for a `CustomerOrder` class, developers first need to determine:

- What is a customer?
- What is an order?
- How are customers and orders related?

Without a stable conceptual model, implementation details are likely to change repeatedly.

Ontology engineering follows the same philosophy.

Before defining:

```
MargheritaPizza
  hasTopping some MozzarellaTopping
```

the ontology must first establish that:

```
Pizza
  NamedPizza
    MargheritaPizza
```

represents a meaningful conceptual structure.

19.5.1.2 Database Design

Database engineers follow a similar discipline.

Before defining:

- foreign keys;
- constraints;
- indexes; and
- transaction rules,

database designers first identify:

- entities;
- attributes; and
- relationships.

For example:

```
graph TD; Customer --> places; places --> Order;
```

must be understood before introducing database constraints.

A database schema with incorrect entities cannot be repaired simply by adding more constraints.

The same applies to ontologies:

A semantically incorrect concept hierarchy cannot be fixed by adding more OWL restrictions.

19.5.1.3 Enterprise Architecture

Enterprise architecture also separate conceptual understanding from detailed specification.

Before analyzing:

- application dependencies;
- integration patters; and
- technology platforms,

architects first establish:

- business capabilities;
- organizational concepts; and
- information domains.

For example:

```

Business Capability
  |
  supports
  |
Business Process

```

must be conceptually understood before detailed architecture decisions can be made.

19.5.1.4 Ontology Engineering

Ontology engineering applies the same engineering discipline:

Engineering Discipline	First Establish	Then Define
Software Engineering	Classes and objects	Behavior and implementation
Database Design	Entities and relationships	Constraints and optimization
Enterprise Architecture	Capabilities and domains	Architecture decisions
Ontology Engineering	Concepts and taxonomy	Semantic restrictions and logical axioms

The pattern is consistent:

Conceptual structure provides the foundation upon which detailed meaning is built.

19.5.2 Why Introducing Semantics Too Early Creates Problems

A common beginner mistake in ontology engineering is attempting to define detailed semantic restrictions immediately after creating a class.

For example, imaging starting with:

```
MargheritaPizza`
```

and immediately adding:

```

MargheritaPizza
  hasTopping some MozzarellaTopping

```

At first glance, this appears reasonable.

However, several important questions remain unanswered:


- Is `MargheritaPizza` correctly positioned in the hierarchy?
- Should it be a subclass of `NamedPizza`?
- Should there be another abstraction such as `ItalianPizza`?

- Are toppings the correct defining characteristics?
- Should ingredients, cooking methods, or regional origins also contribute to the definition?

If the conceptual model later changes, every semantic restriction attached to those concepts may require modification.

This creates unnecessary complexity.

The engineering consequence is:



Therefore, professional ontology development separates:

Conceptual Stability

before:

Semantic Complexity.

19.5.3 Exercise 14 and Exercise 15 Demonstrate This Principle

The design of Michael DeBellis' *Protégé 5 New OWL Pizza Tutorial* intentionally follows this engineering sequence.

Exercise 14

First, the ontology establishes conceptual organization:

```

Pizza
  NamedPizza
    MargheritaPizza
  
```

At this stage, the ontology knows:

- what concepts exist;
- how concepts are related; and
- where each concept belongs in the taxonomy.

However, it does not yet define the detailed meaning of these concepts.

Exercise 15

Only after the hierarchy exists does the ontology introduce semantic meaning:

```

MargheritaPizza
  SubClassOf
  
```

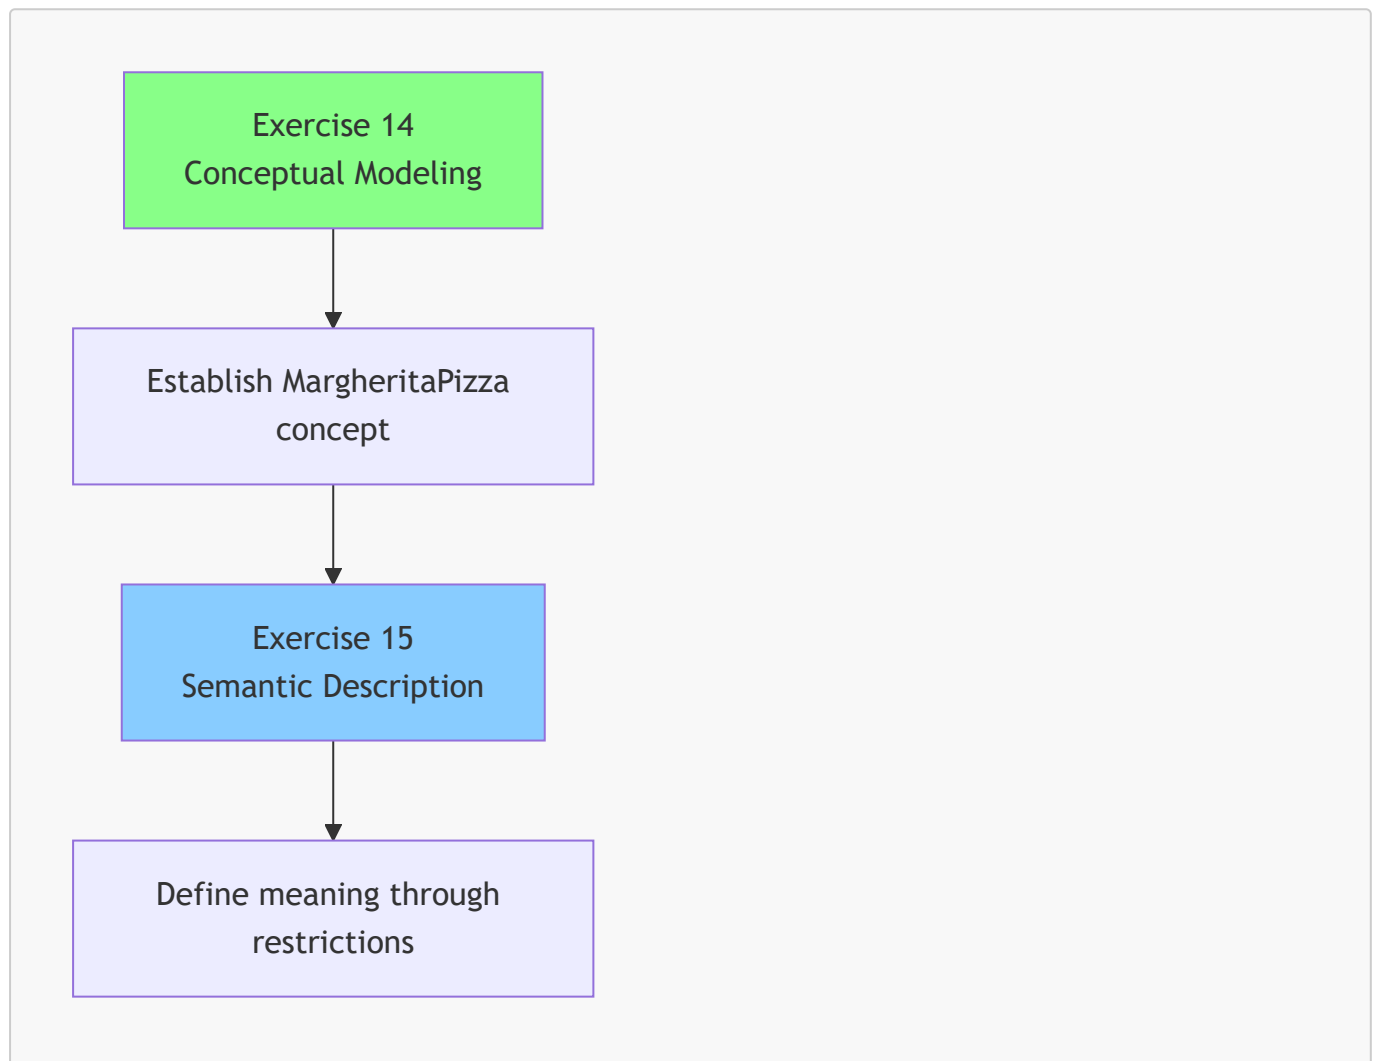
```
hasTopping some MozzarellaTopping

MargheritaPizza
  SubClassOf
    hasTopping some TomatoTopping
```

Now the ontology can express:

A **MargheritaPizza** is not only a named pizza, it is a pizza concept characterized by specific semantic conditions.

The sequence is therefore like below:



This separation is deliberate.


The tutorial is not simply teaching Protégé operations. It is teaching a professional ontology engineering methodology.

19.5.4 Concept Maturity Before Semantic Commitment

A useful way to understand this engineering principle is through the concept of **concept maturity**.

Not every newly discovered concept is immediately ready for detailed semantic definition.

A concept typically evolves through several levels:



Candidate Concept

A possible domain idea has been identified.

Example:

```
SpecialPizza
```

At this stage, its meaning may still be unclear.

Recognized Concept

The concept is accepted as meaningful within the domain.

Example:

```
NamedPizza
```

The ontology community agrees that named pizzas represent a useful category.

Organized Concept

The concept has a defined position in the taxonomy.

```
Pizza
  NamedPizza
    MargheritaPizza
```

Semantically Defined Concept

Formal restrictions are introduced.

Example:

```
MargheritaPizza
  hasTopping some MozzarellaTopping
```

Reasoning-Ready Concept

The ontology contains sufficient formal semantics to support inference (through reasoning).

This maturity model explains why ontology engineering is an incremental process.

19.5.5 The Connection with "Delay Decisions Until Understanding Improves"

This engineering approach also reflects a broader principle used in modern software and system engineering:

Delay irreversible decisions until sufficient understanding exists.


In agile development, architects avoid excessive upfront assumptions because requirements evolve.

Ontology engineering follows the same principle.

A premature semantic definition can become a constraint that limits future modeling choices.

By first establishing conceptual structures (through conceptual modeling), ontology engineers preserve flexibility while gradually increasing semantic precision.

The result is a more adaptable knowledge model:



19.5.6 Engineering Takeaway

Conceptual Modeling and Semantic Description are not competing activities.

They are complementary stages of knowledge engineering.

- Conceptual Modeling creates the **semantic vocabulary**.
- Semantic Description defines the **formal meaning** of that vocabulary.

The relationship can be summarized as:

- **Conceptual Modeling** tells us what exists. - **Semantic Description** tells us what those concepts mean.

Within the Semantic Knowledge Development Lifecycle:

- Stage 1 - Conceptual Modeling: establishes concepts and taxonomy
- Stage 2 - Semantic Description: enriches concepts with formal meaning

Only after both stages have been completed does an ontology become sufficiently expressive for advanced capabilities such as **reasoning**, **validation**, **knowledge reuse**, and **executable intelligence**.

Therefore, the engineering discipline established in this chapter can be summarized by one fundamental rule:

Do not define complex semantics for concepts that have not yet achieved conceptual stability.

This principle is the foundation for building ontologies that remain understandable maintainable, and scalable as semantic knowledge systems evolve.

19.6 Introduction to Description Logic (DL) -- The Formal Language Behind OWL

Having established **why Semantic Description follows Conceptual Modeling** in the previous sections, we are now ready to examine the formal language that makes semantic description precise, machine interpretable, and logically verifiable.

Although Protégé allows use to create *classes*, *properties*, and *restrictions* through an intuitive graphical interface, the ontology is not stored merely as a collection of diagrams or forms. Behind every OWL construct lies a rigorous logical language known as **Description Logic (DL)**.

Understanding this language is not required to operate Protégé effectively, but it greatly improves one's understanding of

- **why OWL behaves as it does,**
- **how reasoners derive new knowledge,** and
- **what semantic restrictions actually mean.**

For ontology engineers, Description Logic provides the conceptual bridge between graphical modeling and formal knowledge representation.

19.6.1 What Is Description Logic?

Description Logic (DL) is a family of **decidable knowledge representation languages** designed for representing structured knowledge using formally defined concepts, relationships, and individuals.

Unlike general-purpose mathematical logic, Description Logic focuses specifically on describing domains of knowledge in a way that is both expressive and computationally manageable.

Its primary objective is to enable computers to answer questions such as:

- Is this ontology logically consistent?
- Which classes does an individual belong to?
- Does one concept imply another?

- Are two concepts logically equivalent?
- Can new knowledge be inferred (derived) automatically?

These capabilities form the foundation of modern ontology engineering and semantic reasoning.

Consequently, every OWL ontology is not merely a collection of XML or RDF statements. It is also a **formal Description Logic knowledge base** that **can be interpreted by automated reasoners**.

19.6.2 Why Does OWL Need Description Logic?

Earlier chapters introduced OWL as a language for modeling semantic knowledge.

However, a modeling language alone is insufficient.

Computer ('Machine') cannot reliably interpret concepts such as **Pizza**, **MargheritaPizza**, or **MozzarellaTopping** unless their meanings are expressed using **precise logical rules**.

Description Logic provides exactly this formal semantics.

Instead of:

relying on natural-language interpretation,

DL:

defines every construct mathematically, ensuring that semantic meaning is unambiguous and independent of human interpretation.

This logical foundation enables different OWL tools and reasoners to interpret the same ontology consistently.

More importantly, it allows automated reasoning to become predictable, reproducible, and mathematically sound.

[!note] When we use "DL" as abbreviation for "Description Logic", do not mix that with "Deep Learning", 😊

19.6.3 The Relationship Between OWL and Description Logic


A common misconception is that OWL and Description Logic are separate technologies.

In reality, they are closely related.

Description Logic provides the formal semantics, while OWL provides a standardized engineering language for expressing those semantics.

From an engineering perspective:

- **Description Logic** defines *what the statements mean*.
- **OWL** provides *how those statements are written and exchanged*.
- **Protégé** provide *how ontology engineers create and manage those statements*.
- **Reasoners** use Description Logic to interpret the ontology and derive logical consequences.



Seen from this perspective, Protégé is not "performing magic". It is constructing Description Logic expressions through a graphical interface, while reasoners apply formal logic to those expressions.

19.6.4 The Three Fundamental Building Blocks of Description Logic

Despite its mathematical foundations, Description Logic is built upon only three fundamental kinds of elements, as below table shows:

Description Logic	OWL	Meaning
Concept	Class	A category of things
Role	Object Property	A relationship between concepts
Individual	Instance	A concrete object belonging to a concept

These correspond directly to the principal modeling elements introduced throughout the previous chapters.

For example, within the **Pizza** ontology:

Description Logic	Pizza Ontology Example
Concept	Pizza , NamedPizza , MargheritaPizza
Role	hasTopping , hasBase
Individual	A particular pizza served in a restaurant

This correspondence demonstrates that OWL terminology is largely an engineering vocabulary built upon Description Logic concepts.

19.6.5 Decidability -- Why Description Logic Matters

One of the defining characteristics of Description Logic is **decidability**.

A logical language is **decidable** if every valid reasoning task is guaranteed to terminate with a definite answer.

In practical terms, this means that a reasoner will always determine whether:

- an ontology is consistent;
- a class is satisfiable;
- one class is a subclass of another; or
- an individual belongs to a particular class.

The reasoning process may require substantial computation for very large ontologies, but it is guaranteed for finish.

This property is critically important in engineering environments.

Without decidability, an ontology reasoner could enter an infinite reasoning process or produce unpredictable behavior, making large semantic systems unsuitable for practical applications.

OWL deliberately balances expressive power with computational tractability so that automated reasoning remains reliable even for complex knowledge bases.

19.6.6 Description Logic in Exercise 15

Although Exercise 15 appears to involve only two simple Protégé operations, it actually introduces the first Description Logic constructor used in this tutorial.

When we add the restriction:

```
hasTopping some MozzarellaTopping
```

we are not merely editing a class description.

We are constructing a formal Description Logic expression.

In standard Description Logic notation, this restriction is written as:

```
 $\exists$  hasTopping . MozzarellaTopping
```

where:

- $\exists$ denotes existential quantification ("there exists");
- **hasTopping** is a role (object property); and
- **MozzarellaTopping** is a concept (class).

This expression states that every instance satisfying the restriction must be related through the **hasTopping** role to **at least one** instance of **MozzarellaTopping**.

Although Protégé hides much of this mathematical notation behind a user-friendly interface, every restriction created in the editor ultimately corresponds to a Description Logic expression interpreted by the reasoner.

19.6.7 Preparing for the First Description Logic Constructor

Description Logic contains numerous logic constructors for describing concepts.

Among the most frequently used are:

- existential restrictions;
- universal restrictions;
- intersections;
- unions;
- complements;
- cardinality restrictions; and
- equivalence definitions.

Exercise 15 introduces the simplest and perhaps the most frequently used of these: the **existential restriction**.

In the next section, we will examine this constructor in detail, explaining not only its OWL syntax (**some**), but also its underlying Description Logic semantics ($\exists R.C$) and its role in enabling machine-understandable semantic descriptions.

19.7 Existential Restrictions (**some**) -- The First Description Logic Constructor

Having introduced **Description Logic** as the formal language underlying OWL in the previous section, we can now examine the first semantic constructor encountered in the **Pizza** tutorial.

Exercise 15 introduces the **existential restriction**, represented in OWL by the keyword **some**.

Although the syntax appears simple, its engineering significance is substantial. This is the first time the ontology begins expressing **conditions that define the meaning of a concept**, rather than merely identifying the concept within a hierarchy.

Unlike the subclass relationships introduced in Chapter (18), which organized concepts into a conceptual taxonomy, existential restrictions begin describing the characteristics that distinguish one concept from another.

The ontology therefore takes an important step forward:

Concepts are no longer defined solely by where they appear in the hierarchy, but also by the semantic conditions they satisfy.

19.7.1 From Naming Concepts to Description Concepts

Before Exercise 15, the class **MargheritaPizza** possessed only a conceptual identity.

The ontology knew that it represented a specialized type of **NamedPizza**, but it did not yet know what characteristics distinguished a Margherita pizza from any other pizza.

Exercise 15 changes this situation fundamentally.

Instead of merely assigning a label to the class, we begin describing the semantic conditions that characterize its members.

This represents an important conceptual transition in ontology engineering.

Rather than asking:

What is this class called?

we now ask:

What must be true for something to belong to this class?

Ontology engineering therefore shifts its emphasis from:

classification by name

to:

classification by meaning.

The class name remains useful for human communication, but it is the semantic description that enables machines to interpret the concept consistently.

19.7.2 Understanding "At Least One"

One of the most common misunderstanding among beginners concerns the meaning of the keyword **some**.

In everyday natural English, the word "some" often implies an unspecified quantity, perhaps several objects or "a few".

Within OWL, however, the meaning is much more precise.

An existential restriction specifies that **at least one** relationship satisfying the stated condition must exist.

It DOES **NOT** specify exactly how many such relationships exist.

For example, if a pizza is required to have a mozzarella topping, the ontology does not distinguish between **one** mozzarella topping, **two** mozzarella toppings, or **several** mozzarella toppings. Any of these situations satisfies the existential condition.

The restriction therefore establishes a **minimum semantic requirement**, rather than a precise quantity.

The distinction between several commonly used cardinality restrictions is summarized as below:

Restriction	Semantic Meaning
some	At least one related individual satisfies the condition.
exactly 1	Exactly one related individual satisfies the condition.
min 2	At least two related individuals satisfy the condition.
only	Every related individual must satisfy the specified condition.

Later chapters will gradually introduce these additional restriction types.

Exercise 15 intentionally begins with the simplest and most widely used semantic constructor.

19.7.3 How Existential Restrictions Supports Automated Classification

The primary purpose of an existential restriction is not to make the ontology more descriptive for human readers.

Its real value lies in enabling automated semantic classification.

Rather than relying on class names, an OWL reasoner evaluates *whether the semantic conditions associated with a class* is satisfied.

When those conditions hold, the reasoner may infer that an individual belongs to the corresponding class.

Consequently, semantic description become the foundation upon which logical inference is later performed.

This illustrated an important engineering principle:

Reasoner classify concepts according to semantic conditions rather textual labels.

The ontology therefore evolves from a static hierarchy into a semantic model capable of supporting intelligent interpretation.

Although Exercise 15 introduces only two simple restrictions, it establishes the modeling pattern that will be used repeatedly throughout the remainder of the **Pizza** ontology.

19.7.4 Semantic Descriptions Grow Incrementally

Professional ontologies rarely define every characteristic of a concept at once.

Instead, semantic descriptions typically evolve through a sequence of incremental refinements.

Exercise 15 demonstrates this engineering philosophy clearly.

At this stage, **MargheritaPizza** is characterized by only a small number of semantic conditions. Numerous additional characteristics could potentially be added later, including

- the type of pizza base,
- vegetarian status,
- country of origin,
- preparation style,
- nutritional properties, or
- commercial classifications.

Rather than attempting to model every possible aspect immediately, the **Pizza** tutorial introduces only the semantic knowledge required to support the current learning objective.

This incremental approach offers several important advantages:

1. It keeps semantic models understandable while they are still evolving.
2. It minimizes unnecessary restructuring as domain knowledge matures.
3. It allows reasoning results to be validated continuously throughout development.
4. It supports gradual enrichment without sacrificing conceptual stability.

This development strategy aligns closely with the philosophy of the **Semantic Knowledge Development Lifecycle (SKDL)**, where semantic knowledge is refined progressively rather than constructed in one single step.

19.7.5 Common Misunderstandings About Existential Restrictions

Because existential restrictions are often introduced early in OWL tutorials, several misconceptions frequently arise.

The table below summarized some of the most common misunderstandings.

Misunderstanding	Correct Interpretation
" some means several."	It means at least one .
" some specifies an exact quantity."	It specifies only a minimum existence requirement.
"Restrictions create new individuals."	Restrictions describe classes, not individual objects.
"Class names determine semantic meaning."	Semantic meaning is determined by logical description rather than labels.
"Existential restrictions perform reasoning."	They provide the semantic condition upon which reasoning operates.

Understanding these distinctions early prevents many modeling mistakes that become increasingly difficult to correct as an ontology grows.

19.7.6 Engineering Takeaway

Existential restrictions represent the first point at which an ontology begins expressing **formal semantic knowledge** rather than simple conceptual organization.

Although adding a restriction in Protégé requires only a few mouse clicks, the underlying engineering significance is much greater.

For the first time, the ontology begins defining concepts according to **their semantic characteristics** instead of relying solely on their position within a hierarchy.

This transition -- from organizing concepts to describing their meaning -- marks one of the most important milestones in ontology engineering.

It establishes the semantic foundation upon which

- richer class expressions,
- logical reasoning,
- semantic validation,
- knowledge reuse, and
- ultimately executable knowledge

will be progressively developed throughout the remaining stages of the **Semantic Knowledge Development Lifecycle (SKDL)**.

19.8 Class Expressions -- Building Semantic Meaning with Description Logic

In the previous sections, we examined the first Description Logic constructor introduced in Exercise 15: the **existential restriction (some)**.

Although existential restrictions allow us to describe individual semantic characteristics of a concept, real-world concepts are rarely defined by a single condition.

A pizza is not characterized solely by possessing mozzarella.

Likewise, an enterprise application is not defined only by the data it manages, and a biological organism is not identified by a single characteristic.

Most concepts are distinguished by a **combination of semantic properties** that together define their meaning.

Ontology engineering therefore requires a language capable of combining multiple semantic statements into coherent logical definitions.

In OWL, this mechanism is provided through **class expressions**.

Rather than describing concepts using isolated restrictions, class expressions allow ontology engineers to compose multiple logical conditions into increasingly expressive semantic descriptions that machines can interpret consistently.

This represents another important milestone in the Semantic Knowledge Development Lifecycle (SKDL). During **Stage 1 -- Conceptual Modeling**, we organized concepts into taxonomies. During **Stage 2 -- Semantic Description**, we begin expressing what those concepts actually mean by constructing logical descriptions that build upon the semantic vocabulary established in Stage 1.

19.8.1 From Individual Restrictions to Complete Semantic Descriptions

Exercise 15 introduced two independent semantic restrictions from **MargheritaPizza**:

```
hasTopping some MozzarellaTopping
```

and

```
hasTopping some TomatoTopping
```


Each restriction contributes one semantic characteristic.

Viewed independently, each statement answers only one single question:

- Does the pizza contain mozzarella?
- Does the pizza contain tomato?

Neither restriction alone adequately characterized a Margherita pizza.

Only when both restrictions are considered **together** does a more meaningful semantic description begin to emerge.



This illustrates an important principle of ontology engineering:

Meaning is rarely expressed through isolated logical statements. It emerges from the combination of multiple semantic constraints.

Description Logic therefore allows multiple restrictions to be composed into larger logical expressions.

19.8.2 What Is a Class Expression?

A **class expression** is a logical description that specifies the conditions required for membership within a class.

Unlike a named class such as `MargheritaPizza`, which simply provides a reusable identifier, a class expression explicitly describes the semantic characteristics that define a concept.

For example, in Exercise 15:

```
hasTopping some MozzarellaTopping
```

is itself a simple class expression.

It does not introduce a new named concept. Instead, it describes the collection of all pizzas that possess at least one mozzarella topping.

As additional logical conditions are combined, increasingly sophisticated semantic descriptions can be constructed without introducing additional class names.

Consequently, class expressions should be viewed as **descriptions of meaning**, rather than merely labels within a hierarchy.

19.8.3 Named Classes and Anonymous Classes

One of the most important concepts in Description Logic is the distinction between **named classes** and **anonymous classes**.

A **named class** possesses a unique identifier (IRI) within the ontology.

Examples include:

- `Pizza`
- `NamedPizza`
- `MargheritaPizza`
- `PizzaTopping`

These classes become part of the ontology's conceptual vocabulary and appear explicitly within Protégé's class hierarchy.

By contrast, a **class expression** usually represents an **anonymous class**.

For example:

```
hasTopping some MozzarellaTopping
```

does not create a new class named `MozzarellaPizza`.


Instead, it describes an unnamed collection of pizzas satisfying that restriction.

Similarly,

```
NamedPizza
and
(hasTopping some MozzarellaTopping)
```

also represents an anonymous class.

Although unnamed, anonymous classes possess complete semantic meaning and may participate in logical reasoning exactly like named classes.



The distinction is summarized in below table:

Named Classes	Anonymous Class
Possesses an IRI	Has no explicit name
Appears within the class hierarchy	Exists only as a logical expression

Named Classes	Anonymous Class
Represents reusable domain concepts	Represents semantic descriptions
Typically created by ontology engineers	Often generated or composed dynamically during modeling

Understanding this distinction helps explain why Protégé frequently displays complex logical expressions that do not correspond to explicit class names.

19.8.4 Description Logic as a Language of Composition

Unlike traditional taxonomies, Description Logic is not limited to simple subclass relationships.

Instead, it provides a small collection of logical constructors that can be combined recursively to describe increasingly complex concepts.

Some of the most frequently used constructors include:

Description Logic Constructor	Manchester OWL Syntax	Purpose
Conjunction	and	All conditions must be satisfied simultaneously.
Disjunction	or	At least one condition must be satisfied.
Negation	not	Excludes concepts satisfying a particular condition.
Existential Restriction	some	Requires at least one relationship satisfying a condition.
Universal Restriction	only	Restricts all related individuals to satisfy a condition.

These logical constructors behave much like grammatical rules in a natural language.

Just as words combine into phrases and sentences, logical constructors combine into class expressions that describe increasingly rich semantic meaning.

The expressive power of OWL therefore arises not from the number of available keywords, but from the ability to compose simple semantic statements into sophisticated knowledge models.

19.8.5 Semantic Composition in the **Pizza** Ontology

Let's return to Exercise 15, the semantic description of **MargheritaPizza** can now be viewed as the composition of several independent semantic statements.

Conceptually, we begin with the existing abstraction:

```
NamedPizza
```

We then enrich that concept through additional semantic restrictions:

```
hasTopping some MozzarellaTopping
and
hasTopping some TomatoTopping
```

Rather than:

representing unrelated statements,

these restrictions:

collectively describe one coherent concept.

As additional semantic characteristics are introduced throughout later chapters, increasingly expressive class expressions will emerge by extending this same compositional principle.

Complex semantic definitions therefore evolve incrementally from smaller, understandable building blocks.

19.8.6 Necessary Conditions Versus Complete Definitions

At this stage of the **Pizza** tutorial, the restrictions introduced for **MargheritaPizza** are attached through **SubClassOf** relationships.

This distinction is extremely important.

A subclass restriction specifies a **necessary condition**.

It states that every instance of **MargheritaPizza** must satisfy the specified restriction.

However, it does **not** imply that satisfying the restriction alone is sufficient to conclude that an arbitrary pizza is a Margherita pizza.

Later chapters will introduce **EquivalentTo**, which represents a much stronger semantic commitment, and also can be configured in the same screen as **SubClassOf** in Protégé:

The screenshot displays the Protégé interface with the 'Classes' tab selected. On the left, the 'Class hierarchy: MargheritaPizza' tree shows a hierarchy starting from 'owl:Thing', followed by 'NoClassDefined', 'Pizza', 'CheesePizza', 'NamedPizza', and then 'MargheritaPizza' (highlighted with a blue box). Other classes like 'AmericanaPizza', 'SohoPizza', 'NonVegetarianPizza', 'VegetarianPizza', 'PizzaBase', and 'PizzaTopping' are also visible. On the right, the 'Annotations: MargheritaPizza' panel shows an 'rdfs:comment' stating 'A pizza that only has Mozzarella and Tomato toppings'. Below this, the 'Description: MargheritaPizza' panel shows an 'Equivalent To' restriction (highlighted with a red box) and a 'SubClass Of' list containing 'hasTopping some MozzarellaTopping', 'hasTopping some TomatoTopping', and 'NamedPizza'. The 'General class axioms' panel is also visible at the bottom.

The engineering distinction can be summarized as follows:

SubClassOf	EquivalentTo
Specifies necessary conditions	Specifies necessary and sufficient conditions
Supports incremental semantic enrichment	Provides a complete logical definition
Appropriate during early ontology development	Appropriate after concepts have stabilized

This design philosophy directly reflects the principles introduced earlier in this chapter.

Professional ontology engineers typically begin by **accumulating necessary semantic descriptions** before attempting to define **complete semantic equivalence**.

Doing so allows the ontology to evolve naturally while reducing the cost of future redesign.

19.8.7 Open World Assumption Enables Incremental Semantic Modeling

Throughout the earlier chapters of this book, we have encountered the **Open World Assumption (OWA)** on several occasions. Rather than repeating its formal definition, it is now useful to examine its engineering significance from the perspective of **semantic description**.

Exercise 15 introduced the first semantic restrictions describing **MargheritaPizza**:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

These restrictions define semantic characteristics that a Margherita pizza must possess.

Importantly, however, they **do not attempt to describe everything about the pizza**.

The ontology does not claim that mozzarella and tomato are the *only* toppings, nor does it assert that no additional ingredients may exist.

Instead, each restriction contributes one additional piece of semantic knowledge.

*This is proved by the **real world knowledge**! Thanks for my wife who is professional in cooking and making pizza, with far beyond my skills, she looks at my ontology and immediately points out the missing knowledge - Basil leaf, as below image, those are the at least three key toppings / ingredients for a pizza named as Margherita pizza, while our ontology only describes two:*



This illustrates an important engineering consequence of the Open World Assumption:

Semantic descriptions are intended to accumulate rather than terminate knowledge.

Unlike traditional database schemas, where developers often attempt to specify complete records from the beginning, ontology engineering encourages semantic models to evolve incrementally.

Concepts gradually become richer as additional knowledge is discovered, validated, and incorporated into the ontology.

This incremental philosophy aligns naturally with the Semantic Knowledge Development Lifecycle (SKDL):

- **Stage 1** establishes the conceptual vocabulary.
- **Stage 2** begins enriching concepts with semantic descriptions.
- Later stages introduce reusable patterns, governance rules, validation constraints, reasoning axioms, and executable behaviors.

Each stage adds semantic knowledge without requiring earlier descriptions to be discarded.

Consequently, ontology engineers should regard OWL restrictions as **incremental semantic contributions**, not as exhaustive specifications.

This perspective also explains why the previous section emphasized the use of `SubClassOf` rather than `EquivalentTo`.

A subclass restriction simply contributes another necessary semantic characteristic.

As understanding of the domain matures, additional restrictions can be introduced naturally without redefining the concept from scratch, for example, adding `Basil` into restrictions to `MargheritaPizza`.

Only when the ontology reaches a sufficient level of semantic completeness does it become appropriate to define concepts using `EquivalentTo`, thereby providing both necessary and sufficient conditions for automated classification.

From this perspective, the Open World Assumption is more than a logical principle.

It is an **engineering strategy for long-lived knowledge systems**.

Enterprise ontologies, biomedical terminologies, digital libraries, and large-scale knowledge graphs continually evolve as new information (knowledge) becomes available.

Because OWL assumes that knowledge is always potentially incomplete, these systems can expand organically while preserving semantic consistency.

This is one of the reasons why ontology engineering scales successfully to domains whose knowledge can never be considered permanently complete.

19.8.8 Engineering Perspective -- Building Meaning Incrementally

Throughout this chapter, we have introduced class expressions not merely as OWL syntax, but as a systematic way of enriching the semantic meaning of concepts.

An important engineering principle emerges from this process:

Semantic knowledge should be developed incrementally rather than all at once.

Beginners often assume that every class should immediately receive a complete logical definition. In practice, however, professional ontology engineering follows a much more disciplined approach.

Concepts are first identified and organized into a stable **conceptual taxonomy** (Stage 1 of the Semantic Knowledge Development Lifecycle). Only after that conceptual foundation has matured do ontology engineers begin introducing **semantic descriptions** through carefully designed class expressions

The evolution of the **Pizza** ontology provides an excellent illustration of this engineering methodology.

In **Exercise 14**, the class **MargheritaPizza** was introduced solely as part of the conceptual hierarchy. At that stage, the ontology answered only one question:

Where does this concept belong within the domain?

No semantic characteristics were yet associated with the class.

Exercise 15 answered a different question:

What semantic characteristics begin to distinguish this concept from other pizzas?

Rather than attempting to produce a complete logical definition immediately, the ontology introduces only two existential restrictions:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

These restrictions enrich the semantic meaning of **MargheritaPizza** without claiming that the description is complete. Additional semantic constraints (e.g. Basil) may be introduced later as the ontology evolves and the domain becomes better understood.

This gradual refinement offers several important engineering advantages:

- It separates conceptual organization from semantic specification.
- It minimizes unnecessary restructuring as domain knowledge evolves.
- It allows semantic definitions to mature alongside the ontology.
- It reduces the likelihood of introducing inconsistent or overly restrictive axioms during the early stages of development.
- It encourages ontology engineers to validate semantic assumptions incrementally rather than attempting to model everything at once.

This philosophy closely resembles engineering practices found in many other disciplines.

Software architects establish class structures before implementing complex behaviors (methods).

Database designers identify entities before defining integrity constraints.

Enterprise architects develop capability maps before modeling operational interactions.

Ontology engineers follow exactly the same progression.

Conceptual organization provides the structural framework, while **semantic description** progressively enriches that framework with machine-understandable meaning.

Consequently, class expressions should not be viewed as isolated pieces of OWL syntax. They are modular semantic building blocks that allow an ontology to evolve steadily from conceptual organization into a rich and formally defined knowledge model.

As the ontology grows, these individual semantic building blocks combine to form increasingly expressive descriptions that enable automated reasoning, semantic validation, knowledge reuse, semantic governance, and ultimately executable intelligence.

19.8.9 Preparing for Future Chapters

Exercise 15 introduces only the first examples of semantic class expressions, yet the modeling principles established here will underpin every subsequent stage of ontology development.

In the chapters that follow, the **Pizza** ontology will gradually acquire richer semantic descriptions through additional Description Logic constructors and increasingly expressive class expressions.

Future modeling activities will introduce concepts such as universal restrictions (**only**), cardinality constraints (**min**, **max**, and **exactly**), equivalent class definitions (**EquivalentTo**), and reusable semantic modeling patterns. Each new construct will build upon the semantic foundation established in this chapter.

As these semantic descriptions become progressively richer, the ontology will evolve from a simple conceptual hierarchy into a formal knowledge model capable of supporting automated inference.

Eventually, the reasoner will no longer rely solely on manually asserted subclass relationships. Instead, it will begin deriving new knowledge automatically by evaluating the logical meaning expressed through class definitions.

This progression mirrors the overall structure of the **Semantic Knowledge Development Lifecycle (SKDL)**.

- **Stage 1 -- Conceptual Modeling** established the conceptual vocabulary and semantic taxonomy of the domain.
- **Stage 2 -- Semantic Description** enriches those concepts with formal meaning through Description Logic and OWL class expressions.
- **Stage 3 -- Knowledge Reuse** will demonstrate how well-designed semantic patterns can be reused consistently across multiple concepts and ontologies.
- Later stages will progressively introduce semantic governance, validation, automated reasoning, and ultimately executable knowledge within the Executable Knowledge Architecture (EKA).

This incremental progression is intentional.

Rather than introducing every OWL construct simultaneously, the tutorial follows a disciplined engineering methodology in which each stage builds upon the stable foundations established by the previous one.

In doing so, the ontology grows

not merely in size,

but

in semantic richness, logical precision, and engineering maturity.

By mastering class expressions in this chapter, you are learning far more than another feature of Protégé.

You are learning the **formal language** through which ontologies communicate meaning to machines -- a language whose theoretical foundations span more than two thousand years of logical thought and whose practical applications now power modern knowledge graphs, semantic web technologies, digital twins, enterprise ontologies, and intelligent AI systems.

The next chapter continues this journey by demonstrating how semantic definitions, once established, can be reused systematically to reduce duplication, improve consistency, and accelerate ontology development -- bringing us to **Stage 3 -- Knowledge Reuse** of the Semantic Knowledge Development Lifecycle (SKDL).

19.9 Mathematical Perspective -- Semantic Description as Formal Knowledge Representation

One of the defining characteristics of ontology engineering is that semantic descriptions are **not merely textual documentation**.

Unlike comments written for human readers, OWL class expressions possess **precise mathematical meaning**.

Every restriction, every logical constructor, and every class expression introduced throughout this chapter can be interpreted formally using the language of **set theory** and **first-order logic (FOL)**.

This mathematical foundation is precisely what enables ontology reasoners to

- verify consistency,
- infer new knowledge,
- classify concepts automatically, and
- guarantee logically sound conditions.

For ontology engineers, understanding this mathematical perspective is not required to use Protégé effectively, but it greatly deepens one's appreciation of **why Description Logic works**.

19.9.1 From Conceptual Sets to Semantic Sets

In **Chapter (18)**, we viewed ontology classes primarily as conceptual categories organized into taxonomies.

Mathematically, each class may also be interpreted as a **set of individuals**.

For example, **Pizza** represents the **set** containing every pizza.

Similarly, **NamedPizza** represents the **subset** containing all pizzas possessing recognized names.

This interpretation naturally extends the mathematical view introduced during **Conceptual Modeling**.

The subclass relationship:

```
NamedPizza SubClassOf Pizza
```

corresponds directly to the subset relation:

$$NamedPizza \subseteq Pizza$$

meaning:

every **NamedPizza** is also a **Pizza**.

Stage 2 - Semantic Description - extends this idea considerably.

Instead of defining sets only through hierarchy, ontology engineers now define sets through **logical conditions**.

Membership within a class is determined not merely by its position in the taxonomy, but by whether an individual satisfies the logical expression associated with that class.

Concepts therefore evolve from:

Named Sets

into:

Logically Defined Sets.

19.9.2 Object Properties as Binary Relations

Classes alone cannot describe semantic meaning.

Meaning emerges through relationships.

Mathematically, every OWL object property represents a **binary relation** between two sets.

For example,

```
hasTopping
```

is interpreted as:

$$hasTopping \subseteq Pizza \times PizzaTopping$$

where

$$Pizza \times PizzaTopping$$

denotes the Cartesian product of the two sets.

Every pair (p, t) contained within this relation states that pizza p possesses topping t .

Unlike ordinary database foreign keys, object properties are treated as **logical relations** rather than **implementation artifacts**.

This distinction allows Description Logic to reason about relationships independently of any physical storage mechanism.

Consequently, semantic descriptions become mathematical statements about relationships between sets rather than records inside database tables.

19.9.3 Existential Restrictions as Set constructors

Exercise 15 introduced the existential restriction:

```
hasTopping some MozzarellaTopping
```

From a mathematical perspective, this expression defines a completely new set.

Using Description Logic notation:

$$\exists hasTopping.MozzarellaTopping$$

is interpreted as

$$\{x \mid \exists y ((x, y) \in hasTopping \wedge y \in MozzarellaTopping)\}$$

This notation can be read as:

the set of all pizzas **x** for which there exists at least one topping **y** such that **x** has topping **y** and **y** belongs to the class **MozzarellaTopping**.

Notice something remarkable!

No new class name was required.

Instead, the class expression itself **constructs a mathematically defined set**.

Description Logic therefore acts as a language for building new semantic concepts directly from existing ones.

19.9.4 Building Meaning Through Set Operations

Once every class expression represents a set, logical constructors become familiar mathematical operations.

For example,

Intersection: $C \sqcap D$

corresponds directly to: $C \cap D$

meaning:

individuals belonging to both sets.

Union: $C \sqcup D$

corresponds to: $C \cup D$

meaning:

individuals belonging to either set.

Negation: $\neg C$

corresponds to the complement: $\overline{C}$

meaning:

every individual not belonging to the set.

Consequently,

```
NamedPizza
and
(hasTopping some MozzarellaTopping)
and
(hasTopping some TomatoTopping)
```

is simply the intersection of three mathematically defined sets.

Each additional restriction progressively narrows the collection of individuals satisfying the class expression.

Semantic enrichment therefore corresponds mathematically to **successively refining a set through logical intersection**.

19.9.5 Interpretation Functions Give Meaning to Syntax

Description Logic carefully distinguishes between **syntax** and **semantics**.

- The OWL expressions entered into Protégé constitute the **syntax**.
- Their mathematical meaning is supplied through an **interpretation function**.

Formally, an interpretation:

$$I = (\Delta^{\mathcal{I}}, \cdot^{\mathcal{I}})$$

consists of

- a domain of individuals: $\Delta^{\mathcal{I}}$, and
- an interpretation function: $\cdot^{\mathcal{I}}$

that maps every class, property, and individual to its mathematical meaning.

For example, **Pizza** is mapped to:

$$Pizza^{\mathcal{I}} \subseteq \Delta^{\mathcal{I}}$$

while, **hasTopping** is mapped to:

$$hasTopping^{\mathcal{I}} \subseteq \Delta^{\mathcal{I}} \times \Delta^{\mathcal{I}}$$

This semantic interpretation ensures that every OWL ontology possesses a precise formal meaning independent of software implementation.

- Protégé stores the syntax.
- Description Logic defines its semantics.

Reasoners operate over that mathematical semantics rather than the textual representation.

19.9.6 Consistency, Satisfiability, and Logical Soundness

The mathematical semantics of Description Logic enable reasoners to answer questions that would otherwise be impossible.

Among the most important formal relationships is **subsumption**, which corresponds directly to the subclass relationship in OWL.

Subsumption: ($\sqsubseteq$) is formally defined as:

$$C \sqsubseteq D \quad \text{if and only if} \quad C^{\mathcal{I}} \subseteq D^{\mathcal{I}} \text{ for every interpretation } \mathcal{I}$$

In other words, class C is subsumed by class D precisely when the set of individuals belonging to C is a subset of the set of individuals belonging to D under every possible valid interpretation of the ontology.

This mathematical definition explains why an OWL reasoner can determine subclass relationships automatically: it verifies that the semantic conditions defining C logically imply the semantic conditions defining D across all interpretations.

From this foundation, reasoners answer four fundamental questions that would otherwise be impossible to resolve through inspection alone:

- **Consistency** -- Does the ontology contain any logical contradiction?
- **Satisfiability** -- Could a class possibly contain at least one instance?
- **Subsumption** -- Is one class logically more general than another?
- **Classification** -- Where should a class appear within the hierarchy?

These questions are answered not through heuristic algorithms but through **formal logical reasoning** based on the mathematical semantics defined for every class expression.

This explains why ontology engineering differs fundamentally from ordinary data modeling.

The ontology is not simply storing information.

It is encoding **logical knowledge** whose properties can be analyzed mathematically.

19.9.7 Decidability Makes Automated Reasoning Possible

Earlier in this chapter, we introduced **Description Logic** as a carefully designed subset of first-order logic.

One of the principal reasons for this design is **decidability**.

A logical language is said to be decidable if every valid reasoning task is guaranteed to terminate after a finite number of computational steps.

This property has profound engineering significance.

Without decidability:

- ontology consistency might never be determined,
- automated classification could continue indefinitely,
- reasoning services could become computationally impractical for real-world applications.

Description Logic deliberately balances **expressive power** against **computational tractability**.

It is expressive enough to model complex domains while remaining sufficiently restricted to guarantee reliable reasoning.

This balance explains why OWL adopts Description Logic rather than unrestricted first-order logic as its formal foundation.

19.9.8 From Mathematics to Engineering

Although ontology engineers rarely write mathematical formulas during everyday modeling, every OWL expression entered into Protégé implicitly carries the formal semantics described throughout this section.

The two restrictions added during **Exercise 15** may appear deceptively simple:

```
hasTopping some MozzarellaTopping  
hasTopping some TomatoTopping
```

Behind these concise expressions lies a rigorous mathematical framework based upon:

- set theory,
- binary relations,
- logical quantifiers,
- interpretation theory, and
- decidable Description Logic.

This mathematical foundation explains why semantic descriptions can be interpreted consistently by machines, shared across organizations, validated automatically, and processed reliably by reasoning engines.

Ultimately, the purpose of ontology engineering is not simply to organize information.

It is to construct **formal knowledge models whose semantics are mathematically precise, computationally interpretable, and logically verifiable.**

This is the foundation upon which semantic reasoning, knowledge reuse, executable knowledge and the remaining stages of the **Semantic Knowledge Development Lifecycle (SKDL)** are progressively built.

19.10 Interesting Reading -- From Aristotle to Description Logic

Throughout this chapter, we have examined **semantic description** from practical, engineering, and mathematical perspectives. We introduced **Description Logic (DL)** as the formal language underlying OWL, explored existential restrictions (**some**), and showed how class expressions describe the meaning of concepts through logical composition.

At first glance, these constructs may appear to be **modern inventions** created specifically for ontology engineering.

In reality, they represent the culmination of a remarkable **long intellectual tradition**.

The logical expressions used in OWL today did not emerge suddenly with the Semantic Web. They evolved through more than two millennia of **philosophical inquiry, mathematical logic, and knowledge representation** research.

Understanding this historical progression provides us an important perspective:

Ontology engineering is not merely learning Protégé syntax -- it is participating in one of humanity's oldest efforts to organize and reason about knowledge.

19.10.1 Aristotle -- the Birth of Conceptual Definition

More than two thousand years before computers existed, **Aristotle (384-322 BCE)** established many of the fundamental intellectual principles that continue to influence ontology engineering today.

Although Aristotle never imagined knowledge graphs, Description Logic, or the Semantic Web, his approach to **classification** and **definition** introduced ideas that remain recognizable in modern ontology languages.

Two ideas proved especially influential.

The **first** was the distinction between **genus** and **differentia**.

A concept should first be placed within a broader category (its *genus*) and then distinguished from related concepts through characteristics that make it unique (its *differentia*).

For example, Aristotle would not define a human being simply by assigning it a name. Instead, he would describe it as:

A human is a rational animal.

Here:

- **Animal** represents the **genus**, placing human within a broader conceptual category.
- **Rational** provides the **differentia**, distinguishing humans from other animals.

This approach closely resembles modern ontology engineering.

When we modeled:

```
MargheritaPizza  
  SubClassOf NamedPizza
```

in Chapter (18), we established the **genus** of the concept.

When Exercise 15 added:

```
hasTopping some MozzarellaTopping  
hasTopping some TomatoTopping
```

we began introducing its **differentia** -- the semantic characteristics that distinguish a Margherita pizza from other named pizzas.

The **second** influential idea was that **knowledge should be expressed through explicit definitions rather than isolated labels**.

For Aristotle, assigning a name to a concept did not explain its meaning. Genuine understanding required identifying the essential characteristics that made the concept what it was.

This philosophical principle is precisely what distinguishes **Conceptual Modeling** from **Semantic Description** in the Semantic Knowledge Development Lifecycle (SKDL).

A class name such as **MargheritaPizza** provides only an identifier. It is the accompanying **semantic restrictions** and **class expressions** that communicate its **meaning** to both humans and machines.

In this sense, Description Logic can be viewed as a modern mathematical language for expressing Aristotle's centuries-old idea that *concepts should be defined by their essential characteristics rather than merely named*.

19.10.2 Porphyry's Tree -- Visualizing Hierarchical Knowledge

Several centuries after Aristotle, the philosopher **Porphyry (c. 234-305 CE)** produced one of the earliest graphical representations of conceptual organization, now known as the **Porphyrian Tree**.

Rather than introducing a new logical theory, Porphyry provided a systematic method for visualizing Aristotle's principle of **genus** and **differentia** through successive specialization.

A simplified representation is shown below:


```
. /  
└ Substance  
  └ Body  
    └ Living Body  
      └ Animal  
        └ Rational Animal  
          └ Human
```

Each level introduces an additional distinguishing characteristic while preserving the properties inherited from its parent.

This idea closely resembles the conceptual hierarchies developed in Chapter (18).

For example:

```
. /  
└ Pizza  
  └ NamedPizza  
    └ MargheritaPizza
```

Each subclass inherits the semantic meaning of its parent while introducing additional distinguishing characteristics.

Although Protégé displays the ontology as a tree, the tree is not merely a visualization.

It represents a semantic hierarchy in which inheritance preserves knowledge across successive levels of abstraction.

Porphyry therefore transformed Aristotle's philosophical principles into one of history's earliest systematic knowledge organization models.

The influence of the Porphyrian Tree can still be recognized in modern taxonomies, biological classification, library classification systems, enterprise architecture meta-models, and ontology engineering.

For ontology engineers, the lesson remains timeless:

A hierarchy is not simply a drawing. It is a semantic organization of knowledge.

19.10.3 Leibniz -- The Dream of a Universal Knowledge Language

During the seventeenth century, **Gottfried Wilhelm Leibniz (1646-1716)** envisioned a remarkable idea that anticipated many goals of modern knowledge representation.

Leibniz imagined a symbolic language capable of expressing all human knowledge with mathematical precision.

He called this vision the **Characteristica Universalis** -- a universal language in which concepts could be represented unambiguously and combined according to logical rules.

He also proposed the **Calculus Ratiocinator** -- a symbolic reasoning system through which disagreements could be resolved by calculation rather than debate.

Leibniz famously imagined two scholars ending an argument by saying:

"Let us calculate."

Although the mathematics and computing technology necessary to realize this vision did not yet exist, the philosophical objective was remarkably similar to that of modern ontology engineering.

Today, OWL ontologies assign every concept a global unique **IRI**, express semantic meaning through formal logic, and allow reasoners to derive conclusions automatically.

The vision is no longer philosophical speculation.

It has become practical engineering.

Modern ontology languages therefore realize, at least in part, Leibniz's centuries-old ambition to **transform knowledge into a form that machines can process objectively and consistently**.

19.10.4 Frege -- Modern Predicate Logic

The next major milestone occurred during the nineteenth century through the work of **Gottlob Frege (1848-1925)**.

Where Aristotle's logic primarily analyzed categorical statements, Frege introduced **predicate logic**, providing a far more expressive mathematical language for representing knowledge.

Most importantly, predicate logic introduced **quantifiers**, allowing logical statements to describe entire collections of objects rather than individual examples.

Two quantifiers became fundamental:

Quantifier	Meaning	Manchester OWL Equivalent
$\forall$	For all	only
$\exists$	There exists	some

The existential restriction introduced in Exercise 15 therefore has deep historical roots.

When we write:

```
hasTopping some MozzarellaTopping
```

we are expressing the Description Logic statement:

$\exists$ hasTopping.MozzarellaTopping

whose logical origin can be traced directly to Frege's formal treatment of existential quantification.

Similarly, later chapters will introduce universal restrictions such as:

```
hasTopping only CheeseTopping
```

which correspond to universal quantification.

Frege therefore supplied the mathematical language that eventually enabled logic to become computational rather than merely philosophical.

19.10.5 Description Logic -- Making Logic Computable

Although predicate logic is extremely expressive, unrestricted first-order is computationally difficult.

Many reasoning tasks become undecidable, meaning that no algorithm can guarantee termination for every possible knowledge base.

During the late twentieth century, researchers in **knowledge representation** sought a practical compromise.

The result was **Description Logic (DL)**.

Rather than attempting to represent every possible logical statement, Description Logic deliberately limits the available language to carefully designed constructors that preserve **decidability**.

This engineering decision allows automated reasoners to perform important tasks such as:

- satisfiability testing;
- consistency checking;
- subsumption computation;
- equivalence detection; and
- automatic classification.

In other words, Description Logic sacrifices unrestricted expressive power in exchange for reliable computation.

This balance explains why OWL is based upon Description Logic rather than unrestricted predicate logic.

Ontology engineering therefore emphasizes not only logical correctness, but also computational feasibility.

A semantic language is valuable only if machines can reason with it effectively.


19.10.6 OWL -- The Modern Realization of a Two-Thousand-Year Journey

The publication of **OWL (Web Ontology Language)** by the World Wide Web Consortium (W3C) represents the modern culmination of this intellectual journey.

OWL did not invent conceptual classification, logical definition, or formal reasoning.

Instead, it integrated ideas developed over more than two thousand years into a practical engineering language for **machine-understandable language**.

The progression can be summarized as follows:



From this perspective, the OWL restriction introduced in Exercise 15 is much more than a piece of Protégé syntax.

The expression:

```
hasTopping some MozzarellaTopping
```

embodies ideas that began with Aristotle's distinction between **genus** and **differentia**, evolved through centuries of philosophical and mathematical development, and finally became a formally computable representation language capable of supporting automated reasoning.

The ontology engineer working in Protégé therefore participates in a tradition that spans philosophy, mathematics, computer science, and knowledge engineering.

The tools have changed dramatically over the centuries.

The fundamental question, however, has remained remarkably constant:

How can knowledge be represented so that its meaning is both precise enough for machines to compute and understandable enough for humans to share?

19.10.7 Reflection -- Why Ontology Engineering Is Different

The historical journey presented in this section reveals an important truth about ontology engineering.

Unlike many areas of software development, ontology engineering is not built solely upon recent advances in computer science.

It draws simultaneously upon **philosophy**, **mathematical**, **logic**, **knowledge organization**, and **software engineering**.

Each discipline contributes a different perspective:

Discipline	asks
Philosophy	<i>What does a concept mean?</i>
Logic	<i>How can that meaning be expressed precisely?</i>
Mathematics	<i>Can the representation be defined formally and unambiguously?</i>
Computer Science	<i>Can machines compute with that representation efficiently?</i>
Knowledge Engineering	<i>Can this knowledge remain understandable, reusable, and maintainable as it evolves?</i>

Ontology engineering brings these perspectives together into a **single engineering discipline**.

This multi-disciplinary nature explains why developing an ontology differs fundamentally from designing a database schema or implementing an object-oriented software model.

- A database primarily stores **facts**.
- An object-oriented model primarily organizes **software structures and behavior**.
- An ontology, however, attempts to capture the **meaning** of a domain in a form that both humans and machines can understand.

For this reason, ontology engineers must think at multiple levels simultaneously.

They must identify appropriate concepts, organize them into coherent taxonomies, define their semantic characteristics using formal logic, and ensure that these definitions remain computationally tractable for automated reasoning.

The work performed in **Exercise 15** illustrates this broader perspective.

Adding the restriction:

```
hasTopping some MozzarellaTopping
```

may appear to be a simple editing operation within Protégé.

In reality, this single statement embodies several complementary ideas developed throughout the history of knowledge representation.

From the perspective of **Aristotle**, it contributes part of the *differentia* that distinguishes one concept from another.

From the perspective of **Description Logic**, it represents an existential restriction of the form $\exists R.C$, expressed using a mathematically defined constructor.

From the perspective of **OWL**, it becomes a machine-interpretable semantic axiom that automated reasoners can evaluate consistently.

From the perspective of the **Semantic Knowledge Development Lifecycle (SKDL)**, it marks the transition from identifying concepts to formally describing their meaning.

Finally, from the perspective of the **Executable Knowledge Architecture (EKA)**, it enriches the *K* - **Knowledge Graph layer** with semantics that will enables reasoning (*R* - Reasoning and Rules layer), governance (*Γ* - Governance layer), validation, and ultimately executable intelligence.

These perspectives are not separate interpretations of the same statement.

They are complementary layers of understanding that coexist within every well-designed ontology.

Professional ontology engineers therefore develop more than technical proficiency with Protégé or OWL syntax.

They cultivate the ability to think simultaneously as philosophers, mathematicians, software engineers, and knowledge architects.

This broader perspective explains why ontology engineering occupies a unique position within modern information systems.

It is the discipline that transforms isolated data into structures knowledge, structured knowledge into formal semantics, and formal semantics into machine-executable intelligence.

The remaining chapters of this book continue that transformation.

Having established the conceptual vocabulary in **Stage 1** and introduced formal semantic descriptions in **Stage 2**, we are now ready to explore how semantic knowledge can be organized, reused, governed,

validated, and eventually reasoned over in a systematic and scalable manner.

Viewed from this perspective, the OWL constructs introduced in this chapter are not isolated language features.

They are the first practical steps in a much larger journey -- one that began over two thousand years ago with philosophical questions about the nature of concepts and continues today through the engineering of executable knowledge systems.

19.11 Engineering Perspective -- Meaning Before Reasoning

Throughout this chapter, we have progressively enriched the **Pizza** ontology with its first formal semantic descriptions. Using Description Logic (DL) class expressions, we transformed **MargheritaPizza** from a simple concept within a taxonomy into a concept possessing machine-understandable semantic characteristics.

Although these semantic descriptions may appear modest, they establish one of the most important engineering principles in ontology development:

Reasoning cannot begin until meaning has first been formally described.

This principle extends far beyond the **Pizza** tutorial.

Whether constructing enterprise knowledge graphs, biomedical ontologies, digital twins, or intelligent AI systems, automated reasoning always depends upon a previously established semantic model.

A reasoner does not invent knowledge!

Instead, it derives new conclusions by applying formal logic to knowledge that has already been explicitly described.

Consequently, semantic description should be viewed not as reasoning itself, but as the engineering activity that prepares an ontology for future reasoning.

19.11.1 Semantics Description Provides the Raw Material for Reasoning

Many beginners imagine that ontology reasoners behave similarly to modern Large Language Models (LLMs), capable of generating knowledge from incomplete information.

This is not how Description Logic reasoners operate.

A DL reasoner performs logical inference over **explicitly defined semantics**.

If no semantic meaning has been modeled, the reasoner has nothing upon which to operate.

For example, immediately after **Exercise 14**, the ontology contained the following conceptual hierarchy:

```
Pizza
  NamedPizza
    MargheritaPizza
```

Although this hierarchy organized the conceptual vocabulary, it contained almost no semantic information.

A reasoner could verify that the taxonomy remained logically consistent, but it could not infer anything interesting about pizzas because the ontology had not yet described what distinguished one pizza from another.

Exercise 15 introduced the first semantic restrictions:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

These restrictions do not perform reasoning.

Instead, they establish the semantic conditions upon which future reasoning will operate.

This distinction is fundamental.

Semantic description defines knowledge. Reasoning derives knowledge.

19.11.2 Every Restriction Should Capture One Independent Semantic Fact

One of the most effective engineering practices in ontology development is to treat each restriction as expressing **one independent semantic statement**.

Rather than attempting to describe an entire concept through a single large logical expression, ontology engineers typically build semantic definitions from multiple small, understandable restrictions.

For example, **Exercise 15** introduces two separate restrictions:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

Although these restrictions eventually contribute to the semantic description of **MargheritaPizza**, each communicate exactly one piece of domain knowledge.

The first states that every Margherita pizza possesses at least one mozzarella topping.

The second states that every Margherita pizza possesses at least one tomato topping.

Keeping these semantic statements separate provides several important engineering advantages:

- Individual restrictions remain easy to read and review.
- Semantic changes affect only one logical statement.
- Domain experts can validate each assertion independently.
- Additional restrictions can be introduced without rewriting existing ones.
- Reasoning results become easier to explain and debug.

Only after several independent semantic facts have been validated should they be viewed collectively as the **semantic description** of the concept.

This modular approach closely resembles good software engineering.

Rather than writing one enormous function, developers divide behavior into smaller methods.

Likewise, ontology engineers divide semantic meaning into reusable logical components:

One restriction should normally express one semantic fact. Complex meaning should emerge through composition rather than monolithic expressions.

Future chapters will continue enriching the ontology through additional Description Logic constructors, progressively increasing its semantic expressiveness while preserving the stable conceptual foundation established in Stage 1.

19.11.3 Engineering Analogy -- Blueprint Before Construction

Every mature engineering discipline distinguishes between **design** and **execution**.

Construction engineers do not begin by pouring concrete before architectural drawings have been completed.

Mechanical engineers do not manufacture components before the product has been designed.

Likewise, ontology engineers should not expect automated reasoning to compensate for concepts whose semantic meaning has never been explicitly modeled.

Semantic Description serves the same role as an engineering blueprint.

A blueprint does not construct the building.

Instead, it specifies enough structural information for construction to proceed safely and consistently.

Similarly, OWL restrictions do not perform reasoning themselves.

Rather, they provide the formal semantic specification upon which the reasoner operates.

This distinction mirrors the progression established throughout the **Semantic Knowledge Development Lifecycle (SKDL)**:

- **Stage 1** establishes the conceptual structure.
- **Stage 2** enriches that structure with semantic meaning.
- Only after sufficient semantic knowledge exists can later stages perform validation, reasoning, governance, and ultimately executable intelligence.

Attempting to reason without semantic description is therefore comparable to attempting construction without architectural plans.

There is simply nothing sufficiently precise for the execution process to interpret.

Semantic descriptions are the blueprints from which reasoning constructs new knowledge.

19.11.4 Incremental Semantic Maturity

One of the most important engineering principles illustrated by **Exercise 15** is that:

semantic models mature gradually.

Beginners often believe an ontology should contain complete logical definitions as soon as classes are created.

Professional ontology development follows different philosophy.

- Semantic meaning accumulates incrementally.
- Each restriction contributes another piece of knowledge.
- Each chapter introduces additional logical precision.
- Each validation cycle improves semantic quality.

Exercise 15 deliberately introduces only two existential restrictions:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

These restrictions do not completely define a Margherita pizza.

Instead, they establish an **initial** semantic description that can be extended later.

Future chapters will introduce:

- universal restrictions (**only**);
- cardinality restrictions (**min**, **max**, **exactly**);
- equivalent class definitions;
- reusable semantic patterns; and
- automated classification through Description Logic.

This gradual refinement reflects an important engineering principle:

Semantic quality improves through iteration rather than completeness on the first attempt.

The Open World Assumption discussed earlier in this chapter & book naturally supports this development strategy.

Knowledge is expected to grow.

Semantic descriptions therefore remain open to future refinement as domain understanding evolves.

19.11.5 Design Semantic Models for Humans First

Although OWL restrictions are ultimately interpreted by machines, they are first created, reviewed, maintained, and governed by people.

Consequently, semantic models should be designed with human understanding in mind.

Clear class names, **consistent modeling patterns**, and **readable class expressions** reduce the effort required to maintain an ontology over time.

For example, `hasTopping some MozzarellaTopping` communicates its meaning - to human - almost immediately.

By contrast, deeply nested class expressions containing numerous logical operators may be mathematically correct while remaining difficult for ontology engineers to understand and maintain.

Professional ontology engineering therefore **balances** two complementary objectives:

- logical precision for automated reasoning; and
- conceptual clarity for human comprehension.

As ontologies expand to thousands or even millions of concepts, **readability** becomes an essential aspect of engineering quality rather than merely a stylistic preference.

A semantic model that cannot be understood by its maintainers is unlikely to remain correct over the long term.

19.11.6 Let the Reasoner Become Your Design Assistant

Many beginners regard the reasoner as a validation tool that is executed only after modeling has finished.

Experienced ontology engineers use it very differently.

The reasoners become **an active design assistant** throughout ontology development.

After introducing each new semantic restriction, the ontology should be synchronized and classified.

This continuous feedback allows engineers to detect:

- inconsistent class definitions;
- unsatisfiable concepts;
- unintended subclass relationships;
- missing semantic constraints; and
- unexpected logical consequences.

Rather than:

waiting until the ontology is complete,

semantic validation:

becomes part of the normal modeling workflow.

Exercise 15 already demonstrates this engineering practice.

Immediately after introducing the new restrictions, the tutorial instructs the reader to synchronize the reasoner and verify that the ontology remains logically consistent.

This iterative workflow greatly reduces debugging effort while improving confidence in the semantic model as it evolves.

19.11.7 Engineering Discipline Throughout the Semantic Knowledge Development Lifecycle

The principles introduced in this chapter extend well beyond **Exercise 15**.

They establish an engineering discipline that will guide every remaining stage of the Semantic Knowledge Development Lifecycle (SKDL).

Throughout SKDL, ontology engineers should consistently apply the following practices:

- establish conceptual structure before semantic definition (meaning);
- introduce semantic descriptions incrementally;
- validate changes continuously using the reasoner;
- prefer readable and reusable class expressions;
- avoid premature logical commitments; and
- allow semantic knowledge to mature alongside the ontology.

These practices enable semantic models to evolve naturally while preserving **logical consistency**, **maintainability**, and **extensibility**.

They also prepare the ontology for the more advanced stages of SKDL, including semantic reuse, governance, automated reasoning, and executable intelligence.

19.11.8 Engineering Takeaway

Exercise 15 introduces only two OWL restrictions, yet it represents a significant transition in ontology engineering.

For the first time, concepts acquire formal semantic meaning that can be interpreted mathematically and processed automatically by Description Logic reasoners.

The engineering lesson extends far beyond Protégé.

Semantic description should not be viewed as writing logical syntax for its own sake.

Instead, it is the disciplined process of transforming conceptual structures into precise, machine-understandable knowledge.

When semantic models are **developed incrementally**, **validated continuously**, and **expressed clearly**, they become a durable foundation upon which increasingly sophisticated knowledge systems can be constructed.

This philosophy will continue to guide every subsequent stage of the Semantic Knowledge Development Lifecycle.

From this point onward, the ontology will evolve not simply by adding new concepts, but by enriching those concepts with progressively

- deeper semantic meaning,
- stronger logical precision, and
- greater capacity

for automated knowledge processing.

19.12 EKA Perspective -- Stage 2 Enriches the Knowledge Layer K and Prepares Reasoning R

One of the central ideas introduced throughout this chapter is that:

knowledge becomes executable ONLY AFTER concepts acquires formal semantic meaning.

In **Chapter (18)**, Stage 1 of the Semantic Knowledge Development Lifecycle (SKDL) established the conceptual vocabulary of the domain. The ontology identified concepts such as **Pizza**, **NamedPizza**, and **MargheritaPizza**, and organized them into a coherent semantic taxonomy.

At this stage, however, these classes functioned primarily as **conceptual placeholders**. They identified *what* concepts existed within the domain, but they did not yet explain *what those concepts meant*.

Exercise 15 fundamentally changes this situation.

By introducing the first OWL restrictions:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

the ontology begins attaching **formal semantic descriptions** to its conceptual vocabulary.

This marks the transition from a taxonomy on names to a **machine-interpretable semantic knowledge model**.

Within the **Executable Knowledge Architecture (EKA)**, this represents the second major milestone in the gradual construction of executable knowledge.

As introduced in Chapter (00), EKA organizes executable knowledge into five complementary architectural components:

- K - Knowledge Graph
- R - Reasoning & Rules
- Θ - Trigger
- Φ - Execution
- Γ - Governance

The work completed in this chapter contributes primarily to the **Knowledge Graph layer (K)** while simultaneously preparing the logical foundation required for the future **Reasoning & Rules layer (R)**.

Unlike Stage 1, which focused on conceptual organization, Stage 2 enriches each knowledge node with formal semantics.

The resulting contribution can be summarized as follows:

EKA Component	Contribution of Chapter (19)
K - Knowledge Graph	Enriches conceptual nodes with formal semantic descriptions through Description Logic class expressions and OWL restrictions, transforming concepts into machine-understandable knowledge objects.
R - Reasoning & Rules	Establishes the logical conditions upon which future automated reasoning and classification will operate. Although no inference is performed yet, the semantic foundations required by the reasoner are now being constructed.

EKA Component	Contribution of Chapter (19)
Θ - Trigger	Not yet involved. Semantic descriptions alone do not initiate events, but they provide the knowledge structures that future triggers will evaluate.
Φ - Execution	No executable behavior is introduced. The ontology continues to model knowledge rather than operational processes.
Γ - Governance	Semantic restrictions become governed knowledge assets. They introduce explicit, verifiable definitions that improve consistency, documentation, and long-term maintainability.

This mapping highlights an important architectural observation.

During Stage 1, the Knowledge Graph consisted primarily of **conceptual nodes connected through subclass relationships**.

After Stage 2, those same nodes begin to acquire **semantic content**.

Conceptually, the transformation can be visualized as from:

Stage 1

MargheritaPizza (Concept only)

toward:

Stage 2

```
MargheritaPizza
|
|-- hasTopping some MozzarellaTopping
|-- hasTopping some TomatoTopping
```

The ontology is no longer storing only *names*.

It is beginning to represent **meaning**.

This distinction is critical because reasoning engines do not operate on names.

A reasoner cannot infer that a pizza belongs to **MargheritaPizza** simply because the class exists in the hierarchy.

Instead, reasoning depends entirely upon the logical descriptions attached to that class.

Without semantic restrictions such as **hasTopping some MozzarellaTopping**, there is nothing for the Description Logic reasoner to evaluate.

From the perspective of EKA, Stage 2 therefore serves as the **bridge between knowledge representation and logical reasoning**.

Stage 1 answered the question:

What concepts exist?

Stage 2 answers the next question:

What do those conceptual formally mean?

Only after these semantic meanings have been established can future stages ask a third question:

What new knowledge can be inferred automatically?

That question belongs to the Reasoning & Rules layer (*R*) and will be explored in later chapters.

Another important architectural consequence concerns the relationship between the Knowledge Graph and semantic governance.

Once concepts acquire formal semantic descriptions, those descriptions themselves become governed assets.

A restriction such as `hasTopping some MozzarellaTopping` is no longer merely documentation.

It becomes part of the **formal contract** defining the concept.

Changes to these restrictions may alter

- reasoning results,
- validation outcomes,
- knowledge reuse patterns, and
- downstream executable behavior.

Consequently, semantic descriptions must eventually be governed with the same discipline applied to source code, enterprise models, or database schemas.

This illustrates another important principle of EKA:

Knowledge becomes governable only after it becomes explicit.

Conceptual names alone provide little opportunity for validation or quality assurance.

Formal semantic definitions, by contrast, can be reviewed, verified, reused, versioned, reasoned over, and governed throughout the lifetime of the ontology.

From the perspective of the Semantic Knowledge Development Lifecycle, Stage 2 therefore represents much more than the introduction of OWL syntax.

It transforms conceptual structures into **formal semantic knowledge** that machines can interpret mathematically, reasoning engines can process logically, and governance processes can manage systematically.

Just as Stage 1 established the structural foundation of the ontology, Stage 2 establishes its semantic foundation.

Together, these two stages complete the first half of the knowledge modeling process.

The chapters that follow will demonstrate how these semantic definitions can be **reused, refined, governed**, and ultimately **reasoned over**, progressively transforming a conceptual ontology into a fully executable knowledge system within the **Executable Knowledge Architecture**.

19.13 Engineering Guidelines for Semantic Description

Semantic Description represents the transition from **identifying concepts** to **formally defining their meaning**.

Although Protégé allows OWL restrictions to be created with only a few mouse clicks, designing semantic descriptions that remain **correct, maintainable**, and **reusable** throughout the lifetime of an ontology requires careful engineering discipline.

Professional ontology engineers recognize that every logical restriction added today becomes part of the knowledge base that future reasoning, validation, governance, and application systems will depend upon.

The following engineering guidelines summarize widely accepted practices for constructing robust semantic descriptions.

19.13.1 Model Meaning Incrementally

One of the most common mistakes made by beginners is attempting to produce complete semantic definitions as soon as a new concept is created.

Professional ontology development follows a far more incremental approach.

Rather than attempting to capture every characteristic immediately, ontology engineers gradually enrich concepts as domain understanding matures.

Exercise 15 demonstrates this philosophy perfectly.

Only two semantic restrictions are introduced for **MargheritaPizza**:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

These restrictions clearly do **not** describe every characteristic of a real Margherita pizza.

Instead, they represent the first steps toward its semantic definition.

As discussed earlier in this chapter, additional restrictions may naturally be introduced later without requiring existing semantic descriptions to be discarded.

This incremental strategy offers several important advantages:

- It reduces unnecessary redesign.
- It supports evolving domain knowledge.
- It simplifies validation.

- It keeps semantic models understandable throughout development.

Remember:

Semantic maturity should evolve through successive refinement rather than one-time specification.

19.13.2 Prefer Necessary Conditions Before Complete Definitions

Another frequent mistake is introducing complete logical definitions too early.

OWL provides two different mechanisms for describing classes.

`SubClassOf` expresses **necessary conditions**.

`EquivalentTo` expresses **necessary and sufficient conditions**.

During the early stages of ontology development, concepts often continue evolving as engineers gain a deeper understanding of the domain.

Under these circumstances, necessary conditions provide a much safer modeling strategy.

They allow semantic knowledge to accumulate naturally without prematurely constraining future development.

Complete definitions should generally be introduced only after:

- conceptual boundaries have stabilized;
- domain terminology has become consistent;
- reusable modeling patterns have emerged; and
- engineers possess sufficient confidence in the semantic definition.

This staged progression mirrors the overall philosophy of the Semantic Knowledge Development Lifecycle (SKDL).

First describe concepts sufficiently. Only later define them completely.

19.13.3 Write Readable Class Expressions

Description Logic provides considerable expressive power.

However, expressive power should never come at the expense of readability.

Ontology engineers do not write class expressions solely for reasoners.

They also write them for future engineers, domain experts, reviewers, and maintainers.

A good class expression should communicate its meaning almost as naturally as ordinary language.

For example,

```
hasTopping some MozzarellaTopping  
and
```

```
hasTopping some TomatoTopping
```

can be understood immediately by both ontology engineers and domain specialists.

By contrast, deeply nested expressions containing multiple negations, intersections, unions, and cardinality restrictions may become difficult to interpret, even when logically correct.

Whenever possible:

- compose expressions from smaller logical building blocks;
- avoid unnecessary nesting;
- group related restrictions together; and
- prioritize semantic clarity over logical cleverness.

Remember:

Readable ontologies are maintainable ontologies.

19.13.4 Reuse Semantic Patterns

As ontologies grow, similar semantic descriptions inevitably appear across multiple concepts.

Rather than repeatedly constructing nearly identical restrictions, professional ontology engineers identify reusable semantic patterns.

For example, many pizzas may require similar topping restrictions.

Instead of redefining the same logical structures repeatedly, later chapters will demonstrate how common semantic descriptions can be reused systematically throughout the ontology.

Semantic reuse offers several important benefits:

- Reduced duplication.
- Improved consistency.
- Easier maintenance.
- Lower probability of modeling errors.

This principle closely parallels software engineering.

Just as reusable libraries replace duplicated code, reusable semantic patterns replace duplicated logical definitions.

19.13.5 Validate Frequently with the Reasoner

Semantic modeling should never be viewed as a purely editing activity.

Each logical restriction changes the knowledge model that the reasoner must interpret.

For this reason, experienced ontology engineers synchronize the reasoner regularly throughout development rather than waiting until the ontology becomes large.

Exercise 15 concludes by synchronizing the reasoner after introducing only two restrictions.

Although the ontology remains relatively small, this engineering discipline should already become routine.

Frequent validation allows engineers to:

- verify that restrictions behave as intended;
- detect inconsistencies early;
- observe inferred classifications; and
- identify modeling mistakes before they propagate throughout the ontology.

Reasoning should therefore become an integral part of the development workflow rather than a final verification step.

19.13.6 Avoid Over-Constraining Early Models

A concept should rarely receive its most restrictive semantic definition during its initial development.

Overly restrictive axioms often reduce flexibility and make future extensions unnecessarily difficult.

Whenever possible, begin with relatively simple existential restrictions.

For example, `hasTopping some MozzarellaTopping` is usually more appropriate during early modeling than immediately introducing combinations of:

- `only`
- `exactly`
- `max`
- complex nested expressions.

As understanding of the domain improves, additional restrictions can gradually refine the semantic description.

This strategy complements the Open World Assumption discussed earlier in this chapter.

Knowledge grows incrementally, so semantic description should evolve in the same manner.

19.13.7 Separate Conceptual Organization from Semantic Description

One of the central engineering principles of the Semantic Knowledge Development Lifecycle is the clear separation between **conceptual organization** and **semantic definition**.

Stage 1 answered the question:

What concepts exist within the domain?

Stage 2 answers a different question:

What do those concepts actually mean?

Although these activities are closely related, they should not be performed simultaneously.

Attempting to redesign conceptual hierarchies while simultaneously introducing complex semantic restrictions often leads to unnecessary restructuring and increased maintenance costs.

A stable taxonomy provides the structural foundation upon which semantic descriptions can evolve confidently.

Separating these responsibilities also makes ontology development

- more predictable,
- easier to review, and
- simpler to maintain over time.

19.13.8 Engineering Takeaway

Although **Exercise 15** introduces only two OWL restrictions, its engineering significance extends far beyond the Protégé user interface.

For the first time, the ontology begins expressing **formal semantic meaning** rather than merely organizing concepts into a taxonomy.

Semantic Description transforms conceptual structures into machine-interpretable knowledge by attaching logical meaning to previously abstract concepts.

Well-designed semantic descriptions should:

- communicate meaning clearly;
- support automated reasoning;
- remain understandable to both engineers and domain experts;
- accommodate future evolution;
- encourage semantic reuse; and
- provide a stable foundation for validation, governance, and executable knowledge.

The principles presented in this section are not unique to the **Pizza** ontology.

They apply equally to enterprise knowledge graphs, biomedical ontologies, digital twins, manufacturing knowledge bases, scientific taxonomies, and large-scale Semantic Web systems.

By following these engineering practices, ontology engineers create semantic models that remain **correct**, **maintainable**, **reusable**, and **extensible** long after the ontology's initial development has been completed.

19.14 Key Concepts

Throughout this chapter, we explored **Semantic Description** from practical, engineering, mathematical, historical, and architectural perspectives.

The following concepts summarize the essential terminology introduced during **Stage 2 -- Semantic Description** of the **Semantic Knowledge Development Lifecycle (SKDL)**.

Concept	Description
Semantic Description	The process of enriching conceptual models with formal semantic meaning through OWL axioms and Description Logic expressions, enabling machines to interpret concepts according to their logical characteristics rather than their names alone.

Concept	Description
Description Logic (DL)	A family of mathematically defined, decidable knowledge representation languages that provide the formal logical foundation of OWL and support automated reasoning over ontologies.
Formal Semantics	A mathematically precise interpretation of ontology constructs that ensures every class, property, and logical expression has an unambiguous meaning independent of software implementation.
Class Expression	A logical description that specifies the conditions required for membership within a class. Class expressions may combine multiple logical constructors to represent increasingly rich semantic definitions.
Named Class	A class identified by a globally unique IRI that forms part of the ontology's conceptual vocabulary and appears explicitly within the class hierarchy.
Anonymous Class	A class defined by a logical expression without introducing a new named concept. Anonymous classes are widely used within restrictions and logical definitions.
Existential Restriction (<i>some</i>)	A Description Logic constructor ($\exists R.C$) stating that an individual must participate in at least one relationship R whose target belongs to class C . In OWL Manchester Syntax, this is expressed as Property some Class .
Necessary Condition	A semantic condition that every instance of a class must satisfy. Necessary conditions are commonly introduced through SubClassOf axioms and support incremental semantic enrichment.
Sufficient Condition	A semantic condition that completely characterizes a concept. When specified using EquivalentTo , it enables reasoners to classify individuals and classes automatically.
Semantic Constructor	A logical operator such as and , or , not , some , or only used to build increasingly expressive class expressions within Description Logic.
Semantic Composition	The process of combining multiple logical restrictions into a single coherent class expression that collectively describes the meaning of a concept.
Open World Assumption (OWA)	The assumption that knowledge may always be incomplete. The absence of information does not imply falsity, allowing ontologies to evolve incrementally while remaining logically consistent.
Decidability	A mathematical property of Description Logic ensuring that reasoning algorithms always terminate and produce a definitive result for supported reasoning tasks.
Knowledge Representation	The discipline of modeling domain knowledge in a formal, machine-interpretable form so that it can be validated, queried, shared, and reasoned over automatically.
<i>K</i> - Knowledge Graph Layer	The first component of the Executable Knowledge Architecture (EKA) . During Stage 2, conceptual nodes established in Stage 1 are enriched with formal semantic meaning through Description Logic and OWL class expressions.
Stage 2 - Semantic Description	The second stage of the Semantic Knowledge Development Lifecycle (SKDL) , responsible for transforming conceptual taxonomies into formal semantic

Concept	Description
	knowledge models by introducing logical restrictions, class expressions, and mathematically defined semantics before reasoning is introduced.

19.15 Chapter Summary

This chapter introduced **Stage 2 -- Semantic Description** of the **Semantic Knowledge Development Lifecycle (SKDL)**, making the transition from organizing concepts to formally defining their semantic meaning.

In **Chapter (18)**, we established the conceptual vocabulary of the **Pizza** ontology by identifying classes and organizing them into a coherent taxonomy. Although that hierarchy described the structural organization of the domain, it did not explain what distinguished one concept from another.

This chapter addressed that limitation by introducing the first semantic descriptions through **Exercise 15**.

Using Protégé, we enriched **MargheritaPizza** with two existential restrictions:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

Although these additions required only a few mouse clicks within the Protégé interface, they represented a major conceptual advancement.

For the first time, the ontology expressed **knowledge** in a form that machines could interpret **logically** rather than merely **structurally**.

Throughout the chapter, we gradually expanded this practical exercise into a broader understanding of semantic knowledge representation.

We first examined why semantic description must follow conceptual modeling, emphasizing that stable conceptual structures provide the foundation upon which semantic meaning can be added systematically.

We then introduced **Description Logic (DL)** as the formal mathematical language underlying OWL, demonstrating that the Manchester Syntax used throughout the tutorial is not simply application-specific notation but a human-readable representation of well-defined logical expressions.

Building upon this foundation, we explored **existential restriction (some)** as the first Description Logic constructor encountered in the tutorial. Rather than viewing **some** as merely another OWL keyword, we interpreted it as the logical existential quantifier ($\exists$), capable of expressing the existence of semantic relationships between concepts.

The chapter then broadened this perspective through **class expressions**, illustrating how multiple logical constructors can be composed to define increasingly rich semantic descriptions. We distinguished between named and anonymous classes, discussed the difference between necessary and sufficient conditions, and examined how semantic meaning accumulates incrementally through logical composition.

To reinforce the theoretical foundation, we examined semantic description from a mathematical perspective. Using the language of set theory and first-order logic, we showed that

- classes correspond to sets,
- object properties correspond to binary relation, and
- OWL restrictions correspond to formally defined set-theoretic operations.

These mathematical semantic guarantee that ontology reasoning is **precise**, **consistent**, and **independent** of any particular software implementation.

The historical development of these ideas further demonstrated that modern ontology engineering represents the continuation of a much longer intellectual tradition. Beginning with Aristotle's concepts of **genus** and **differentia**, continuing through Porphyry's hierarchical classifications, Leibniz's vision of a universal symbolic language, Frege's predicate logic, and finally twentieth-century Description Logic, the chapter illustrated how contemporary OWL inherits and formalizes ideas developed over more than two millennia.

From an engineering perspective, we emphasized that semantic description should be viewed as an **incremental modeling process** rather than an attempt to capture complete knowledge immediately. Guided by the Open World Assumption, ontology engineers progressively enrich concepts as domain understanding evolves, allowing semantic models to remain adaptable while preserving logical consistency.

Within the **Executable Knowledge Architecture (EKA)**, Stage 2 significantly enriches the **Knowledge Graph (*K*)** layer. Concepts that previously existed only as nodes within a taxonomy now acquire formal semantic descriptions, transforming the ontology into a machine-interpretable knowledge model. At the same time, these logical descriptions establish the prerequisites for the **Reasoning & Rules (*R*)** layer, where automated inference will later derive new knowledge from the semantic defined here.

The progression established across the first two stages of SKDL can therefore be summarized as below:

SKDL Stage	Primary Outcome
Stage 1 - Conceptual Modeling	Identify concepts and organize them into a coherent semantic taxonomy.
Stage 2 - Semantic Description	Enrich those concepts with formal semantic meaning through Description Logic and OWL class expressions.

Together, these two stages establish both the **structure** and the **meaning** of the ontology.

Only after concepts have been identified and their semantics formally described can knowledge be

- reused consistently,
- validated automatically,
- reasoned over logically, and
- ultimately transformed into executable intelligence.

This progression illustrates one of the central principles of ontology engineering:

Concepts organize knowledge. Semantic description explain knowledge. Together, they create the formal foundation upon which intelligent knowledge systems are built.

19.16 Looking Ahead -- Toward Knowledge Reuse

By the end of this chapter, the **Pizza** ontology contains both a stable conceptual hierarchy and its first formal semantic descriptions. Concepts are no longer represented merely as names within a taxonomy; they now carry explicit meaning expressed through Description Logic and OWL class expressions.

As the ontology continues to expand, however, a new engineering challenge naturally emerges.

Many concepts begin to **share similar semantic structures**.

Different pizza varieties often inherit common toppings, restrictions, annotations, and modeling patterns. Constructing every new class from the beginning would quickly become repetitive, time-consuming, and increasingly difficult to maintain. More importantly, repeated manual modeling introduces opportunities for inconsistency, where semantically similar concepts evolve differently simply because they were modeled independently.

Professional ontology engineering therefore emphasizes **knowledge reuse** as the next stage of semantic development.

Rather than repeatedly recreating similar semantic structures, ontology engineer identify reusable patterns that can be extended, refined, and specialized as new concepts are introduced.

Reuse not only

- improves modeling efficiency,

but also

- strengthens semantic consistency across the entire ontology.

This engineering philosophy is demonstrated directly in **Exercise 16** of the original *Protégé 5 New OWL Pizza Tutorial*.

Instead of constructing a new pizza definition from scratch, the exercise begins by duplicating the existing **MargheritaPizza** class and renaming it **AmericanaPizza**. The previously established semantic descriptions are preserved, after which only a single additional restriction is introduced:

```
hasTopping some PepperoniTopping
```

Although this appears to be a simple editing operation with Protégé, it represents a significant shift in engineering methodology.

Rather than rebuilding knowledge, we being **reusing existing semantic assets**.

The new concept inherits an established semantic foundation and is extended only where its meaning differs. This mirrors practices found throughout modern engineering disciplines.

Software engineers extend existing classes rather than rewriting identical code.

Database designers normalize schemas to eliminate redundant structures.

Enterprise architects develop reference architectures that can be adapted across multiple solutions.

Likewise, ontology engineers construct reusable semantic models that promote consistency while reducing duplication.

Within the **Semantic Knowledge Development Lifecycle (SKDL)**, this marks the beginning of **Stage 3 -- Knowledge Reuse**.

Having established conceptual structures in **Stage 1** and enriched them with formal semantics in **Stage 2**, we now focus on systematically reusing and extending those semantic definitions across the ontology. Reuse transforms isolated semantic descriptions into coherent modeling patterns capable of supporting larger and more maintainable knowledge bases.

From the perspective of the **Executable Knowledge Architecture (EKA)**, Stage 3 further strengthens the **Knowledge Graph (*K*)** layer by improving semantic consistency and reducing redundant modeling effort. At the same time, reusable semantic patterns provide a richer and more uniform foundation for the **Reasoning & Rules (*R*)** layer, enabling future inference to operate over knowledge that is both logically consistent and systematically organized.

The next chapter therefore explores how professional ontology engineers build ontologies that scale not by repeatedly creating new knowledge, but by **reusing**, **extending**, and **refining** the semantic knowledge they have already established.

This progression reflects another enduring engineering principle:

The maturity of an ontology is measured not by how much knowledge it contains, but by how effectively that knowledge can be reused.

Last Updated at 2026-08-15

Chapter 20 -- Knowledge Reuse: Stage 3 of the Semantic Knowledge Development Lifecycle

Reusing Semantic Knowledge Through OWL Design Patterns

"Knowledge becomes truly valuable not when it is created once, but when it can be reused many times."

The first two stages of the **Semantic Knowledge Development Lifecycle (SKDL)** established the semantic foundations of ontology engineering.

During **Stage 1 -- Conceptual Modeling**, we identified the concepts that define a knowledge domain and organized them into a coherent semantic taxonomy.

During **Stage 2 -- Semantic Description**, we enriched those concepts with formal meaning using **Description Logic (DL)**, OWL restrictions, and class expressions

Chapter 18 established the conceptual vocabulary. Chapter 19 gave that vocabulary formal semantic meaning. By the end of **Chapter (19)**, our ontology was no longer simply a collection of class names.

It had become a **formal semantic knowledge model** whose concepts possess machine-understandable meaning.

At this point, however, another important engineering question naturally arises:

Once semantic knowledge has been created, how should it be managed as the ontology continues to grow?

Should every new concept independently redefine knowledge that already exists?

Should ontology engineers repeatedly recreate similar semantic descriptions?

Or should previously established semantic knowledge itself become a reusable engineering asset?

Professional ontology engineering adopts the **third approach**.

Rather than viewing semantic definitions as isolated pieces of modeling work, mature ontologies treat them as **reusable knowledge assets** that can be shared, extended, specialized, and maintained systematically across the entire knowledge base.

This philosophy introduces **Stage 3 -- Knowledge Reuse** of the Semantic Knowledge Development Lifecycle (SKDL).

The objective of this stage is not to create additional semantic knowledge simply for its own sake.

Instead, it is to organize semantic knowledge so that it can be **reused consistently**, reducing duplication while improving maintainability, scalability, and long-term semantic quality.

Within the **Executable Knowledge Architecture (EKA)**, this transition represents another important increase in engineering maturity.

Knowledge is no longer viewed merely as information stored within an ontology.

Instead, semantic knowledge begins evolving into a reusable engineering asset capable of supporting


- multiple concepts,
- multiple ontologies,
- multiple applications, and
- eventually multiple intelligent systems.

Beginning with **Exercise 16** from the *Protégé 5 New OWL Pizza Tutorial*, this chapter explores how OWL enables semantic reuse through inheritance, logical abstraction, reusable modeling patterns, and carefully designed class definitions.

Along the way, we will examine knowledge reuse from multiple complementary perspectives, including software engineering, Enterprise Architecture, scientific research, biological evolution, mathematical abstraction, and modern Artificial Intelligence.

By the end of this chapter, you will recognize that **knowledge reuse is far more than another OWL modeling technique**.

It represents one of the universal engineering principles through which human knowledge continuously grows, evolves, and accumulates over time.



Whether knowledge is transmitted through scientific publications, enterprise architecture repositories, biological genomes, AI foundation models, or semantic ontologies, the underlying principle remains the same:

Progress is achieved not by repeatedly creating knowledge, but by continuously reusing, refining, and extending the knowledge that already exists.

Table of Contents

- 20.1 Introduction -- From Individual Knowledge to Reusable Knowledge
- 20.2 Why Knowledge Reuse Is a Universal Engineering Principle
 - 20.2.1 Reuse in Software Engineering
 - 20.2.2 Reuse in Enterprise Architecture
 - 20.2.3 Reuse in Scientific Research
 - 20.2.4 Reuse in Biological Evolution
 - 20.2.5 Reuse in Artificial Intelligence
 - 20.2.5.1 From Knowledge Reuse to Recursive Improvement
 - 20.2.6 From Universal Principle to Ontology Engineering
- 20.3 Exercise 16 -- Reusing Semantic Knowledge in Practice
 - 20.3.1 Engineering Objective of Exercise 16
 - 20.3.2 From Copying Classes to Reusing Knowledge
 - 20.3.3 Reuse Before Redesign
 - 20.3.4 Knowledge Reuse versus Knowledge Duplication
 - 20.3.5 Looking Beyond the User Interface
- 20.4 Reuse Before Rebuild -- An Engineering Pattern
 - 20.4.1 The Pattern in Software Engineering
 - 20.4.2 The Pattern in Enterprise Architecture
 - 20.4.3 The Pattern in Design and Manufacturing
 - 20.4.4 From Copying to Knowledge Reuse
 - 20.4.5 The Role of Variation in Knowledge Reuse
 - 20.4.6 The Limits of Cloning
 - 20.4.7 The Engineering Attitude
- 20.5 Reusable Semantic Patterns in OWL
 - 20.5.1 Inheritance (`SubClassOf`)
 - 20.5.2 Equivalent Class Definitions (`EquivalentTo`)
 - 20.5.3 Reusable Restriction Patterns
 - 20.5.4 Ontology Imports (`owl:imports`)
 - 20.5.5 Design Patterns
 - 20.5.6 Semantic Alignment
- 20.6 Reuse Across Disciplines - Mechanisms
 - 20.6.1 Software Engineering -- Inheritance and Composition
 - 20.6.2 Enterprise Architecture -- Capability Reuse and Reference Architecture
 - 20.6.3 Scientific Research -- Citation and Systematic Review
 - 20.6.4 Biology -- Gene Duplication and Regulatory Review
 - 20.6.5 Artificial Intelligence -- Transfer Learning and Model Adaptation
- 20.7 Mathematical Perspective -- Abstraction, Generalization and Reuse
 - 20.7.1 Generalization as Subset Inclusion
 - 20.7.2 Abstraction and Dimension Reduction
 - 20.7.3 Posets and Lattices
 - 20.7.4 Reuse and Logical Consistency
- 20.8 Interesting Reading -- From Plato to Design Patterns
- 20.9 EKA Perspective -- Stage 3 Strengthens *K* and Prepares *Gamma*
- 20.10 Engineering Guidelines for Knowledge Reuse
- 20.11 Key Concepts
- 20.12 Chapter Summary
- 20.13 Looking Ahead -- Knowledge Governance

20.1 Introduction -- From Individual Knowledge to Reusable Knowledge

One of the defining characteristics of human civilization is that **knowledge accumulates through reuse rather than continual reinvention**.

Very few ideas are created in complete isolation.

Instead, almost every advance builds upon knowledge that already exists.

Scientists extend previous discoveries through research and citation.

Engineers improve proven designs rather than redesigning every component from first principles.

Software developers reuse libraries, frameworks, and design patterns instead of repeatedly implementing identical functionality.

Enterprise architects organize reusable business capabilities, reference architectures, and architectural building blocks that can support numerous initiatives across an organization.

Even biological evolution follows the same principle.

Rather than redesigning living organisms from the beginning, evolution preserves successful genetic information across generations while gradually introducing variation through mutation and natural selection.

Modern Artificial Intelligence demonstrates a remarkably similar pattern.

Foundation models learn from enormous collections of accumulated human knowledge, while techniques such as **transfer learning**, **Retrieval-Augmented Generation (RAG)**, and **knowledge distillation** continually reuse previously acquired knowledge to solve new problems more effectively.

Although these disciplines appear unrelated, they all reflect the same underlying engineering principle:

Progress is achieved not by repeatedly creating knowledge, but by continuously reusing, refining, and extending existing knowledge.

Ontology engineering is no exception.

After completing **Stage 2 -- Semantic Description**, the **Pizza** ontology now contains its first meaningful semantic definitions.

For example, **MargheritaPizza** is no longer merely a concept within a taxonomy.

It possesses formal semantic descriptions through restrictions such as:

```
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

These restrictions successfully describe one particular pizza.

However, an important engineering challenge immediately appears.

As the ontology expands from a handful of concepts to hundreds -- or even millions -- of concepts, should ontology engineers repeatedly recreate similar semantic definitions?

Or should those semantic descriptions themselves become reusable knowledge assets?

This question motivates **Stage 3 -- Knowledge Reuse** of the Semantic Knowledge Development Lifecycle (SKDL).

Knowledge reuse extends the engineering principles established in the previous chapters.

- **Stage 1** introduced reusable conceptual abstractions through semantic taxonomies.
- **Stage 2** introduced reusable logical constructors through Description Logic (DL).
- Now **Stage 3** goes one step further by creating **semantic knowledge itself as reusable**.

Rather than viewing each class definition as an isolated modeling exercise, ontology engineers identify semantic patterns that can be shared consistently across many concepts.

As ontologies continue to grow, this shift becomes increasingly important.

Without systematic reuse, semantic models gradually accumulate duplicated restrictions, inconsistent definitions, unnecessary maintenance effort, and increasingly governance complexity.

Small educational ontologies may tolerate such duplication

Large enterprise knowledge graphs containing tens of thousands of concepts **CANNOT**.

Knowledge reuse therefore should not be viewed simply as a convenient modeling technique.

It represents a fundamental engineering discipline for controlling semantic complexity while enabling ontologies to evolve sustainably over long periods of time.

This engineering philosophy is familiar across many mature disciplines.

Enterprise Architecture promotes the principle of **Reuse before Buy before Build**, encouraging organizations to maximize the value of existing architectural assets before acquiring or constructing new solutions.

Scientific research measures impact not only by creating new knowledge, but also by how frequently that knowledge is reused and extended through subsequent publications.

Likewise, successful ontologies are characterized not by the number of concepts they contain, but by the extent to which their semantic knowledge can be reused consistently throughout the knowledge base.

Exercise 16 introduces this principle through a seemingly simple modeling activity.

Rather than creating entirely new semantic definitions from scratch, we begin constructing new ontology concepts by reusing semantic structures that already exist.

Although the practical exercise requires only a small number of operations within Protégé, its engineering significance extends far beyond the user interface.

It introduces one of the defining characteristics of every mature engineering discipline:

Well-designed knowledge should be created once, reused many times, and refined only when necessary.

Throughout this chapter, we will examine knowledge reuse from practical, engineering, mathematical, architectural, biological, and artificial intelligence perspectives.

By the end of the chapter, you will recognize that reusable semantic knowledge is more than a modeling convenience.

It is a **reusable knowledge asset** -- one that enables ontologies, enterprise architectures, scientific knowledge, AI systems, and even biological life to grow in complexity while remaining coherent, maintainable, and sustainable.

20.2 Why Knowledge Reuse Is a Universal Engineering Principle

One of the most remarkable characteristics shared by mature engineering disciplines is that:

progress depends far more on reuse than on continual reinvention.

Although innovation often receives the greatest attention, sustained progress is rarely achieved by repeatedly creating everything from the beginning.

Instead, successful systems grow by preserving proven knowledge while extending it to address new challenges.

This principle appears so consistently across science and engineering that it can reasonably be regarded as a universal engineering law.

Ontology engineering is no exception.

Once semantic knowledge has been carefully **designed, validated, and understood**, it becomes an engineering asset that should be reused rather than recreated. **Exercise 16** introduces this principle within the **Pizza** ontology, but its significance extends far beyond OWL or Protégé. It reflects a pattern that underlies software engineering, enterprise architecture, scientific research, biology, mathematics, and modern artificial intelligence.

The remainder of this chapter explores these perspectives to demonstrate that **knowledge reuse is not merely an ontology technique -- it is a fundamental mechanism through which complex knowledge systems evolve.**

20.2.1 Reuse in Software Engineering

Software engineering has long recognized that **maintainability depends upon systematic reuse.**

Modern applications are rarely developed entirely from scratch.

Instead, developers build upon existing programming languages, operating systems, libraries, frameworks, APIs, design patterns, and open-source ecosystems.

When creating a new application, an experienced engineer does not begin by writing a new compiler, networking stack, or database engine. Instead, existing components are reused whenever they have already demonstrated correctness, reliability, and maintainability.

This philosophy is reflected in principle such as:

- Design Patterns
- Don't Repeat Yourself (DRY)
- Convention over Configuration
- Framework-based Development
- Package Management Ecosystems

Knowledge reuse therefore reduces both development effort and engineering risk.

Ontology engineering follows exactly the same philosophy.

Instead of repeatedly redefining identical semantic structures, reusable ontology patterns become the semantic equivalent of software libraries.

20.2.2 Reuse in Enterprise Architecture

Enterprise Architecture extends this philosophy from software components to organizational knowledge.

Rather than designing every capability independently, mature EA practices encourage organizations to reuse proven architectural assets whenever possible.

A widely recognized principle is:

Reuse before Buy before Build

The sequence reflects increasing cost and decreasing reuse --

1. Reuse existing enterprise assets.
2. Purchase proven commercial capability (Commercial Off-The-Shelf, or COTS) when suitable assets are unavailable.
3. Build new solutions only when neither reuse nor acquisition satisfies the business need.

This principle minimizes duplication, reduces operational risk, improves governance, and accelerates delivery.

Semantic knowledge should be treated in exactly the same way.

Well-designed ontology modules, reusable class expressions, common vocabularies, and shared semantic patterns become enterprise assets that can be leveraged across projects rather than recreated independently.

Within the EKA framework, reusable semantic knowledge represents reusable organizational intelligence.

20.2.3 Reuse in Scientific Research

Scientific knowledge advances through cumulative refinement rather than isolated discovery.

Every research paper begins with a literature review because no scientific contribution exists in isolation.

Researchers study previous work, identify existing theories, reproduce experiments, extend established results, and contribute only the incremental knowledge that advances the field.

Consequently, one indicator of a publication's long-term impact is the number of times it is cited by subsequent research.

A highly cited paper demonstrates that its knowledge has been successfully reused by the scientific community.

Knowledge therefore achieves influence not simply by being published, but by becoming reusable.

Ontology engineering follows the same principle.

A reusable ontology contributes value far beyond its original project because its semantic structures can support future ontologies, knowledge graphs, and intelligent applications.

20.2.4 Reuse in Biological Evolution

Biological inheritance provides a useful analogy for cumulative information preservation and variation.

The genomes of modern organisms are not created independently for every generation.

Instead, genetic information is inherited, reused, recombined, and gradually refined through evolutionary processes.

Successful biological structures persist because they continue to prove useful under changing environmental conditions.

Evolution therefore preserves accumulated biological knowledge while permitting incremental adaptation.

Complex life emerges through billions of years of reuse rather than continual recreation.

Ontology evolution follows a remarkably similar pattern.

Successful semantic structures should likewise be inherited, extended, specialized, and adapted instead of repeatedly reconstructed.

In short:

- Biology: inheritance -> variation -> selection
- Ontology: reuse -> specialization -> refinement

20.2.5 Reuse in Artificial Intelligence

Recent advances in Artificial Intelligence further illustrate the importance of reusable knowledge.

Large Language Models are trained on vast collections of accumulated human knowledge rather than beginning with an empty understanding of the world.

Subsequent techniques continue this philosophy:

- Transfer Learning
- Fine-Tuning
- Retrieval-Augmented Generation (RAG)
- Knowledge Distillation
- Model Adaptation

Each technique seeks to reuse previously learned knowledge instead of rediscovering it independently.

Knowledge distillation provides an especially compelling example.

Rather than retraining a smaller model from the beginning, knowledge learned by a larger model is transferred into a more efficient representation.

The objective is not to recreate intelligence but to reuse it.

Ontology engineering pursues the same objective.


Semantic definitions, once established, should become reusable knowledge assets capable of supporting future applications.

20.2.5.1 From Knowledge Reuse to Recursive Improvement

Recent discussions surrounding **Artificial General Intelligence (AGI)** have introduced another perspective on knowledge reuse.

One frequently discussed hypothesis proposes that once an AI system becomes capable of contributing to AI research itself, it may begin improving the very models from which it originated.

Conceptually, the process can be viewed as:



Whether recursive self-improvement eventually produces AGI or even Artificial Super-Intelligence (ASI) remains an open scientific question.

Researchers continue to debate both its feasibility and its expected timescale. Nevertheless, one engineering observation is already evident:

each new generation of AI systems is built primarily by reusing the accumulated knowledge, architectures, and capabilities of previous generations.

Modern foundation models do not rediscover mathematics, language, software engineering, medicine, or ontology engineering independently.

They inherit enormous amounts of previously created human knowledge through trained corpora, reusable model architectures, transfer learning, fine-tuning, retrieval mechanisms, and knowledge distillation.

Even when new capabilities emerge, they typically represent **extensions of existing knowledge rather than complete reinvention**.

Consequently, **knowledge reuse** is becoming not merely an optimization technique, but one of the central mechanisms through which modern AI systems continue to evolve.

Ontology engineering follows the same engineering philosophy.

Every reusable ontology, semantic pattern, and knowledge graph contributes to a growing body of machine-interpretable knowledge that future intelligent systems may build upon rather than recreate.

From the perspective of the **Executable Knowledge Architecture (EKA)**, reusable ontologies can be viewed as persistent organizational knowledge assets. As AI systems increasingly consume structured knowledge rather than isolated documents, well-governed ontologies and knowledge graphs become reusable semantic foundations upon which future intelligent systems can continuously build.

In this sense, **Stage 3 -- Knowledge Reuse** is concerned not only with reducing modeling effort today, but also with enabling the accumulation of semantic intelligence for tomorrow.

From this perspective, knowledge reuse is one important mechanism through which machine capabilities can accumulate and improve.

20.2.6 From Universal Principle to Ontology Engineering

Across all these disciplines, a common pattern emerges as below:

Discipline	Reusable Asset
Software Engineering	Libraries, frameworks, APIs
Enterprise Architecture	Capability maps, reference architectures, reusable building blocks
Scientific Research	Publications, theories, experimental evidence
Biology	Genetic information
Artificial Intelligence	Learned representations, foundation models, distilled knowledge
Ontology Engineering	Ontologies, semantic patterns, reusable class expressions

Although the artifacts differ, the engineering philosophy remains remarkably consistent.

Knowledge becomes increasingly valuable as its ability to be reused increases.

Exercise 16 now introduces this same principle within OWL ontology engineering.

The objective is no longer merely to define semantic knowledge correctly, but to organize that knowledge so it can be

- reused consistently,
- governed effectively, and
- extended systematically

throughout the remainder of the **Semantic Knowledge Development Lifecycle (SKDL)**.

However, reuse also introduces new responsibilities. Once semantic knowledge is shared across multiple concepts and applications, changes made to one reusable definition may affect many dependent models. Consequently, semantic reuse naturally raises questions of **ownership**, **versioning**, **validation**, and **governance**. These engineering concerns become the focus of the next stage of the SKDL as well.

Reuse does not merely preserve existing knowledge. It increases its value! Every successful reuse validates previous knowledge, extends its applicability, reduces future engineering, and creates a stronger foundation for subsequent innovation.

In this sense, the value of knowledge compounds through reuse in much the same way that financial capital compounds through investment.

Knowledge survives civilizations not because it is repeatedly rediscovered, but because it is successfully transmitted, reused, and refined across generations.

In short:

Human civilization advances because knowledge accumulates through systematic reuse. Ontology engineering simply makes this principle explicit and machine-interpretable.

20.3 Exercise 16 -- Reusing Semantic Knowledge in Practice

In the previous section we examined why knowledge reuse represents one of the most fundamental engineering principles across science, software engineering, enterprise architecture, biology, and artificial intelligence.

Exercise 16 -- Create *AmericanaPizza* by Cloning *MargheritaPizza* and Adding Additional Restrictions now demonstrates how this principle is applied within an ontology.

At first glance, the exercise appears surprisingly simple.

Instead of creating a new pizza from scratch, we duplicate an existing class and introduce one additional semantic restriction.

Although this operation requires only a few clicks inside Protégé, it illustrates an engineering principle that scales to ontologies containing thousands -- or even millions -- of concepts.

The Protégé operation begins by copying an existing class definition. The engineering objective, however, is to reuse the semantic knowledge embodied in that definition while introducing only the differences required by the new concept.

This is precisely how mature knowledge systems evolve.

20.3.1 Engineering Objective of Exercise 16

Exercise 16 asks us to create *AmericanaPizza* by duplicating *MargheritaPizza* and adding one additional topping restriction.

Conceptually the exercise can be summarized as:

```
Existing Semantic Definition + Small Semantic Difference  
= New Knowledge
```

Rather than redefining every characteristic, we preserve the semantic knowledge that has already been validated.


Only the distinguishing features is added.

This reflects one of the oldest engineering principles:

Never recreate what has already been proven correct.


20.3.2 From Copying Classes to Reusing Knowledge

Protégé performs the operation using **Duplicate Selected Class**, either from Edit menu or simply right click on the class name:




In the "Duplicate Class" dialogue, change the name from `MargheritaPizza` to `AmericanaPizza`, leave all the other options as they are and select **OK**:

From:




to:



Notice: when you change **Name**, the IRI identifier is also changed.

After clicking **OK**, new named pizza is added:




Switch to the **Entities** tab, you may see this new named pizza **AmericanaPizza** contains the copied **SubClassOf** axioms from **MargheritaPizza**:

The screenshot shows the Protege ontology editor interface. On the left, the 'Class hierarchy' panel displays a tree structure starting from `owl:Thing`, with `Pizza` as a subclass, and `AmericanaPizza` as a subclass of `NamedPizza`. `AmericanaPizza` is highlighted with a red box. On the right, the 'Annotations' tab for `AmericanaPizza` is active, showing an `rdfs:comment` 'A pizza that only has Mozzarella and Tomato toppings'. Below this, the 'Description' tab shows the class expression editor. The 'SubClass Of' section contains three entries: `hasTopping some MozzarellaTopping`, `hasTopping some TomatoTopping`, and `NamedPizza`. A red box highlights these entries, and a red arrow points from the `AmericanaPizza` class in the hierarchy to this box.


Now add one more restriction to make `AmericanaPizza` distinguished from the copied one: `hasTopping some PepperoniTopping`. Note that if you see a dotted red line under any concept (for example, under `PepperoniTopping`), it means you haven't yet added that topping to the ontology. Add it first, then return to Class Expressions Editor:

The screenshot shows the 'Class expression editor' tab for `AmericanaPizza`. The text input field contains the expression `hasTopping some PepperoniTopping`. A red box highlights the text, and a red dotted line is visible under the `PepperoniTopping` concept, indicating it has not yet been added to the ontology.

After `PepperoniTopping` is added, you should see below:



Also modifying comment annotation, then the end state of **Exercise 16** should be like below:



From a software perspective this appears to be cloning.

From an ontology perspective something more interesting is occurring.

We are not merely copying a class.

We are reusing:

- conceptual placement,
- inherited semantics,
- logical restrictions,
- annotations,
- documentation, and
- modeling decisions.

Every one of these represents accumulated knowledge.

The duplicated class therefore becomes the starting point for further semantic specialization rather than an independent creation.

Knowledge evolves through **modification** rather than **repeated invention**.

20.3.3 Reuse Before Redesign

When `AmericanaPizza` is created, almost everything already exists:

```
NamedPizza
hasTopping some MozzarellaTopping
hasTopping some TomatoTopping
```

Only one new semantic statement is introduced:

```
hasTopping some PepperoniTopping
```

The resulting definition becomes:

```
MargheritaPizza + PepperoniTopping
= AmericanaPizza
```

From an engineering perspective, most of the semantic structure is retained, while a small number of additional axioms distinguish the new concept.

This ratio appears repeatedly throughout software engineering, enterprise architecture, and scientific modeling.

Innovation typically occurs by extending existing knowledge rather than replacing it.

20.3.4 Knowledge Reuse versus Knowledge Duplication

Although **Exercise 16** begins with duplication, duplication itself is not the engineering objective.

There is an important distinction.

Duplication	Reuse
Copies information	Preserves validated knowledge
Produces independent artifacts	Maintains semantic lineage
Encourages divergence	Encourages consistency
Increases maintenance	Reduces maintenance effort

The duplicate serves only as an intermediate step.

Its purposes is to accelerate semantic reuse while preserving correctness.

As the ontology matures, increasingly sophisticated reuse mechanisms -- such as reusable superclass patterns, equivalent class definitions, ontology imports, and modular ontologies -- will gradually replace simple cloning.

Exercise 16 therefore introduces the simplest manifestation of a much broader engineering discipline.

20.3.5 Looking Beyond the User Interface

You may initially regard **Duplicate Selected Class** as merely a convenient Protégé feature.

It is considerably more important than that.

The command embodies a universal engineering philosophy.

Knowledge that has already been

- understood,
- validated,
- documented, and
- successfully applied

should become a **reusable asset** rather than disposable effort.

This principle explains:

- why software libraries exist,
- why enterprise architecture repositories exist,
- why scientific literature is cited,
- why biological evolution conserves successful genes, and
- why ontologies continue growing through semantic specialization instead of continual reinvention.

Exercise 16 therefore represents the first practical demonstration of **Stage 3 -- Knowledge Reuse** within the Semantic Knowledge Development Lifecycle (SKDL).

20.4 Reuse Before Rebuild -- An Engineering Pattern

The simple act of duplicating **MargheritaPizza** to create **AmericanaPizza** in **Exercise 16** illustrates a much broader engineering principle that extends far beyond Protégé or OWL. It is a concrete manifestation of a pattern that appears across virtually every mature engineering discipline:

clone before create -- or more formally, **reuse before build**.

This pattern is not merely a productivity shortcut.

It embodies a fundamental insight about how complex systems are constructed and maintained over time.

Engineering disciplines that have successfully scaled to manage extraordinary complexity -- from software engineering to aerospace systems, from enterprise architecture to semiconductor design -- all rely on some form of this principle.

Rather than beginning each new design from a blank canvas, engineers identify existing solutions that closely approximate the desired outcome, duplicate them, and then introduce only the modifications necessary to address the new requirements.

This approach is not an admission of limited creativity. It is a recognition that the vast majority of engineering effort is spent not on novel invention, but on the **adaptation**, **extension**, and **combination** of proven solutions.

By beginning with an existing artifact, engineers inherit the accumulated knowledge, testing, validation, and documentation embedded within it. Every hour spent designing a new solution from scratch is an hour that could have been spent refining a proven solution to fit a new context -- or, even more productively, improving the shared foundation so that everyone benefits.

20.4.1 The Pattern in Software Engineering

In software engineering, the **clone before create** pattern appears in numerous forms:

- **Class inheritance**: Subclasses extend the behavior of parent class, inheriting their methods and attributes while adding or overriding specific features. This is not duplication; it is specialization.
- **Copy-to-write**: Data structures are shared until modification is required, reducing memory consumption and improving performance.
- **Forking**: Open-source projects are forked to create new projects that share a common codebase but evolve independently.
- **Scaffolding**: Code generators produce initial versions of applications or components based on established templates, which are then customized.
- **Design Patters**: Reusable solutions to common problems, such as the Factory pattern or Observer pattern, are applied across projects without being re-implemented from scratch.

Each of these mechanisms shares a common logic:

start with a working solution, modify only what is necessary, and preserve what already works.

This logic is not accidental. It emerges from the observation that most engineering problems are variations of problems that have already been solved. To solve a new problem by reinventing the whole wheel is to ignore the accumulated knowledge of the engineering community. To solve it by adapting an existing solution is to build upon that knowledge.

The principle of **Don't Repeat Yourself (DRY)** formalizes this idea:

every piece of knowledge should have a single, unambiguous representation within a system.

When knowledge is duplicated, changes must be made in multiple places, increasing the risk of inconsistency and the cost of maintenance. DRY is not merely a coding guideline -- it is a reuse discipline.

Similarly, **Convention over Configuration** reduces the need for repetitive specification by establishing default behaviors that can be overridden only when necessary. This allows engineers to reuse established patterns without explicitly redefining them for each new context.

Package management ecosystems -- such as **npm**, **Maven**, **PyPI**, and others -- represent reuse at an industrial scale. Rather than implementing common functionality themselves, developers declare dependencies on existing libraries that have been tested, documented, and maintained by others. This dramatically reduces development effort while improving reliability.

Ontology engineering follows exactly the same philosophy. Instead of repeatedly redefining identical semantic structure, reusable ontology patterns become the semantic equivalent of software libraries. Our **Pizza** ontology itself can be viewed as a reusable module that could be imported into larger food ontologies, menu systems, or nutritional knowledge graphs.

20.4.2 The Pattern in Enterprise Architecture

Enterprise Architecture formalizes the same principle through the **Reuse before Buy before Build** framework.

This principle is not merely a suggestion; it is a governance rle that guides investment decisions across the enterprise.

Decision	Implication
Reuse	Leverage existing capabilities, reducing cost and risk
Buy	Acquire proven commercial solutions when reuse is not feasible
Build	Construct solutions only when neither reuse nor buy suffices

The discipline enforced by this framework ensures that enterprise assets are not continually recreated.

Instead, architectural building blocks -- capabilities, applications, data models, and reference architectures -- are designed explicitly for reuse.

When a new project begins, the first question is not "*What should be build?*" but rather "*What can we reused?*"

This is not merely an economic optimization. It is a strategy for managing complexity. Every time an existing asset is reused, the organization avoids introducing new variability into its architecture. Consistency improves. Governance becomes simpler. Maintenance effort declines.

Reference Architectures provide a particularly clear example. Rather than designing a new solution architecture for every project, organizations develop reference architectures that embody proven patterns, standards, and best (good) practices. These reference architectures are then adapted -- not recreated -- for each specific context.

Capability-based planning extends this thinking to the business level. Business capabilities are defined once and then realized through multiple projects and initiatives. The capability itself becomes a reusable asset that provides a stable foundation for investments decisions, roadmap planning, and architectural governance.

In enterprise architecture, reuse is not merely a technical practice -- it is a strategic discipline that aligns technology investment with business value.

20.4.3 The Pattern in Design and Manufacturing

Beyond software and architecture, the **clone before create** pattern appears in physical engineering as well:

- **Product line engineering:** Automobile manufacturers produce multiple vehicle models from shared platforms. The underlying chassis, powertrain, and electrical architecture are reused across models; only exterior styling and interior features vary. Volkswagen's MQB platform (*Modularer QuerBaukasten*, https://en.wikipedia.org/wiki/Volkswagen_Group_MQB_platform), for example, underpins dozens of models across multiple brands, sharing approximately 60% of components.
- **Design reuse in aerospace:** Aircraft designers are refined over decades. The Boeing 737, first introduced in 1967, has been continuously extended and modified, but its fundamental architecture remains recognizable. The 737 MAX shares the same basic airframe design as earlier models, enabling airlines to reuse pilot training, maintenance procedures, and spare parts.
- **Semiconductor design:** Integrated circuits (ICs) are built from reusable intellectual property (IP) blocks -- processors, memory controllers, interface modules -- that are combined and configured rather than designed from scratch. Modern system-on-chip (SoC) designs typically integrate dozens of pre-verified IP blocks, dramatically reducing design time and risk.
- **Mechanical engineering:** Standardized fasteners, bearing, motors, and other components are also reused across products. Engineers do not design a new screw for every application; instead, they select from standard catalogs.

In each case, the pattern is identical:

establish a proven base, maximize the reuse of the elements in this base, extend it incrementally, and minimize the introduction of unproven elements.

The alternative -- designing every system from the ground up -- is economically infeasible at scale.

Crucially, these physical engineering disciplines also demonstrate the limits of reuse. Components that are too tightly coupled to their original context cannot be reused without significant modification. Successful reuse requires thoughtful design for adaptability -- a theme we will return to in [Section 20.4.6](#).

20.4.4 From Copying to Knowledge Reuse

The software and design examples above share one subtle but important characteristic:

reuse in these disciplines is not merely about **copying**.

It is about:

preserving the knowledge embedded in the original artifact.

When a software engineer extends a class, they inherit not only *the code* but also **the design decisions, testing, and documentation** that accompanied the original implementation.

When an automobile manufacturer reuses a platform, they inherit **crash-test results, manufacturing processes, and supply-chain relationships**.

When an ontology engineer duplicates a class, they inherit **the semantic decisions, logical restrictions, and validation history** attached to the original concept.

This is the key insight that distinguishes **reuse** from mere **copying**:

	Copying	Knowledge Reuse
Preserves	Text, Code, or Structure	Validated Knowledge, Design Rationale, and Engineering Decisions
Reduces	Typing Effort	Error Risk and Design Uncertainty
Requires	No Understanding	Understanding of What is being Reused
Supports	Independent Evolution	Consistent Evolution across Dependent Artifacts
Outcome	Redundant Artifacts	Coherent, Maintainable Knowledge Base

The duplicate operation in **Exercise 16** is therefore not a mechanical copy operation. It is a deliberate act of knowledge transmission.

By duplicating **MargheritaPizza**, the ontology engineer inherits:

- **Conceptual Placement** -- the class's position within the **NamedPizza** hierarchy
- **Inherited Semantics** -- all restrictions inherited from **NamedPizza** and **Pizza**
- **Logical Restrictions** -- the existential restrictions on **MozzarellaTopping** and **TomatoTopping**
- **Annotation** -- labels, comments, and documentation
- **Modeling decisions** -- the choice of toppings and the semantic structure

The new class does not start from scratch; it starts from a position of existing semantic definitions --

The ontology engineer's task is no longer to **recreate** but to **extend**.

This distinction becomes particularly important when considering large ontologies. In a knowledge base containing thousands of concepts, the difference between "copying and pasting" and "reusing previous established semantic knowledge" is not merely semantic -- it is the difference between a **maintainable ontology** and an **unmaintainable one**.

20.4.5 The Role of Variation in Knowledge Reuse

Critically, knowledge reuse does not imply **stagnation**.

Every reuse introduces variation. The new class is modified to accommodate new requirements.

In **Exercise 16**, **AmericanaPizza** receives an additional topping restriction that distinguishes it from **MargheritaPizza**.

This variation is not a defect; it is the mechanism through which knowledge evolves.

Biological evolution provides a useful analogy. Genetic information is inherited from parent to offspring, but mutations introduce variation. Most mutations are neutral or harmful; occasionally, one proves beneficial and is preserved. Over time, the population evolves -- not by reinventing the genome, but by accumulating beneficial variations on a shared genetic foundation.

Ontology engineering follows a similar pattern. Semantic knowledge is inherited, but variations are introduced to address new requirements. Most variations are minor and have limited impact; occasionally, one proves valuable and becomes part of the reusable foundation for future concepts. The ontology evolves -- not by reinventing its concepts, but by accumulating useful semantic variations.

In this sense, **Section 20.3.3**'s observation that approximately 90% of the semantic definition is reused is not a limitation of the exercise; it is a feature of how knowledge evolves in any mature discipline. Innovation typically occurs by extending existing knowledge rather than replacing it.

Variation without reuse leads to fragmentation. Each new concept becomes a silo island, unrelated to what came before. The ontology becomes difficult to navigate, govern, or reason over.

Reuse without variation leads to stagnation. The ontology cannot adapt to new requirements or domains. It becomes brittle and unresponsive.

The balance between reuse and variation is essential. It is the mechanism through which knowledge systems grow while maintaining coherence. This balance is not achieved automatically; it requires conscious engineering decisions about when to extend, when to specialize, and when to create something genuinely new.

20.4.6 The Limits of Cloning

While **cloning** is a useful starting point for knowledge reuse, it is not the ultimate mechanism.

As ontologies grow, simple duplication becomes insufficient for several reasons:

1. **Maintenance burden:** Changes to the original class are not automatically propagated to cloned classes. If `MargheritaPizza` is refined -- for example, if a new restriction is added or an annotation is corrected -- `AmericanaPizza` must be manually updated. In a large ontology, this manual synchronization becomes a significant maintenance cost.
2. **Semantic lineage:** The ontological relationship between cloned classes is not captured. The ontology knows that `AmericanaPizza` was copied from `MargheritaPizza` in a modeling session, but it does NOT know *why* they are related or what semantic relationship exists between them. This information is lost unless explicitly documented.
3. **Scalability:** Cloning is effective for small-scale reuse but becomes unwieldy when hundreds of classes must be derived from shared foundations. The manual effort of duplicating, renaming, and modifying each class scales linearly with the number of variants, while the potential for error grows super-linearly.
4. **Reasoning limitations:** The reasoner does not recognize the lineage between cloned classes. It treats `MargheritaPizza` and `AmericanaPizza` as unrelated concepts unless additional axioms are introduced. The semantic connection exists in the engineer's mind, but not in the ontology.
5. **Governance challenges:** When cloned classes diverge, they may become inconsistent over time. The conceptual relationship between them -- originally a shared foundation -- becomes obscured. Governing such an ontology requires additional effort to maintain consistency.

For these reasons, cloning is best understood as an **introductory mechanism** for knowledge reuse. As ontology engineers become more experienced, they adopt more sophisticated patterns:

- **Shared superclasses:** Common restrictions and properties are abstracted into parent classes that multiple subclasses inherit. For example, all named pizzas might inherit a common set of base restrictions.
- **Equivalent class definitions:** Classes are defined by logical conditions that can be independently applied and reused. A `VegetarianPizza` definition, once established, can be reused across multiple contexts.

- **Ontology modules:** Self-contained ontologies that can be imported and extended across projects. The **Pizza** ontology itself could become a module imported into larger food ontologies.
- **Design patterns:** Proven modeling templates that capture recurring semantic structures, such as the **hasTopping** restriction pattern.

Exercise 16 introduces the simplest of these mechanisms because it establishes the fundamental principle:

The goal is not to copy; the goal is to reuse validated knowledge while introducing only necessary variation.

Once this principle is understood, the more sophisticated mechanisms become natural extensions of the same practice.

20.4.7 The Engineering Attitude

Ultimately, the **clone before create** pattern is not a technique but an **engineering attitude**.

It is the disposition to ask, before beginning any new modeling activity:

- What existing knowledge can I reuse?
- How can I introduce only the minimum necessary novelty?
- What has already been validated that I can build upon?
- What will be the maintenance cost of creating something entirely new?

This attitude is not limited to ontology engineering. It is the attitude that has enabled every successful engineering discipline to scale. It is the recognition that progress is achieved not by repeatedly creating knowledge, but by continuously reusing, refining, and extending the knowledge that already exists.

Exercise 16 serves as the first practical demonstration of this attitude within the Semantic Knowledge Development Lifecycle (SKDL). It is a small operation, but it embodies a large principle -- one that will reappear throughout the remaining stages of SKDL and across every mature engineering discipline that has ever managed to grow without collapsing under its own complexity.

The engineering attitude also requires humility.

- It acknowledges that most problems have been encountered before and that someone else has likely developed a solution worth studying.
- It does not devalue creativity; it redirects it toward the genuinely novel aspects of a problem rather than its routine elements.

In this sense, knowledge reuse is not merely a technique for managing complexity. It is a recognition of our interdependence as engineers -- an acknowledgement that we build upon the work of those who came before us and that our own work will, if we are fortunate, provide a foundation for those who follow.

20.5 Reusable Semantic Patterns in OWL

Exercise 16 demonstrated reuse through the simplest possible mechanism: **cloning**.

But cloning, as **Section 20.4.6** established, is only the **entry point**.

OWL itself provides several formal mechanisms through which semantic knowledge can be reused -- not by copying axioms, but by **referencing** them.

This distinction matters enormously.

When knowledge is copied, two definitions now exist. When knowledge is referenced, one definition exists -- used in many places.

The remainder of this section catalogs the primary reuse mechanisms available in OWL, moving from the simplest (inheritance) to the most structurally significant (semantic alignment).

When possible, each mechanism is tied back to a specific exercise from the *Protege 5 New OWL Pizza Tutorial* (M. DeBellis), so you can see exactly where in your own ontology you already built, and will build, the pattern being described.

20.5.1 Inheritance (SubClassOf)

The most fundamental reuse mechanism in OWL is one you have already used many times: **SubClassOf**.

When **VegetarianPizza** is declared **SubClassOf** **Pizza**, it does not redefine what a **Pizza** is.

Instead, it **reuses** the entire semantic definition of **Pizza** -- every restriction, every annotation, every axiom that applies to **Pizza** -- and add only what is genuinely new:

Class: VegetarianPizza
SubClassOf: Pizza

The screenshot shows the Protege 5 interface. On the left, the 'Classes' tab is active, displaying a class hierarchy. 'VegetarianPizza' is highlighted with a red box, and a red arrow points from it to the 'SubClass Of' section on the right. The right pane shows the 'Annotations' and 'Usage' tabs for 'VegetarianPizza'. The 'Annotations' section shows 'rdfs:label' as 'Vegetarian Pizza' and 'rdfs:comment' as 'A pizza containing only vegetable toppings'. The 'Description' section shows 'Equivalent To' as 'hasTopping only VegetableTopping'. The 'SubClass Of' section shows 'Pizza' as the superclass.

The relationship is often described as inheritance because it has a useful analogy with object-oriented inheritance. However, OWL inheritance is fundamentally logical rather than behavioral.

OOP	OWL
Reuses behavior	Reuses semantic consequences
Methods / fields	Class axioms
Runtime behavior	Logical entailment
Overriding	No direct equivalent
Object-oriented type system	Description Logic semantics

A subclass is subject to the logical consequences of the axioms that apply to its superclass -- everything the reasoner knows about **Pizza** becomes, *automatically and without restatement*, something the reasoner also knows about **VegetarianPizza**.

This is why [Section 20.2.6](#) described class hierarchy itself as the **first and most pervasive reuse pattern** in ontology engineering. Every hierarchy you have built since **Chapter (04)** has been a reuse structure, whether it was labeled as one or not.

There is, however, a precise logical limit to what `SubClassOf` reuses, and it is worth naming explicitly.

`SubClassOf` tells the reasoner that every member of the subclass must also satisfy the superclass definition. It does not say that satisfying the superclass definition is enough to become a member of the subclass.

A class defined only through `SubClassOf` axioms is called a **primitive class** -- and, as the tutorial's **Section 4.11** puts it directly, a primitive class is one where *"if something is a member of this class then it is necessary to fulfil these conditions,"* but the reverse can never be inferred.

Protégé will never automatically move an individual *into* a primitive class through classification, no matter how well that individual matches the class' intuitive description, because the ontology has never stated what would be *sufficient* to qualify.

Demonstrated in the tutorial: Nearly every class built through **Exercises 4-19** -- `Pizza`, `PizzaTopping`, `NamedPizza`, `MargheritaPizza`, and their siblings -- is a primitive class, defined purely through `SubClassOf`.

Exercise 20 will build `CheesyPizza` the same way, as a primitive class, specifically so that **Exercise 21** can then demonstrate what changes when it becomes a defined class instead.

[!Note] If you inspect the Pizza ontology's **Asserted** hierarchy against its **Inferred** hierarchy in Protégé (a UI mechanic explained in the `Pizza` tutorial's **Section 4.11**, building on the general asserted-vs-inferred distinction **Chapter (06)** first raised), primitive classes are exactly the ones where the two hierarchies never diverge. Their position in the tree is something you declared, not something the reasoner discovered.

20.5.2 Equivalent Class Definitions (`EquivalentTo`)

Inheritance reuses a class **as a starting point**.

`EquivalentTo` reuses a class **as a complete, standalone definition**.

Recall the classic pattern from **Chapter (15)**:

```
Class: VegetarianPizza
  EquivalentTo: Pizza
    and (hasTopping only (CheeseTopping or VegetableTopping))
```

The screenshot shows the Protégé interface with the `VegetarianPizza` class selected. On the left, the 'Class hierarchy: VegetarianPizza' panel shows a tree where `VegetarianPizza` is a subclass of `Pizza`. On the right, the 'Annotations: VegetarianPizza' panel shows the `rdfs:label` as 'Vegetarian Pizza' and the `rdfs:comment` as 'A pizza containing only non-meat toppings, suitable for vegetarian diets.' Below this, the 'Description: VegetarianPizza' panel shows the `Equivalent To` constraint: `hasTopping only (CheeseTopping or VegetableTopping)`. A red arrow points from the `VegetarianPizza` class in the hierarchy to its description panel.

This is not merely a restriction attached to a class. It is a **logical definition** -- a statement that anything satisfying the right-hand side of the expression **is**, by definition, a **VegetarianPizza**.

Once defined this way, the expression becomes reusable in a much stronger sense than inheritance allows. Any other class, anywhere in the ontology, can reference the same logical pattern -- **hasTopping only (CheeseTopping or VegetableTopping)** -- without needing to inherit from **VegetarianPizza** at all. The **pattern itself** becomes a reusable unit of semantic knowledge, independent of where it first appeared.

Inheritance reuses **classes**. Equivalent class definitions reuse **logical conditions**.

A class defined through **EquivalentTo** is called a **defined class**, in direct contrast to the **primitive class** from **Section 20.5.1**.

Because the right-hand side of an **EquivalentTo** axiom states a necessary-**and**-sufficient condition, the reasoner can do something it cannot do for primitive classes:

it can **automatically classify** any individual satisfying that condition as a member of the class, even if no engineer ever explicitly asserted the membership.

Demonstrated in the tutorial:

- Defined Class: **Exercise 21**, which takes the primitive **CheesyPizza** from **Exercise 20** and run **Edit > Convert to defined class** -- the exact moment the tutorial shows the reasoner beginning to auto-classify **MargheritaPizza** and the other named pizzas as subclasses of **CheesyPizza** without being told to.
- **Exercise 22** will repeat the pattern for **VegetarianPizza**, built directly as a defined class using a universal restriction.
- **Exercise 24** shows a second flavor of the same mechanism: the enumerated class **Spiciness**, defined as **EquivalentTo {Hot, Medium, Mild}** -- proof that **EquivalentTo** reuses not just restriction patterns but explicit value sets.

Two syntax variants are worth recognizing, since you will encounter both in Protégé and in the wider OWL ecosystem:

```
# Manchester Syntax (used throughout this book and the tutorial's Class expression editor)
Class: VegetarianPizza
  EquivalentTo: Pizza and (hasTopping only (CheeseTopping or VegetableTopping))

# OWL Functional-Style Syntax (used in the underlying RDF/XML and most OWL APIs)
EquivalentClasses(
  :VegetarianPizza
  ObjectIntersectionOf(
    :Pizza
    ObjectAllValuesFrom(
      :hasTopping
      ObjectUnionOf(
        :CheeseTopping :VegetableTopping
      )
    )
  )
)
```

One pitfall deserves a direct warning, since it is common enough to have its own name in ontology engineering discussions: a class may legally carry **more than one EquivalentTo axiom in OWL**.

When it does, the reasoner treats every listed expression as equivalent not only to the class, but to **each other**. If those expressions were not intended to be logically identical, this silently over-constrains the ontology -- and because it is valid OWL syntax, no reasoner error will warn you.

This is a reuse mechanism that punishes carelessness precisely because it is powerful.

20.5.3 Reusable Restriction Patterns

Chapter (14) and (15) introduced existential and universal restrictions individually. What was not emphasized at the time is that **the restrictions themselves are reusable patterns**, independent of any single class -- and the **Pizza** tutorial's Section 4.10.1 (page 30) actually names all three restriction families explicitly:

- **quantifier restrictions** (**some** / **only**),
- **cardinality restrictions** (**min** / **max** / **exactly**), and
- **hasValue restrictions**.

Consider the general shape of a quantifier restriction:

```
<property> only (<ClassA> or <ClassB> or ...)
```

This is not specific to **VegetarianPizza**. It is a **Template**. The same template, applied to different properties and different classes, produces an entire family of definitions.

Demonstrated in the tutorial:

- Quantifier restriction: **Exercise 13** (**hasBase some PizzaBase**, the existential template), and **Exercise 22** (**hasTopping only (VegetableTopping or CheeseTopping)**, the universal template).

Cardinality restrictions generalize the same idea toward counting rather than membership:

```
<property> min <n> <Class>      -- at least n relationships to Class
<property> max <n> <Class>      -- at most n relationships to Class
<property> exactly <n> <Class>  -- exactly n relationships to Class
```

Demonstrated in the tutorial:

- Cardinality restriction: **Exercise 26**, which will define **InterestingPizza** as **hasTopping min 3 PizzaTopping** -- a reusable "at least N fillers of this type" template that has nothing to do with pizza specifically and could be applied to any property and class pair.

The third family, hasValue restrictions, is **different** in kind: rather than requiring *some* or *only* individuals of a given class, it requires a relationship to one **specific** individual.

```
<property> value <specific individual>
```

Demonstrated in the tutorial:

hasValue restriction: **Exercise 25**, which will assert **hasSpiciness value Hot** on **JalapenoPepperTopping**, then reuse that exact pattern one level up by defining **SpicyPizza** as **hasTopping some (hasSpiciness value Hot)** -- a restriction nested inside a restriction, and a good illustration of how these templates compose.

Recognizing all three restriction families as templates -- rather than one-off constructors -- is what separates ontology engineers who model efficiently from those who rebuild the same logical shape from scratch every time a new class is needed.

This is the same principle Section 20.2.1 identified in software engineering: a well-designed function is written once and called many times.

A well-designed restriction pattern is written once and **instantiated many times**.

One caution belongs here, because it connects directly to the Open World Assumption first introduced in **Chapter 14** and revisited in the **Pizza** tutorial's Section 4.13: a cardinality or quantifier restriction constrains how many **asserted or inferable** relationships an individual has, not how many the ontology engineer happens to have typed into Protégé.

The same OWA reasoning that governed restrictions in **Chapter (14) and (15)** governs every restriction family identically.

20.5.4 Ontology Imports (**owl:imports**)

Every reuse mechanism discussed so far operates **within** a single ontology file.

owl:imports extends reuse **across** ontology files.

```
Ontology: <http://example.org/pizza-extended.owl>
Import: <http://example.org/pizza.owl>
```

When one ontology imports another, every class, property, and axiom defined in the imported ontology becomes available -- without copying a single line into the importing file. This is the OWL equivalent of a software library import, and it carries the same engineering benefit: the imported ontology can be maintained, corrected, and extended **independently**, and every ontology that imports it inherits the improvement automatically.

[!Note] Unlike every other mechanism in this section, **owl:imports** is **not** demonstrated anywhere in the **Pizza** Tutorial exercises, and this book's **Pizza.owl** does not use it either. The closest the tutorial comes is its Chapter 8, where a second, pre-populated version of the ontology is *loaded in place of* the original -- a replacement, not an import. **owl:imports** becomes essential once an ontology engineer begins working across multiple, independently maintained ontologies, a scenario more typical of enterprise-scale EKA deployments than of a single teaching ontology. It is included here because it is one of OWL's core reuse mechanisms, even though neither the tutorial nor this book has yet occasion to use it.

Three technical details are worth knowing before you use **owl:imports** outside a teaching context, because each one is a common source of confusion the first time an ontology engineer encounters it.

First, the IRI in an import statement is a logical identifier, not necessarily a network address.

- **owl:imports** names *which* ontology to bring in; it does not, by itself, say *where* the file physically lives.
- Protégé resolves this mapping through a **catalog file** (typically named **catalog-v001.xml**), which pairs each logical IRI with a local file path or a remote URI.

Second, imports are transitive.

- If Ontology A imports Ontology B, and Ontology B imports Ontology C, then Ontology A's reasoner sees the full contents of C as well, whether or not A's author has ever opened C directly.
- This transitive closure is powerful, but it also means that importing one ontology can silently pull in far more axioms, and far more reasoning cost, than the size of the imported file suggests.

Third, **owl:versionIRI exists specifically to keep reuse stable over time.**

- Importing against a version IRI, rather than the bare ontology IRI, protects a dependent ontology from silently inheriting un-reviewed changes the next time the imported ontology is updated -- the same discipline that version-pinning a software dependency provides, and for the same reason.

20.5.5 Design Patterns

The reuse mechanisms discussed so far -- *inheritance*, *equivalent classes*, *restriction templates*, and *imports* -- are OWL **language features**.

Design patterns are something different: they are **proven combinations** of those language features, applied to recurring modeling problems. The **Pizza** tutorial's Section 4.14 (page 44) introduces the idea directly, noting that OWL design patterns play the same role that patterns like **Model-View-Controller** play in object-oriented programming:

reusable shapes at a higher level of abstraction than any single class library.

Two design patterns in this section come straight from the tutorial itself, fully named and demonstrated:

Demonstrated in the tutorial:

- Enumerated Class pattern: **Exercise 24** -- **Spiciness** defined as {Hot, Medium, Mild}, exactly the pattern Section 4.14 introduces by name.
- Probe Class pattern: **Exercise 19**, which builds **ProbeInconsistentTopping** specifically to test whether **CheeseTopping**, and **VegetableTopping** were correctly declared disjoint. The tutorial names this explicitly: "*such a class is called a Probe Class*" -- a reusable diagnostic pattern, built, used once to confirm a design assumption, and then deleted, rather than kept as a permanent part of the ontology.

The ontology engineering community more broadly distinguishes two categories of pattern beyond these two:

- **Logical Ontology Design Patterns (Logical ODPs)** operate purely at the level of OWL constructs, domain-independent -- the Enumerated Class and Probe Class patterns above are both logical patterns.
- **Content Ontology Design Patterns (Content ODPs)** capture a recurring **domain modeling** shape, not just a logical structure -- for example, **Participation** (an agent taking part in an event) or **Time-Indexed Situation** (a fact that holds only during a specific interval), documented in a public catalog maintained by the Semantic Web research community.

Recognizing that a modeling problem matches a documented pattern is itself a form of reuse: it reuses not the axioms, but the **prior analysis** of how the problem should be structured.

An ontology engineer who recognizes a design pattern reuses not just a class or an axiom, but

an entire **way of thinking about a problem**.

20.5.6 Semantic Alignment

The final reuse mechanism in this section operates at the largest scale of all:

reuse **across ontologies that were never designed to work together**.

Semantic alignment is the practice of establishing formal correspondences between concepts in independently developed ontologies -- typically through properties such as `owl:equivalentClass`, `owl:equivalentProperty`, or `owl:sameAs` used **between** ontologies rather than within one.

```
pizza:VegetarianPizza owl:equivalentClass foodOntology:MeatFreeDish
```

[!Note] As with `owl:imports`, semantic alignment is not something the `Pizza` tutorial exercises demonstrate directly -- the tutorial works within a single, self-contained ontology throughout. The closest real-world illustration comes from outside the exercise set entirely: the tutorial's Section 9.22 mentions a companion example (built by DeBellis with Franz Inc.) that federates a SPARQL query across the Pizza ontology's sibling "`People`" ontology and external public knowledge graphs such as DBpedia and Geonames. That example is closer to federated querying than to formal `owl:equivalentClass` alignment, but it demonstrates the same underlying motivation: ***making independently built knowledge sources work together without merging them.***

`owl:equivalentClass` is the strongest alignment statement OWL offers, and that strength is not always appropriate.

Declaring two classes fully equivalent means the reasoner will treat every inference licensed by one class' axioms as equally licensed by the other's --

- **correct** when two teams have genuinely modeled the same concept,
- but actively **harmful** when the concepts merely overlap.

For this reason, ontology engineers reaching across independently governed ontologies frequently reach for a **deliberately weaker** vocabulary:

the mapping properties from **SKOS** (Simple Knowledge Organization System), a W3C standard built specifically for this situation.

```
pizza:VegetarianPizza skos:closeMatch foodOntology:MeatFreeDish
pizza:VegetarianPizza skos:broadMatch foodOntology:PlantBasedDish
```

`skos:exactMatch`, `skos:closeMatch`, `skos:broadMatch`, and `skos:narrowMatch` express degrees of correspondence without asserting the full logical machinery of `owl:equivalentClass`.

`owl:equivalentClass` says: **these two classes are logically the same thing.** `skos:closeMatch` says: **these two classes are close enough to be useful together, and no closer claim is being made.**

Choosing between them is itself a governance decision, not merely a technical one -- which is exactly why semantic alignment, of every reuse mechanism in this section, is the one most directly connected to Γ - **Governance**, previewed briefly here and developed further once this chapter reaches its EKA Perspective in Section 20.9.

Aligning two ontologies asserts that two independently maintained models agree on **meaning**, at a stated strength, and that agreement itself must be maintained, reviewed, and, where necessary, revised as either ontology evolves.

20.6 Reuse Across Disciplines - Mechanisms

Section 20.2 argued that reuse is a **universal engineering principle** -- present, in some form, in

- software,
- enterprise architecture,
- science,
- biology, and
- artificial intelligence.

Section 20.5 then catalogued the **specific mechanisms** through which OWL implements that principle:

- `SubClassOf` inheritance,
- `EquivalentTo` defined classes,
- reusable restriction templates,
- `owl:imports`,

- design patterns, and
- semantic alignment.

This section closes the loop between the two.

Rather than asking *why reuse* matters, each subsection below asks a narrower, more useful question:


what is the closest mechanical analog, in each discipline, to a specific OWL reuse mechanism from Section 20.5?

The mapping is not perfect -- no cross-disciplinary analogy ever is -- but it is close enough to be genuinely useful the next time you are deciding which OWL mechanism fits a modeling problem.

20.6.1 Software Engineering -- Inheritance and Composition

Software engineering achieves reuse through two primary architectural mechanisms: **inheritance** and **composition**.

Using PlantUML, you may illustrate in Class Diagram's Parent-Child inheritance relationship , with syntax `Parent <|-- Child1`, and composition relationship with syntax `Parent *-- Child2`:



In comparison:

- **Inheritance** establishes an **IS-A** relationship where a subclass inherits state and behavior from a superclass, specializing its capabilities while reusing its foundational design (analogous to OWL's `SubClassOf`).
 - Object-oriented inheritance is the closest software analog to `SubClassOf` ([Section 20.5.1](#)): a subclass inherits behavior from its parent without restating it, exactly as `VegetarianPizza` inherits every axiom that applies to `Pizza`.
- **Composition** establishes **HAS-A** relationship where complex objects are constructed by assembling independent, pre-tested component modules (analogous to OWL object properties and existential restrictions like `hasTopping some MozzarellaTopping`).
 - Composition -- assembling a new object out of several existing, independently reusable components -- is the closer analog to restriction templates ([Section 20.5.3](#)). A composed object does not inherit a single ancestor; it combines several independently defined pieces, the same way `hasTopping some (hasSpiciness value Hot)` composes a `hasValue` restriction inside an existential one.

Software engineers have a well-known preference between the two: "*favor composition over inheritance*", because deep inheritance hierarchies become brittle as a codebase grow.

Ontology engineers face a version of the same tension -- a taxonomy that grows too deep, or restriction expressions that grow too nest, become just as hard to maintain as a ten-level class hierarchy in code.

20.6.2 Enterprise Architecture -- Capability Reuse and Reference Architecture

Section 20.2.2 introduced "*Reuse before Buy before Build*" as an EA principle. Its mechanical implementation is the **reference architecture**:

a vetted, published design that other teams instantiate rather than redesign.

This is the closest EA analog to Design Patterns ([Section 20.5.5](#)). Just as an Enumerated Class or a Probe Class is a proven shape an ontology engineer reaches for rather than reinvents, a reference architecture is a proven shape an enterprise architect reaches for rather than redesigns.

`owl:imports` ([Section 20.5.4](#)) has an EA analog too: the **shared capability catalog**, maintained centrally and referenced -- not copied -- by every project that depends on it.

The version-pinning discipline urged for `owl:versionIRI` has a direct EA counterpart: capability catalogs that change without versioning break every downstream project that depends on them exactly as an unpinned `owl:imports` does.

20.6.3 Scientific Research -- Citation and Systematic Review

A citation is the closest research analog to `owl:equivalentClass` and the SKOS mapping properties from Semantic Alignment ([Section 20.5.6](#)):

it formally asserts that one piece of work relates to another, at a stated strength, without merging the two papers into one.

The analogy runs deeper than it first appears.

- A citation that says "*this result confirms X*" behaves like `owl:equivalentClass` -- a strong, direct claim.
- A citation that says "*this result is consistent with X*" behaves more like `skos:closeMatch` -- a weaker, ore cautious correspondence.

Researchers choose between these citation strengths for exactly the reason [Section 20.5.6](#) gave for choosing between `owl:equivalentClass` and `skos:closeMatch`:

overclaiming alignment is a scholarly integrity problem in research, just as it is a reasoning-integrity problem in OWL.

A **systematic review and meta-analysis** -- which synthesizes many independent studies into one structured account -- is the closer analog to Design Patters:

it does not create new findings, but documents a **reusable shape** for how existing findings fit together.

20.6.4 Biology -- Gene Duplication and Regulatory Review

Gene duplication is the closest biological analog to `Edit > Convert to defined class` from [Section 20.5.2](#):

a sequence that already works is copied, and the copy is then free to accumulate mutations under **relaxed selective pressure**, exactly as cloning `MargheritaPizza` into `AmericanaPizza` in **Exercise 16** preserved a working definition while introducing one deliberate variation.

The disjointness checking behind **Chapter (07)** and **Exercise 19**'s Probe Class pattern ([Section 20.5.5](#)) has a biological counterpart too:

regulatory review -- the network of biochemical checks that prevent a mutated gene from producing a nonviable organism. A reasoner flagging `ProbeInconsistentTopping` as equivalent to `owl:Nothing` is doing, at the level of logic, roughly what a regulatory checkpoint does at the level of biochemistry: catching an invalid combination before it propagates.

20.6.5 Artificial Intelligence -- Transfer Learning and Model Adaptation

Transfer learning -- taking a model trained on one task and adapting it to a related one, rather than training from nothing -- is the closest AI analog to `SubClassOf` inheritance ([Section 20.5.1](#)):

the adapted model does not discard what the base model already knows, and only the delta needs to be learned.

Fine-tuning and **model adaptation** are closer analogs to [EquivalentTo](#) defined classes ([Section 20.5.2](#)):

where transfer learning reuses a starting point, fine-tuning changes what the model is **defined to do**, the same way converting [CheesyPizza](#) from a primitive class to a defined class in **Exercise 21** changed what the class was, not merely where it sat in the hierarchy.

Retrieval-Augmented Generation (RAG), mentioned briefly in [Section 20.2.5](#), is the closest AI analog to [Semantic Alignment](#) ([Section 20.5.6](#)):

it does not merge external knowledge into the model's weights, but references it at query time, at a started degree of trust -- the same posture `skos:closeMatch` takes toward an external ontology it does not fully absorb.

20.7 Mathematical Perspective -- Abstraction, Generalization and Reuse

While software patters, enterprise architecture frameworks, and biological analogies provide **intuitive models** for reuse, the foundational guarantees of semantic knowledge reuse rest upon formal mathematics.

Description Logics (DL), set theory, and lattice theory provide the rigorous framework that explains **why** semantic reuse is logically sound, safe, and computationally tractable.

This section examines the mathematical anatomy of **abstraction**, **generalization**, and **structural reuse** within ontological models.

20.7.1 Generalization as Subset Inclusion

In model-theoretic semantics, an ontology interpretation $\mathcal{I} = (\Delta^{\mathcal{I}}, \cdot^{\mathcal{I}})$ consists of a non-empty domain $\Delta^{\mathcal{I}}$ and an interpretation function $\cdot^{\mathcal{I}}$ that maps:

- Concept names C to subsets of the domain ($C^{\mathcal{I}} \subseteq \Delta^{\mathcal{I}}$)
- Role name R to binary relations ($R^{\mathcal{I}} \subseteq \Delta^{\mathcal{I}} \times \Delta^{\mathcal{I}}$)

Generalization and **inheritance** correspond directly to subset inclusion.

When a class C is declared a subclass of D via the axiom $C \sqsubseteq D$, the semantic constraint enforced by the interpretation is:

$$C^{\mathcal{I}} \subseteq D^{\mathcal{I}}$$

Reusing a semantic definition through inheritance meaning that every individual that satisfies the more specialized concept C automatically inherits all membership criteria, property restrictions, and logical consequences of the generalized concept D .

Mathematically, any property holding for all elements in $D^{\mathcal{I}}$ trivially holds for all elements in $C^{\mathcal{I}}$ because $C^{\mathcal{I}}$ is embedded within $D^{\mathcal{I}}$.

When we created [AmericanaPizza](#) by cloning [MargheritaPizza](#) and adding `hasTopping some PepperoniTopping` in **Exercise 16**, the mathematical operation restricted the set extension:

$$\text{AmericanaPizza}^{\mathcal{I}} = \text{MargheritaPizza}^{\mathcal{I}} \cap \{x \in \Delta^{\mathcal{I}} \mid \exists y : (X, y) \in \text{hasTopping}^{\mathcal{I}} \wedge y \in \text{PepperoniTopping}^{\mathcal{I}}\}$$

Thus, reuse via specialization is formally a **constrained intersection** over an inherited base set.

20.7.2 Abstraction and Dimension Reduction

Abstraction in ontology engineering can be viewed as moving from a high-dimensional semantic space (characterized by numerous specific role restrictions and value assignments) to a lower-dimensional conceptual projection that capture

only essential shared characteristics.

Let an individual or class be characterized by a vector of semantic features or role restrictions $x = (r_1, r_2, \dots, r_n)$, where each r_i represents an existential or universal restriction along a specific property. An abstraction function f maps this detailed feature space onto a coarser, reusable template:

$$f : \mathbb{R}^n \rightarrow \mathbb{R}^k \quad \text{where } k < n$$

In Description Logic terms, dropping specific existential or value restrictions (for example, moving from a fully specified pizza back to the generic **Pizza** concept) broadens the concept extension:

$$\Delta^{\mathcal{I}} \supseteq D^{\mathcal{I}} \supseteq C^{\mathcal{I}}$$

This upward movement in the hierarchy discards local variation while retaining structural invariants. By identifying these invariants, ontology engineers create reusable patterns that can be instantiated across multiple distinct sub-domains without duplicating the underlying structural logic.

20.7.3 Posets and Lattices

The taxonomic backbone of an ontology forms a **Partially Ordered Set (Poset)**.

Let $\mathcal{T}$ be the set of all classes in an ontology, and let $\sqsubseteq$ denote the subsumption relation.

The pair $(\mathcal{T}, \sqsubseteq)$ forms a poset because the relation is:

1. **Reflexive**: For every class C , $C \sqsubseteq C$.
2. **Anti-symmetric**: If $C \sqsubseteq D$ and $D \sqsubseteq C$, then $C \equiv D$ (the classes are semantically equivalent).
3. **Transitive**: If $C \sqsubseteq D$ and $D \sqsubseteq E$, then $C \sqsubseteq E$.

When we introduce defined classes and class expressions using logical intersections ($\sqcap$), unions ($\sqcup$), and complements ($\neg$), the structure extends from a mere poset into a bounded **lattice** (specifically, a **Heyting** algebra or **Boolean** lattice depending on the expressivity of the underlying DL fragment, such as **SHIQ** or **SROIQ**).

In a lattice, any two concepts C and D possess a unique:

- **Infimum** (Greatest Lower Bound), corresponding to $C \sqcap D$ (the most specific common specialization).
- **Supremum** (Least Upper Bound), corresponding to $C \sqcup D$ (the most general common abstraction)

These algebraic operations provide the mathematical guarantee that when we combine or reuse semantic patterns, a unique, logically optimal position exists within the hierarchy for the resulting concept, ensuring that the reasoner can consistently position and classify new assets.

20.7.4 Reuse and Logical Consistency

A critical concern in knowledge reuse is whether incorporating existing models and axioms preserves the **logical consistency** of the overall knowledge base.

Let $\mathcal{O}_1$ be an established, verified ontology (such as the foundational **Pizza** ontology), and let $\mathcal{O}_2$ be a set of new axioms or imported modules representing specialized knowledge.

The combined ontology is represented as:

$$\mathcal{O} = \mathcal{O}_1 \cup \mathcal{O}_2$$

Because standard Description Logics are **monotonic** logics, adding new axioms can never invalidate existing entailments derived from $\mathcal{O}_1$ alone:

$$\text{If } \mathcal{O}_1 \models \phi, \text{ then } \mathcal{O}_1 \cup \mathcal{O}_2 \models \phi$$

However, monotonicity does *NOT* imply risk-free consistency.

If $\mathcal{O}_2$ introduces axioms that contradict $\mathcal{O}_1$ (e.g., asserting that `MargheritaPizza` is disjoint from `NamedPizza`, or creating an unsatisfiable class where $C^I = \emptyset$), then combined ontology becomes **inconsistent** ($\mathcal{O} \models \perp$).

Mathematically, safe knowledge reuse requires verifying that the extension of new concepts does not collapse into the empty set under the constraints of the inherited model:

$$C^I \neq \emptyset$$

This formal requirement underpins why automated reasoners (such as Pellet, HermiT, or FcCT++) are indispensable in **Stage 3**:

they execute consistency-checking algorithms (such as tableau expansion algorithms) to mathematically guarantee that reused and extended knowledge remains logically coherent before it is deployed into production systems.

20.8 Interesting Reading -- From Plato to Design Patterns

20.9 EKA Perspective -- Stage 3 Strengthens K and Prepares Γ

20.10 Engineering Guidelines for Knowledge Reuse

20.11 Key Concepts

20.12 Chapter Summary

20.13 Looking Ahead -- Knowledge Governance

Last updated at: 2026-08-18

Appendix A -- Chapter Mapping with Exercises

Chapter	Exercise	Tutorial Section	Video
Chapter 00~03	N/A	N/A	N/A
Chapter 04	Exercise 1, Exercise 2, Exercise 3	Section 4.0	04
Chapter 05	Exercise 4	Section 4.1	05
Chapter 06	Exercise 5	Section 4.2	06
Chapter 07	Exercise 6	Section 4.3	07
Chapter 08	N/A	N/A	
Chapter 09	Exercise 7	Section 4.4	09
Chapter 09	Exercise 8	Section 4.5	09
Chapter 10	Exercise 9	Section 4.6	10
Chapter 11	Exercise 10	Section 4.7	11
Chapter 12	(Added by Author)	N/A	
Chapter 13	Exercise 11, Exercise 12	Section 4.9	13
Chapter 14	Exercise 13	Section 4.10.2	14
Chapter 15	N/A (continues Chapter 14 topic)	14	
Chapter 16	N/A (continues Chapter 14 topic)	14	
Chapter 17	N/A	N/A	
Chapter 18	Exercise 14	Section 4.10.3	15
Chapter 19	Exercise 15	Section 4.10.3	16
Chapter 20	Exercise 16	Section 4.10.3	17
Chapter 21	Exercise 17	Section 4.10.3	18
Chapter 21	Exercise 18	Section 4.10.3	19
Chapter 22	Exercise 19	Section 4.10.4	20
Chapter 23	Exercise 20	Section 4.11	21
	Exercise 21	Section 4.11	
	Exercise 22	Section 4.12	
	Exercise 23	Section 4.13	
	Exercise 24	Section 4.14	

Chapter	Exercise	Tutorial Section	Video
	Exercise 25	Section 4.15	
	Exercise 26	Section 4.15	
	Exercise 27	Section 5.1	
	Exercise 28, 29, 30, 31, 32	Section 5.2	
	Exercise 33	Section 9.1	
	Exercise 34, 35, 36	Section 10.0	

Last Updated at: 2026-08-15

Contributors

Name / GitHub Handle	Contribution Type	Issue/Discussion	Chapter(s)	Date	Memo
LaraAcuna	Question / Content Improvement	#26	13	2026-08-12	Raised a thoughtful question about modeling OR semantics in Domain and Range (Chapter 13). Her question prompted the addition of a complete guide to owl:unionOf , separate properties, and local restriction patterns, including new screenshots, RDF examples, and a comparison table. This contribution significantly improved the practical guidance of Chapter 13. Fixed in v0.44.0.
Michael DeBellis	Refinement Suggestion	#19	15	2026-07-01	Proposed revised definition for VegetarianPizza pattern to align with common real-world interpretation (include cheese/dairy). Pending.
Michael DeBellis	Terminology Clarification	#20	Multiple	2026-07-01	Highlighted need to clearly distinguish "Annotation Properties" (reasoner-invisible) from logical property axioms. Pending.
Michael DeBellis	Critical Distinction / Error Report	#21	00, 08	2026-07-01	Identified fundamental conflation between Property Graphs (e.g., Neo4j) and Native RDF Graphs (e.g., AllegroGraph, GraphDB, Stardog). Major architectural correction required. Pending.
Michael DeBellis	Question / Technical Review	#9	00, 06, 14	2026-06-13	Clarified reasoner inference materialization distinction (Protégé UI behavior vs.

Name / GitHub Handle	Contribution Type	Issue/Discussion	Chapter(s)	Date	Memo
					production triplestores). Fixed in v0.39.0.
Michael DeBellis	Error Report	#10, #11, #12, #13, #14, #15	Various	2026-06-16	Multiple error reports including OWL modeling errors, contributor attribution, proofreading, and other corrections. See individual issues for details.
mlungsta89	Bug Report	#3	—	2025-09-15	Reported reasoner getting stuck issue. Closed.
nikokaoya	Feedback	#2	—	2024-06-26	Positive feedback and encouragement. Closed.
Your name here	Error report / Pull request / Question	—	X	YYYY-MM-DD	Brief Memo

Last updated: 2026-08-12

Error Tracker

Error E-001 | Severity: Must Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

The EKA table in Chapter 00 claims that an OWL reasoner "does not modify a graph or database," labeling it "No (read-only inference)." This is overly broad. While true for the normal Protégé workflow (where inferred axioms are displayed rather than asserted into the saved ontology), it is not true as a general statement about semantic systems or graph databases. Triplestores may materialize inferred triples, maintain inferred closures, or expose inferences virtually depending on configuration.

Suggested Fix:

Revise to: "In Protégé, reasoner results are normally shown as inferred axioms rather than automatically saved as asserted axioms. In graph databases and rule-enabled triple stores, inferred triples may be materialized, persisted, or exposed virtually depending on the system and configuration."

Status: Pending | Chapters: 00, 14 | Target Version: v0.37.0

Error E-002 | Severity: Must Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

The current list of graph database and knowledge graph tools (Neo4j, GraphDB, Stardog, RDFox) omits AllegroGraph by Franz Inc. AllegroGraph is one of the most powerful RDF graph databases, with significant advantages in performance, security, and system integration. The original Pizza.owl tutorial explicitly acknowledges Franz/AllegroGraph/Gruff and uses Gruff in the SHACL discussion. Additionally, Amazon Neptune should be added as it natively supports both Property Graph and RDF graph models.

Suggested Fix:

Add AllegroGraph by Franz Inc. and Amazon Neptune wherever graph databases and knowledge graph tools are discussed. Clarify that OWL ontologies can be loaded into RDF triplestores (AllegroGraph, GraphDB, Stardog, RDFox, Amazon Neptune, Apache Jena/Fuseki) and can also be transformed into property graph systems such as Neo4j, but doing so requires choices about how much OWL semantics, reasoning behavior, and RDF identity structure to preserve.

Status: Pending | Chapters: 08, 15.3 | Target Version: v0.37.0

Error E-003 | Severity: Must Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

The VegetarianPizza definition contains serious OWL modeling errors. The book uses `VegetarianPizza ≡ Pizza AND NOT (hasTopping some MeatTopping)`, but this is not equivalent to the tutorial's definition—seafood

toppings are separate from MeatTopping, so "not meat" would still allow seafood. More critically, the hands-on inconsistency example defines VegetarianPizza as hasTopping only VegetableTopping, which omits both Pizza and CheeseTopping and, due to vacuous truth, any individual with no toppings automatically satisfies this definition, leading to strange over-inferences.

Suggested Fix:

Use the original tutorial definition: VegetarianPizza $\equiv$ Pizza and (hasTopping only (CheeseTopping or VegetableTopping)). Then separately explain that named pizzas need closure axioms such as hasTopping only (MozzarellaTopping or TomatoTopping) before the reasoner can classify them as vegetarian under OWA.

Status: Pending | Chapters: 15.1, 15.2 | Target Version: v0.37.0

Error E-004 | Severity: Suggested Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

The definition of "executable semantics" is somewhat vague. Chapter 00 provides a useful distinction: EKA requires triggers and actions; without those, a system is a semantic repository, not full EKA. But elsewhere the book says things like "the ontology stops being passive documentation and becomes an executable semantic system" simply because reasoning is involved.

Suggested Fix:

Distinguish between "reasoning-enabled" and "executable" throughout the book. Reserve "executable" for cases where Θ (triggers) and Φ (execution actions) are actually present.

Status: Pending | Chapters: 00, 01, 16.1 | Target Version: v0.37.0

Error E-005 | Severity: Must Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

The eBook does not include an explicit license statement. Since the original tutorial is under Creative Commons Attribution-ShareAlike 4.0 (CC BY-SA 4.0), a derivative work should clearly state the license and preserve compatible ShareAlike terms.

Suggested Fix:

Add clear CC BY-SA 4.0 license statements to the GitHub repository, eBook file metadata (dcterms:license), and the documentation itself. Being explicit and redundant with license information is good open-source practice.

Status: Pending | Chapters: All | Target Version: v0.37.0

Error E-006 | Severity: Must Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

The original authors' contributor chain has not been preserved. Michael DeBellis's tutorial states on its first page: "...this version is a revised version of the Protégé 4 Tutorial by Matthew Horridge," and lists previous contributors: Holger Knublauch, Alan Rector, Robert Stevens, Chris Wroe, Simon Jupp, Georgina Moulton, Nick Drummond, and Sebastian Brandt. These names do not appear in the eBook materials.

Suggested Fix:

Add the complete contributor chain in the Acknowledgments section and/or a dedicated Contributors page. Ensure Matthew Horridge and all listed previous contributors receive proper attribution.

Status: Pending | Chapters: Acknowledgments, Contributors | Target Version: v0.37.0

Error E-007 | Severity: Suggested Fix | Reporter: Michael DeBellis | Date: 2026-06-16

Issue Description:

There are numerous visible typos and awkward expressions throughout the book, including but not limited to: "Classis," "Mearning," "Pallet reasonder," "EAK," "ekecutable," "Standford," "Courese," "concsistency," and so on.

Suggested Fix:

Run the entire manuscript through an LLM or professional proofreading tool and generate a fix report. Consider using spell-checking tools or asking Michael to assist with proofreading.

Status: Pending | Chapters: All | Target Version: v0.37.0

Last updated at 2026-08-08

Error Found: A Living Document Invitation

- [Chapter Introduction](#)
- [What Kinds of Errors Might You Find?](#)
- [How to Report an Error](#)
 - [Step 1: Verify the Error](#)
 - [Step 2: Choose Your Contribution Method](#)
 - [Step 3: Use the Error Report Template](#)
- [Error Report](#)
 - [Step 4: A Note on Tone](#)
- [What Happens After Your Report?](#)
 - [For Issues](#)
 - [For Pull Requests](#)
 - [Versioning and Updates](#)
- [Contributors](#)
- [The Spirit of Open Technical Writing](#)
- [Quick Reference: Error Reporting Links](#)
- [Final Request](#)

Chapter Introduction

No technical book is perfect -- including this one.

Despite countless hours of writing, reviewing, testing, and refining, errors inevitably remain still. Some may be typographical. Others may be technical inaccuracies. A few might be conceptual oversights or outdated patterns that better approaches have since clarified.

This chapter exists for one reason:

to invite you, the reader, to help make this book better!

Ontology engineering is a living discipline. Semantic technologies evolve. Best Practices mature. New tools introduce. And with your help, this book can evolve alongside them.

This chapter serves three purposes:

1. **Acknowledging imperfection** -- No author has all the answers. I have done my best, but I'm certain errors remain.
2. **Providing a structured error reporting process** -- So you know exactly how to contribute.
3. **Celebrating community contributions** -- Because open knowledge grows stronger when many minds refine it together.

What Kinds of Errors Might You Find?

Before you begin searching ("hunting"), let me help you understand what to look for.

Errors in a book could like this typically fall into below categories:

Error Category	Description	Sample Types
Technical Errors	These are factual mistakes about OWL, Protégé, RDF, reasoning, or the EKA framework itself	• Incorrect OWL syntax • Misleading logical statements • Reasoning behavior errors
Conceptual Errors	These are misunderstandings or oversimplifications that might mislead readers about deeper semantic principles.	• Oversimplified distinctions • Missing nuance in OWL • Confusing concepts
Typographical and Grammatical Errors	These are the "small stuff" -- but small fixes improve readability and professionalism.	• Misspellings • Missing words • Inconsistent terminology
Outdated or Deprecated Patterns	Ontology engineering tools and standards evolve. What was best practice in 2025 may be outdated by 2028.	• Deprecated reasoner features • Outdated plugin names • Evolving OWL 2 features •
Structural and Pedagogical Issues	These are not "errors" per se, but opportunities to improve learning effectiveness.	• Unclear explanations • • Missing prerequisites • Exercise ambiguity • Broken cross-reference

How to Report an Error

Great, you catch one 😊! Your contribution are sincerely appreciated.

Please follow this structured process so I can address your feedback efficiently.

Step 1: Verify the Error

Before reporting, please:

1. **Check the latest version** -- The error may have already been fixed in a new commit. Visit the repository: https://github.com/yasenstar/protege_pizza

2. **Re-read the relevant section carefully** -- Ensure you haven't misinterpreted the text. OWL semantics can be counter-intuitive, especially around *Open World Assumption*.
3. **Test in Protégé** -- If the error is technical, actually try it in Protégé with a reasoner running. Your real-world testing is invaluable.

Step 2: Choose Your Contribution Method

You have two options:

Method	Best for	Link
GitHub Issue	Reporting errors without proposing specific fixes	Open an issue →
Pull Request	Proposing direct corrections (typos, code examples, clarifications)	Open a Pull Request →

Step 3: Use the Error Report Template

For Issues (no fix proposed), I've configured this template in repository already:

=== Template Start===

Error Report

Chapter: [e.g., Chapter 14]

Section: [e.g., 14.5.2]

Page/Line: [if applicable]

Error Type: (Technical / Conceptual / Typographical / Outdated / Pedagogical)

Current Text:

[Quote the exact text that contains the error]

Suggested Correction:

[Describe what should replace it, or why it is wrong]

Reasoning:

[Explain why this is an error — cite OWL specification, reasoner behavior, or logical principles]

Additional Context:

[Protégé version, reasoner used, screenshots if helpful]

=== Template End===

For Pull Request (fix proposed):

If you are comfortable with Git and Markdown, welcome directly edit the `.md` file in the `ebook/` directory.

Your pull request should:

1. Target the **main** branch
2. Describe what you changed and why
3. Reference any related issue (if applicable)

I will review all pull requests as quickly as possible.

Step 4: A Note on Tone

All feedback is welcome and appreciated -- including critical feedback. However, please follow the **principle of charitable interpretation**.

- Assume good faith.
- Assume the error was accidental, not intentional.
- Assume I genuinely want to fix it.

Rude or dismissive reports will be closed. Thoughtful, respectful reports will be treasured.

What Happens After Your Report?

For Issues

1. **Acknowledgment** -- I will respond within 7 days (try my best) confirming receipt.
2. **Validation** -- I will verify the error independently (re-testing in Protégé, reviewing OWL specifications, or any other tools).
3. **Resolution** -- Depending on severity, I will either:
 - Fix immediately (for minor typos or clear technical errors)
 - Schedule for next revision (for conceptual clarifications or structural improvements)
 - Mark as "won't fix" with explanation (if the report is incorrect or based on misunderstanding)

For Pull Requests

1. **Review** -- I will test your proposed changes locally.
2. **Feedback** -- If adjustments are needed, I will comment inline.
3. **Merge** -- Once approved, your contribution becomes part of the book.
4. **Attribution** -- Your name (or GitHub handle) will be added to the **Contributors** section below.

Versioning and Updates

The book is a living document. When errors are fixed:

- A new commit will be pushed to the repository
- The PDF and EPUB will be regenerated (periodically)
- A **CHANGELOG.md** will track significant corrections

You can always find the latest version at **ebook** folder of https://github.com/yasenstar/protege_pizza.

Contributors

The following individuals have contributed error reports, corrections, or improvements to this book. (also in the dedicated contributors file)

Name / GitHub Handle	Contribution Type	Chapter(s)	Date
<i>Your name here</i>	<i>Error report / Pull request</i>	<i>X</i>	<i>YYYY-MM-DD</i>

This space is reserved for you.

When you report an error that leads to a correction, your name will be added here with my sincere thanks.

The Spirit of Open Technical Writing

This chapter exists because I believe in a specific philosophy of technical education:

Books are not monuments. Books are conversations.

When you read a chapter and find an error, you are not "criticizing" the author. You are **contributing the work** of making knowledge clearer, more accurate, and more useful.

Every technical book ever written contains errors. The only question is whether those errors remain hidden or become opportunities for improvement.

By opening this chapter, by inviting you to report errors, by publishing the entire book on GitHub (and fine formatted to LeanPub), I am choosing the second path.

This book is not finished. It will never be finished. And that is exactly as it should be.

Because ontology engineering itself is a living discipline. The Semantic Web evolves continuously. Protégé development keep updates. Reasoners improves. New enterprise challenges emerge. And this book will evolve alongside them -- with your help!

You are not just a reader. You are a co-creator.

Quick Reference: Error Reporting Links

Action	Link
Report an error (no fix)	Open an Issue →
Propose a correction	Open a Pull Request →
View the latest version	Repository Main →
Browse the eBook directory	ebook/ →

Final Request

Before you close this book and move on with your day, please consider:

Did you find anything in this book that confused you?

That confusion -- even if not a "technical error" -- is worth reporting. Unclear explanations are a kind of error too. If something took you three readings to understand, that section could be improved.

Did you find a type, a missing word, or an inconsistent term?

Those are trivial to fix and make the book more professional. Please report them.

Did you try an example in Protégé and get a different result than the book described?

That is the most valuable report of all. Real-world testing trumps theoretical writing every time.

Thank you for reading. Thank you for learning. And thank you, in advance, for helping make this book better.

Last updated: 2026-08-08

Repository: https://github.com/yasenstar/protege_pizza

Author contact: GitHub Issues (preferred) or via YouTube channel

The best books are those that acknowledge they are never truly finished -- only temporarily improved!